

## Threads

Programs consist of single sequence of instructions. A sequence is called a **thread** (for "thread of control") in

Programs containing **multiple** threads, which (conceptually) run concurrently.

On a uniprocessor, only one thread at a time actually runs, but this is largely invisible.

For program access to threads, Java provides the type `Thread`. Each `Thread` contains information about, and controls,

Access to data from two threads can cause chaos, so Java constructs for controlled communication, allowing threads to **wait** to be notified of events, and to **interrupt** others.

56:28 2019

CS61B: Lecture #37 2

## Java Mechanics

Two actions "walking" and "chewing gum":

```
1 implements Runnable { // Walk and chew gum
    id run()              Thread clomp
    (true) ChewGum(); }   = new Thread(new
                        Chewer1());
1 implements Runnable { Thread clomp
    id run()              = new Thread(new
    (true) Walk(); }      Walker1());
                        clomp.start(); clomp.start();
```

Alternative (uses fact that `Thread` implements `Runnable`):

```
extends Thread {
    run()
    (true) ChewGum(); } Thread clomp = new Chewer2(),
                        clomp = new Walker2();
                        clomp.start(); clomp.start();

extends Thread {
    run()
    (true) Walk(); }
```

56:28 2019

CS61B: Lecture #37 4

## Communicating the Hard Way

Sharing data is tricky: the faster party must wait for the

Proxies for sending data from thread to thread don't

```
changer {
    lue = null;
    ceive() {
        r; r = null;
        (r == null)
        = value; }
    = null;
    r;

    sit(Object data) {
        (value != null) { }
        = data;

    DataExchanger exchanger
        = new DataExchanger();

    -----

    // thread1 sends to thread2 with
    exchanger.deposit("Hello!");

    -----

    // thread2 receives from thread1 with
    msg = (String) exchanger.receive();
```

Thread can monopolize machine while waiting; two threads deposit or receive simultaneously cause chaos.

56:28 2019

CS61B: Lecture #37 6

## Lecture #37

Excursions into nitty-gritty stuff: Threads, storage man-

56:28 2019

CS61B: Lecture #37 1

## But Why?

Programs always have  $> 1$  thread: besides the main thread, others clean up garbage objects, receive signals, update other stuff.

Systems deal with asynchronous events, is sometimes convenient to organize into subprograms, one for each independent, receive of events.

How to insulate one such subprogram from another.

Organized like this: application is doing some computation, another thread waits for mouse clicks (like 'Stop'), pays attention to updating the screen as needed.

Search engines like search engines may be organized this way, with server request.

Otherwise, sometimes we **do** have a real multiprocessor.

56:28 2019

CS61B: Lecture #37 3

## Avoiding Interference

Thread has data for another, one must wait for the other

Two threads use the same data structure, generally only modify it at a time; other must wait.

Could happen if two threads simultaneously inserted an unlinked list at the same point in the list?

Could conceivably execute

```
new ListCell(x, p.next);
```

Update values of `p` and `p.next`; one insertion is lost.

For only one thread at a time to execute a method on an object with either of the following equivalent definitions:

```
) {
    synchronized (this) {
        f f


    synchronized void f(...) {
        body of f
    }
```

56:28 2019

CS61B: Lecture #37 5

## Message-Passing Style

primitives very error-prone. Wait until CS162.  
re higher-level, and allow the following program struc-



Player is a thread that looks like this:

```
ameOver() {  
    ve()  
    .deposit(computeMyMove(lastMove));  
  
    ve = inBox.receive();  
}
```

56:28 2019

CS61B: Lecture #37 8

## Coroutines

is a kind of synchronous thread that explicitly hands  
o other coroutines so that only one executes at a time,  
generators. Can get similar effect with threads and

ursive inorder tree iterator:

```
or extends Thread {  
    lbox r;  
    ree T, Mailbox r) {  
T; this.dest = r;    void treeProcessor(Tree T) {  
                        Mailbox m = new QueuedMailbox();  
                        new TreeIterator(T, m).start();  
                        while (true) {  
                            Object x = m.receive();  
                            if (x is end marker)  
                                break;  
                            do something with x;  
                        }  
                    }  
    Tree t) {  
        l return;  
        eft);  
        label);  
        ight);  
    }
```

56:28 2019

CS61B: Lecture #37 10

## Highlights of a GUI Component

```
hat draws multi-colored lines indicated by mouse. */  
tends JComponent implements MouseListener {  
    <Point> lines = new ArrayList<Point>();  
  
    / Main thread calls this to create one  
    redSize(new Dimension(400, 400));  
    stener(this);  
  
    ronized void paintComponent(Graphics g) { // Paint thread  
        (Color.white);    g.fillRect(0, 0, 400, 400);  
        x = y = 200;  
        Color.black;  
        p : lines)  
        or(c); c = chooseNextColor(c);  
        ne(x, y, p.x, p.y); x = p.x; y = p.y;  
  
    ronized void mouseClicked(MouseEvent e) // Event thread  
        d(new Point(e.getX(), e.getY())); repaint(); }
```

56:28 2019

CS61B: Lecture #37 12

## Primitive Java Facilities

on Object makes thread wait (not using processor) un-  
y notifyAll, unlocking the Object while it waits.

.util.mailbox has something like this (simplified):

```
Mailbox {  
    posit(Object msg) throws InterruptedException;  
    eceive() throws InterruptedException;  
  
    edMailbox implements Mailbox {  
        List<Object> queue = new LinkedList<Object>();  
  
        ynchronized void deposit(Object msg) {  
            add(msg);  
            otifyAll(); // Wake any waiting receivers  
        }  
  
        ynchronized Object receive() throws InterruptedException {  
            (queue.isEmpty()) wait();  
            queue.remove(0);  
        }  
    }  
}
```

56:28 2019

CS61B: Lecture #37 7

## More Concurrency

imple can be done other ways, but mechanism is very

you want to think during opponent's move:

```
ameOver() {  
    ve()  
    .deposit(computeMyMove(lastMove));  
  
    kAheadALittle();  
    Move = inBox.receiveIfPossible();  
    e (lastMove == null);  
  
    ssible(written receive(0) in our actual package) doesn't  
    ; null if no message yet, perhaps like this:  
  
    ynchronized Object receiveIfPossible()  
    InterruptedException {  
        e.isEmpty()  
        null;  
        ueue.remove(0);  
    }
```

56:28 2019

CS61B: Lecture #37 9

## Use In GUIs

e library uses a special thread that does nothing but  
nts like mouse clicks, pressed keys, mouse movement,

ignate an object of your choice as a *listener*; which  
Java's event thread calls a method of that object when-  
t occurs.

your program can do work while the GUI continues to  
uttons, menus, etc.

cial thread does all the drawing. You don't have to be  
his takes place; just ask that the thread wake up when-  
nge something.

56:28 2019

CS61B: Lecture #37 11

## Note Mailboxes (A Side Excursion)

The Method Interface allows one program to refer to objects in another program.

To allow mailboxes in one program to be received from or to in another.

You define an *interface* to the remote object:

```
import java.rmi.*;
interface Mailbox extends Remote {
    void deposit(Object msg)
        throws InterruptedException, RemoteException;
    Object receive()
        throws InterruptedException, RemoteException;
}
```

That actually will contain the object, you define

```
class RemoteMailbox ... implements Mailbox {
    // implementation as before, roughly
}
```

## Scope and Lifetime

Declaration is portion of program text to which it applies

Can be contiguous.

Can be static: independent of data.

*Extent* of storage is portion of program execution during which it exists.

Can be contiguous

Can be dynamic: depends on data

Can be static: extent:

Can be static: lifetime duration of program

*Automatic*: duration of call or block execution (local variables)

From time of allocation statement (*new*) to deallocation.

## Under the Hood: Allocation

References (references) are represented as integer addresses.

How to convert integers to machine's own practice.

How to convert integers ↔ pointers,

Parts of Java runtime implemented in C, or sometimes in assembly, where you can.

How to do in C:

```
char* buf = new char[STORAGE_SIZE]; // Allocated array
int* p = buf; // pointer to anything
```

```
void* malloc(size_t n) { // void*: pointer to anything
    if (n < 0) return 0;
    int remainder = n % 8;
    int r = (remainder - n) & ~0x7; // Make multiple of 8
    return (void*) (store + remainder);
}
```

## Interrupts

An interrupt is an event that disrupts the normal flow of control of a program.

In systems, interrupts can be totally *asynchronous*, occurring at arbitrary points in a program, the Java developers considered that interrupts would occur only at controlled points.

In programs, one thread can interrupt another to inform it that it has unusual needs attention:

```
Thread t = ...;
t.interrupt();
```


A thread does not receive the interrupt until it waits: methods like `sleep` (wait for a period of time), `join` (wait for thread to finish), and `mailbox` deposit and receive.

When a thread uses these methods to throw `InterruptedException`, the thread is like this:

```
try {
    mailbox.receive();
} catch (InterruptedException e) { HandleEmergency(); }
```

## Remote Objects Under the Hood

```
// On Machine #1:
Mailbox m1 = new Mailbox();
// On Machine #2:
Mailbox m2 = getOutBoxFromMachine1(m1);
```



A mailbox is an interface, hides fact that on Machine #2, the mailbox actually has direct access to it.

Method calls are relayed by I/O to machine that has the mailbox.

When a method call or return type OK if it also implements `Remote` or `Serializable`—turned into stream of bytes and back, as can be for objects and `String`.

When a method call is involved, expect failures, hence every method can throw an exception (subtype of `IOException`).

## Explicit vs. Automatic Freeing

Explicit means to free dynamic storage.

When no expression in any thread can possibly be influenced by an object, it might as well not exist:

```
void free(Object o) { ... }
```

```
int* c = new int[3];
// ...
// c is now deallocated, so no way to first cell of list
```

But, Java runtime, like Scheme's, recycles the object with *garbage collection*.

## Explicit Deallocating

ly require explicit deallocation, because of  
in-time information about what is array  
of converting pointers to integers.

in-time information about **unions**:

```
Various {
Int;
* Pntr;
le Double;
// X is either an int, char*, or double
```

all three problems; automatic collection possible.

ing can be somewhat faster, but rather error-prone:

orruption

aks

56:28 2019

CS61B: Lecture #37 20

## Free List Strategies

ests generally come in multiple sizes.

ks on the free list are big enough, and one may have to  
chunk and break it up if too big.

egies to find a chunk that fits have been used:

**l fits:**

cks in LIFO or FIFO order, or sorted by address.

e adjacent blocks.

for **first fit** on list, **best fit** on list, or **next fit** on list  
st-chosen chunk.

**ed fits:** separate free lists for different chunk sizes.

**items:** A kind of segregated fit where some newly ad-  
ze blocks of one size are easily detected and combined  
r chunks.


locks reduces **fragmentation** of memory into lots of lit-  
d chunks.

56:28 2019


CS61B: Lecture #37 22

## rbage Collection: Mark and Sweep

atics)




1. Traverse and **mark** graph of objects.
2. **Sweep** through memory, freeing unmarked objects.



56:28 2019

CS61B: Lecture #37 24

## Example of Storage Layout: Unix



y to turn chunks of unallocated region into heap.  
omatically for stack.

56:28 2019

CS61B: Lecture #37 19

## Free Lists

cator grabs chunks of storage from OS and gives to

ycled storage, when available.


ge is freed, added to a **free list** data structure to be

or explicit freeing and some kinds of automatic garbage

variables  
(to program)

ne Heap

ree List

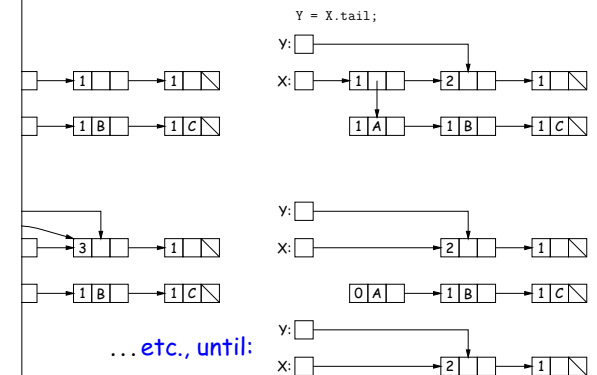


56:28 2019

CS61B: Lecture #37 21

## page Collection: Reference Counting

count of number of pointers to each object. Release  
oes to 0.



56:28 2019

CS61B: Lecture #37 23

## Copying Garbage Collection

reach: *copying garbage collection* takes time proportional to amount of active storage:

traverse the graph of active objects breadth first, *copying* them to a contiguous area (called "to-space").

Copy each object, mark it and put a *forwarding pointer* at the old location; it points to where you copied it.

Next time you have to copy an already marked object, just use the forwarding pointer instead.

When done, the space you copied from ("from-space") becomes to-space; in effect, all its objects are freed in constant time.

## Objects Die Young: Generational Collection

Most objects stay active, and need not be collected.

Need to avoid copying them over and over.

*Generational garbage collection* schemes have two (or more) spaces: one for newly created objects (*new space*) and one for objects that have survived garbage collection (*old space*).

Garbage collection collects only in new space, ignores pointers to old space, and moves objects to old space.

Has usual roots plus pointers in old space that have changed (might be pointing to new space).

When new space full, collect all spaces.


This leads to much smaller *pause times* in interactive systems.

## Cost of Mark-and-Sweep

Mark-and-sweep algorithms don't move any existing objects—pointers are updated.

Amount of work depends on the amount of memory swept—total amount of active (non-garbage) storage + amount of garbage. It's not necessarily a big hit: the garbage had to be active at some point, and hence there was always some "good" processing in the memory byte of garbage scanned.

## Copying Garbage Collection Illustrated



## There's Much More

Most highlights.

Learn how to implement these ideas efficiently.

*Generational garbage collection*: What if objects scattered over many spaces?

*Incremental garbage collection*: where predictable pause times are important, doing a little at a time.