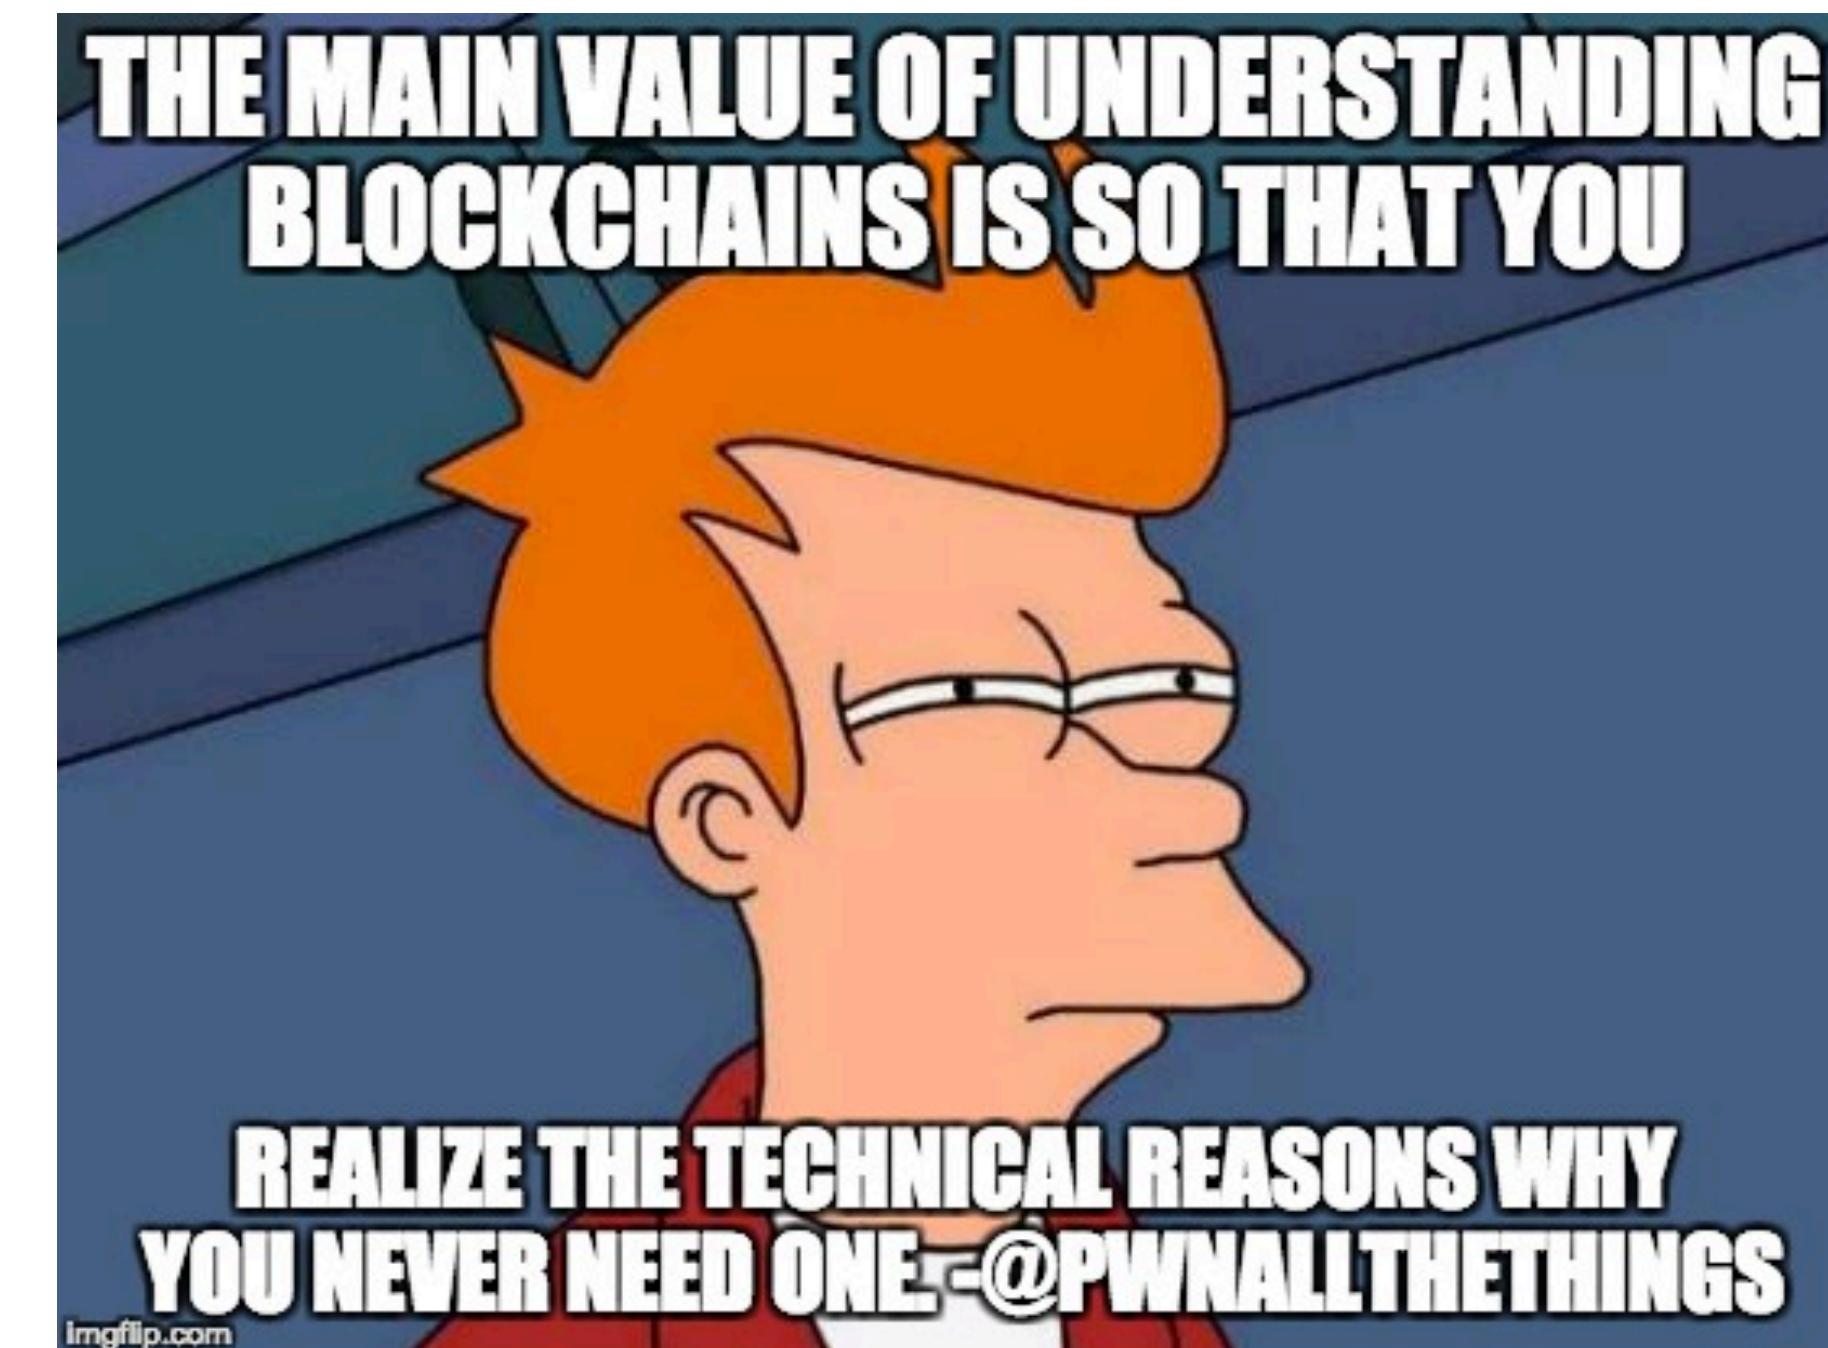


# Crypto 3



# Administrivia!

- Project 1 due Friday the 25th
- Homework 2 due Monday the 28th
- Reminder:
  - Zoom chat for conversation
  - Zoom Q&A for Questions & Answers

# Message Authentication Codes (MACs)

- Symmetric-key approach for integrity
  - Uses a shared (secret) key  $K$
- Goal: when Bob receives a message, can confidently determine it hasn't been altered
  - In addition, whomever sent it must have possessed  $K$   
( $\Rightarrow$  message authentication, sorta...)
- Conceptual approach:
  - Alice sends  $\{M, T\}$  to Bob, with tag  $T = \text{MAC}(K, M)$ 
    - Note,  $M$  could instead be  $C = E_K(M)$ , but not required
  - When Bob receives  $\{M', T'\}$ , Bob checks whether  $T' = \text{MAC}(K, M')$ 
    - If so, Bob concludes message untampered, came from Alice
    - If not, Bob discards message as tampered/corrupted

# Requirements for Secure MAC Functions

- Suppose MITM attacker Mallory intercepts Alice's  $\{M, T\}$  transmission ...
  - ... and wants to replace  $M$  with altered  $M^*$
  - ... but doesn't know shared secret key  $K$
- We have secure integrity if MAC function  $T = \text{MAC}(M, K)$  has two properties:
  - Mallory can't compute  $T^* = \text{MAC}(M^*, K)$ 
    - Otherwise, could send Bob  $\{M^*, T^*\}$  and fool him
  - Mallory can't find  $M^{**}$  such that  $\text{MAC}(M^{**}, K) = T$ 
    - Otherwise, could send Bob  $\{M^{**}, T\}$  and fool him
- These need to hold even if Mallory can observe many  $\{M_i, T_i\}$  pairs, including for  $M_i$ 's she chose

# MAC then Encrypt or Encrypt then MAC

- You should **never** use the same key for the MAC and the Encryption
  - Some MACs will break completely if you reuse the key
  - Even if it is **probably** safe (eg, AES for encryption, HMAC for MAC) its still a bad idea
- MAC then Encrypt:
  - Compute  $T = \text{MAC}(M, K_{\text{mac}})$ , send  $C = E(M||T, K_{\text{encrypt}})$
- Encrypt then MAC:
  - Compute  $C = E(M, K_{\text{encrypt}})$ ,  $T = \text{MAC}(M, K_{\text{mac}})$ , send  $C||T$
- Theoretically they are the same, but...
  - Once again, its time for ...



# HTTPS Authentication in Practice

- When you log into a web site, it sets a "cookie" in your browser
  - All subsequent requests include this cookie so the web server knows who you are
  - If an attacker can get your cookie...
    - They can impersonate you on the "Secure" site
  - And the attacker can create multiple tries
    - On a WiFi network, inject a bit of JavaScript that repeatedly connects to the site
    - While as a man-in-the-middle to manipulate connections



# The TLS 1.0 "Lucky13" Attack: "F-U, This is Cryptography"

- HTTPS/TLS uses MAC then Encrypt
  - With CBC encryption
- The Lucky13 attack changes the cipher text in an attempt to discover the state of a byte
  - But can't predict the MAC
  - The TLS connection retries after each failure so the attacker can try multiple times
    - Goal is to determine the status each byte in the authentication cookie which is in a known position
- It detects the ***timing*** of the error response
  - Which is different if the guess is right or wrong
    - Even though the underlying algorithm was "***proved***" secure!
- So always do Encrypt then MAC since, once again, it is more mistake tolerant



# The best MAC construction: HMAC

- Idea is to turn a hash function into a MAC
  - Since hash functions are often much faster than encryption
  - While still maintaining the properties of being a cryptographic hash
- Reduce/expand the key to a single hash block
- XOR the key with the i\_pad
  - 0x363636... (one hash block long)
- Hash ( $(K \oplus i\_pad) \parallel message$ )
- XOR the key with the o\_pad
  - 0x5c5c5c...
- Hash ( $(K \oplus o\_pad) \parallel first\ hash$ )

```
function hmac (key, message) {  
    if (length(key) > blocksize) {  
        key = hash(key)  
    }  
    while (length(key) < blocksize) {  
        key = key || 0x00  
    }  
    o_key_pad = 0x5c5c... ⊕ key  
    i_key_pad = 0x3636... ⊕ key  
    return hash(o_key_pad ||  
               hash(i_key_pad || message))  
}
```

# Why This Structure?

- `i_pad` and `o_pad` are slightly arbitrary
  - But it is necessary for security for the two values to be different
    - So for paranoia chose very different bit patterns
- Second hash prevents appending data
  - Otherwise attacker could add more to the message and the HMAC and it would still be a valid HMAC for the key
    - Wouldn't be a problem with the key at the *end* but at the start makes it easier to capture intermediate HMACs
- Is a Pseudo Random Function if the underlying hash is a PRF
  - AKA if you can break this, you can break the hash!

```
function hmac (key, message) {  
    if (length(key) > blocksize) {  
        key = hash(key)  
    }  
    while (length(key) < blocksize) {  
        key = key || 0x00  
    }  
    o_key_pad = 0x5c5c... ⊕ key  
    i_key_pad = 0x3636... ⊕ key  
    return hash(o_key_pad ||  
                hash(i_key_pad || message))  
}
```

# Great Properties of HMAC...

- It is still a hash function!
  - So all the good things of a cryptographic hash:  
An attacker or **even the recipient** shouldn't be able to calculate **M** given **HMAC(M,K)**
  - An attacker who doesn't know **K** can't even verify if **HMAC(M,K) == M**
    - Very different from the hash alone, and potentially very useful:  
Attacker can't even brute force try to find **M** based on **HMAC(M,K)**!
- Its probably safe if you screw up and use the same key for both MAC and Encrypt
  - Since it is a different algorithm than the encryption function...
  - **But you shouldn't do this anyway!**

# Considerations when using MACs

- Along with messages, can use for data at rest
  - E.g. laptop left in hotel, providing you don't store the key on the laptop
  - Can build an efficient data structure for this that doesn't require re-MAC'ing over entire disk image when just a few files change
- MACs in general provide no promise not to leak info about message
  - Compute MAC on ciphertext if this matters
  - Or just use HMAC, which **does** promise not to leak info if the underlying hash function doesn't
- **NEVER** use the same key for MAC and Encryption...
  - Known "FU-this-is-crypto" scenarios reusing an encryption key for MAC in some algorithms when its the same underlying block cipher for both



# Plus AEAD Encryption Modes...

- The latest block cipher modes are "AEAD":
  - Authenticated Encryption with Additional Data
- Provides both **integrity** *and* confidentiality over the data
  - With **integrity** also provided for the "Additional Data"
- Used right, these are great
  - Assuming you use a library...
- Used wrong...
  - The AEAD modes are built for "performance", which means parallelization, which means CTR mode, which means IV reuse is a disaster!

# A Lot Of Uses for Random Numbers...

- The key foundation for all modern cryptographic systems is often not encryption but these "random" numbers!
- So many times you need to get something random:
  - A random cryptographic key
  - A random initialization vector
  - A "nonce" (use-once item)
  - A unique identifier
  - Stream Ciphers
- If an attacker can *predict* a random number things can catastrophically fail

# Breaking Slot Machines

- Some casinos experienced unusual bad "luck"
  - The suspicious players would wait and then all of a sudden try to play
- The slot machines have ***predictable*** pRNG
  - Which was based on the current time & a seed
- So play a little...
  - With a cellphone watching
  - And now you know when to press "spin" to be more likely to win
- Oh, and this ***never*** effected Vegas!
  - ***Evaluation standards*** for Nevada slot machines specifically designed to address this sort of issue

BRENDAN KOERNER SECURITY 02.06.17 07:00 AM  
**RUSSIANS ENGINEER A BRILLIANT SLOT MACHINE**  
IN EARLY JUNE 2014, accountants at the Lumiere Place

Casino in St. Louis noticed that several of their slot machines had—just for a couple of days—gone haywire. The government-approved software that powers such machines gives the house a fixed mathematical edge, so that casinos can be certain of how much they'll earn over the long haul—say, 7.129 cents for every dollar played. But on June 2 and 3, a number of Lumiere's machines had spit out far more money than they'd consumed, despite not awarding any major jackpots—an aberration known in industry parlance as a



# Breaking Bitcoin Wallets

- blockchain.info supports "web wallets"
  - Javascript that protects your Bitcoin
- The private key for Bitcoin needs to be random
  - Because otherwise an attacker can spend the money
- An "Improvement" [sic] to the RNG reduced the entropy (the actual randomness)
  - Any wallet created with this improvement was brute-forceable and could be stolen

## Improvements to RNG



zootreeves committed on Dec 7, 2014

1 parent b0d5639

Showing 1 changed file with 26 additions and 28 deletions.

54 bitcoinjs-lib/src/jsbn/rng.js

```
@@ -8,15 +8,16 @@ var rng_state;
 8   8   var rng_pool;
 9   9   var rng_pptr;
10  10
11 // Mix in a 32-bit integer into the pool
12 -function rng_seed_int(x) {
13 -  rng_pool[rng_pptr++] ^= x & 255;
14 -  rng_pool[rng_pptr++] ^= (x >> 8) & 255;
15 -  rng_pool[rng_pptr++] ^= (x >> 16) & 255;
16 -  rng_pool[rng_pptr++] ^= (x >> 24) & 255;
```



# *TRUE* Random Numbers

- True random numbers generally require a physical process
- Common circuit is an unusable ring oscillator built into the CPU
  - It is then sampled at a low rate to generate true random bits which are then fed into a pRNG on the CPU
- Other common sources are human activity measured at very fine time scales
  - Keystroke timing, mouse movements, etc
    - "Wiggle the mouse to generate entropy for a key"
  - Network/disk activity which is often human driven
- More exotic ones are possible:
  - Cloudflare has a wall of lava lamps that are recorded by a HD video camera which views the lamps through a rotating prism: It is just one source of the randomness



# Combining Entropy

- Many physical entropy sources are biased
  - Some have significant biases: e.g. a coin that flips "heads" 90% of the time!
  - Some aren't very good: e.g. keystroke timing at a microsecond granularity
- The general procedure is to combine various sources of entropy
- The goal is to be able to take multiple crappy sources of entropy
  - Measured in how many bits:  
A single flip of a fair coin is 1 bit of entropy
  - And combine into a value where the entropy is the minimum of the sum of all entropy sources (maxed out by the # of bits in the hash function itself)
  - **N-1** bad sources and **1** good source -> good pRNG state

# Pseudo Random Number Generators (aka Deterministic Random Bit Generators)

- Unfortunately one needs a *lot* of random numbers in cryptography
  - More than one can generally get by just using the physical entropy source
- Enter the pRNG or DRBG
  - If one knows the state it is entirely predictable
  - If one doesn't know the state it should be indistinguishable from a random string
- Three operations
  - Instantiate: (aka Seed) Set the internal state based on the real entropy sources
  - Reseed: Update the internal state based on both the previous state and *additional entropy*
    - The big different from a simple stream cipher
  - Generate: Generate a series of random bits based on the internal state
    - Generate can also optionally add in additional entropy
- **instantiate(entropy)**  
**reseed(entropy)**  
**generate(bits, {optional entropy})**

# Properties for the pRNG

- Can a pRNG be truly random?
  - No. For seed length  $s$ , it can only generate at most  $2^s$  distinct possible sequences.
- A cryptographically strong pRNG “looks” truly random to an attacker
  - Attacker ***cannot distinguish*** it from a random sequence:  
If the attacker can tell a sufficiently long bitstream was generated by the pRNG instead of a truly random source it isn't a good pRNG

# Prediction and Rollback Resistance

- A pRNG should be predictable only if you know the internal state
  - It is this predictability which is why its called "pseudo"
- If the attacker does not know the internal state
  - The attacker should not be able to distinguish a truly random string from one generated by the pRNG
- It ***should*** also be rollback-resistant
  - Even if the attacker finds out the state at time T, they should not be able to determine what the state was at T-1
  - More precisely, if presented with two random strings, one truly random and one generated by the pRNG at time T-1, the attacker should not be able to distinguish between the two
  - Rollback resistance isn't specifically required in a pRNG...  
But it should be

# Why "Rollback Resistance" is Essential

- Assume attacker, at time  $T$ , is able to obtain all the internal state of the pRNG
  - How? E.g. the pRNG screwed up and instead of an IV, released the internal state, or the pRNG is bad...
- Attacker observes how the pRNG was used
  - $T_{-1}$  = Random Session key
  - $T_0$  = Nonce/IV
- Now if the pRNG doesn't resist rollback, and the attacker gets the state at  $T_0$ , attacker can know the session key! And we are back to...



# More on Seeding and Reseeding

- Seeding should take all the different physical entropy sources available
  - If one source has 0 entropy, it ***must not*** reduce the entropy of the seed
  - We can shove a whole bunch of low-entropy sources together and create a high-entropy seed
- Reseeding ***adds*** in even more entropy
  - **F(internal\_state, new material)**
  - Again, even if reseeding with 0 entropy, it ***must not*** reduce the entropy of the seed

# Probably the best pRNG/DRBG: HMAC\_DRBG

- Generally believed to be the best
  - ***Accept no substitutes!***
- Two internal state registers,  $V$  and  $K$ 
  - Each the same size as the hash function's output
- $V$  is used as (part of) the data input into HMAC, while  $K$  is the key
- If you can break this pRNG you can ***either break the underlying hash function or break a significant assumption about how HMAC works***
  - Yes, security proofs sometimes are a very good thing and actually do work
  - So as long as the security proof for HMAC is correct, the security proof for HMAC\_DRBG is correct!

# HMAC\_DRBG

## Generate

- The basic generation function
- Remarks:
  - It requires one HMAC call per blocksize-bits of state
  - Then two more HMAC calls to update the internal state
- Prediction resistance:
  - If you can distinguish new **K** from random when you don't know old **K**:  
You've distinguished HMAC from a random function!  
Which means you've either broken the hash or the HMAC construction
- Rollback resistance:
  - If you can learn old **K** from new **K** and **V**:  
**You've reversed the hash function!**

```
function hmac_drbg_generate (state, n) {  
    tmp = ""  
    while(len(tmp) < N) {  
        state.v = hmac(state.k, state.v)  
        tmp = tmp || state.v  
    }  
    // Update state with no input  
    state.k = hmac(state.k, state.v || 0x00)  
    state.v = hmac(state.k, state.v)  
    // Return the first N bits of tmp  
    return tmp[0:N]  
}
```

# HMAC\_DRBG

## Update

- Used instead of the "no-input update" when you have additional entropy on the generate call
- Used standalone for both instantiate (**state.k = state.v = 0**) and reseed (keep **state.k** and **state.v**)
- Designed so that even if the attacker controls the input but doesn't know **k**:  
The attacker should not be able to predict the new **k**

```
function hmac_drbg_update (state, input) {  
    state.k = hmac(state.k, state.v || 0x00  
                  || input)  
    state.v = hmac(state.k, state.v)  
    state.k = hmac(state.k, state.v || 0x01  
                  || input)  
    state.v = hmac(state.k, state.v)  
}
```

# Generating **true** random numbers

- Modern CPUs have true random number generators
  - Sample a noisy circuit at a low rate or similar tricks
- These sources are biased...
  - They are also **slow**
- So use this as an entropy source to feed a pRNG on the chip
  - Now you can get random numbers quickly
- Very fast
- Vulnerable to tampering!
  - You can't actually test that the pRNG circuit is 100% correct without adding paths that could potentially sabotage the pRNG circuit
  - Sabotage that can reduce effective entropy to 32b are possible

# Stream ciphers

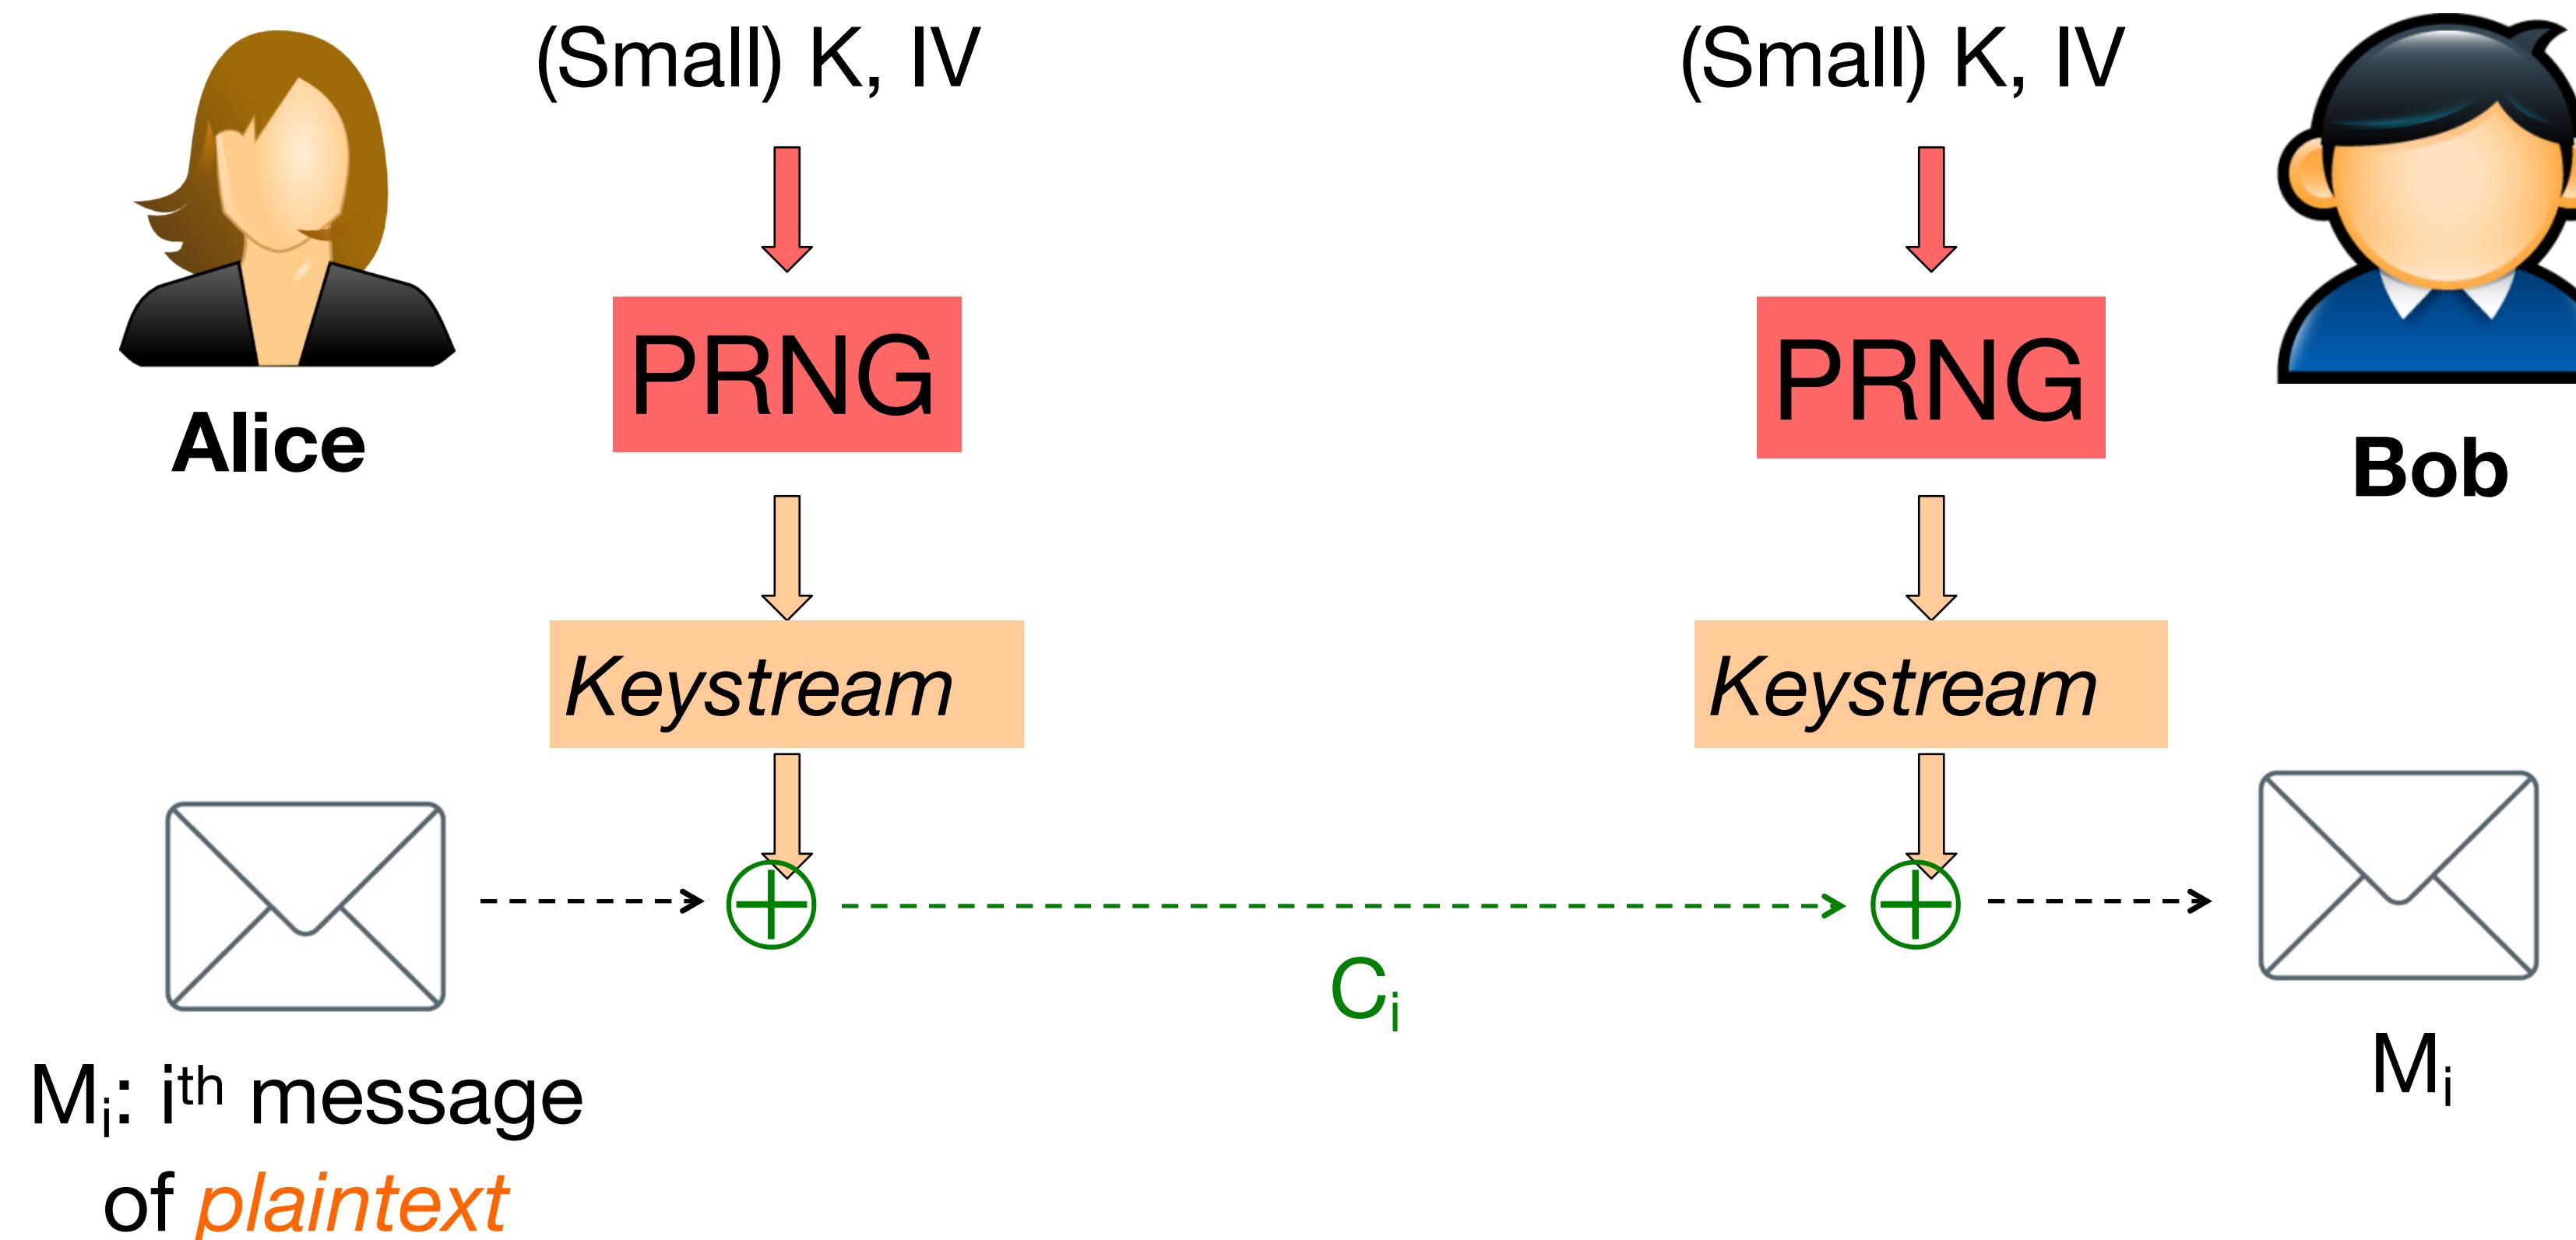
- Block cipher: fixed-size, stateless, requires “modes” to securely process longer messages
- Stream cipher: keeps state from processing past message elements, can continually process new elements
- Common approach: “one-time pad on the cheap”:
  - XORs the plaintext with some “random” bits
  - But: random bits  $\neq$  the key (as in one-time pad)
    - Instead: output from cryptographically strong pseudorandom number generator (pRNG)
    - Anyone who actually calls this a "One Time Pad" is selling snake oil!

# Building Stream Ciphers

- Encryption, given key **K** and message **M**:
  - Choose a random value **IV**
  - $E(M, K) = \text{pRNG}(K, IV) \oplus M$
- Decryption, given key **K**, ciphertext **C**, and initialization vector **IV**:
  - $D(C, K) = \text{PRNG}(K, IV) \oplus C$
- Can encrypt message of any length because pRNG can produce any number of random bits...
  - But in practice, for an n-bit seed pRNG, stop at  $2^{n/2}$ . Because, of course...



# Using a pRNG to Build A Stream Cipher



# CTR mode is (mostly) a stream cipher

- $E(\text{ctr}, K)$  should look like a series of pseudo random numbers...
  - But after a large amount it is *slightly* distinguishable!
- Since it is actually a pseudo-random **permutation**...
  - For a cipher using 128b blocks, you will never get the same 128b number until you go all the way through the  $2^{128}$  possible entries on the counter
  - Reason why you want to stop after  $2^{64}$ 
    - If you use CTR mode in the first place
- Also very minor information leakage:
  - If  $C_i = C_j$ , for  $i \neq j$ , it follows that  $M_i \neq M_j$

# UUID: Universally Unique Identifiers

- You got to have a "name" for something...
  - EG, to store a location in a filesystem
- Your name ***must*** be unique...
  - And your name ***must*** be unpredictable!
- Just chose a ***random*** value!
  - UUID: just chose a 128b random value
    - Well, it ends up being a 122b random value with some signaling information
    - A good UUID library uses a cryptographically-secure pRNG that is properly seeded
- Often written out in hex as:
  - 00112233-4455-6677-8899-aabbccddeeff

# What Happens When The Random Numbers Goes Wrong...

- Insufficient Entropy:
  - Random number generator is seeded without enough entropy
- Debian OpenSSL CVE-2008-0166
  - In "cleaning up" OpenSSL (Debian 'bug' #363516), the author 'fixed' how OpenSSL seeds random numbers
    - Because the code, as written, caused Purify and Valgrind to complain about reading uninitialized memory
  - Unfortunate cleanup reduced the pRNG's seed to be *just* the process ID
    - So the pRNG would only start at one of ~30,000 starting points
- This made it easy to find private keys
  - Simply set to each possible starting point and generate a few private keys
  - See if you then find the corresponding public keys anywhere on the Internet



# And Now Lets Add Some RNG Sabotage...

- The Dual\_EC\_DRBG
  - A pRNG pushed by the NSA behind the scenes based on Elliptic Curves
  - It relies on two parameters,  $P$  and  $Q$  on an elliptic curve
  - The person who generates  $P$  and selects  $Q = eP$  can predict the random number generator, regardless of the internal state
- It also **sucked!**
  - It was horribly slow and even had subtle biases that shouldn't exist in a pRNG:  
You could distinguish the upper bits from random!
  - Now this was spotted fairly early on...
  - Why should anyone use such a horrible random number generator?

# Well, anyone not paid that is...

- RSA Data Security accepted 30 pieces of silver \$10M from the NSA to implement Dual\_EC in their RSA BSAFE library
  - And *silently* make it the default pRNG
  - Using RSA's support, it became a NIST standard
    - And inserted into other products...
  - And then the Snowden revelations
    - The initial discussion of this sabotage in the NY Times just vaguely referred to a Crypto talk given by Microsoft people...
    - That everybody quickly realized referred to Dual\_EC



# But this is insanely powerful...

- It isn't just forward prediction but being able to run the generator backwards!
  - Which is why Dual\_EC is so nasty:  
Even if you know the internal state of HMAC\_DRBG it has rollback resistance!
- In TLS (HTTPS) and Virtual Private Networks you have a motif of:
  - Generate a random session key
  - Generate some other random data that's **public visible**
    - EG, the IV in the encrypted channel, or the "random" nonce in TLS
    - Oh, and an NSA sponsored "standard" to spit out even more "random" bits!
- If you can run the random number generator **backwards**, you can find the session key



# It Got Worse: Sabotaging Juniper

- Juniper also used Dual\_EC in their Virtual Private Networks
  - "But we did it safely, we used a different **Q**"
- Sometime later, someone else noticed this...
  - "Hmm, **P** and **Q** are the keys to the backdoor... Lets just hack Juniper and rekey the lock!"
    - And whoever put in the first Dual\_EC then went "Oh crap, we got locked out but we can't do anything about it!"
- Sometime later, someone else goes...
  - "Hey, lets add an ssh backdoor"
- Sometime later, Juniper goes
  - "Whoops, someone added an ssh backdoor, lets see what else got F'ed with, oh, this # in the prng"
- And then everyone else went
  - "Ohh, patch for a backdoor. Lets see what got fixed. Oh, these look like Dual\_EC parameters..."



# Sabotaging "Magic Numbers" In General

- Many cryptographic implementations depend on "magic" numbers
  - Parameters of an Elliptic curve
  - Magic points like  $P$  and  $Q$
  - Particular prime  $p$  for Diffie/Hellman
  - The content of S-boxes in block cyphers
- Good systems should cleanly describe how they are generated
  - In some sound manner (e.g. AES's S-boxes)
  - In some "random" manner defined by a pRNG with a specific seed
    - Eg, seeded with "Nicholas Weaver Deserves Perfect Student Reviews"...
    - Needs to be very low entropy so the designer can't try a gazillion seeds

# Because Otherwise You Have Trouble...

- Not only Dual-EC's **P** and **Q**
- Recent work: 1024b Diffie/Hellman moderately impractical...
  - But you can create a sabotaged prime that is 1/1,000,000 the work to crack!  
And the most often used "example" **p**'s origin is lost in time!
- It can cast doubt **even when a design is solid:**
  - The DES standard was developed by IBM but with input from the NSA
    - Everyone was suspicious about the NSA tampering with the S-boxes...
    - They did: The NSA made them **stronger** against an attack they knew but the public didn't
  - The NSA-defined elliptic curves P-256 and P-384
    - I trust them because they are in CNSA so the NSA uses them for TS communication:  
A backdoor here would be absolutely unacceptable...  
but **only because I actually believe the NSA wouldn't go and try to shoot itself in the head!**



# So What To Use?

- AES-128-CFB or AES-256-CFB:
  - Robust to screwups encryption
  - Alternately, AES-128-GCM (Galios Counter Mode):  
An AEAD mode, but is ***NOT resistant to screwups***
- SHA-2 or SHA-3 family (256b, 384b, or 512b):
  - Robust cryptographic hashes, SHA-1 and MD5 are broken
- HMAC-SHA256 or HMAC-SHA3:
  - Different function than the encryption:  
Prevents screwups on using the same key & is a hash if not using an AEAD mode
  - ***Always Encrypt Then MAC!***
- HMAC-SHA256-DRBG or HMAC-SHA3-DRBG:
  - The best pRNG available
  - Seed using ***both*** the processor random number generator AND other entropy sources!
    - Don't use the processor RNG bare when building a software cryptosystem:  
Those are potentially sabotage able and use designs without rollback resistance.

# Public Key...

- All our previous primitives required a "miracle":
  - We somehow have to have Alice and Bob get a shared  $k$ .
- Enter Public Key cryptography: the miracle of modern cryptography
  - How starting Friday, but ***what*** today
- Three primitives:
  - Public Key Agreement
  - Public Key Encryption
  - Public Key Signatures
- Based on some families of magic math...
  - For us, we will use some group-theory based primitives

# Public Key Agreement

- Alice and Bob have a channel...
  - There may be an eavesdropper ***but not a manipulator***
  - The goal: Alice & Bob agree on a ***random*** value
    - This will be  $k$  for all subsequent communication
  - When done, the key is thrown away
    - Designed to prevent an attacker who later recovers Alice or Bob's long lived secrets from finding  $k$ .

# Public Key Encryption

- Alice has ***two*** keys:
  - $K_{pub}$ : Her public key, anyone can know
  - $K_{priv}$ : Her private key, a deep dark secret
- Anyone has access to Alice's public key
- For anyone to send a message to Alice:
  - Create a random session key  $k$ 
    - Used to encrypt the rest of the message
  - Encrypt  $k$  using Alice's  $K_{pub}$ .
- Only Alice can ***decrypt*** the message
  - The decryption function only works with  $K_{priv}$ !

# Public Key Signatures

- Once again, Alice has *two* keys:
  - $K_{pub}$ : Her public key, anyone can know
  - $K_{priv}$ : Her private key, a deep dark secret
- She can sign a message
  - Calculate  $H(M)$
  - $S(K_{priv}, H(M))$ : Sign  $H(M)$  with  $K_{priv}$ .
- Anyone can now verify
  - Recalculate  $H(M)$
  - $V(K_{pub}, S(K_{priv}, H(M)), H(M))$ : Verify that the signature was created with  $K_{priv}$

# Things To Remember...

- Public key is ***slow!***
  - Orders of magnitude slower than symmetric key
- Public key is based on delicate magic math
  - Discrete log in a group is the most common
  - RSA
  - Some new "post-quantum" magic...
- Some systems in particular are easy to get wrong
  - We will get to some of the epic crypto-fails later

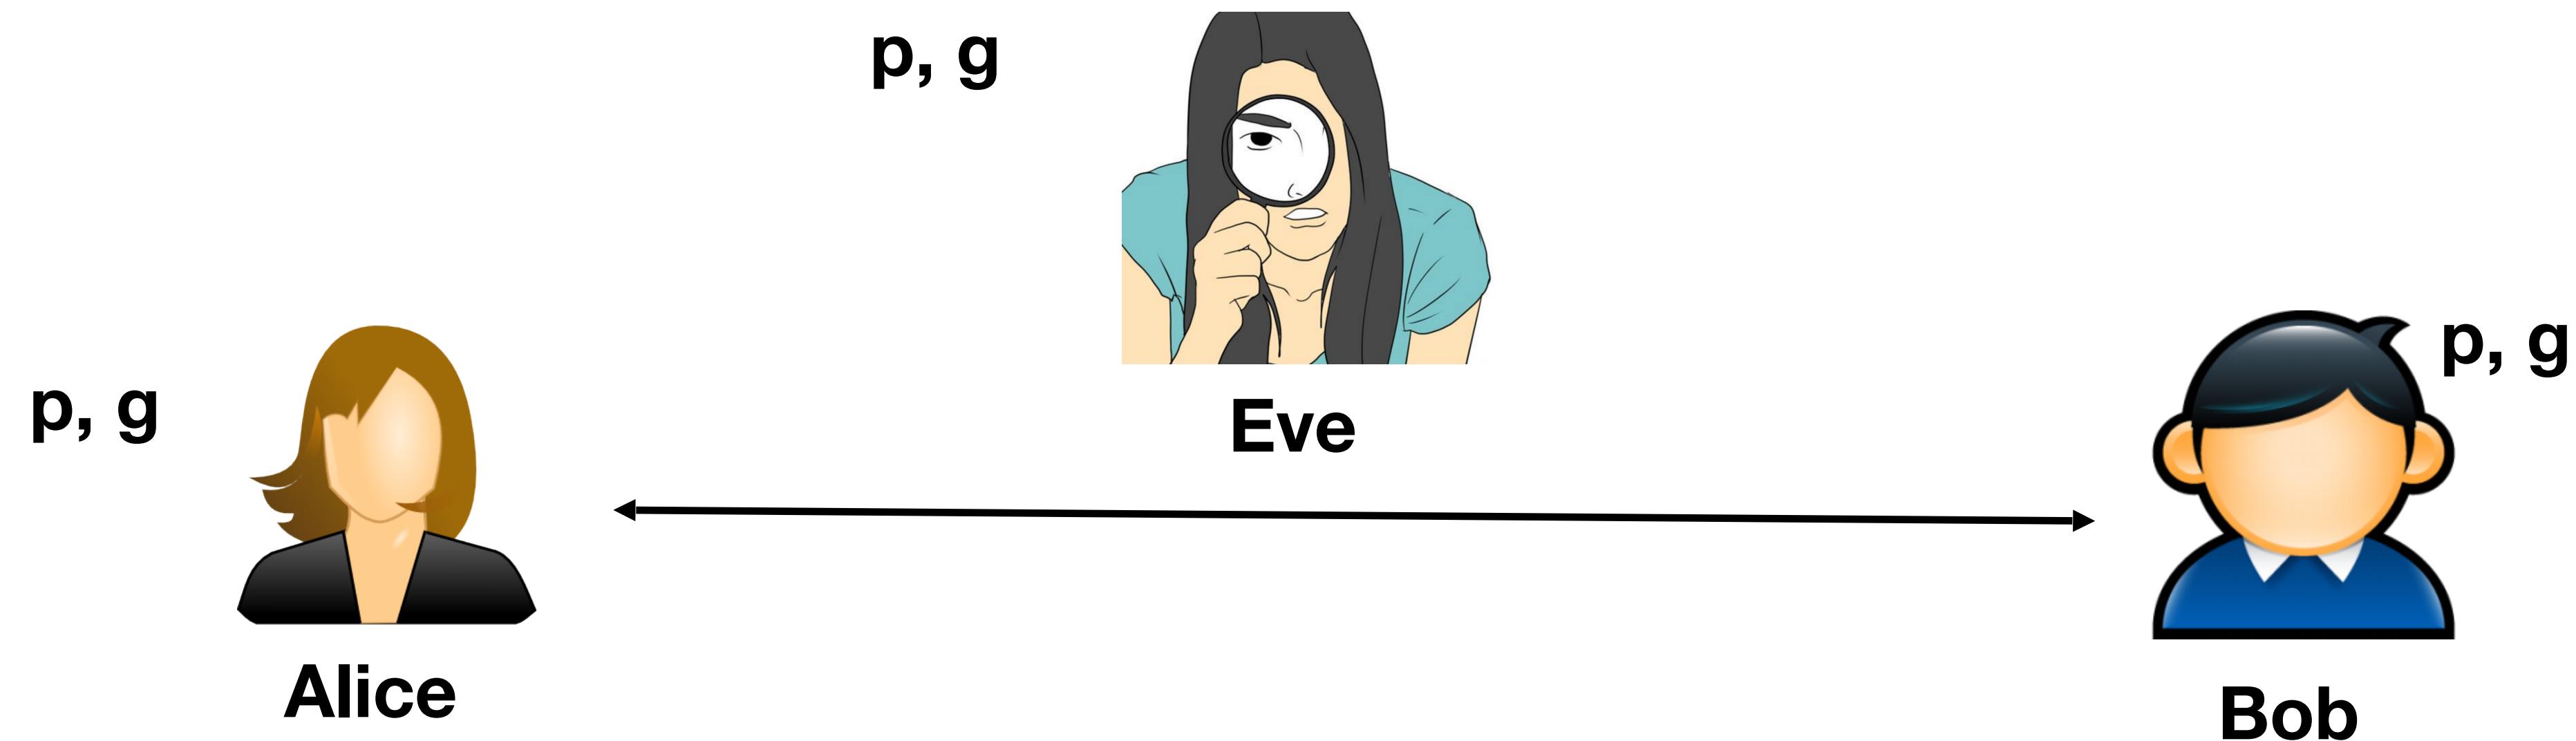
# Our Roadmap For Public Key...

- Public Key:
  - Something **everyone** can know
- Private Key:
  - The secret belonging to a specific person
- Diffie/Hellman:
  - Provides key exchange with no pre-shared secret
- ElGamal & RSA:
  - Provide a message to a recipient only knowing the recipient's **public key**
- DSA & RSA signatures:
  - Provide a message that anyone can prove was generated with a **private key**

# Diffie-Hellman Key Exchange

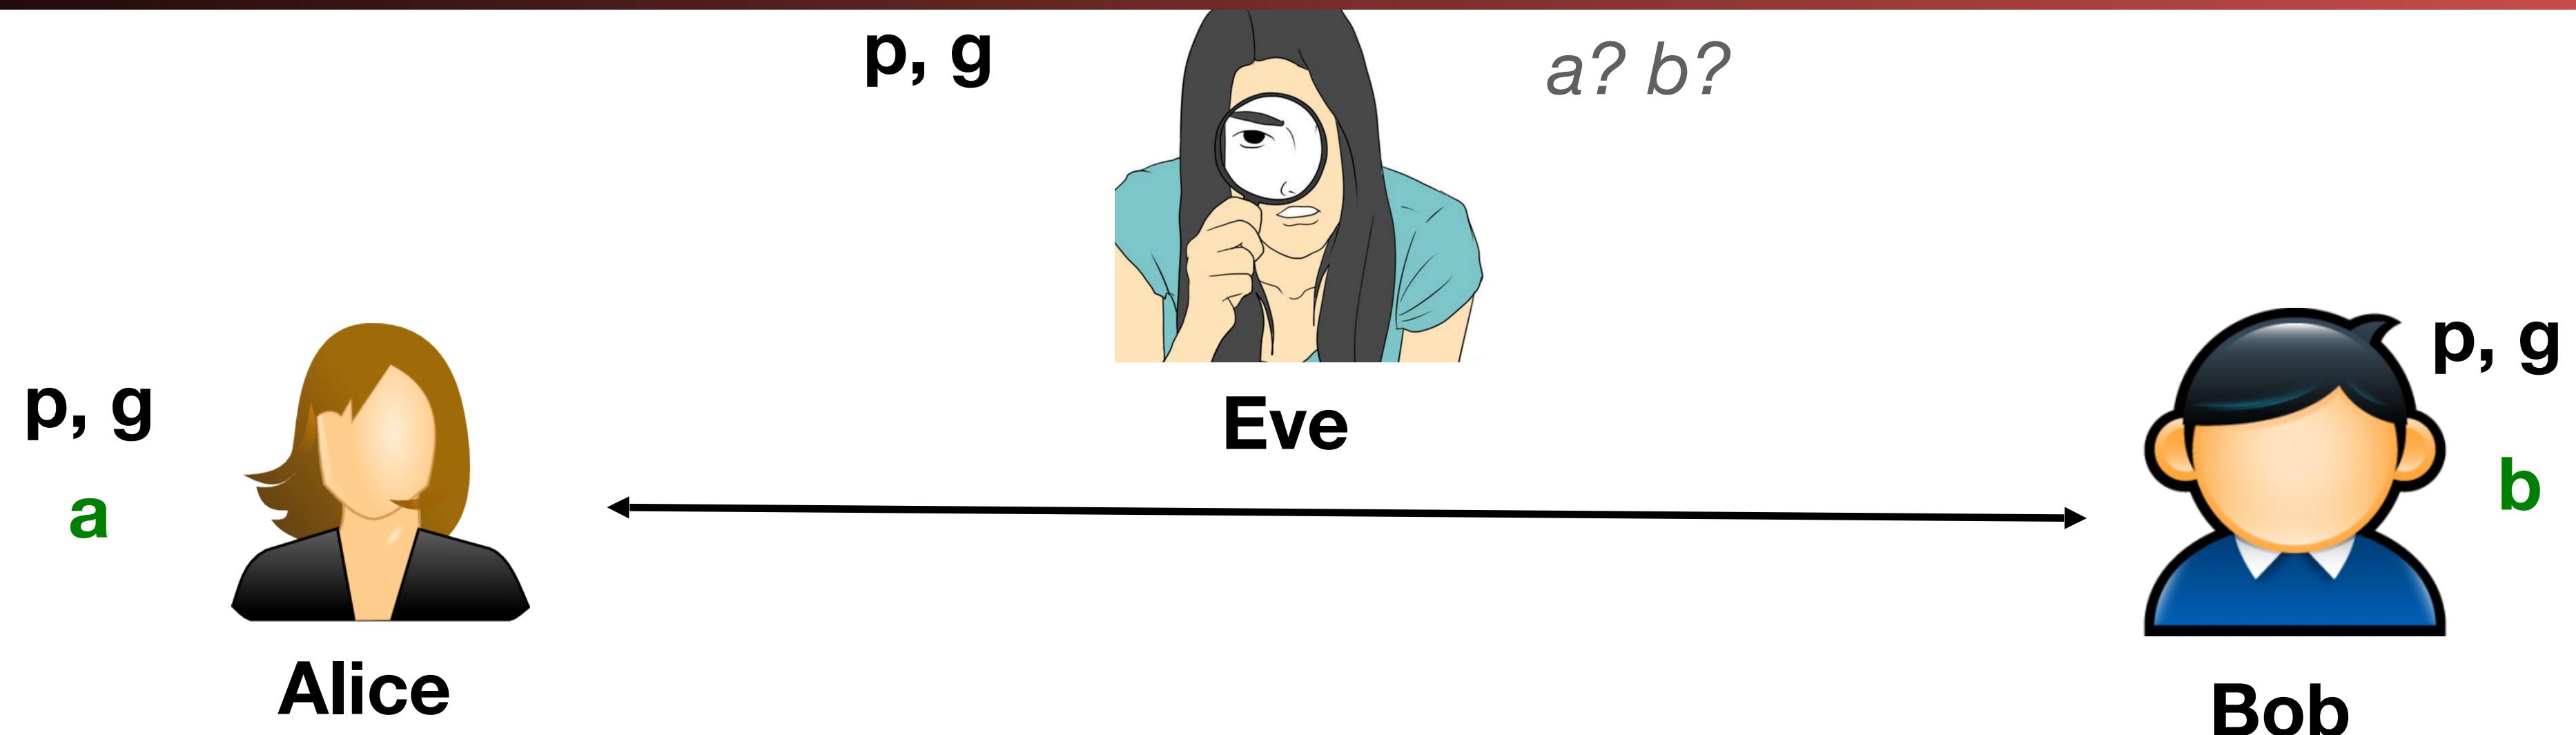
- What if instead they can somehow generate a random key when needed?
- Seems impossible in the presence of Eve observing all of their communication ...
  - How can they exchange a key without her learning it?
  - But: actually is possible using public-key technology
    - Requires that Alice & Bob know that their messages will reach one another without any meddling
  - Protocol: Diffie-Hellman Key Exchange (DHE)
    - The E is "Ephemeral", we use this to create a temporary key for other uses and then forget about it

# Diffie-Hellman Key Exchange



1. Everyone agrees in advance on a well-known (large) prime  $p$  and a corresponding  $g$ :  $1 < g < p-1$

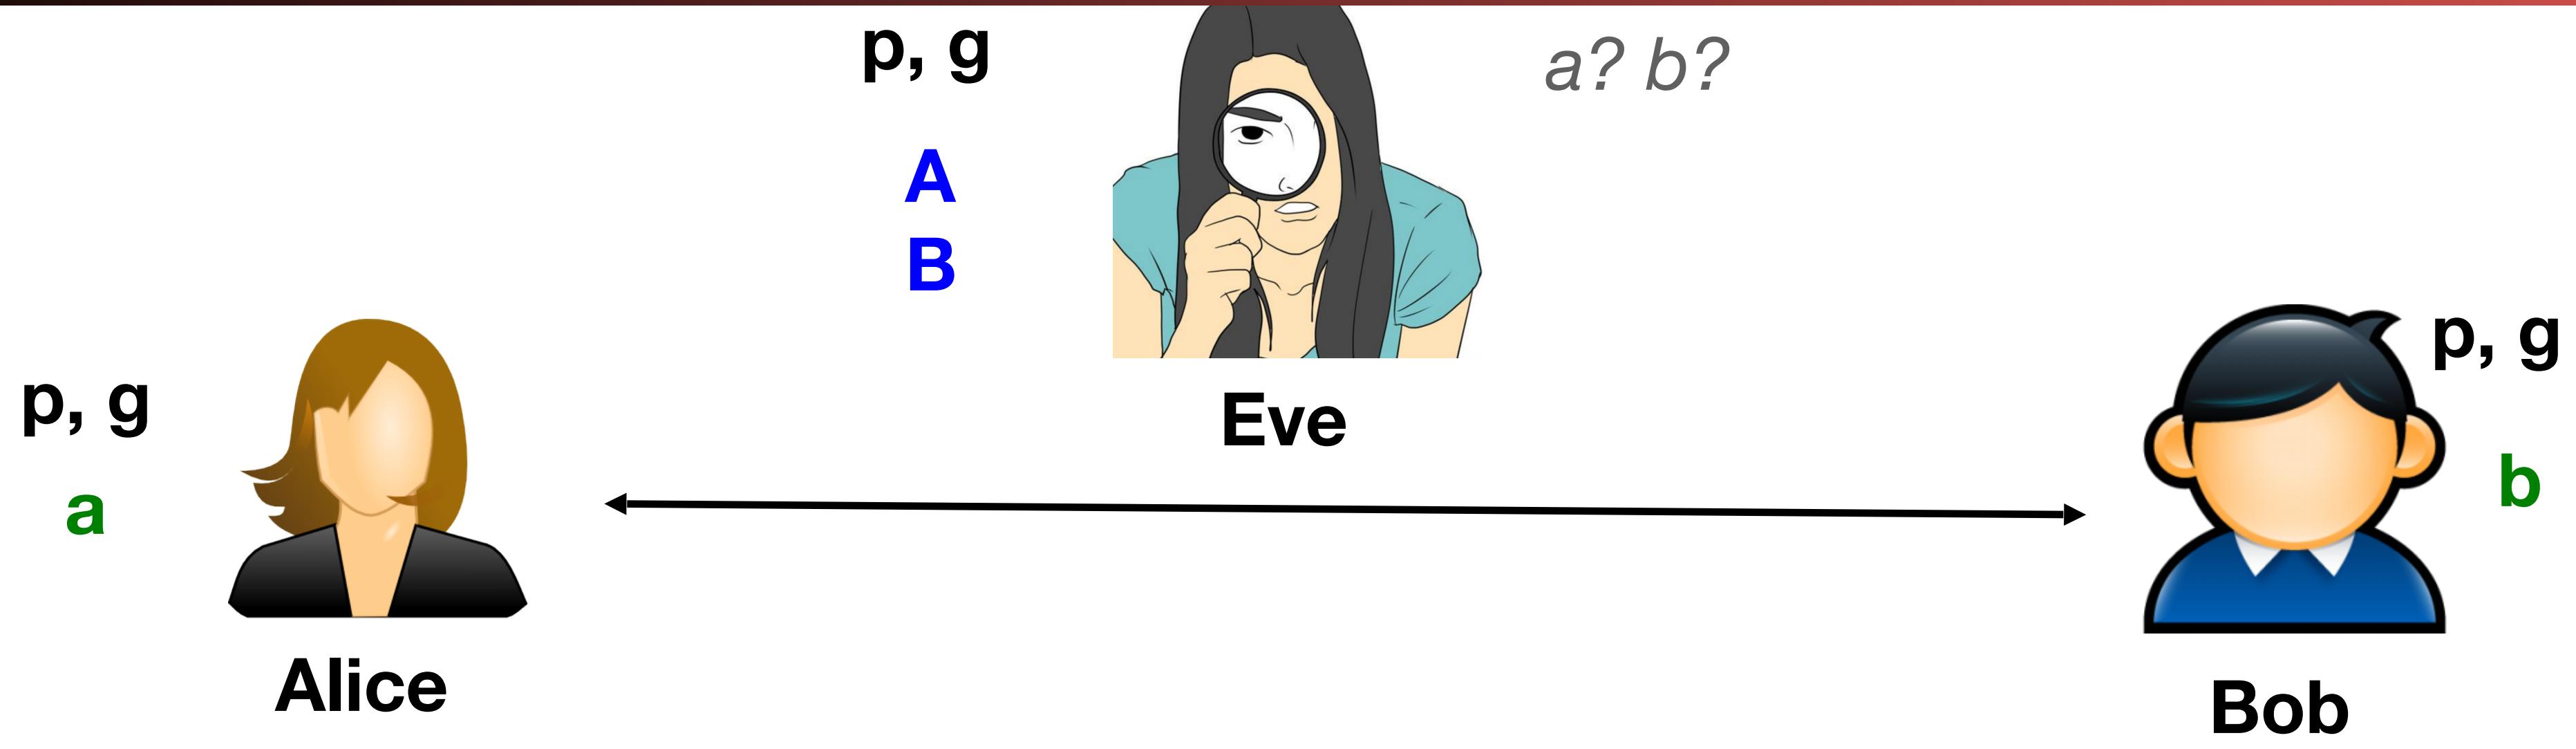
# DHE



2. Alice picks random secret ' $a1 < a < p-1$

3. Bob picks random secret ' $b1 < b < p-1$

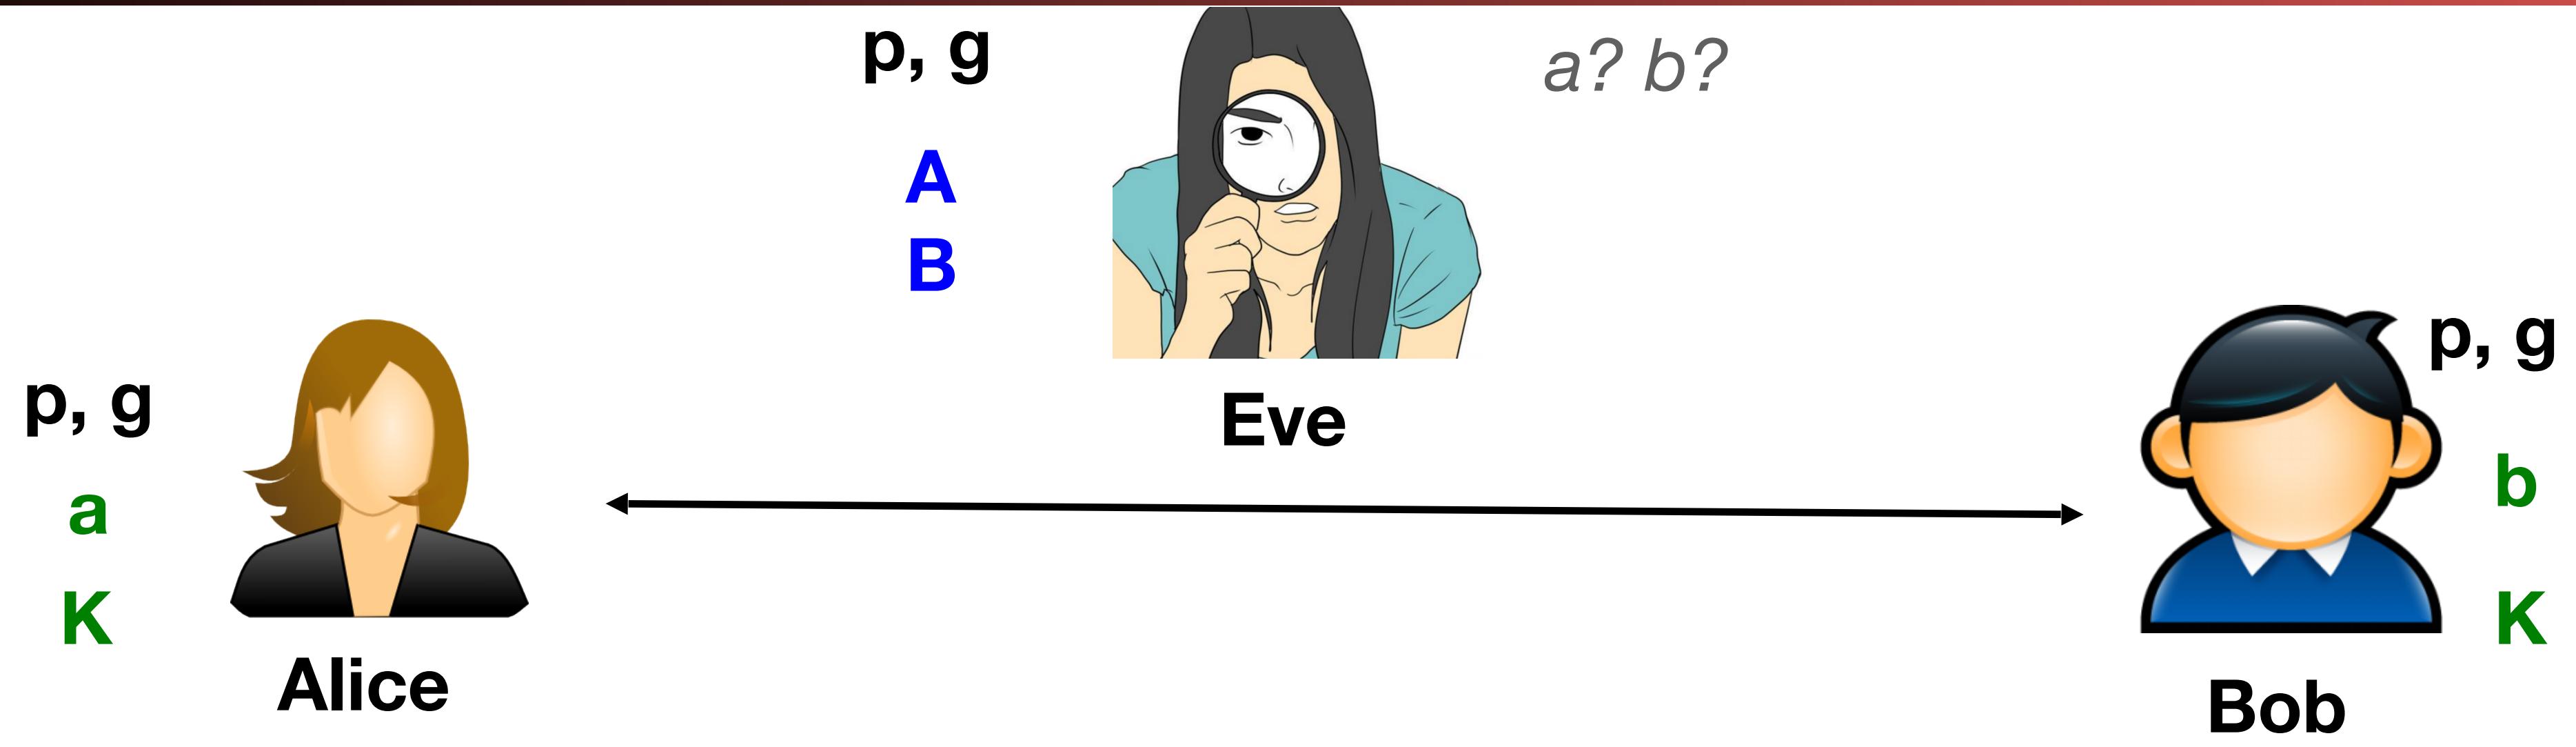
# DHE



$$A = g^a \text{ mod } p$$

4. Alice sends  $A = g^a \text{ mod } p$  to Bob     $g^b \text{ mod } p = B$
5. Bob sends  $B = g^b \text{ mod } p$  to Alice

# DHE



$$A = g^a \bmod p$$

$$B$$

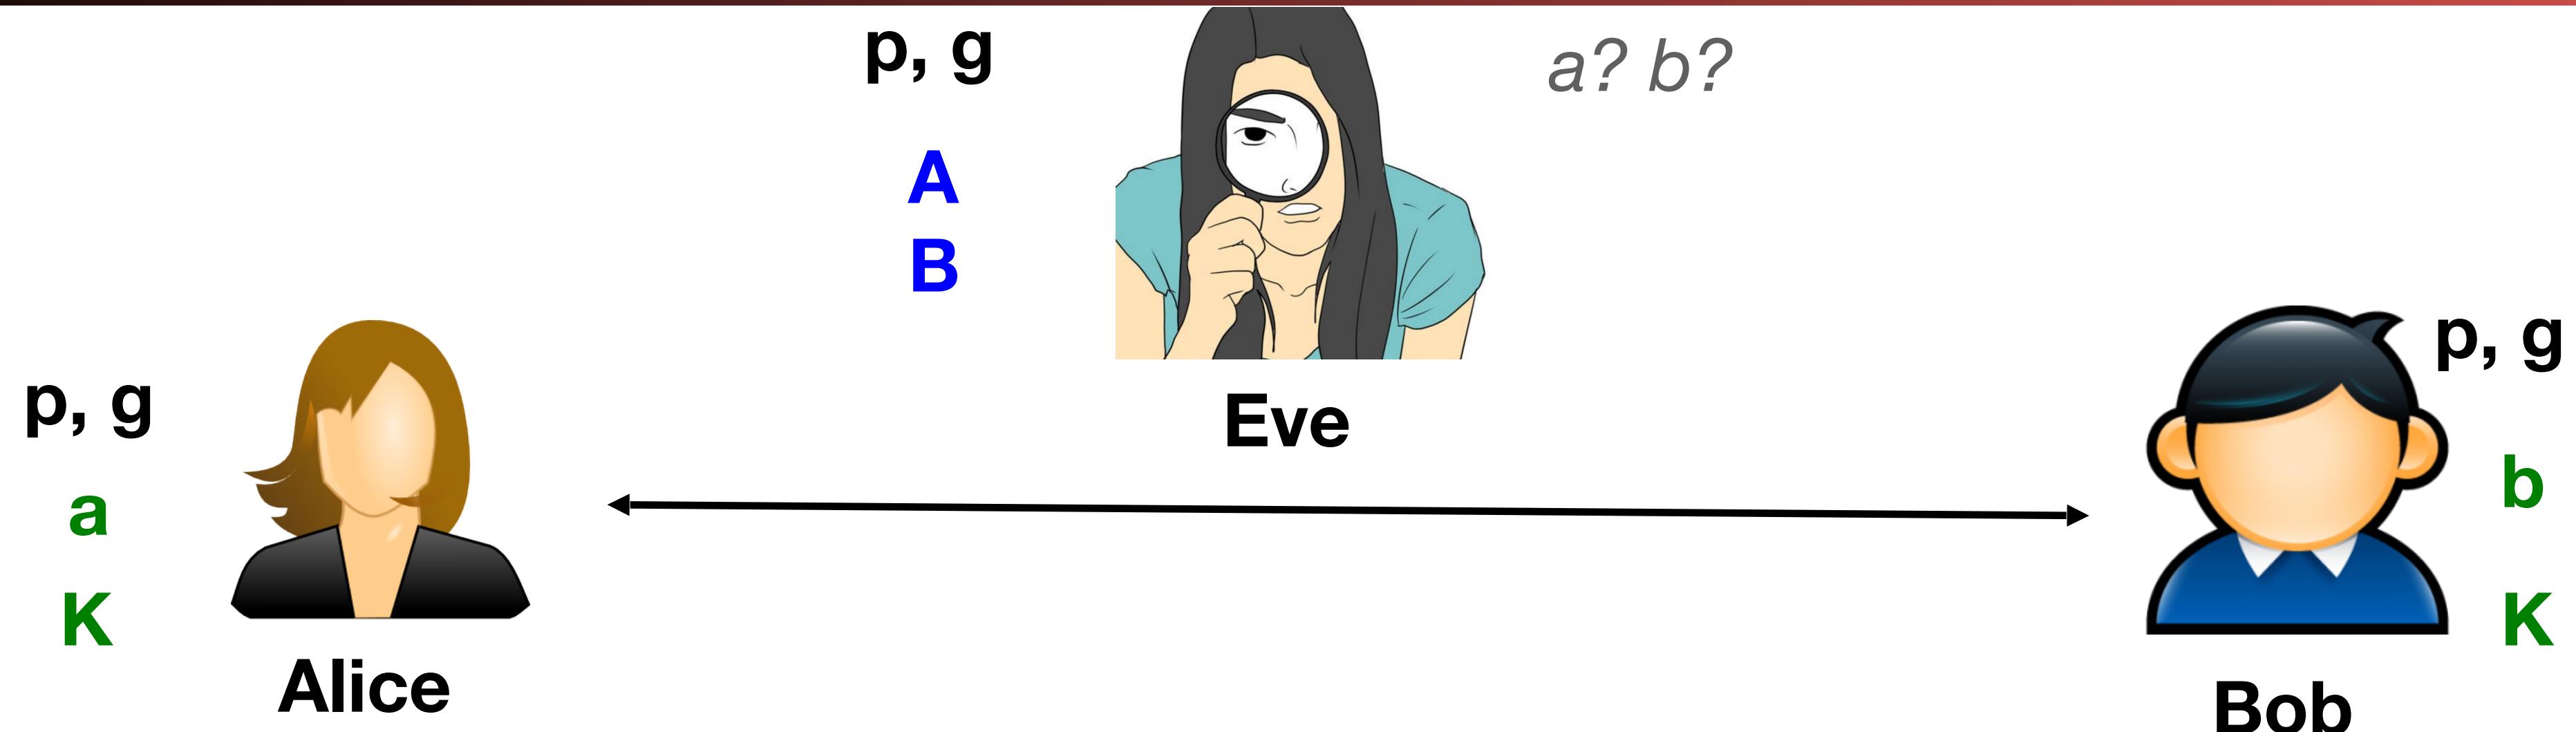
6. Alice knows  $\{a, A, B\}$ , computes  
 $K = B^a \bmod p = (g^b)^a = g^{ba} \bmod p$

$$g^b \bmod p = B$$

7. Bob knows  $\{b, A, B\}$ , computes  
 $K = A^b \bmod p = (g^a)^b = g^{ab} \bmod p$

8.  $K$  is now the shared secret key.

# DHE



While Eve knows  $\{p, g, g^a \text{ mod } p, g^b \text{ mod } p\}$ , believed to be **computationally infeasible** for her to then deduce  $K = g^{ab} \text{ mod } p$ .

She can easily construct  $A \cdot B = g^a \cdot g^b \text{ mod } p = g^{a+b} \text{ mod } p$ .  
But computing  $g^{ab}$  requires ability to take *discrete logarithms* mod  $p$ .  
Discrete log over the group defined by  $p$  and  $g$  **presumed** to be hard

# This is Ephemeral Diffie/Hellman

- $K = g^{ab} \bmod p$  is used as the basis for a "session key"
  - A symmetric key used to protect subsequent communication between Alice and Bob
  - In general, public key operations are vastly more expensive than symmetric key, so it is mostly used just to agree on secret keys, transmit secret keys, or sign hashes
  - If either  $a$  or  $b$  is random,  $K$  is random
- When Alice and Bob are done, they discard  $K, a, b$ 
  - This provides ***forward secrecy***: Alice and Bob don't retain any information that a later attacker who can compromise Alice or Bob's secrets could use to decrypt the messages exchanged with  $K$ .

# Diffie Hellman is part of more generic problem

- This involved deep mathematical voodoo called "Group Theory"
  - Its actually done under a group  $G$
- Two main groups of note:
  - Numbers mod  $p$  with generator  $g$
  - Point addition in an elliptic curve  $C$ 
    - Usually identified by number, eg. p256, p384 (NSA-developed curves) or Curve25519 (developed by Dan Bernstein, also 256b long)
- So EC (Elliptic Curve) == different group
  - Thought to be harder so fewer bits: 384b ECDHE ?= 3096b DHE
  - But otherwise, its "add EC to the name" for something built on discrete log