

Applied Crypto: Passwords & Signal & Tor



Administrivia: Exam Logistics

- <https://cs161.org/exam>
 - READ IT: This is just the tl;dr summary
 - And do Homework 3 now...
At least the part about the exam logistics!
 - Scope is everything up to and including this lecture
- Basic concept
 - You have arbitrary ***hand-written*** paper notes
 - You can compose them on a tablet with a stylus interface but you need to print them out
 - You do the exam on your computer from a pre-distributed encrypted PDF
 - Proctoring is through your phone/second device over Zoom
 - Possible of a "Trust but verify" quick oral quiz afterwards
 - Explain how to solve a variant of a question you successfully solved on the exam

Why this structure?

- We need to do what we can to ensure academic honesty
 - We also have hidden techniques we are using as well
 - If I told all of you "The exam is unproctored and we have no mechanism to detect, respond, or validate that you don't cheat"...
 - I would be ***personally insulted and a total failure if 100% of you DIDN'T CHEAT!***
 - I want my students to be rational, and under those conditions it would be irrational not to cheat
- At the same time, there is a non-linear relationship between effectiveness and burden
 - I doubt the aggressive CS162-style staring on reddit 17 page protocol is significantly better at dissuading dishonesty

Concerns about student privacy

- You will be notified in advance of which TA will be proctoring you
 - For any reason you can switch proctors if you reach out in advance
- You can join using your SID rather than your name
 - Plus zoom on phone sucks for looking at other people's rooms
- If you are not comfortable with this setup, reach out for arranging off-line recording
- If you temporarily lose connectivity, don't worry
 - Focus on the exam, not your Internet connection
- If you need to get up and go to the bathroom, stretch, etc...
Go ahead

Reminder: Cryptographic Hashes...

- We love ourselves some cryptographic hashes
 - SHA_256, SHA_384, SHA3_256, SHA3_384
- Reminder on the properties:
 - Irreversible:
Given $H(X)$, it is infeasible to find X short of simply trying all possibilities
 - First preimage resistant:
Given $H(X)$, it is infeasible to find any X' such that $H(X) = H(X')$
 - Second preimage resistant:
It is infeasible to find X and Y such that $X \neq Y$ and $H(X) = H(Y)$

A Couple Other Hash Properties...

- They accept arbitrarily large inputs
- They "look" random
 - Change a single bit on the input and each output bit has a 50% chance of flipping
 - And until you change the input, you can't predict which output bits are going to change
- The ones we talked about are ***fast***
 - Can operate at many many MB/s:
 - Faster at processing data than block ciphers

A Hash Problem: Proof of work...

- Alice wants Bob to waste a bunch of CPU resources
 - But wants to quickly **check** that Bob wasted that much CPU
- Alice -> Bob: "Here is a message **M** and a factor **x** "
 - Make sure **M** has a nonce in it
- Now Bob needs to provide **M'** such that it starts with **M** and $H(M')$ starts with **x** zero bits
- Alice computes $H(M')$ and verifies that it starts with **x** zero bits
 - Alice now knows that Bob is expected to have had to create **2^x** separate M 's and hash them until he found one that matched

What this provides

- You can use it in a protocol where the user has to waste something...
 - EG, proposals for sending mail as a way of reducing spam
 - It wouldn't: Bad guys can get lots of CPU resources
- Have other options too
- CAPTCHAs:
 - Those "prove your human" web puzzles:
It is a proof you wasted a few seconds of a human's time
(Or that you paid \$.01 to waste a few seconds of a human's time)
- Proof of ***wait***
 - Alice has a secret key k
 - Alice to Bob sends "Don't contact me until time T , here is $\text{HMAC}(k, T)$ "
 - When Bob gets back, he says " T , $\text{HMAC}(k, T)$ "
 - Alice then verifies T is in the past and $\text{HMAC}(k, T)$

Passwords

- The password problem:
 - User Alice authenticates herself with a password P
 - How does the site verify later that Alice knows P ?
- Classic:
 - Just store $\{\text{Alice}, P\}$ in a file...
 - But what happens when the site is hacked?
 - The attacker now knows Alice's password!
 - Enter "Password Hashing"

Password Hashing

- Instead of storing $\{\text{Alice}, P\}$...
 - Store $\{\text{Alice}, H(P)\}$
- To verify Alice, when she presents P
 - Compute $H(P)$ and compare it with the stored value
- Problem: Brute Force tables...
 - Most people chose bad passwords...
And these passwords are known
 - Bad guy has a huge file...
 - $H(P_1), P_1$
 $H(P_2), P_2$
 $H(P_3), P_3\dots$
 - Ways to make this more efficient ("Rainbow Tables")

A Sprinkle of Salt...

- Instead of storing $\{\text{Alice}, \text{H}(\text{P})\}$, also have a user-specific string, the "Salt"
 - Now store $\{\text{Alice}, \text{Salt}, \text{H}(\text{P}||\text{Salt})\}$
 - The salt ideally should be both long and random, but it isn't considered "secret": rather it is a *nonce*
- As long as the salt is unique...
 - An attacker who captures the password file has to **brute force** Alice's password on its own
- It's still an "off-line attack" (Attacker can do all the computation he wants) but...
 - At least the attacker can't **precompute** possible solutions

Slower Hashes...

- Most cryptographic hashes are designed to be ***fast***
 - After all, that is the point: they should not only turn **H(🍔)** to hamburger... they need to do it quickly
- But for password hashes, we ***want*** it to be slow!
 - Its OK if it takes a good fraction of a second to ***check*** a password
 - Since you only need to do it once for each legitimate usage of that password
 - But the attacker needs to do it for each password he wants to try
 - Slower hashes don't change the ***asymptotic difficulty*** of password cracking but can have huge practical impact
 - Slow rate by a factor of 10,000 or more!

PBKDF2

- "Password Based Key Derivation Function 2"
 - Designed to produce a long "random" bitstream derived from the password
 - Used for both a password hash and to generate keys derived from a user's password
 - PBKDF(PRF, P, S, c, len):
 - **PRF** == Pseudo Random Function (e.g. HMAC-SHA256)
 - **P** == Password
 - **S** == Salt
 - **c** == Iteration count
 - **len** == Number of bits/bytes requested
 - **DK** == Derived Key

```
PBKDF(PRF, P, S, c, len) {  
    DK = ""  
    for i = 1, range(len/blocksize)+1) {  
        DK = DK || F(PRF, P, S, c, i)  
    }  
    return DK[0:len]  
}  
  
F(PRF, P, S, c, i) {  
    UR = U = PRF(P, S || INT_32(i))  
    for j = 2; j <= c; ++j {  
        U = PRF(P, U)  
        UR = UR ^ U  
    }  
    return UR  
}
```

Comments on PBKDF2

- Allows you to get effectively an arbitrary long string from a password
 - ***Assuming*** the user's password is strong/high entropy
 - Very good for getting a bunch of symmetric keys from a single password
 - You can also use this to seed a pRNG for generating a "random" public/private key pair
 - Designed to be slow in computation...
 - But it does ***not*** require a lot of memory:
Other functions are also expensive in memory as well, e.g. scrypt & argon2

Passwords...

- If an attacker can do an ***offline*** attack, your password must be ***really good***
 - Attacker simply tries a huge number of passwords in parallel using a GPU-based computer: buy a bunch of used Nvidia 2080 supers from all those upgrading to 3080s
 - So you need a ***high entropy*** password:
 - Even xkcd-style is only 10b/word with a 1000 word dictionary, so need a 7 or more ***random word*** passphrase to resist a determined attacker
 - Life is far better if the attacker can only do ***online*** attacks:
 - Query the device and see if it works
 - Now limited to a few tries per second and ***no parallelism!***



... and iPhones

- Apple's security philosophy:
 - In your hands, the phone should be everything
 - In anybody else's, it should (ideally) be an inert "brick"
- Apple uses a small co-processor in the phone to handle the cryptography
 - The "Secure Enclave"
- The rest of the phone is untrusted
 - Notably the memory: *All* data must be encrypted:
The CPU requests that the Secure Enclave unencrypt data and some data (e.g., your credit card for ApplePay) is only readable by the Secure Enclave
- They also have an ability to effectively erase a small piece of memory
 - "Effaceable Storage": this takes a good amount of EE trickery

Crypto and the iPhone Filesystem

- A lot of keys encrypted by keys...
 - But there is a random master key, k_{phone} , that is the root of all the other keys
- Need to store k_{phone} encrypted by the user's password in the flash memory
 - $\text{PBKDF2}(P, \dots) = k_{\text{user}}$
- But how to prevent an off-line brute-force attack?
 - Also have a 256b *random* secret burned into the Secure Enclave that you can use for encryption
 - Need to take apart the chip to get this!
 - Even the secure enclave can't *read* this secret, only *use* this secret as a key for hardware cryptographic engines
- Now the user key is not just a function of P , but $E(K_{\text{secret}}, P)$
 - Without the secret, *can not* do an offline attack
- All *online* attacks have to go through the secure enclave
 - After 5 tries, starts to slow down
 - After 10 tries, can (optionally) nuke k_{phone} !
 - Erase just that part of memory -> effectively erases the entire phone!
 - Even compromising the secure enclave limits guessing to 10 per second!

Backups...

- Of course there is a ***necessary*** weakness:
 - Backing up the phone copies all the data off in a form not encrypted using the in-chip secret
 - After all, you need to be able to recover it onto a new phone!
- So someone who can get your phone...
And can somehow managed to have it unlocked
 - Thief, abusive boyfriend, cop...
 - Hold it up to your face (iPhone X) or Fingerprint (5s or beyond)
 - And then sync it with a new computer
- Change of policy for iOS-11:
 - Now you also need to put in the passcode to trust a new computer:
Can't create a backup without knowing the passcode

Signal and Tor

- Signal is a messenger protocol and implementation
 - Signal (the company) is a 501(c)3 nonprofit
 - The protocol is also used by WhatsApp, Facebook Messenger, etc...
- Tor is an anonymity tool
 - Designed to provide anonymous but real-time network connectivity in the face of an **aggressive but local adversary**
- Common (bad) information security advice is "Use Signal, Use Tor"
 - In reality, Signal is a great protocol, but some security compromises are annoying in the implementation, so for most, WhatsApp is about as good
 - While Tor is often not just a placebo but **poison!**

End-To-End Messengers

- We love ***end to end*** cryptographic protocols...
 - After all, we just saw why we might want them
- We love ***forward secrecy***...
 - After all, we want things to stay secret even if our keys are compromised
- Forward secrecy is "easy" for online protocols
 - Just make sure to do a DHE/ECDHE key exchange, and throw away the session key when done
- Forward secrecy is ***much more annoying*** for an offline protocol
 - Alice wants to share data with Bob, but Bob is ***not online***
 - Like in project 2...
 - Or any messenger system!

Signal Requirements For Key Agreement

- Three parties: Alice, Bob, and a messenger server
 - The messenger server is like the file store in project 2, an ***untrusted*** entity
 - A **separate** mechanism is used to provide ***key transparency***
- Bob is ***offline***:
 - He has prearranged data stored on the messenger server
- Alice and Bob want to create an ephemeral (DH) key...
 - To use for then encrypting messages
- They need ***mutual authentication***
 - Assuming Alice and Bob have the correct public keys, ***only*** Alice and Bob could have agreed on a key
- They also need ***deniability***
 - Alice or Bob can't create a record ***proving*** the other side participated in creating the key:
So no "Alice just signs her DH..." design

Extended Triple Diffie-Hellman

- Key idea:
 - Lets use multiple Diffie-Hellman exchanges combined into one
 - Some to perform mutual authentication
 - Some to generate an ephemeral key
 - Shove them ALL into a hash-based key derivation function
- They use elliptic curves, but the design would be the same for conventional DH, so we will use the former
 - We will use $DH(A,B)$ as $DH(g^a, g^b)$ where we know a but not b .
(So A is our private value, B is someone else's public value)
 - Also have $Sign(K,M)$ for signing and $KDF(KM)$ which derives a bunch of session keys for a hash-based key derivation function (e.g. $PBKDF2$ with only a couple iterations)

Lots of Keys!

- All keys have both a public & private component
 - Private components always stay with Alice and Bob
 - Anything broadcast is always the public component
- Alice:
 - IK_A : Alice's identity key: for both DH and signatures
 - EK_A : Alice's ephemeral key: Created randomly just to talk to Bob.
- Bob:
 - IK_B : Bob's identity key, long lived
 - SPK_B : Bob's signed rekey, rotates ~weekly/monthly
 - Has corresponding signature $\text{Sign}(IK_b, SPK_b)$
 - OPK_B : Bob's one time use keys (One Time Prekey)
 - Can run out, but designed to increase security when available

Before We Start: Bob to Server, Server to Alice

- Bob uploads:
 - IK_B , SPK_B , $\text{Sign}(IK_B, SPK_B)$, $\{OPK_B^1, OPK_B^2, OPK_B^3 \dots\}$
- Now when Alice wants to talk to Bob...
- Gets from the server:
 - IK_B , SPK_B , $\text{Sign}(IK_B, SPK_B)$, OPK_B ?
 - Told which OPK it is or "There are no $OPKs$ left"
 - $OPKs$ are designed to prevent replay attacks:
Bob will *never* allow any particular OPK to be used twice
- This is now the input into Alice's DH calculations

Alice now does a lot of DH...

- $DH1 = DK(IK_A, SPK_B)$
 - Acts as authentication for Alice when Bob does the same
- $DH2 = DK(EK_A, IK_B)$
 - Forces Bob to do mutual authentication
- $DH3 = DK(EK_A, SPK_B)$
 - Adds in ephemeral EK_A to short lived SPK_B
- $DH4 = DK(EK_A, OPK_B)$
 - Adds in one-time used OPK_B , if available
- $SK = HKDF(DH1 \parallel DH2 \parallel DH3 \parallel DH4)$
 - Skip DH4 if no one time pre-keys are available
 - Now discard the private part of EKA and the intermediate DH calculations

Now Alice Sends To Bob

- IK_A , EK_A , which OPK used (if any), and
 $E(SK, M, IK_A \parallel IK_B)$
 - Using an AEAD encryption mode:
Authenticated Encryption with Additional Data modes allow additional data to be protected by the MAC but sent in the clear:
In this case IK_A and IK_B
 - Bob can do the same DH calculations to generate SK
 - Since Bob knows the private keys corresponding to the public values Alice used
 - If it fails to verify the AEAD data abort

Key Transparency

- For now, Alice and Bob are trusting the server to report IK_A and IK_B correctly
 - If the server lies, 🤡
 - Fortunately there is an answer:
If Alice and Bob are **ever** together:
 - One person's phone displays $H(IK_A // IK_B)$ as a QR Code
 - Other person's phone verifies that it is the same
 - Plus the voice channel...
 - Display "Two Words" on screen:
 $F(H(IK_A // IK_B // SK))$
 - Assumption is a MitM attacker can't fake a voice conversation quickly enough, so if each person says one of the words...

Considerations

- Authentication requires the out-of-channel methods
 - Otherwise no guarantees
- Replay attacks
 - Only if no OPK is available: Can be potentially bad
- Deniability
 - No cryptographic proofs available as to the sender/receiver!

And Then Ratchets...

- A "ratchet" is a one-way function for message keys
 - $\text{Ratchet}(K_i) \rightarrow K_{i+1}, MK_i$
 - But can't take K_{i+1} and MK_i to find K_i
- A symmetric key ratchet is easy
 - We've seen these already:
Any secure PRNG with rollback resistance is a ratchet
 - Can do it slightly more efficiently with HMAC:
 $HMAC(K_i, 0x01) \rightarrow MK_i$
 $HMAC(K_i, 0x02) \rightarrow K_{i+1}$
- It's OK to keep around the intermediate session keys
 - Thanks to HMAC we can't go backwards with them anyway:
Needed for out of order messages

Signal adds in DH ratchets too...

- So for a few messages in a chain you use a symmetric key ratchet...
 - You gain forward secrecy by discarding the old internal state
- But occasionally you rekey with an additional DH
 - Used to add into the ratchet internal state: update K_i to $H(K_{i-1} \parallel DH)$
 - Acts to reset everything with even more randomness
 - So even if you compromise Bob's device at time T and steal all the keys...
 - You can't decrypt old messages that aren't on Bob's device:
can't run the symmetric ratchet backwards
 - You can't decrypt subsequent messages once Bob & Alice use a DH ratchet

The Protocol is Great... BUT!

- The app itself does some eh thing in the usability/security tradeoff...
 - ***No mechanism to back-up messages!***
If your phone is toast, your messages are gone!
 - ***No mechanism to migrate to a new phone!***
If you upgrade to a new phone, your messages are gone!
 - ***Auto-notifies all those where you are in their contacts that they join***
- This is where WhatsApp has a huge competitive advantage
 - They allow backup of messages, message migration etc...

Tor: The Onion Router

Anonymous Websurfing

- Tor actually encompasses many different components
- The Tor network:
 - Provides a means for anonymous Internet connections with low(ish) latency by relaying connections through multiple Onion Router systems
- The Tor Browser bundle:
 - A copy of FireFox extended release with privacy optimizations, configured to only use the Tor network
- Tor Hidden Services:
 - Services only reachable though the Tor network
- Tor bridges with pluggable transports:
 - Systems to reach the Tor network using encapsulation to evade censorship
- Tor provides three separate capabilities in one package:
 - Client anonymity, censorship resistance, server anonymity

The Tor Threat Model: Anonymity of content against *local* adversaries

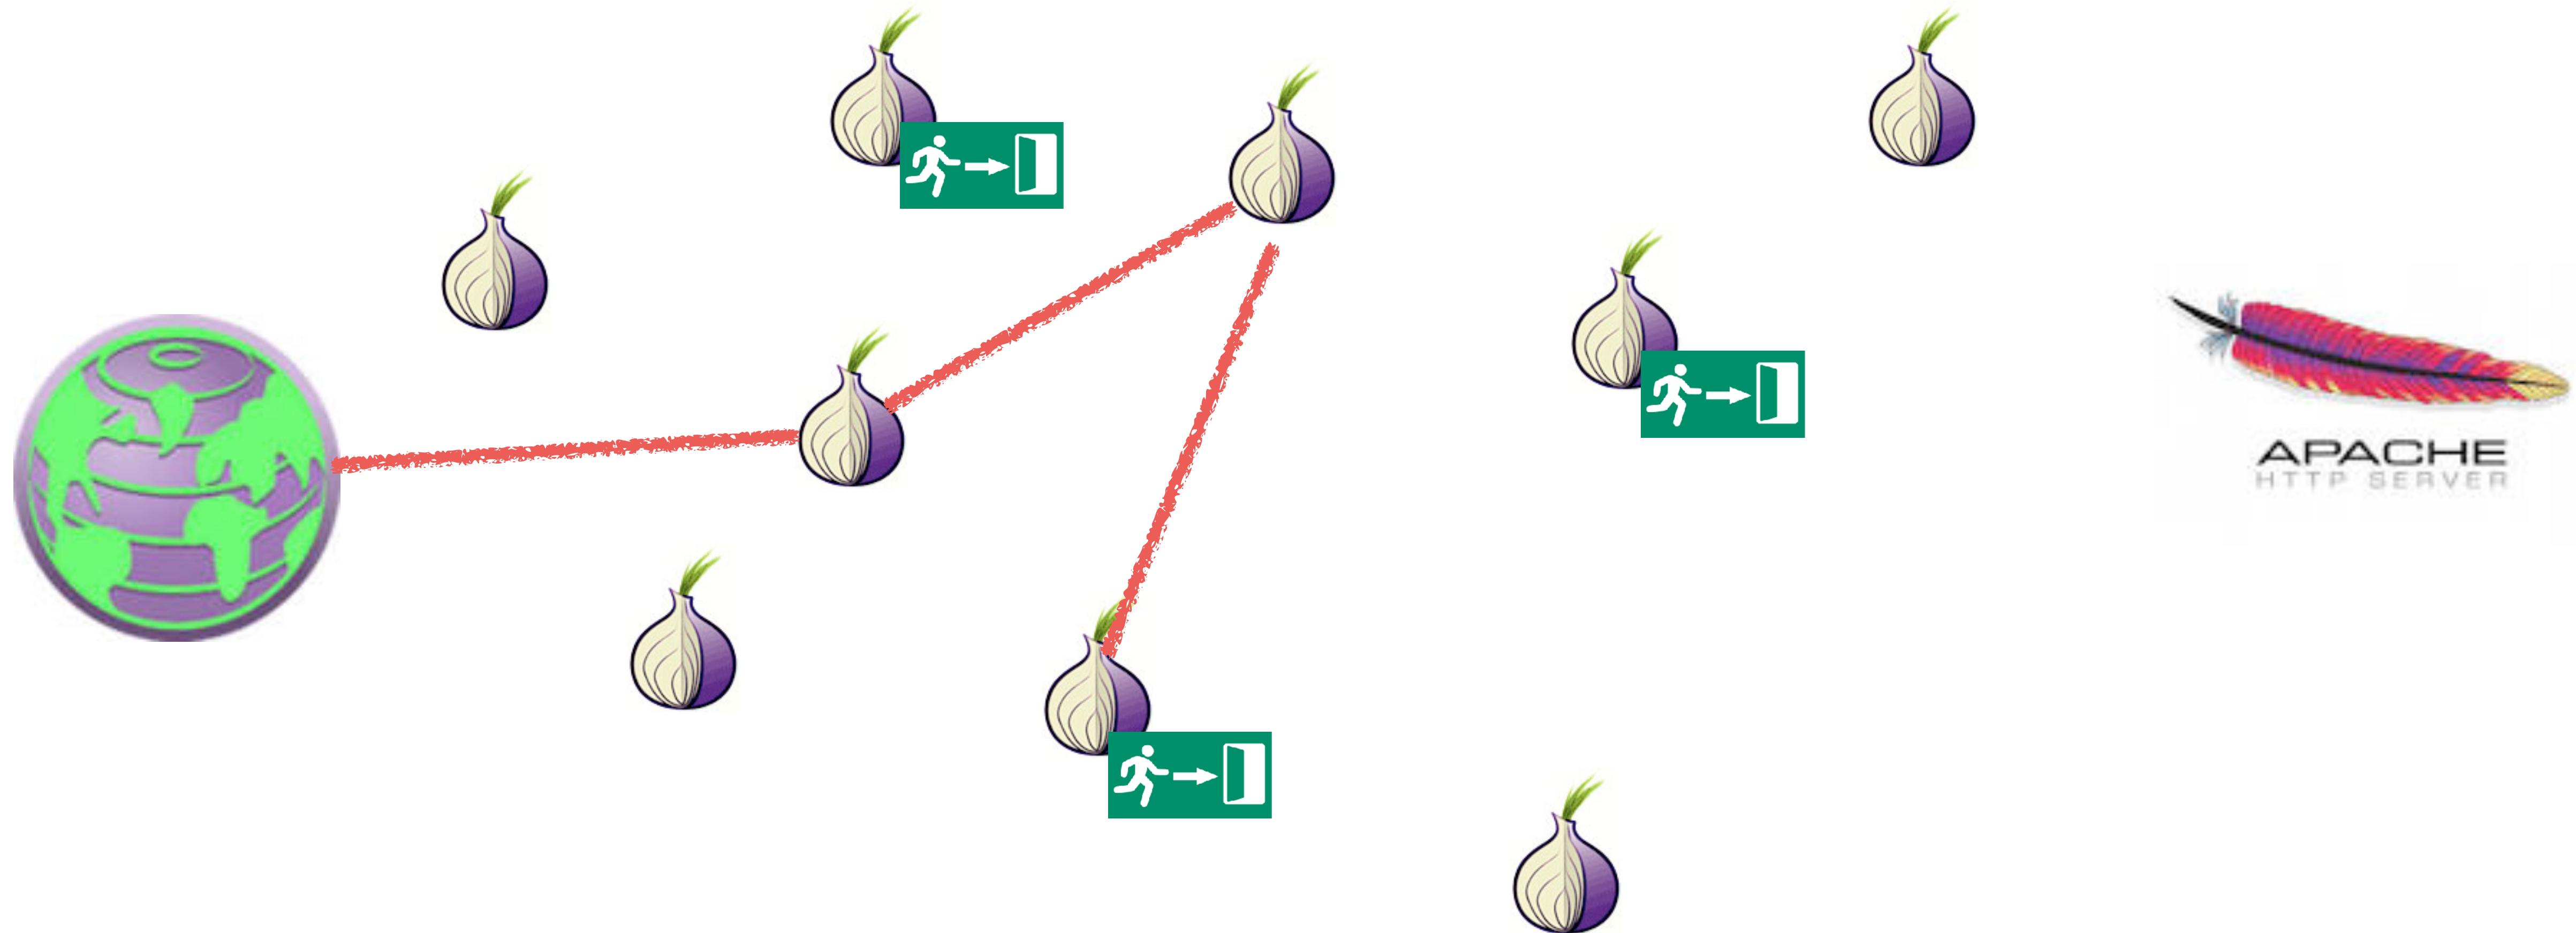
- The goal is to enable users to connect to other systems “anonymously” but with low latency
 - The remote system should have no way of knowing the IP address originating traffic
 - The local network should have no way of knowing the remote IP address the local user is contacting
- Important what is excluded:
The *global* adversary
 - Tor does not even attempt to counter someone who can see *all* network traffic:
It is probably *impossible* to do so and be low latency & efficient



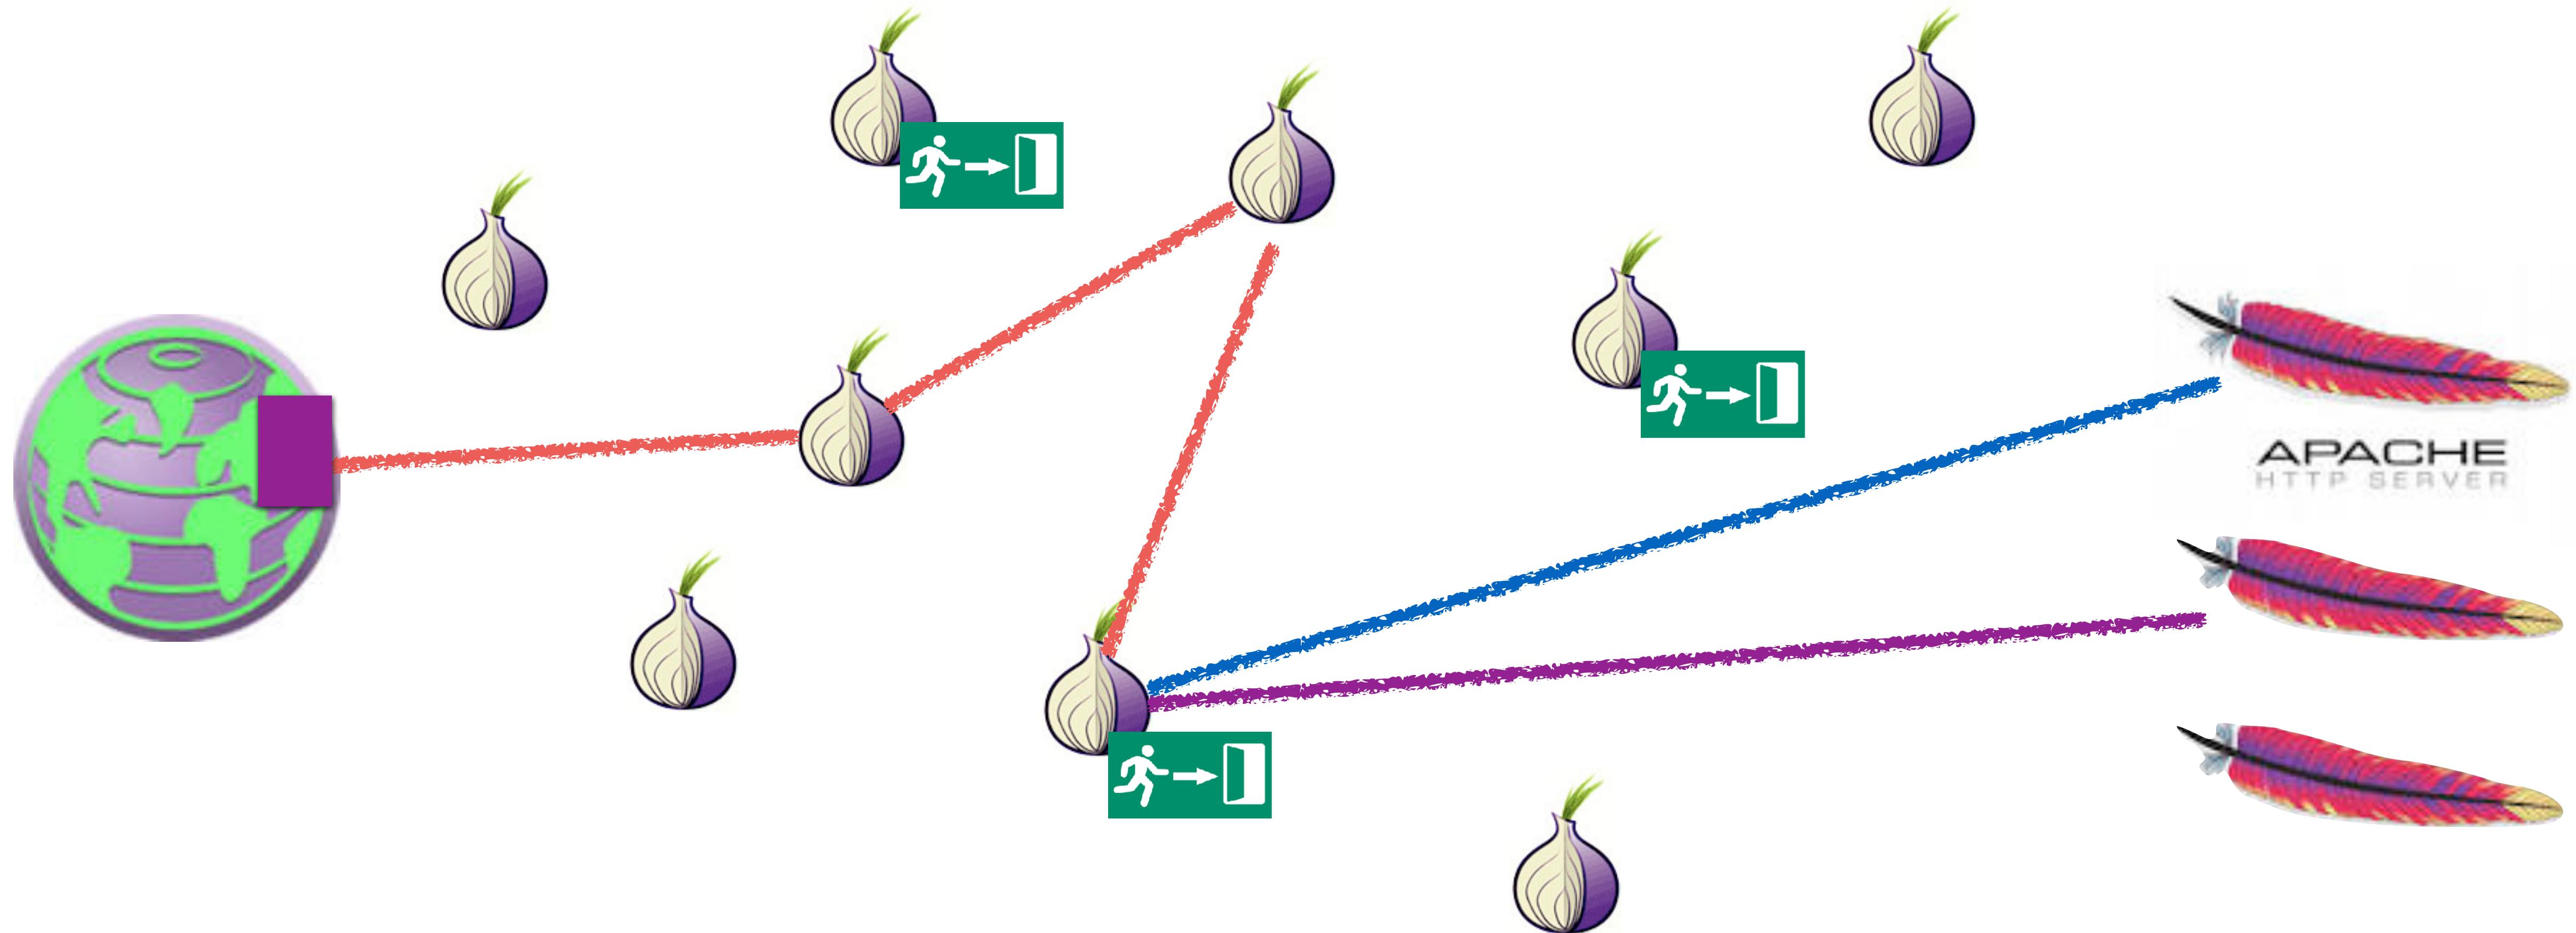
The High Level Approach: Onion Routing

- The Tor network consists of thousands of independent Tor nodes, or “Onion Routers”
 - Each node has a distinct public key and communicates with other nodes over TLS connections
 - A Tor circuit encrypts the data in a series of layers
 - Each hop away from the client removes a layer of encryption
 - Each hop towards the client adds a layer of encryption
- During circuit establishment, the client establishes a session key with the first hop...
 - And then with the second hop through the first hop
- The client has a ***global*** view of the Tor Network:
The directory servers provide a list of all Tor relays and ***their public keys***

Tor Routing In Action



Tor Routing In Action



Creating the Circuit Layers...

- The client starts out by using an authenticated DHE key exchange with the first node...
 - OR1 creates g^a , signs it with its private key, sends g^a , $\text{Sign}(K_{priv_or1}, g^a)$ to client
 - Client creates g^b , sends it to OR1
 - Client does $\text{Verify}(K_{pub_or1}, g^a)$
 - Creating a session key K_{OR1} as $H(g^{ab})$
 - This first hop is commonly referred to as the “guard node”
- It then tells OR1 to extend this circuit to OR2
 - Through that, creating a session key for the client to talk to OR2 that OR1 **does not know**
 - And OR2 doesn't know what the client is, just that it is somebody talking to OR1 requesting to extend the connection...
- It then tells OR2 to extend to OR3...
 - And OR1 won't know where the client is extending the circuit to, only OR2 will

Unwrapping the Onion

- Now the client sends some data...
 - $E(K_{or1}, E(K_{or2}, E(K_{or3}, \text{Data})))$
- OR1 decrypts it and passes on to OR2
 - $E(K_{or2}, E(K_{or3}, \text{Data}))$
- OR2 then passes it on...
- Generally go through at least 3 hops...
 - Why 3? So that OR1 can't call up OR2 and link everything trivially
 - Messages are a fixed-sized payload

The Tor Browser...

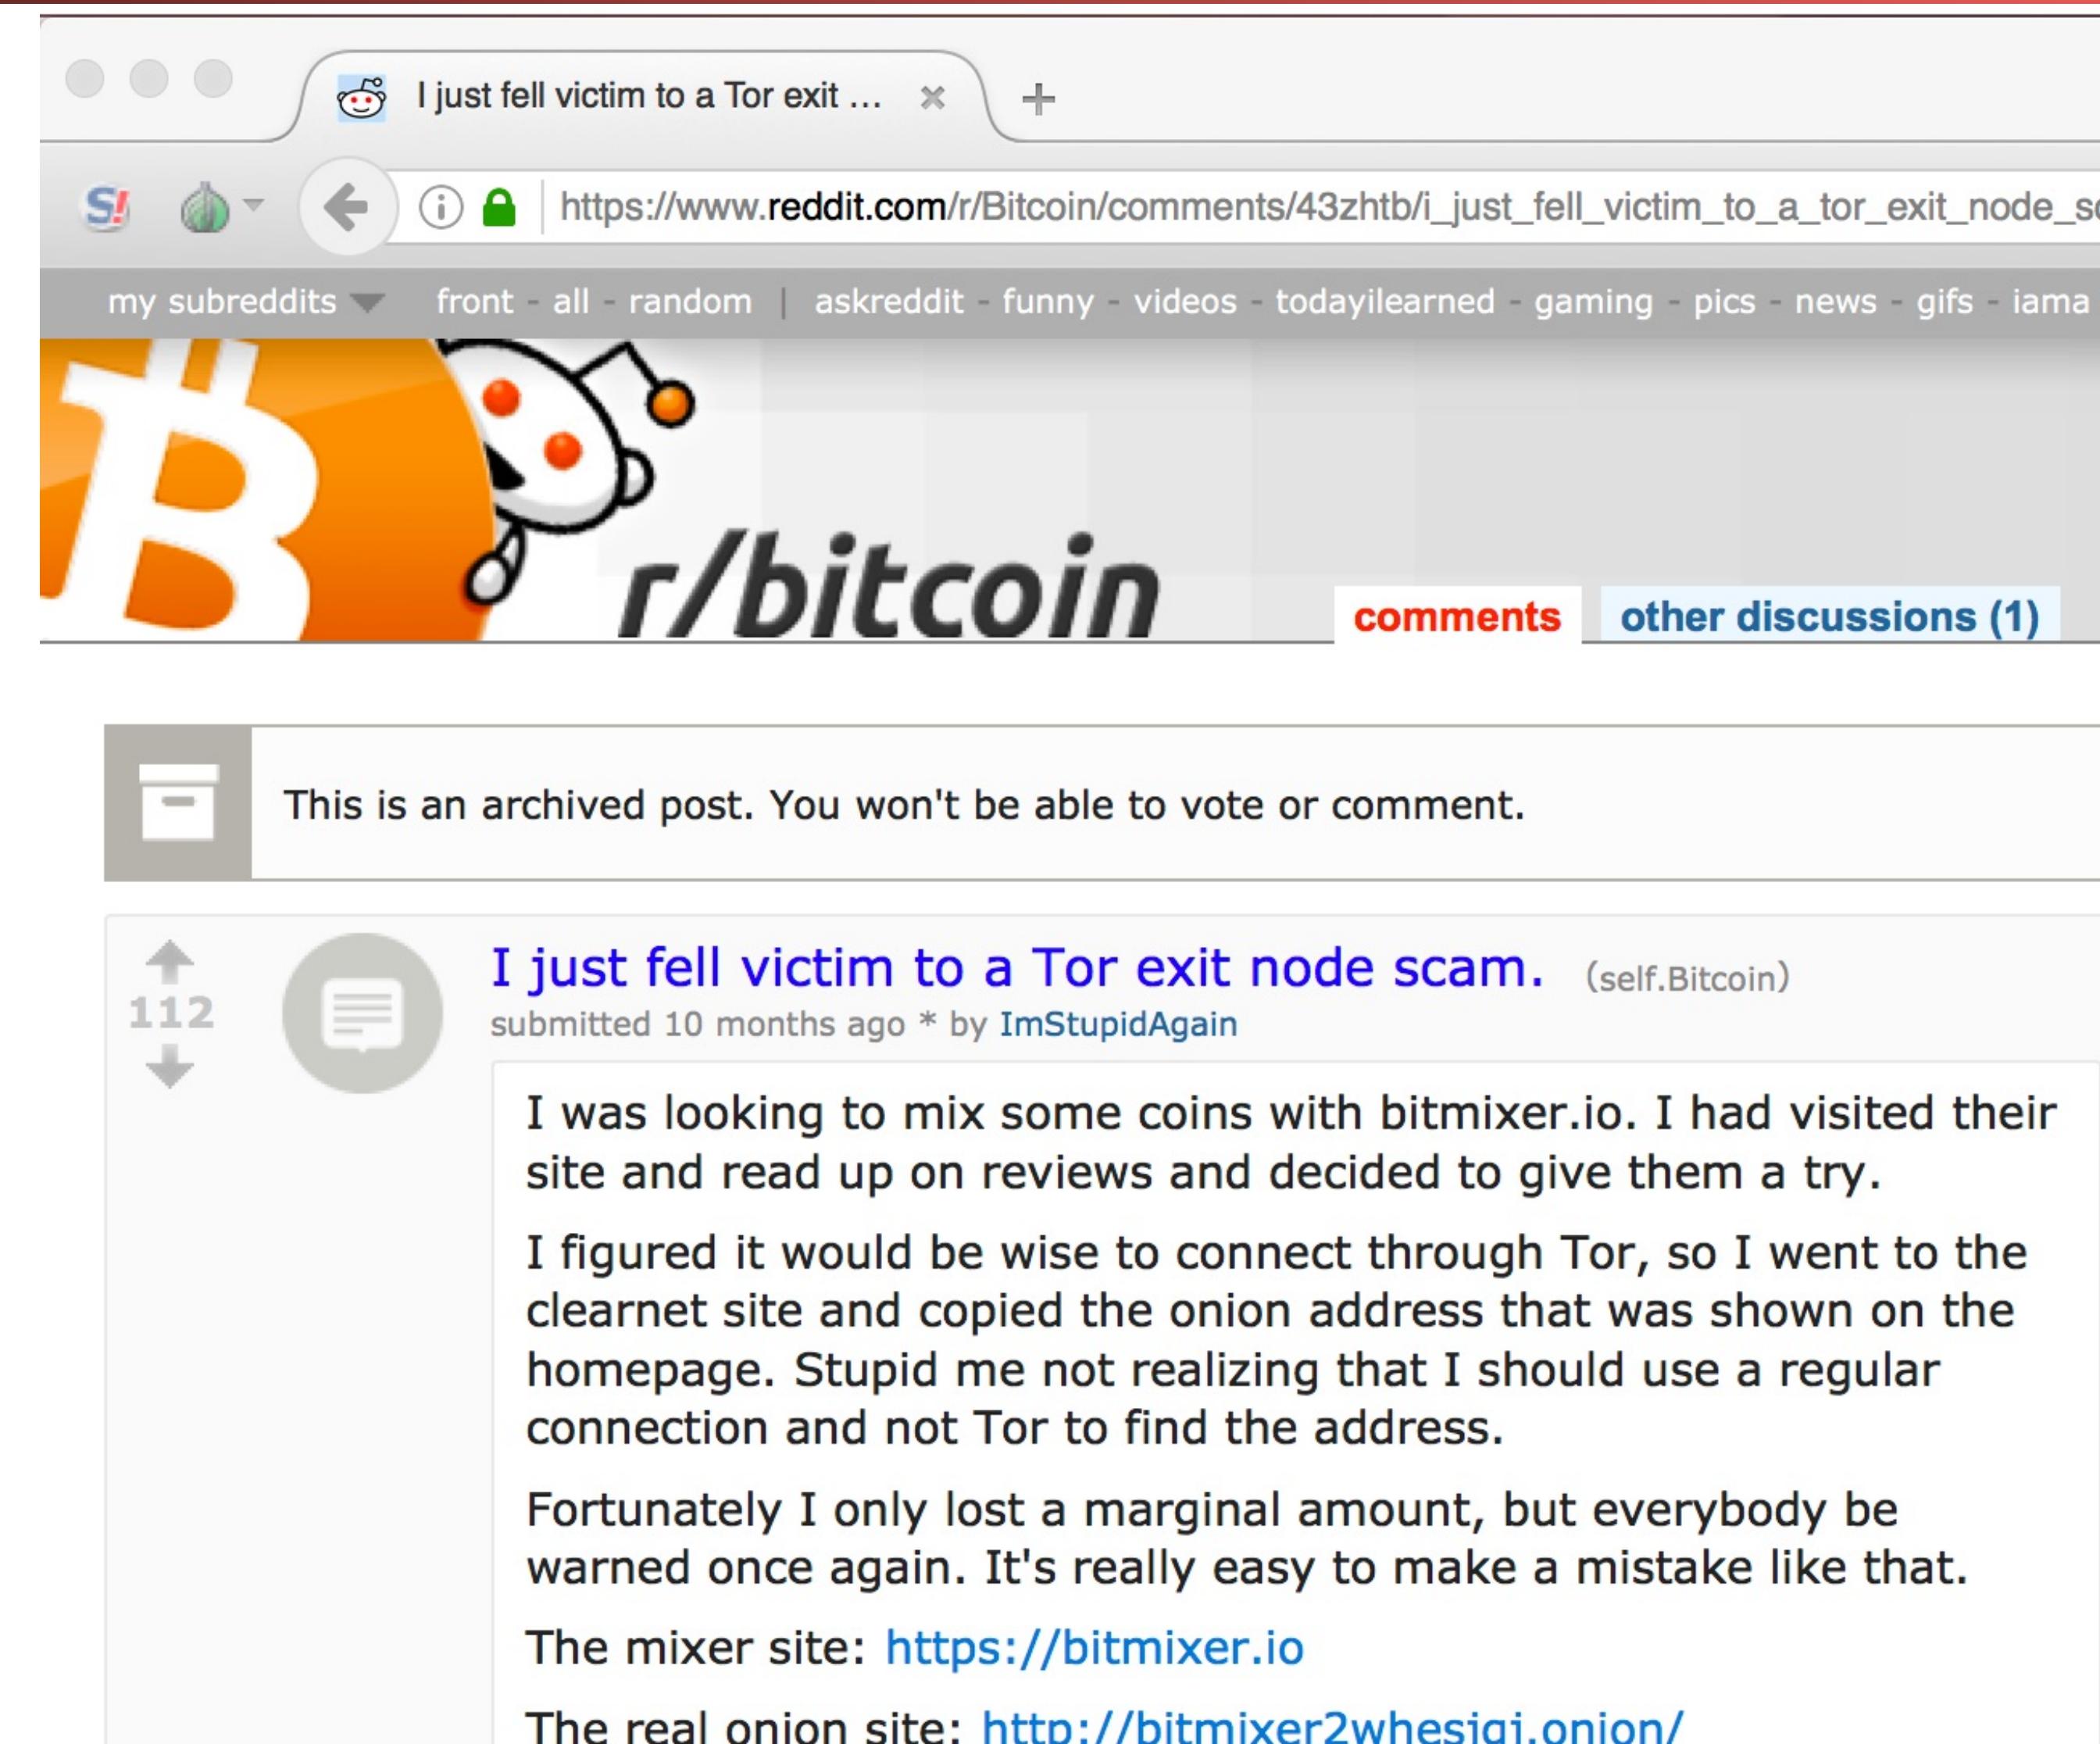
- Surfing “anonymously” doesn’t simply depend on hiding your connection...
- But also configuring the browser to make sure it resists tracking
 - No persistent cookies or other data stores
 - ***No deviations from other people*** running the same browser
- Anonymity ***only works in a crowd...***
 - So it really tries to make it all the same
 - But by default it makes it easy to say “this person is using Tor”

But You Are Relying On Honest Exit Nodes...

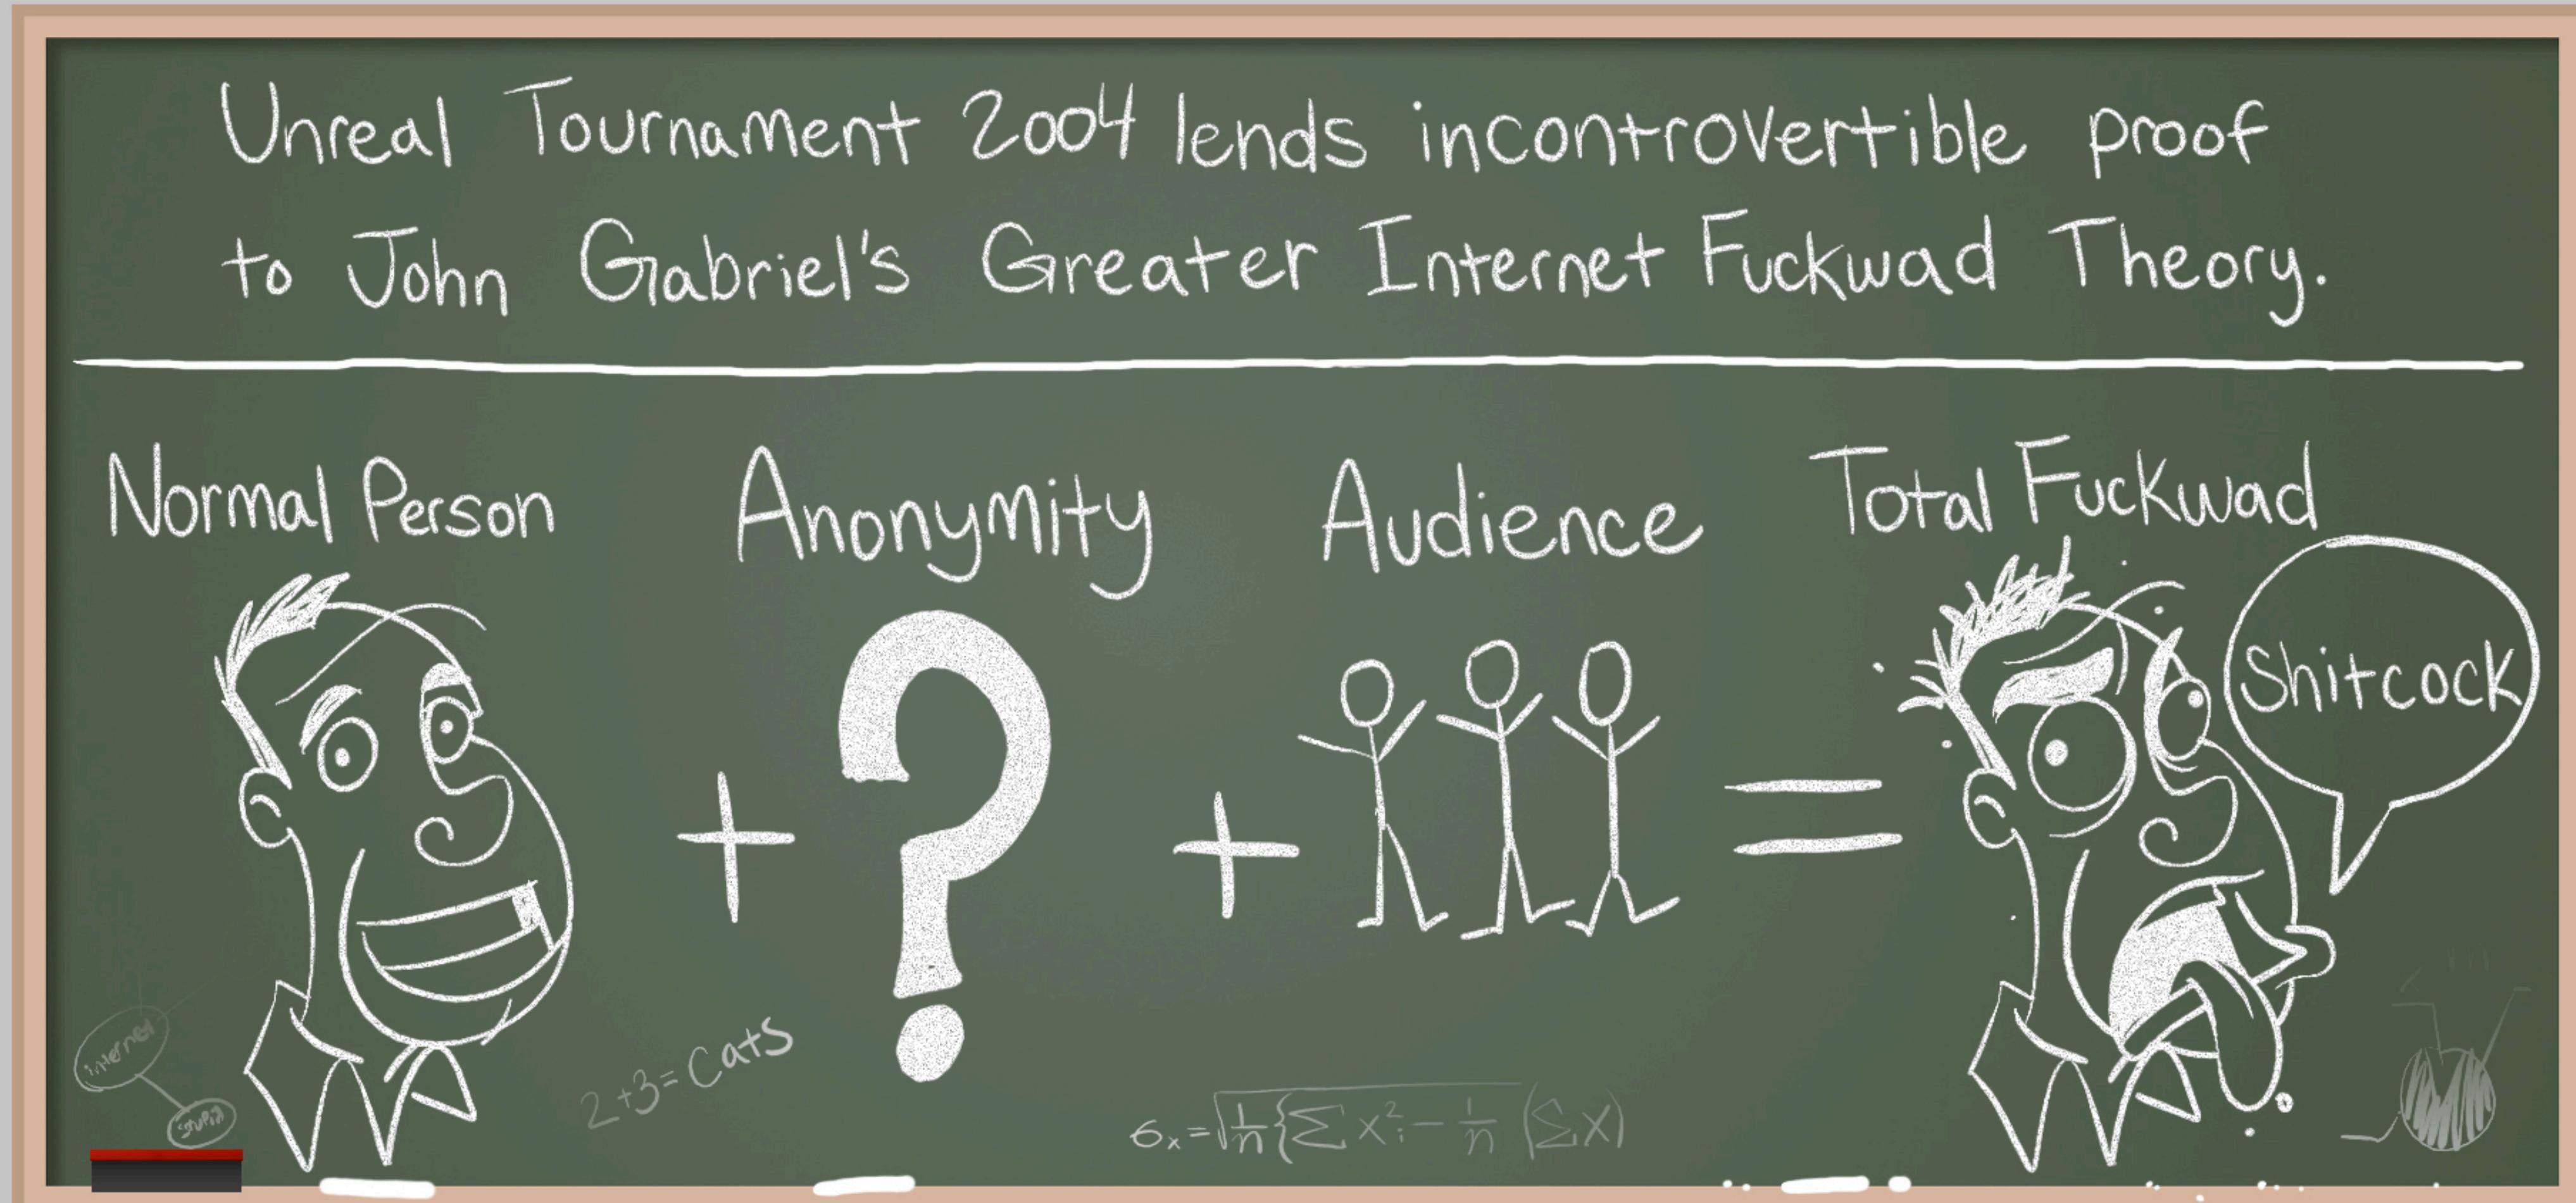
Computer Science 161 Fall 2019

Weaver

- The exit node, where your traffic goes to the general Internet, is a man-in-the-middle...
- Who can see and modify all non-encrypted traffic
- The exit node also does the DNS lookups
- Exit nodes have not always been honest...



Anonymity Invites Abuse... (Stolen from Penny Arcade)



This Makes Using Tor Browser Painful...

Computer Science 161 Fall 2019

Weaver



And Also Makes Running Exit Nodes Painful...

- If you want to receive abuse complaints...
 - Run a Tor Exit Node
- Assuming your ISP even allows it...
 - Since they don't like complaints either
- Serves as a large limit on Tor in practice:
 - Internal bandwidth is plentiful, but exit node bandwidth is restricted
- Know a colleague who ran an exit node for research...
 - And got a *visit from the FBI!*

One Example of Abuse: The Harvard Bomb Threat...

- On December 16th, 2013, a Harvard student didn't want to take his final in "Politics of American Education" ...
 - So he emailed a bomb threat using Guerrilla Mail
 - But he was "smart" and used Tor and Tor Browser to access Guerrilla Mail
- Proved easy to track
 - "Hmm, this bomb threat was sent through Tor..."
 - "So who was using Tor on the Harvard campus..." (look in Netflow logs..)
 - "So who is this person..." (look in authentication logs)
 - "Hey FBI agent, wanna go knock on this guy's door?!"
- There is no magic Operational Security (OPSEC) sauce...
 - And again, anonymity only works if there is a crowd

Censorship Resistance: Pluggable Transports

- Tor is really used by separate communities
 - Anonymity types who want anonymity in their communication
 - Censorship-resistant types who want to communicate despite government action
 - The price for "free" censorship evasion is that your traffic acts to hide other anonymous users
- Vanilla Tor fails the latter ***completely***
- So there is a framework to deploy bridges that encapsulate Tor over some other protocol
 - So if you are in a hostile network...
 - Lots of these, e.g. OBS3 (Obfuscating Protocol 3), OBS4, Meek...
 - But its an arm's race