

Captchas & The Net



**"Some speak of
an Armageddon;
A time when humans
will build machines
they neither understand
nor control."**

**To myself I whisper,
'We already have.'"**

- Taylor Swift


And We Call It "Machine Learning"

Announcements:

- Midterm grades released
- HW4 due Friday, October 23, 11:59 PM PT
- Project 2 design doc draft due Wednesday, October 28, 11:59 PM PT
- Really happy with the high grades on the midterm...
 - But for those unhappy, we ***will allow clobbering***:
Your scaled final grade will replace the midterm
 - The final is expected to be harder, so we want to scale your final grade as appropriate for midterm clobbering purposes
- Lets also "lock in" some additional constraints on the final curve...
 - Already present: Target GPA of 3.5...
If it goes over, 
 - Add in: Minimum bin size: 3.333% per step:
So 90% = A- guaranteed, 80% = B- guaranteed, 70% = C- guaranteed
In shifting down the bin size won't shrink, only expand

Election Week Planning...

- Yes, things are still shaping up for a 🦇💩-crazy week...
 - Prepare now: **VOTE NOW!**
 - Tension in the class: Maintain "normalcy" vs "omg we all want to curl up in a ball..."
 - Compromise:
 - Lecture will still happen, ***but it will be the dealers-choice lectures (not tested on the final)***
 - Discussion will still happen, ***but they will be just vent/chat sessions***
 - Attendance ***will not be taken***



From Last Year: Bug Of The Day...

- Yet Another Buffer overflow...
 - You think we'd be bored of them by now
- But...
 - The operating system and device drivers are special...
- They need very low level access
 - As they are working in a world where everything *is* just a big pile of bits!
- Perhaps you could use rust...
 - But you would need to rewrite a huge amount of code:
So most drivers are in C/C++



ars TECHNICA

NOTICE OF ABSENCE —

Unpatched Linux bug may open devices to serious attacks over Wi-Fi

Buffer overflow can be triggered in Realtek Wi-Fi chips, no user interaction needed.


DAN GOODIN - 10/17/2019, 2:35 PM

The Problem: Automation...

- You host some website...
- It is intended for ***human*** usage
 - One person, one mouse, one clickstream of behavior...
- But you want to lock out ***robot*** usage
- Why?
 - Selling something
 - Offering something for free
 - Dealing with load from an attack
- Enter the CAPTCHA:
A way to go “Is this a human?”

CAPTCHAs: How Lazy Cryptographers Do AI

- The whole point of CAPTCHAs is *not* just to solve "is this human" ...
 - But leverage bad guys to force them to solve hard problems
 - Primarily focused on machine vision problems



Visual code | [Audio code](#)

Help



Type the code shown


 [Try a new code](#)

By clicking the "Create My Account" button below, I certify that I have read and agree to the [Yahoo! Terms of Service](#), [Yahoo! Privacy Policy](#) and [Communication Terms of Service](#), and to receive account related communications from Yahoo! electronically. Yahoo! automatically identifies items such as words, links, people, and subjects from your Yahoo! communications services to deliver product features and relevant advertising.

Create My Account

CAPTCHAs

- *Reverse Turing Test*: present “user” a challenge that’s easy for a human to solve, hard for a program to solve
- One common approach: distorted text that’s difficult for character-recognition algorithms to decipher






(a) AOL.



(b) mail.ru



(c) phpBB 3.0



(d) Simple Machines Forum



(e) Yahoo!



(f) youku

Figure 1: Examples of CAPTCHAs from various Internet properties.

Problems?

vatinKes nipyous



Verify Your Registration

- Enter the code shown:

 [More info](#)

This helps prevent automated registrations.



Please enter the code you see below. [what's this?](#)

Qualifying question

Just to prove you are a human, please answer the following math challenge.

Q: Calculate:

$$\frac{\partial}{\partial x} \left[4 \cdot \sin \left(7 \cdot x - \frac{\pi}{2} \right) \right] \Big|_{x=0}$$

A:

mandatory

Note: If you do not know the answer to this question, reload the page and you'll get another question.

fuck Cloudflare

Issues with CAPTCHAs


- Inevitable arms race: as solving algorithms get better, defense erodes



Figure 4: Examples of images from the hard CAPTCHA puzzles dataset.

Issues with CAPTCHAs

- Inevitable arms race: as solving algorithms get better, defense erodes, or gets harder for humans



Asirra

Asirra is a human interactive proof that asks users to identify photos of cats and dogs. It's powered by over two million photos from our unique partnership with Petfinder.com. Protect your web site with Asirra — free!

Please click on the images that show cats:

The grid contains 16 images:

- Row 1: Dog (black and white), Dog (brown and tan), Dog (brown), Dog (brown and white)
- Row 2: Dog (yellow), Cat (orange tabby), Cat (orange tabby), Dog (black and white)
- Row 3: Dog (brown), Cat (white), Cat (orange tabby), Dog (black and white)
- Row 4: Dog (brown), Dog (brown and tan), Cat (grey), Dog (white)

Each image has a blue "adopt me" link below it.

Issues with CAPTCHAs


Computer Science 161 Fall 2020

- Inevitable arms race: as solving algorithms get better, it gets harder for humans

Security Check
Enter both words below, separated by a space.
Can't read the words below? Try different words or an audio cap




Text in the box:



- *Accessibility*: not all humans can see
- *Granularity*: not all bots are bad (e.g., crawlers)
- *Ambiguity*: No clear solution!

Issues with CAPTCHAs, con't

- Deepest problem: CAPTCHAs are inherently vulnerable to *outsourcing* attacks
 - Attacker gets real humans to solve them



captcha solving

X

Computer Science 161 Fall 2020

All

Videos

Images

News

Books

More

Settings

T

Weaver

About 5,360,000 results (0.40 seconds)

Ad · www.2captcha.com/fast/recognition ▾

Captcha solving service - Always Cheap

Stable quality service. Even for difficult to recognize images. Start use now! Try the quality of our service right now and enjoy. API in all languages. **Solution** time 9 seconds. Steps: Register, Implement Our API, Send Us Your CAPTCHAs, Get Your Answer As Text.

Ad · www.anti-captcha.com/ ▾

Captcha Solving Service - API

GitHub/npm/pip3 code packages, education tutorials, browser plugin, local payment methods.

\$0.5 per 1000 image **captchas**, \$1.5 for Recaptcha v2/v3. Huge amounts of trained workers.

[View Documentation](#).

[View Tutorials](#) · [Read The FAQs](#)

prowebscraper.com › blog › top-10-captcha-solving-servi...

Top 10 Captcha Solving Services Compared – ProWebScraper

1. Anticaptcha. Anticaptcha. Powered by 99% success rate and 7 seconds response time, Anticaptcha provides ...

Dec 19, 2017 · Uploaded by ProWebScraper

The image shows two browser tabs side-by-side. The left tab displays the homepage of Anti-Captcha (<https://anti-captcha.com/mainpage>). It features a superhero character in a blue and yellow suit with 'AC' on the chest, flexing his bicep. The page highlights four main benefits: 'Cheapest price on the market' (starting from 0.5USD per 1000 images), 'Pay as you go' (pay-per-captcha basis, minimum refill of 1 USD), '99.99% uptime since 2007' (vast amount of workers and premium infrastructure), and 'Solving Google Recaptcha since 2016' (stable solution). A 'Create Account' button is visible. The right tab shows the API documentation for the Anti-Captcha service (<https://api.anti-captcha.net/wiki/spaces/API/>). The title is 'Documentation in English'. It includes a profile picture of the administrator, last updated on Oct 17, 2018. Below this, it provides instructions for solving captchas via the API, listing three steps: creating a task, waiting for assignment, and requesting the solution. A flowchart at the bottom illustrates the process: 'createTask' leads to '...wait 5-10s...', which then leads to 'getTaskResult'.

createTask

- NoCaptchaTask
- NoCaptchaTaskProxyless
- ImageToTextTask

getTaskResult

Language	Example			AG	BC	BY	CB	DC	IT	All
English	one	two	three	51.1	37.6	4.76	40.6	39.0	62.0	39.2
Chinese (Simp.)	一	二	三	48.4	31.0	0.00	68.9	26.9	35.8	35.2
Chinese (Trad.)	一	二	三	52.0	24.4	0.00	62.9	30.2	22.0	34.1
Spanish	uno	dos	tres							
Italian	uno	due	tre							
Tagalog	isá	da	ta							
Portuguese	um	dois	três							
Russian	один	два	три							
Tamil	ஒன்று	இரண்டு	மூன்று							
Dutch	een	two	three							
Hindi	एक	दो	तीन							
German	eins	zwei	drei							
Malay	satu	dua	tiga	0.00	1.42	0.00	0.00	0.55	29.4	5.23
Vietnamese	một	hai	ba	0.46	2.07	0.00	0.00	1.74	18.1	3.72
Korean	일	이	삼	0.00	0.00	0.00	0.00	0.00	20.2	3.37
Greek	ένα	δύο	τρία	0.45	0.00	0.00	0.00	0.00	15.5	2.65
Arabic	واحد	اثنين	ثلاثة	0.00	0.00	0.00	0.00	0.00	15.3	2.56
Bengali	এক	দুই	তিনি	0.45	0.00	9.89	0.00	0.00	0.00	1.72
Kannada	ಒಂದು	ಎರಡು	ಮೂರು	0.91	0.00	0.00	0.00	0.55	6.14	1.26
Klingon	✓	<	↖	0.00	0.00	0.00	0.00	0.00	1.12	0.19
Farsi	یک	دو	سه	0.45	0.00	0.00	0.00	0.00	0.00	0.08

Re: CAPTCHAs – Understanding CAPTCHA-Solving Services in an Economic Context

Marti Motoyama, Kirill Levchenko, Chris Kanich, Damon McCoy,
 Geoffrey M. Voelker and Stefan Savage
 University of California, San Diego

{mmotoyam, klevchen, ckanich, dlmccoy, voelker, savage}@cs.ucsd.edu

Table 2: Percentage of responses from the services with correct answers for the language CAPTCHAs.

These Days: CAPTCHAs are ways of *training* AI systems

- Plus are all about an economic protection
 - Even the best CAPTCHA doesn't say "Is this a human or a bot"...
 - but...
 - "Is this a human or a bot willing to spend a couple pennies?"
- Acts as a hard limit on what a CAPTCHA can really protect!

TO COMPLETE YOUR REGISTRATION, PLEASE TELL US WHETHER OR NOT THIS IMAGE CONTAINS A STOP SIGN:



NO YES

ANSWER QUICKLY—OUR SELF-DRIVING CAR IS ALMOST AT THE INTERSECTION.

SO MUCH OF "AI" IS JUST FIGURING OUT WAYS TO OFFLOAD WORK ONTO RANDOM STRANGERS.

Network Security

- Why study network security?
 - Networking greatly extends our **overall attack surface**
 - Networking = the [Internet](#)
 - Opportunity to see *how large-scale design affects security issues*
 - Protocols a great example of *mindless agents* in action
- This lecture + next: sufficient background in networking to then explore security issues in next ~8 lectures
- Complex topic with many facets
 - We will omit concepts/details that aren't very security-relevant
 - But to no small extent we are speed running about 1/2 a dozen worth of "networking" lectures!
 - By all means, ask questions when things are unclear

Protocols

- A protocol is an **agreement** on how to communicate
- Includes **syntax** and **semantics**
 - How a communication is specified & structured
 - Format, order messages are sent and received
 - What a communication means
 - Actions taken when transmitting, receiving, or timer expires
- E.g.: making a comment in lecture in the Before Times?
 1. Raise your hand.
 2. Wait to be called on.
 3. Or: wait for speaker to **pause** and vocalize
 4. If unrecognized (after **timeout**): vocalize w/ “excuse me”

So Let's Do A Google Search...


- Walk into a coffee shop
- Open a laptop
- Search google...

Coffee Shop





1. Join the wireless network





1. Join the wireless network




Wireless access point(s)
continually shout:
**HEY, I'M WIRELESS
NETWORK Y, JOIN ME!**






1. Join the wireless network




If either match up, your laptop joins the network. Optionally performs a cryptographic exchange.






2. Configure your connection






2. Configure your connection





2. Configure your connection



192.168.1.14


The configuration includes:

- (1) An Internet address (**IP address**) your laptop should use; typ. 32 bits (IPv4). May also include 64b of the 128b IPv6 address
- (2) The address of a “**gateway**” system to use to access **hosts** beyond the local network
- (3) The address of a **DNS server (“resolver”)** to map names like google.com to IP addresses






3. Find the address of **google . com**




Your laptop sends a **DNS** request
asking: “***address for google . com?***”

It’s transmitted using the **UDP** protocol
(lightweight, unreliable).

The **DNS resolver** might not be on the
local network.




3. Find the address of google.com




3. Find the address of google.com

*The Rest of
the Internet*



3. Find the address of `google.com`


*The Rest of
the Internet*



3. Find the address of google.com


(The resolver now itself uses DNS queries to other DNS servers to figure out the address associated with google.com.)

The Rest of
the Internet




3. Find the address of google.com

The Rest of
the Internet



4. Connect to google.com server

*The Rest of
the Internet*




4. Connect to google.com server

The Rest of
the Internet

Your laptop now **establishes a connection** with the web server at 172.217.6.78. It uses **TCP** for this rather than UDP, to obtain reliability.


172.217.6.78



4. Connect to google.com server

The Rest of
the Internet

The first step of establishing the connection is to send a TCP connection request ("SYN") to the server.




4. Connect to google.com server

The Rest of
the Internet

If the server accepts the connection, it
replies with a "SYN ACK".


172.217.6.78



4. Connect to google.com server

Your laptop completes the connection establishment by likewise sending an acknowledgement.

The Rest of
the Internet




4. Connect to google.com server

The Rest of
the Internet


At this point the connection is established and data can be (reliably) exchanged.

172.217.6.78

The Rest of the Internet



The Rest of the Internet




5. Establish a secure
connection using **TLS**
([https](https://))

Coffee Shop

Well if you really possess the corresponding private key, prove it by decrypting this blob which we'll use to establish shared secret keys (RSA key exchange)



192.168.1.14



The Rest of
the Internet


5. Establish a secure
connection using TLS
(https)

Coffee Shop

Well if you really possess the corresponding private key, prove it by signing your Diffie/Hellman half of a DHE using the private key?
(DHE/ECDHE key exchange)




192.168.1.14




The Rest of
the Internet

5. Establish a secure
connection using TLS
(https)

The Rest of the Internet




The Rest of the Internet




6. Finally, your laptop can send along
your query!
(Using HTTP inside the **TLS channel**)

Layering

- Internet design is strongly partitioned into **layers**
 - Each layer relies on services provided by next layer below ...
 - ... and provides services to layer above it
- Analogy:
 - Consider structure of an application you've written and the “services” each layer relies on / provides



Internet Layering (“Protocol Stack”)/“OSI Model”)



Note on a point of potential confusion: these diagrams are always drawn with lower layers **below** higher layers ...

But diagrams showing the layouts of packets are often the *opposite*, with the lower layers at the **top** since their headers precede those for higher layers


(And nobody remembers what layers 5 and 6 are for (“Session” and “Presentation) for the trivia buffs because they aren’t really used)

(also, layer 8 is a “joke”, but really is important)


Packets and The Network

- Modern networks break communications up into packets
 - For our purposes, packets contain a variable amount of data up to a maximum specified by the particular network
- The sending computer breaks up the message and the receiving computer puts it back together
 - So the software doesn't actually see the packets per-se
 - Network itself is **packet switched**: sending each packet on towards its next destination
- Other properties:
 - Packets are received **correctly** or not at all in the face of **random** errors
 - The network does not enforce correctness in the face of adversarial inputs:
They are checksums not cryptographic MACs.
 - Packets may be **unreliable** and “dropped”
 - Its up to higher-level protocols to make the connection Reliable


Horizontal View of a Single Packet




Vertical View of a Single Packet




Internet Layering (“Protocol Stack”)



Layer 1: Physical Layer



Layer 2: Link Layer




Framing and transmission of a collection of bits into individual **messages** sent across a single “subnetwork” (one physical technology)

Might involve multiple *physical links* (e.g., modern Ethernet)

Often technology supports **broadcast** transmission (**every** “node” connected to subnet receives)

Layer 3: (Inter)Network Layer (*IP*)




Bridges multiple “subnets” to provide *end-to-end* **internet** connectivity between **nodes**

- Provides global addressing

Works across **different link technologies**

***Different* for each Internet “hop”**

Layer 4: Transport Layer




End-to-end communication between processes

Different services provided:
TCP = reliable byte stream
UDP = unreliable *datagrams*

(Datagram = single packet message)

Layer 7: Application Layer



Communication of
whatever you wish


Can use whatever
transport(s) is convenient

Freely structured

E.g.:

Skype, SMTP (email),
HTTP (Web), Halo, BitTorrent

4.5: Some Crypto...




TLS cryptography
(aka the 's' in HTTPS)


Often basically used as a
“layer 4.5” transport layer to
encrypt otherwise
unencrypted network
connections

Other times crypto may be at
the application layer (e.g. ssh


Internet Layering (“Protocol Stack”)



Internet Layering (“Protocol Stack”)




Internet Layering (“Protocol Stack”)




Hop-By-Hop vs. End-to-End Layers

Host A communicates with Host D




Hop-By-Hop vs. End-to-End Layers

Host A communicates with Host D



Hop-By-Hop vs. End-to-End Layers


Host A communicates with Host D



E.g., HTTP over TCP over IP

Same Network / Transport / Application Layers (3/4/7)
(Routers **ignore** Transport & Application layers)

Layer 3: (Inter)Network Layer (*IP*)




Bridges multiple “subnets” to provide *end-to-end* **internet** connectivity between **nodes**


- Provides global addressing

Works across **different link technologies**


IPv4 Packet Structure (IP version 6 is different)




IP Packet Structure




IP Packet Structure



IP Packet Structure



IP Packet Structure



IP Packet Header (Continued)

- Two IP addresses
 - Source IP address (32 bits in main IP version, IPv4)
 - Destination IP address (32 bits, likewise)
- Destination address
 - Unique **identifier/locator** for the receiving host
 - Allows each node to make forwarding decisions
- Source address
 - Unique identifier/locator for the sending host
 - Recipient can decide whether to accept packet
 - Enables recipient to send reply back to source

The Basic Ethernet Packet: The near-universal Layer 2

- An Ethernet Packet contains:
 - A preamble to synchronize data on the wire
 - We normally ignore this when talking about Ethernet
 - 6 bytes of destination MAC address
 - In this case, MAC means media access control address, not message authentication code!
 - 6 bytes of source MAC address
 - Optional 4-byte VLAN tag
 - 2 bytes length/type field
 - 46-1500B of payload

DST MAC	SRC MAC	VLAN	Type	PAYLOAD
---------	---------	------	------	---------

The MAC Address

- The MAC acts as a device identifier
 - The upper 3 bytes are assigned to a manufacturer
 - Can usually identify product with just the MAC address
 - The lower 3 bytes are assigned to a specific device
 - Making the MAC a de-facto serial #
- Usually written as 6 bytes in hex:
 - e.g. 13:37:ca:fe:f0:0d
- A device ***should ignore*** all packets that aren't to itself or to the broadcast address (**ff:ff:ff:ff:ff:ff**)
 - But almost all devices can go into ***promiscuous mode***
 - This is also known as "sniffing traffic"
- A device generally should only send with its own address
 - But this is enforced with software and can be trivially bypassed when you need to write "raw packets"