


Crypto 3



Administrivia...

- Project 1 question 3 starter code has a bug
 - See Piazza
- HW2 due Friday
- Project 1 due Friday the 19th

A Lot Of Uses for Random Numbers...


- The key foundation for all modern cryptographic systems is often not encryption but these "random" numbers!
- So many times you need to get something random:
 - A random cryptographic key
 - A random initialization vector
 - A "nonce" (use-once item)
 - A unique identifier
 - Stream Ciphers
- If an attacker can *predict* a random number things can catastrophically fail

Breaking Slot Machines

Computer Science 161

Weaver

- Some casinos experienced unusual bad "luck"
 - The suspicious players would wait and then all of a sudden try to play
- The slot machines have ***predictable*** pRNG
 - Which was based on the current time & a seed
- So play a little...
 - With a cellphone watching
 - And now you know when to press "spin" to be more likely to win
- Oh, and this ***never*** effected Vegas!
 - ***Evaluation standards*** for Nevada slot machines specifically designed to address this sort of issue



Breaking Bitcoin Wallets

Computer Science 161

Weaver

- blockchain.info supports "web wallets"
 - Javascript that protects your Bitcoin
- The private key for Bitcoin needs to be random
 - Because otherwise an attacker can spend the money
- An "Improvement" [sic] to the RNG reduced the entropy (the actual randomness)
 - Any wallet created with this improvement was brute-forceable and could be stolen

Improvements to RNG

zootreeves committed on Dec 7, 2014

1 parent b0d5639

Showing 1 changed file with 26 additions and 28 deletions.

S4 bitcoinjs-lib/src/jspb/rng.js

```
@@ -8,15 +8,16 @@ var rng_state;
 8   8   var rng_pool;
 9   9   var rng_pptr;
10  10
11 // Mix in a 32-bit integer into the pool
12 -function rng_seed_int(x) {
13 -  rng_pool[rng_pptr++] ^= x & 255;
14 -  rng_pool[rng_pptr++] ^= (x >> 8) & 255;
15 -  rng_pool[rng_pptr++] ^= (x >> 16) & 255;
16 -  rng_pool[rng_pptr++] ^= (x >> 24) & 255;
```



TRUE Random Numbers

- True random numbers generally require a physical process
- Common circuit is an unusable ring oscillator built into the CPU
 - It is then sampled at a low rate to generate true random bits which are then fed into a pRNG on the CPU
- Other common sources are human activity measured at very fine time scales
 - Keystroke timing, mouse movements, etc
 - "Wiggle the mouse to generate entropy for a key"
 - Network/disk activity which is often human driven
- More exotic ones are possible:
 - Cloudflare has a wall of lava lamps that are recorded by a HD video camera which views the lamps through a rotating prism: It is just one source of the randomness



Combining Entropy

- Many physical entropy sources are biased
 - Some have significant biases: e.g. a coin that flips "heads" 90% of the time!
 - Some aren't very good: e.g. keystroke timing at a microsecond granularity
- The general procedure is to combine various sources of entropy
- The goal is to be able to take multiple crappy sources of entropy
 - Measured in how many bits:
A single flip of a fair coin is 1 bit of entropy
 - And combine into a value where the entropy is the minimum of the sum of all entropy sources (maxed out by the # of bits in the hash function itself)
 - **N-1** bad sources and **1** good source -> good pRNG state

Pseudo Random Number Generators (aka Deterministic Random Bit Generators)

- Unfortunately one needs a *lot* of random numbers in cryptography
 - More than one can generally get by just using the physical entropy source
- Enter the pRNG or DRBG
 - If one knows the state it is entirely predictable
 - If one doesn't know the state it should be indistinguishable from a random string
- Three operations
 - Instantiate: (aka Seed) Set the internal state based on the real entropy sources
 - Reseed: Update the internal state based on both the previous state and **additional entropy**
 - The big different from a simple stream cipher
 - Generate: Generate a series of random bits based on the internal state
 - Generate can also optionally add in additional entropy
- **instantiate(entropy)**
reseed(entropy)
generate(bits, {optional entropy})

Properties for the pRNG

- Can a pRNG be truly random?
 - No. For seed length s , it can only generate at most 2^s distinct possible sequences.
- A cryptographically strong pRNG “looks” truly random to an attacker
 - Attacker ***cannot distinguish*** it from a random sequence:
If the attacker can tell a sufficiently long bitstream was generated by the pRNG instead of a truly random source it isn't a good pRNG

Prediction and Rollback Resistance

- A pRNG should be predictable only if you know the internal state
 - It is this predictability which is why its called "pseudo"
- If the attacker does not know the internal state
 - The attacker should not be able to distinguish a truly random string from one generated by the pRNG
- It **should** also be rollback-resistant
 - Even if the attacker finds out the state at time T, they should not be able to determine what the state was at T-1
 - More precisely, if presented with two random strings, one truly random and one generated by the pRNG at time T-1, the attacker should not be able to distinguish between the two
 - Rollback resistance isn't specifically required in a pRNG...
But it should be

Why "Rollback Resistance" is Essential

- Assume attacker, at time T , is able to obtain all the internal state of the pRNG
 - How? E.g. the pRNG screwed up and instead of an IV, released the internal state, or the pRNG is bad...
 - Attacker observes how the pRNG was used
 - T_{-1} = Random Session key
 T_0 = Nonce/IV
 - Now if the pRNG doesn't resist rollback, and the attacker gets the state at T_0 , attacker can know the session key! And we are back to...



More on Seeding and Reseeding

- Seeding should take all the different physical entropy sources available
 - If one source has 0 entropy, it ***must not*** reduce the entropy of the seed
 - We can shove a whole bunch of low-entropy sources together and create a high-entropy seed
- Reseeding ***adds*** in even more entropy
 - **F(internal_state, new material)**
 - Again, even if reseeding with 0 entropy, it ***must not*** reduce the entropy of the seed

Probably the best pRNG/DRBG: HMAC_DRBG

- Generally believed to be the best
 - ***Accept no substitutes!***
- Two internal state registers, V and K
 - Each the same size as the hash function's output
- V is used as (part of) the data input into HMAC, while K is the key
- If you can break this pRNG you can ***either break the underlying hash function or break a significant assumption about how HMAC works***
 - Yes, security proofs sometimes are a very good thing and actually do work
 - So as long as the security proof for HMAC is correct, the security proof for HMAC_DRBG is correct!

HMAC_DRBG

Generate

- The basic generation function
- Remarks:
 - It requires one HMAC call per blocksize-bits of state
 - Then two more HMAC calls to update the internal state
- Prediction resistance:
 - If you can distinguish new **K** from random when you don't know old **K**:
You've distinguished HMAC from a random function!
Which means you've either broken the hash or the HMAC construction
- Rollback resistance:
 - If you can learn old **K** from new **K** and **V**:
You've reversed the hash function!

```
function hmac_drbg_generate (state, n) {  
    tmp = ""  
    while(len(tmp) < N) {  
        state.v = hmac(state.k, state.v)  
        tmp = tmp || state.v  
    }  
    // Update state with no input  
    state.k = hmac(state.k, state.v || 0x00)  
    state.v = hmac(state.k, state.v)  
    // Return the first N bits of tmp  
    return tmp[0:N]  
}
```

HMAC_DRBG

Update

- Used instead of the "no-input update" when you have additional entropy on the generate call
- Used standalone for both instantiate (**state.k = state.v = 0**) and reseed (keep **state.k** and **state.v**)
- Designed so that even if the attacker controls the input but doesn't know **k**:
The attacker should not be able to predict the new **k**

```
function hmac_drbg_update (state, input) {  
    state.k = hmac(state.k, state.v || 0x00  
                  || input)  
    state.v = hmac(state.k, state.v)  
    state.k = hmac(state.k, state.v || 0x01  
                  || input)  
    state.v = hmac(state.k, state.v)  
}
```

Generating ***true*** random numbers

- Modern CPUs have true random number generators
 - Sample a noisy circuit at a low rate or similar tricks
- These sources are biased...
 - They are also ***slow***
- So use this as an entropy source to feed a pRNG on the chip
 - Now you can get random numbers quickly
- Very fast
- Vulnerable to tampering!
 - You can't actually test that the pRNG circuit is 100% correct without adding paths that could potentially sabotage the pRNG circuit
 - Sabotage that can reduce effective entropy to 32b are possible

Stream ciphers


- Block cipher: fixed-size, stateless, requires “modes” to securely process longer messages
- Stream cipher: keeps state from processing past message elements, can continually process new elements
- Common approach: “one-time pad on the cheap”:
 - XORs the plaintext with some “random” bits
 - But: random bits \neq the key (as in one-time pad)
 - Instead: output from cryptographically strong pseudorandom number generator (pRNG)
 - Anyone who actually calls this a “One Time Pad” is selling snake oil!

Building Stream Ciphers

- Encryption, given key **K** and message **M**:
 - Choose a random value **IV**
 - $E(M, K) = \text{pRNG}(K, IV) \oplus M$
- Decryption, given key **K**, ciphertext **C**, and initialization vector **IV**:
 - $D(C, K) = \text{PRNG}(K, IV) \oplus C$
- Can encrypt message of any length because pRNG can produce any number of random bits...
- But in practice, for an n-bit seed pRNG, stop at $2^{n/2}$. Because, of course...



Using a pRNG to Build A Stream Cipher



CTR mode is (mostly) a stream cipher

- $E(ctr, K)$ should look like a series of pseudo random numbers...
 - But after a large amount it is *slightly* distinguishable!
 - Since it is actually a pseudo-random ***permutation***...
 - For a cipher using 128b blocks, you will never get the same 128b number until you go all the way through the 2^{128} possible entries on the counter
 - Reason why you want to stop after 2^{64}
 - If you use CTR mode in the first place
 - Also very minor information leakage:
 - If $C_i = C_j$, for $i \neq j$, it follows that $M_i \neq M_j$

UUID: Universally Unique Identifiers

- You got to have a "name" for something...
 - EG, to store a location in a filesystem
- Your name ***must*** be unique...
 - And your name ***must*** be unpredictable!
- Just chose a ***random*** value!
 - UUID: just chose a 128b random value
 - Well, it ends up being a 122b random value with some signaling information
 - A good UUID library uses a cryptographically-secure pRNG that is properly seeded
- Often written out in hex as:
 - 00112233-4455-6677-8899-aabbccddeeff

What Happens When The Random Numbers Goes Wrong...

- Insufficient Entropy:
 - Random number generator is seeded without enough entropy
- Debian OpenSSL CVE-2008-0166
 - In "cleaning up" OpenSSL (Debian 'bug' #363516), the author 'fixed' how OpenSSL seeds random numbers
 - Because the code, as written, caused Purify and Valgrind to complain about reading uninitialized memory
 - Unfortunate cleanup reduced the pRNG's seed to be **just** the process ID
 - So the pRNG would only start at one of ~30,000 starting points
- This made it easy to find private keys
 - Simply set to each possible starting point and generate a few private keys
 - See if you then find the corresponding public keys anywhere on the Internet



<http://blog.dieweltistgarnichtso.net/Caprica,-2-years-ago> 22

And Now Lets Add Some RNG Sabotage...

- The Dual_EC_DRBG
 - A pRNG pushed by the NSA behind the scenes based on Elliptic Curves
 - It relies on two parameters, P and Q on an elliptic curve
 - The person who generates P and selects $Q=eP$ can predict the random number generator, regardless of the internal state
 - It also **sucked!**
 - It was horribly slow and even had subtle biases that shouldn't exist in a pRNG:
You could distinguish the upper bits from random!
 - Now this was spotted fairly early on...
 - Why should anyone use such a horrible random number generator?

Well, anyone not paid that is...

- RSA Data Security accepted 30 pieces of silver \$10M from the NSA to implement Dual_EC in their RSA BSAFE library
 - And *silently* make it the default pRNG
 - Using RSA's support, it became a NIST standard
 - And inserted into other products...
 - And then the Snowden revelations
 - The initial discussion of this sabotage in the NY Times just vaguely referred to a Crypto talk given by Microsoft people...
 - That everybody quickly realized referred to Dual_EC



But this is insanely powerful...

- It isn't just forward prediction but being able to run the generator backwards!
 - Which is why Dual_EC is so nasty:
Even if you know the internal state of HMAC_DRBG it has rollback resistance!
- In TLS (HTTPS) and Virtual Private Networks you have a motif of:
 - Generate a random session key
 - Generate some other random data that's **public visible**
 - EG, the IV in the encrypted channel, or the "random" nonce in TLS
 - Oh, and an NSA sponsored "standard" to spit out even more "random" bits!
- If you can run the random number generator **backwards**, you can find the session key



It Got Worse: Sabotaging Juniper

Computer Science 161

Weaver

- Juniper also used Dual_EC in their Virtual Private Networks
 - "But we did it safely, we used a different **Q**"
- Sometime later, someone else noticed this...
 - "Hmm, **P** and **Q** are the keys to the backdoor... Lets just hack Juniper and rekey the lock!"
 - And whoever put in the first Dual_EC then went "Oh crap, we got locked out but we can't do anything about it!"
- Sometime later, someone else goes...
 - "Hey, lets add an ssh backdoor"
- Sometime later, Juniper goes
 - "Whoops, someone added an ssh backdoor, lets see what else got F'ed with, oh, this # in the pRNG"
- And then everyone else went
 - "Ohh, patch for a backdoor. Lets see what got fixed. Oh, these look like Dual_EC parameters..."



Sabotaging "Magic Numbers" In General

- Many cryptographic implementations depend on "magic" numbers
 - Parameters of an Elliptic curve
 - Magic points like **P** and **Q**
 - Particular prime **p** for Diffie/Hellman
 - The content of S-boxes in block ciphers
- Good systems should cleanly describe how they are generated
 - In some sound manner (e.g. AES's S-boxes)
 - In some "random" manner defined by a pRNG with a specific seed
 - Eg, seeded with "Nicholas Weaver Deserves Perfect Student Reviews" ...
Needs to be very low entropy so the designer can't try a gazillion seeds

Because Otherwise You Have Trouble...

- Not only Dual-EC's **P** and **Q**
- Recent work: 1024b Diffie/Hellman moderately impractical...
 - But you can create a sabotaged prime that is 1/1,000,000 the work to crack!
And the most often used "example" **p**'s origin is lost in time!
- It can cast doubt **even when a design is solid:**
 - The DES standard was developed by IBM but with input from the NSA
 - Everyone was suspicious about the NSA tampering with the S-boxes...
 - They did: The NSA made them **stronger** against an attack they knew but the public didn't
 - The NSA-defined elliptic curves P-256 and P-384
 - I trust them because they are in CNSA so the NSA uses them for TS communication:
A backdoor here would be absolutely unacceptable...
but only because I actually believe the NSA wouldn't go and try to shoot itself in the head!



So What To Use?

- AES-128-CFB or AES-256-CFB:
 - Robust to screwups encryption
 - Alternately, AES-128-GCM (Galios Counter Mode):
An AEAD mode, but is ***NOT resistant to screwups***
- SHA-2 or SHA-3 family (256b, 384b, or 512b):
 - Robust cryptographic hashes, SHA-1 and MD5 are broken
- HMAC-SHA256 or HMAC-SHA3:
 - Different function than the encryption:
Prevents screwups on using the same key & is a hash if not using an AEAD mode
 - ***Always Encrypt Then MAC!***
- HMAC-SHA256-DRBG or HMAC-SHA3-DRBG:
 - The best pRNG available
 - Seed using ***both*** the processor random number generator AND other entropy sources!
 - Don't use the processor RNG bare when building a software cryptosystem:
Those are potentially sabotage able and use designs without rollback resistance.

Public Key...

- All our previous primitives required a "miracle":
 - We somehow have to have Alice and Bob get a shared k .
- Enter Public Key cryptography: the miracle of modern cryptography
 - How starting Friday, but ***what*** today
- Three primitives:
 - Public Key Agreement
 - Public Key Encryption
 - Public Key Signatures
- Based on some families of magic math...
 - For us, we will use some group-theory based primitives

Public Key Agreement

- Alice and Bob have a channel...
 - There may be an eavesdropper ***but not a manipulator***
- The goal: Alice & Bob agree on a ***random*** value
 - This will be ***k*** for all subsequent communication
- When done, the key is thrown away
 - Designed to prevent an attacker who later recovers Alice or Bob's long lived secrets from finding ***k***.

Public Key Encryption

- Alice has **two** keys:
 - K_{pub} : Her public key, anyone can know
 - K_{priv} : Her private key, a deep dark secret
- Anyone has access to Alice's public key
- For anyone to send a message to Alice:
 - Create a random session key k
 - Used to encrypt the rest of the message
 - Encrypt k using Alice's K_{pub} .
- Only Alice can **decrypt** the message
 - The decryption function only works with K_{priv} !

Public Key Signatures

- Once again, Alice has **two** keys:
 - K_{pub} : Her public key, anyone can know
 - K_{priv} : Her private key, a deep dark secret
- She can sign a message
 - Calculate $H(M)$
 - $S(K_{priv}, H(M))$: Sign $H(M)$ with K_{priv} .
- Anyone can now verify
 - Recalculate $H(M)$
 - $V(K_{pub}, S(K_{priv}, H(M)), H(M))$: Verify that the signature was created with K_{priv}

Things To Remember...

- Public key is **slow!**
 - Orders of magnitude slower than symmetric key
- Public key is based on delicate magic math
 - Discrete log in a group is the most common
 - RSA
 - Some new "post-quantum" magic...
- Some systems in particular are easy to get wrong
 - We will get to some of the epic crypto-fails later

Our Roadmap For Public Key...

- Public Key:
 - Something **everyone** can know
- Private Key:
 - The secret belonging to a specific person
- Diffie/Hellman:
 - Provides key exchange with no pre-shared secret
- ElGamal & RSA:
 - Provide a message to a recipient only knowing the recipient's **public key**
- DSA & RSA signatures:
 - Provide a message that anyone can prove was generated with a **private key**

Diffie-Hellman Key Exchange


- What if instead they can somehow generate a random key when needed?
- Seems impossible in the presence of Eve observing all of their communication ...
 - How can they exchange a key without her learning it?
 - But: actually is possible using public-key technology
 - Requires that Alice & Bob know that their messages will reach one another without any meddling
 - Protocol: Diffie-Hellman Key Exchange (DHE)
 - The E is "Ephemeral", we use this to create a temporary key for other uses and then forget about it

Diffie-Hellman Key Exchange



1. Everyone agrees in advance on a well-known (large) prime p and a corresponding g : $1 < g < p-1$


DHE



2. Alice picks random secret ' a ': $1 < a < p-1$

3. Bob picks random secret ' b ': $1 < b < p-1$


DHE



$$\mathbf{A} = g^a \bmod p$$

4. Alice sends $A = g^a \bmod p$ to Bob $g^b \bmod p = \mathbf{B}$
5. Bob sends $B = g^b \bmod p$ to Alice

DHE



$$A = g^a \text{ mod } p$$

B

6. Alice knows $\{a, A, B\}$, computes

$$K = B^a \text{ mod } p = (g^b)^a = g^{ba} \text{ mod } p$$

7. Bob knows $\{b, A, B\}$, computes


$$K = A^b \text{ mod } p = (g^a)^b = g^{ab} \text{ mod } p$$

8. K is now the shared secret key.

A

$$g^b \text{ mod } p = \mathbf{B}$$

DHE



While Eve knows $\{p, g, g^a \text{ mod } p, g^b \text{ mod } p\}$, believed to be **computationally infeasible** for her to then deduce $K = g^{ab} \text{ mod } p$.

She can easily construct $A \cdot B = g^a \cdot g^b \text{ mod } p = g^{a+b} \text{ mod } p$.

But computing g^{ab} requires ability to take *discrete logarithms* mod p .

Discrete log over the group defined by p and g **presumed** to be hard

This is Ephemeral Diffie/Hellman

- $K = g^{ab} \text{ mod } p$ is used as the basis for a "session key"
 - A symmetric key used to protect subsequent communication between Alice and Bob
 - In general, public key operations are vastly more expensive than symmetric key, so it is mostly used just to agree on secret keys, transmit secret keys, or sign hashes
 - If either a or b is random, K is random
 - When Alice and Bob are done, they discard K, a, b
 - This provides ***forward secrecy***: Alice and Bob don't retain any information that a later attacker who can compromise Alice or Bob's secrets could use to decrypt the messages exchanged with K .

Diffie Hellman is part of more generic problem


- This involved deep mathematical voodoo called "Group Theory"
 - Its actually done under a group G
- Two main groups of note:
 - Numbers mod p with generator g
 - Point addition in an elliptic curve C
 - Usually identified by number, eg. p256, p384 (NSA-developed curves) or Curve25519 (developed by Dan Bernstein, also 256b long)
- So EC (Elliptic Curve) == different group
 - Thought to be harder so fewer bits: 384b ECDHE ?= 3096b DHE
 - But otherwise, its "add EC to the name" for something built on discrete log

But Its Not That Simple

- What if Alice and Bob aren't facing a passive eavesdropper
 - But instead are facing Mallory, an **active** Man-in-the-Middle
- Mallory has the ability to change messages:
 - Can remove messages and add his own
- Lets see... Do you think DHE will still work as-is?




Attacking DHE as a MitM




What happens if instead of Eve watching, Alice & Bob face the threat of a hidden Mallory (MITM)?

The MitM Key Exchange




2. Alice picks random secret ' a ': $1 < a < p-1$
3. Bob picks random secret ' b ': $1 < b < p-1$




$$\mathbf{A} = g^a \text{ mod } p$$

4. Alice sends $A = g^a \text{ mod } p$ to Bob
5. Mallory prevents Bob from receiving A




$$A = g^a \text{ mod } p$$

6. Mallory generates her own a', b'
7. Mallory sends $A' = g^{a'} \text{ mod } p$ to Bob



8. The same happens for Bob and B/B'
- $$g^b \text{ mod } p = B$$
- $$A' \quad \quad \quad$$




8. The same happens for Bob and B/B'

9. Alice and Bob now compute keys they share with ...
Mallory!

10. Mallory can relay encrypted traffic between the two ...

10'. Modifying it or making stuff up however she wishes



So We Will Want More...

- This is online:
 - Alice and Bob actually need to be active for this to work...
- So we want offline encryption:
 - Bob can send a message to Alice that Alice can read at a later date
- And authentication:
 - Alice can publish a message that Bob can verify was created by Alice later
 - Can also be used as a building-block for eliminating the MitM in the DHE key exchange:
Alice authenticates **A**, Bob verifies that he receives **A** not **A'**.