

The Net


Part 3




Addressing on the Layers On The Internet

- Ethernet:
 - Address is 6B MAC address, Identifies a machine on the local LAN
- IP:
 - Address is a 4B (IPv4) or 16B (IPv6) address, Identifies a system on the Internet
- TCP/UDP:
 - Address is a 2B port number, Identifies a particular listening server/process/activity on the system
 - Both the client and server have to have a port associated with the communication
 - Ports 0-1024 are for privileged services
 - Must be root to accept incoming connections on these ports
 - Any thing can do an outbound request to such a port
 - Port 1025+ are for anybody
 - And high ports are often used ephemerally


IP Packet Structure




IP Packet Structure




IP Packet Structure




IP Packet Structure



IP Packet Structure



IP Packet Structure



IP Packet Header (Continued)

- Two IP addresses
 - Source IP address (32 bits)
 - Destination IP address (32 bits)
- Destination address
 - Unique identifier/locator for the receiving host
 - Allows each node to make forwarding decisions
- Source address
 - Unique identifier/locator for the sending host
 - Recipient can decide whether to accept packet
 - Enables recipient to send a reply back to source
- Checksum is arithmetic, not CRC...
 - To allow easily modification of the packet by the network

IP: “Best Effort” Packet Delivery


- Routers inspect destination address, locate “next hop” in forwarding table
 - Address = ~unique identifier/locator for the receiving host
- Only provides a “*I'll give it a try*” delivery service:
 - Packets may be lost
 - Packets may be corrupted (but that is 'assume drop' based on layer 2 error detection)
 - Packets may be delivered out of order



IP Routing: Autonomous Systems

- Your system sends IP packets to the gateway...
 - But what happens after that?
- Within a given network its routed internally
 - Identified by its ASN (Autonomous System Number)
- But the key is the Internet is a network-of-networks
 - Each "autonomous system" (AS) handles its own internal routing
 - The AS knows the next AS to forward a packet to
- Primary protocol for communicating in between ASs is BGP:
 - Each router announces what networks it can provide and the path onward
 - **Most precise** route with the shortest path and no loops preferred


Packet Routing on the Internet: Border Gateway Protocol & Routing Tables




IP Spoofing And Autonomous Systems

- The edge-AS where a user connects **should** restrict packet spoofing
 - Sending a packet with a different sender IP address
 - But about 25% of them don't...
 - So a system can simply lie and say it comes from someplace else
 - This enables blind-spoofing attacks
 - Such as the Kaminski attack on DNS which we will see in a second
 - It also enables "reflected DOS attacks"
 - Send a small request...
That sends a large reply...
To the fake "sender" of the packet

On-path Injection vs Off-path Spoofing



Lying in BGP



UDP:

Datagrams on the Internet


- UDP is a protocol built on the Internet Protocol (IP)
- It is an "unreliable, datagram protocol"
 - Messages may or may not be delivered, in any order
 - Messages can be larger than a single packet (but probably shouldn't)
 - IP will fragment these into multiple packets (mostly... Single digit %-age of hosts can't receive fragmented traffic)
- Programs create a socket to send and receive messages
 - Just create a datagram socket for an ephemeral port
 - Bind the socket to a particular port to receive traffic on a specified port
 - Basic recipe for Python:
<https://wiki.python.org/moin/UdpCommunication>

DNS Overview

- DNS translates `www.google.com` to `74.125.25.99`
 - Turns a human abstraction into an IP address
 - Can also contain other data
- It's a performance-critical distributed database.
- DNS security is critical for the web.
(Same-origin policy **assumes** DNS is secure.)
 - Analogy: If you don't know the answer to a question, ask a friend for help (who may in turn refer you to a friend of theirs, and so on).
 - Based on a notion of hierarchical trust:
 - You trust `.` for everything, `com.` for any com, `google.com.` for everything google...

DNS Lookups via a *Resolver*

Host at **xyz.poly.edu** wants IP address for **eecs.mit.edu**



Security risk #1: malicious DNS server

- Of course, if *any* of the DNS servers queried are malicious, they can lie to us and fool us about the answer to our DNS query
- (In fact, they used to be able to fool us about the answer to other queries, too. We'll come back to that.)

Security risk #2: on-path eavesdropper

- If attacker can eavesdrop on our traffic... we're hosed.
- Why? We'll see why.

Security risk #3: off-path attacker

- If attacker can't eavesdrop on our traffic, can he inject spoofed DNS responses?
- This case is especially interesting, so we'll look at it in detail.

DNS Threats

- DNS: path-critical for just about everything we do
 - Maps hostnames ⇔ IP addresses
 - Design only **scales** if we can minimize lookup traffic
 - #1 way to do so: **caching**
 - #2 way to do so: return not only answers to queries, but **additional info** that will likely be needed shortly
 - The "glue records"
- What if attacker eavesdrops on our DNS queries?
 - Then similar to DHCP, ARP, AirPwn etc, can spoof responses
- Consider attackers who **can't** eavesdrop - but still aim to manipulate us via *how the protocol functions*
- Directly interacting w/ DNS: **dig** program on Unix
 - Allows querying of DNS system
 - Dumps each field in DNS responses

dig eecs.mit.edu A

Use Unix “dig” utility to look up IP address (“A”) for hostname eecs.mit.edu via DNS

```
; ; <>> DiG 9.6.0-APPLE-P2 <>> eecs.mit.edu a
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 19901
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 3, ADDITIONAL: 3

;; QUESTION SECTION:
;eecs.mit.edu.           IN      A

;; ANSWER SECTION:
eecs.mit.edu.        21600    IN      A      18.62.1.6

;; AUTHORITY SECTION:
mit.edu.            11088    IN      NS      BITSY.mit.edu.
mit.edu.            11088    IN      NS      W20NS.mit.edu.
mit.edu.            11088    IN      NS      STRAWB.mit.edu.

;; ADDITIONAL SECTION:
STRAWB.mit.edu.     126738    IN      A      18.71.0.151
BITSY.mit.edu.      166408    IN      A      18.72.0.3
W20NS.mit.edu.     126738    IN      A      18.70.0.160
```

```
dig eecs.mit.edu A
```

```
; ; <>> DiG 9.6.0-APPLE-P2 <>> eecs.mit.edu a
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 19901
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 3, ADDITIONAL: 3

;; QUESTION SECTION:
;eecs.mit.edu.          IN      A

;; ANSWER SECTION:
eecs.mit.edu.        21600   IN      A      18.62.1.6

;; AUTHORITY SECTION:
mit.edu.            11088   IN      NS     BITSY.mit.edu.
mit.edu.            11088   IN      NS     W20NS.mit.edu.
mit.edu.            11088   IN      NS     STRAWB.mit.edu.

;; ADDITIONAL SECTION:
STRAWB.mit.edu.    126738   IN      A      18.71.0.151
BITSY.mit.edu.     166408   IN      A      18.72.0.3
W20NS.mit.edu.    126738   IN      A      18.70.0.160
```

The question we asked the server

```
dig eecs.mit.edu A
```

```
; ; <>> DiG 9.6.0-APPLE-P2 <>> eecs.mit.edu a
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 19901
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 3, ADDITIONAL: 3

;; QUESTION SECTION:
;eecs.mit.edu.           IN      A

;; ANSWER SECTION:
eecs.mit.edu.          21600   IN      A          18.71.0.151
                        IN      NS      BITSY.mit.edu.
                        IN      NS      W20NS.mit.edu.
                        IN      NS      STRAWB.mit.edu.

;; AUTHORITY SECTION:
mit.edu.               11088   IN      NS      BITSY.mit.edu.
mit.edu.               11088   IN      NS      W20NS.mit.edu.
mit.edu.               11088   IN      NS      STRAWB.mit.edu.

;; ADDITIONAL SECTION:
STRAWB.mit.edu.        126738  IN      A       18.71.0.151
BITSY.mit.edu.         166408  IN      A       18.72.0.3
W20NS.mit.edu.         126738  IN      A       18.70.0.160
```

A 16-bit transaction identifier that enables the DNS client (`dig`, in this case) to match up the reply with its original request

```
dig eecs.mit.edu A
```

```
; ; <>> DiG 9.6.0-APPLE-P2 <>> eecs.mit.edu a  
;; global options: +cmd  
;; Got answer:  
;; ->>HEADER<<- opcode: 0x00  
;; flags: qr rd ra; QUERY  
;;QUESTION SECTION:
```

```
;eecs.mit.edu.  
  
;; ANSWER SECTION:
```


```
;; AUTHORITY SECTION:
```

mit.edu.	11088	IN	NS	BITSY.mit.edu.
mit.edu.	11088	IN	NS	W20NS.mit.edu.
mit.edu.	11088	IN	NS	STRAWB.mit.edu.

```
;; ADDITIONAL SECTION:
```

STRAWB.mit.edu.	126738	IN	A	18.71.0.151
BITSY.mit.edu.	166408	IN	A	18.72.0.3
W20NS.mit.edu.	126738	IN	A	18.70.0.160

“Answer” tells us the IP address associated with eecs.mit.edu is 18.62.1.6 and we can cache the result for 21,600 seconds



```
dig eecs.mit.edu A
```

```
; ; <>> DiG 9.6.0-APPLE-P2 <>> eecs.mit.edu a
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 19901
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 3, ADDITIONAL: 3

;; QUESTION SECTION:
;eecs.mit.edu.           IN      A

;; ANSWER SECTION:
eecs.mit.edu.        21600   IN      A      18.62.1.6

;; AUTHORITY SECTION:
mit.edu.                 11088   IN      NS     BITSY.mit.edu.
mit.edu.
mit.edu.

;; ADDITIONAL SECTION
STRAWB.mit.edu.          106408  IN      A      18.72.0.5
BITSY.mit.edu.            126738  IN      A      18.70.0.160
W20NS.mit.edu.            126738  IN      A      18.70.0.160
```


In general, a single Resource Record (RR) like this includes, left-to-right, a DNS name, a time-to-live, a family (IN for our purposes - ignore), a type (A here), and an associated value

dig eecs.mit.edu A

```
; ; <>> DiG 9.6.0-APPLE-P2 <>> eecs.mit.edu a
;; global options: +cm
;; Got answer:
;; ->>HEADER<<- opcode
;; flags: qr rd ra; QU
;;
;; QUESTION SECTION:
;eecs.mit.edu.
;;
;; ANSWER SECTION:
eecs.mit.edu.
;;
;; AUTHORITY SECTION:
mit.edu.
mit.edu.
mit.edu.
;;
;; ADDITIONAL SECTION:
STRAWB.mit.edu.          126738  IN      A       18.71.0.151
BITSY.mit.edu.           166408  IN      A       18.72.0.3
W20NS.mit.edu.          126738  IN      A       18.70.0.160
```

“Authority” tells us the name servers responsible for the answer. Each RR gives the hostname of a different name server (“NS”) for names in mit.edu. We should cache each record for 11,088 seconds.

If the “Answer” had been empty, then the resolver’s next step would be to send the original query to one of these name servers.



3

```
dig eecs.mit.edu A
```

```
; ; <>> DiG 9.6.0-APPLE-P2 <>> eecs.mit.edu a
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 19901
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 3, ADDITIONAL: 3
```

```
; ; QUESTION SECTION:
```

```
;eecs.mit.edu.
```

```
; ; ANSWER SECTION
```

```
eecs.mit.edu.
```

```
; ; AUTHORITY SECTION
```

```
mit.edu.
```

```
mit.edu.
```

```
mit.edu.
```

```
; ; ADDITIONAL SECTION:
```

```
STRAWB.mit.edu.
```

```
BITSY.mit.edu.
```

```
W20NS.mit.edu.
```

“Additional” provides extra information to save us from making separate lookups for it, or helps with bootstrapping.

Here, it tells us the IP addresses for the hostnames of the name servers. We add these to our cache.

11088	IN	NS	BITSY.mit.edu.
11088	IN	NS	W20NS.mit.edu.
11088	IN	NS	STRAWB.mit.edu.
			18.71.0.151
			18.72.0.3
			18.70.0.160

DNS Protocol

Lightweight exchange of *query* and *reply* messages, both with **same** message format


Primarily uses UDP for its transport protocol, which is what we'll assume

Servers are on port 53 always

Frequently, clients used to use port 53 but can use any port


UDP Header

UDP Payload



Message header:

- **Identification:** 16 bit # for query, reply to query uses same #
- Along with repeating the Question and providing Answer(s), replies can include “**Authority**” (name server responsible for answer) and “**Additional**” (info client is likely to look up soon anyway)
- Each Resource Record has a **Time To Live** (in seconds) for **caching** (not shown)



```
dig eecs.mit.edu A
```

```
; ; <>> DiG 9.6.0-APPLE-P2 <>> eecs.mit.edu a
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 19901
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 0, ADDITIONAL: 3
;
;; QUESTION SECTION:
;eecs.mit.edu.

;; ANSWER SECTION:
eecs.mit.edu.          216    IN      A       18.71.0.151

;; AUTHORITY SECTION:
mit.edu.               11088   IN      NS     BITSY.mit.edu.
mit.edu.               11088   IN      NS     W20NS.mit.edu.
mit.edu.               11088   IN      NS     STRAWB.mit.edu.

;; ADDITIONAL SECTION:
STRAWB.mit.edu.        126738   IN      A      18.70.0.160
BITSY.mit.edu.          166408   IN      A      18.72.0.3
W20NS.mit.edu.         126738   IN      A      18.71.0.151
```

What if the mit.edu server
is untrustworthy? Could
its operator steal, say, all
of our web surfing to
berkeley.edu's main web
server?

```
dig eecs.mit.edu A
```

```
; ; <>> DiG 9.6.0-APPLE-P2 <>> eecs.mit.edu a
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 19901
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 3, ADDITIONAL: 3

;; QUESTION SECTION:
;eecs.mit.edu.

;; ANSWER SECTION:
eecs.mit.edu.          216    IN      A       18.71.0.151

;; AUTHORITY SECTION:
mit.edu.               11088   IN      NS      BITSY.mit.edu.
mit.edu.               11088   IN      NS      W20NS.mit.edu.
mit.edu.               11088   IN      NS      STRAWB.mit.edu.

;; ADDITIONAL SECTION:
STRAWB.mit.edu.        126738   IN      A       18.70.0.160
BITSY.mit.edu.          166408   IN      A       18.72.0.3
W20NS.mit.edu.         126738   IN      A       18.71.0.151
```

Let's look at a flaw in the
original DNS design
(since fixed)

```
dig eecs.mit.edu A
```

```
; ; <>> DiG 9.6.0-APPLE-P2 <>> eecs.mit.edu a
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 19901
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 3, ADDITIONAL: 3

;; QUESTION SECTION:
;eecs.mit.edu.

;; ANSWER SECTION:
eecs.mit.edu.          21600   IN      A       18.62.1.6

;; AUTHORITY SECTION:
mit.edu.                11088   IN      NS      BITSY.mit.edu.
mit.edu.                11088   IN      NS      W20NS.mit.edu.
mit.edu.                11088   IN      NS      www.berkeley.edu.

;; ADDITIONAL SECTION:
www.berkeley.edu.     100000   IN      A       18.6.6.6
BITSY.mit.edu.         166408   IN      A       18.72.0.3
W20NS.mit.edu.        126738   IN      A       18.70.0.160
```

What could happen if the mit.edu server returns the following to us instead?

```
dig eecs.mit.edu A
```

```
; ; <>> DiG 9.6.0-APPLE-P2 <>> eecs.mit.edu a
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 19901
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 3, ADDITIONAL: 3

;; QUESTION SECTION:
;eecs.mit.edu.           IN      A

;; ANSWER SECTION:
eecs.mit.edu.          11088   IN      NS      BITSY.mit.edu.
eecs.mit.edu.          11088   IN      NS      W20NS.mit.edu.

;; AUTHORITY SECTION:
mit.edu.               11088   IN      NS      www.berkeley.edu.

;; ADDITIONAL SECTION:
www.berkeley.edu.      100000  IN      A       18.6.6.6
BITSY.mit.edu.          166408  IN      A       18.72.0.3
W20NS.mit.edu.          126738  IN      A       18.70.0.160
```

We'd dutifully store in our cache a mapping of
www.berkeley.edu to an IP address under MIT's
control. (It could have been any IP address they
wanted, not just one of theirs.)

```
dig eecs.mit.edu A
```

```
; ; <>> DiG 9.6.0-APPLE-P2 <>> eecs.mit.edu a
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 19901
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 3, ADDITIONAL: 3

;; QUESTION SECTION:
;eecs.mit.edu.           IN      A

;; ANSWER SECTION:
eecs.mit.edu.          11088   IN      NS      BITSY.mit.edu.
                        11088   IN      NS      W20NS.mit.edu.
                        11088   IN      NS      www.berkeley.edu.

;; AUTHORITY SECTION:
mit.edu.               11088   IN      NS      BITSY.mit.edu.
mit.edu.               11088   IN      NS      W20NS.mit.edu.
mit.edu.               11088   IN      NS      www.berkeley.edu.

;; ADDITIONAL SECTION:
www.berkeley.edu.      100000  IN      A       18.6.6.6
BITSY.mit.edu.          166408  IN      A       18.72.0.3
W20NS.mit.edu.          126738  IN      A       18.70.0.160
```

In this case they chose to make the mapping last a long time. They could just as easily make it for just a couple of seconds.

6

```
dig eecs.mit.edu A
```

```
; ; <>> DiG 9.6.0-APPLE-P2 <>> eecs.mit.edu a
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 19901
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 3, ADDITIONAL: 3

;; QUESTION SECTION:
;eecs.mit.edu.           IN      A

;; ANSWER SECTION:
eecs.mit.edu.          11088   IN      NS      BITSY.mit.edu.
                        11088   IN      NS      W20NS.mit.edu.
mit.edu.               30      IN      NS      www.berkeley.edu.

;; AUTHORITY SECTION:
mit.edu.                11088   IN      NS      BITSY.mit.edu.
mit.edu.                11088   IN      NS      W20NS.mit.edu.
mit.edu.                 30      IN      NS      www.berkeley.edu.

;; ADDITIONAL SECTION:
www.berkeley.edu.       30      IN      A       18.6.6.6
BITSY.mit.edu.          166408   IN      A       18.72.0.3
W20NS.mit.edu.          126738   IN      A       18.70.0.160
```

How do we fix such **cache poisoning**?

```
dig eecs.mit.edu A
```

```
; ; <>> DiG 9.6.0-APPLE-DO-<>>  
;; global options: +cr  
;; Got answer:  
;; ->>HEADER<<- opcode: 0x00  
;; flags: qr rd ra; qd=1, qdcount=1  
;; QUESTION SECTION:  
;eecs.mit.edu.  
  
;; ANSWER SECTION:  
eecs.mit.edu.
```

```
; ; AUTHORITY SECTION:
```

```
mit.edu. 11088 IN  
mit.edu. 11088 IN  
mit.edu. 11088 IN
```

```
; ; ADDITIONAL SECTION:
```

```
www.berkeley.edu. 100000 IN  
BITSY.mit.edu. 166408 IN  
W20NS.mit.edu. 126738 IN
```

Don't accept **Additional** records unless they're for the domain we're looking up

E.g., looking up eecs.mit.edu \Rightarrow only accept additional records from *.mit.edu

No extra risk in accepting these since server could return them to us directly in an **Answer** anyway.

This is called "**Bailiwick** checking"

bail·i·wick
/bālē'wɪk/

noun

1. one's sphere of operations or particular area of interest.
"you never give the presentations—that's my bailiwick"
2. LAW
the district or jurisdiction of a bailie or bailiff.

DNS Resource Records and RRSETs

- DNS records (Resource Records) can be one of various types
 - Name TYPE Value
 - Also a “time to live” field: how long in seconds this entry can be cached for
 - Addressing:
 - A: IPv4 addresses
 - AAAA: IPv6 addresses
 - CNAME: aliases, “Name X should be name Y”
 - MX: “the mailserver for this name is Y”
 - DNS related:
 - NS: “The authority server you should contact is named Y”
 - SOA: “The operator of this domain is Y”
 - Other:
 - text records, cryptographic information, etc....
- Groups of records of the same type form RRSETs:
 - E.g. all the nameservers for a given domain.

The Many Moving Pieces In a DNS Lookup of www.isc.org



? A www.isc.org



- User's ISP's ? A www.isc.org
Recursive Resolver



? A www.isc.org Authority
Answers: (the “root”
Authority:
org. NS a0.afilias-nst.info
Additional:
a0.afilias-nst.info A 199.19.5

The Many Moving Pieces In a DNS Lookup of www.isc.org



User's ISP's ? A ~~www.isc.org~~

Recursive Resolver



org. Authority Server

? A www.isc.org

Answers:

Authority:

isc.org. NS sfba.sns-pb.isc.org.

isc.org. NS ns.isc.afilias-nst.info.

Additional:

stba.sns-pb.isc.org. A 199.6.1.30
ns-isc-afiliis.net.info A 100.254.63.254

HS.TSC.all

The Many Moving Pieces In a DNS Lookup of www.isc.org



- User's ISP's ? A www.isc.org
Recursive Resolver



isc.org.
Authority Server

? A www.isc.org

Answers:

www.isc.org. A 149.20.64.42

Authority:

isc.org. NS sfba.sns-pb.isc.org.

isc.org. NS ns.isc.afilias-nst.info.

Additional:

sfba.sns-pb.isc.org. A 199.6.1.30

ns.isc.afilias-nst.info. A 199.254.63.254

The Many Moving Pieces In a DNS Lookup of **www.isc.org**



User's ISP's

? A **www.isc.org**

Recursive Resolver **Answers: www.isc.org A 149.20.64.42**

Name	Type	Value	TTL
org.	NS	a0.afiliias-nst.info	172800
a0.afiliias-nst.info.	A	199.19.56.1	172800
isc.org.	NS	sfba.sns-pb.isc.org.	86400
isc.org.	NS	ns.isc.afiliias-net.info.	86400
sfbay.sns-pb.isc.org.	A	199.6.1.30	86400
www.isc.org	A	149.20.64.42	600

Stepping Through This With `dig`

- Some flags of note:
 - `+norecurse`: Ask directly like a recursive resolver does
 - `+trace`: Act like a recursive resolver without a cache

```
nweaver% dig +norecurse slashdot.org @a.root-servers.net

; <>> DiG 9.8.3-P1 <>> +norecurse slashdot.org @a.root-servers.net
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 26444
;; flags: qr; QUERY: 1, ANSWER: 0, AUTHORITY: 6, ADDITIONAL: 12

;; QUESTION SECTION:
;slashdot.org.           IN      A

;; AUTHORITY SECTION:
org.                    172800  IN      NS      a0.org.afilia-nst.info.
...
;ADDITIONAL SECTION:
a0.org.afilia-nst.info. 172800  IN      A      199.19.56.1
```

So in `dig` parlance

- So if you want to recreate the lookups conducted by the recursive resolver:
 - `dig +norecurse www.isc.org @a.root-servers.net`
 - `dig +norecurse www.isc.org @199.19.56.1`
 - `dig +norecurse www.isc.org @199.6.1.30`

Security risk #1: malicious DNS server

- Of course, if *any* of the DNS servers queried are malicious, they can lie to us and fool us about the answer to our DNS query...
- and they used to be able to fool us about the answer to other queries, too, using *cache poisoning*. Now fixed (phew).

Security risk #2: on-path eavesdropper

- If attacker can eavesdrop on our traffic... we're hosed.
- Why?

Security risk #2: on-path eavesdropper

- If attacker can eavesdrop on our traffic... we're hosed.
- Why? They can see the query and the 16-bit transaction identifier, and race to send a spoofed response to our query.
 - China does this operationally:
 - `dig www.benign.com @www.tsinghua.edu.cn`
 - `dig www.facebook.com @www.tsinghua.edu.cn`

Security risk #3: off-path attacker


- If attacker can't eavesdrop on our traffic, can he inject spoofed DNS responses?
- Answer: It used to be possible, via *blind spoofing*. We've since deployed mitigations that makes this harder (but not totally impossible).

Blind spoofing

- Say we look up `mail.google.com`; how can an **off-path** attacker feed us a **bogus A answer** before the legitimate server replies?
- How can such a **remote** attacker even know we are looking up `mail.google.com`?

Suppose, e.g., we visit a web page under their control:

... ...



Blind spoofing


- Say we look up mail.google.com; how can an **off-path** attacker feed us a bogus A answer before the legitim

This HTML snippet causes our browser to try to fetch an image from

- How do we get mail.google.com. To do that, our even legitimate browser first has to look up the IP address associated with that name.


Suppose, e.g., we visit a web page under their control:

```
.... ...
```



Blind spoofing

Fix?



Once they know we're looking it up, they just have to guess the Identification field and reply before legit server.

How hard is that?

Originally, identification field incremented by 1 for each request. How does attacker guess it?

 They observe ID k here
 So this will be k+1


DNS Blind Spoofing, cont.

Once we **randomize** the Identification, attacker has a 1/65536 chance of guessing it correctly.

Are we pretty much safe?

Attacker can send lots of replies, not just one ...

However: once reply from legit server arrives (with correct Identification), it's **cached** and no more opportunity to poison it. Victim is innoculated!



Unless attacker can send 1000s of replies before legit arrives, we're likely safe – phew! ?

Enter Kaminski... Glue Attacks

- Dan Kaminski noticed something strange, however...
 - Most DNS servers would **cache** the in-bailiwick glue...
 - And then **promote** the glue
 - And will also **update** entries based on glue
- So if you first did this lookup...
 - And then went to query for **a0.org.afiliias-nst.info**
 - there would be no other lookup!

```
nweaver% dig +norecurse slashdot.org @a.root-servers.net
; <>> DiG 9.8.3-P1 <>> +norecurse slashdot.org @a.root-servers.net
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 26444
;; flags: qr; QUERY: 1, ANSWER: 0, AUTHORITY: 6, ADDITIONAL: 12
;; QUESTION SECTION:
;slashdot.org.           IN      A
;; AUTHORITY SECTION:
org.                      172800  IN      NS      a0.org.afiliias-nst.info
...
;; ADDITIONAL SECTION:
a0.org.afiliias-nst.info. 172800  IN      A      199.19.56.1
...
;; Query time: 128 msec
;; SERVER: 198.41.0.4#53(198.41.0.4)
;; WHEN: Tue Apr 16 09:48:32 2013
;; MSG SIZE  rcvd: 432
```

The Kaminski Attack In Practice

- Rather than trying to poison `www.google.com`...
- Instead try to poison `a.google.com`...
And state that "`www.google.com`" is an authority
And state that "`www.google.com A 133.7.133.7`"
 - If you succeed, great!
 - But if you fail, just try again with `b.google.com`!
 - Turns "Race once per timeout" to "race until win"
 - So now the attacker may still have to send lots of packets
 - In the 10s of thousands
 - The attacker can keep trying until success

Defending Against Kaminski: Up the Entropy

- Also randomize the UDP source port
 - Adds ~16 bits of entropy
- Observe that most DNS servers just copy the request directly
 - Rather than create a new reply
- So caMeLcase the NamE ranDomly
 - Adds only a few bits of entropy however, but it does help

Defend Against Kaminski: Validate Glue

- Don't blindly accept glue records...
 - Well, you **have** to accept them for the purposes of resolving a name
- But if you are going to cache the glue record...
- Either only use it for the context of a DNS lookup
 - No more promotion
- Or explicitly validate it with another fetch
- Unbound implemented this, bind did not
 - Largely a **political** decision:
bind's developers are heavily committed to DNSSEC (an upcoming topic)

Oh, and Profiting from Rogue DNS

Computer Science 161

- Suppose you take over a lot of home routers...
 - How do you make money with it?
- Simple: Change their DNS server settings
 - Make it point to yours instead of the ISPs
- Now redirect all advertising
 - And instead serve up ads for "Vimax" pills...
 - Can only do this for unencrypted sites, but....

