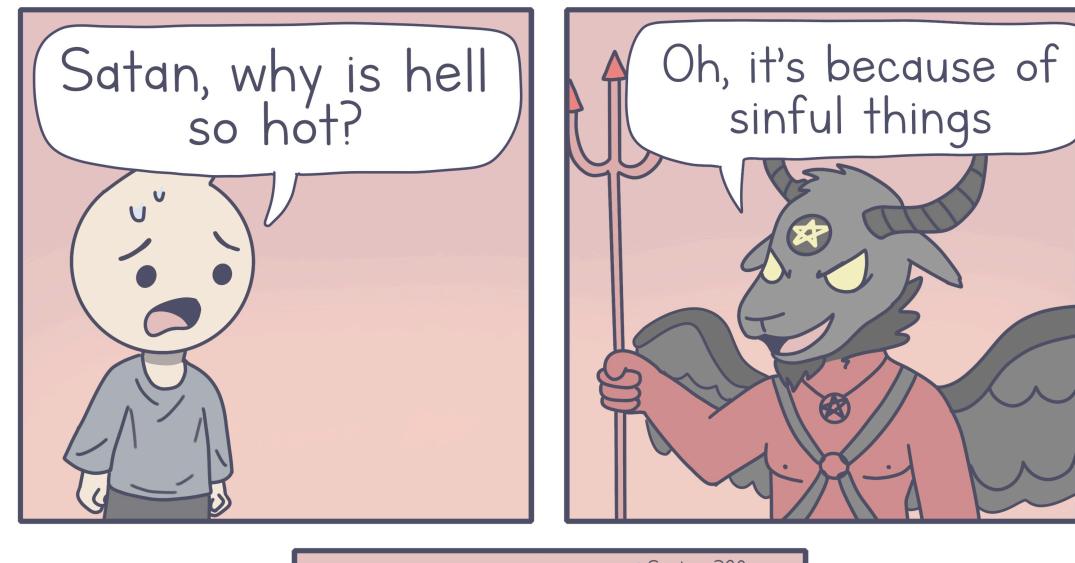


# More Dealer's Choice



# Announcements for the Last Lecture...

- The last lecture will be special...
- The first fraction will ***not*** be recorded
  - Instead it is how I would design a solution for Project 3:  
Based on my previous solution and tweaks for new requirements
- The rest will be an "Ask Me Anything"

# Funky Hardware SideChannels...

- The Meltdown and Spectre Intel bugs...
  - Both were effectively side-channels
- The key idea:
  - You could trick the speculative execution engine to compute on memory that you don't own
  - And that computation will take a different amount of time depending on the memory contents
- So between the two, you could read past isolation barriers
  - Meltdown: Read operating system (and other) memory from user level
  - Spectre: Read in JavaScript from other parts of the web browser

# How Meltdown Works...

- In a CPU, precise exceptions are hard: that is, stopping things when something "happens" at a specific instruction
- x86 actually provides two page table hardware pointers
  - One for the current user program, one for supervisor mode
  - Allows the OS to have virtual memory for the interrupt handler and other things
- Concept behind meltdown:
  - x86 allows "load whatever that memory location points to + base register"
- Do a bunch of loops that are always taken
  - Now the CPU will predict that the next time this loop is taken...
- Now do a load of memory you aren't supposed to read belonging to the OS
  - CPU guesses branch will taken, so is just going to do it speculatively.  
Only when it finally writes to a register will the exception be checked
- Now have the results of that load do a load to memory you are supposed to read
  - But dependent on what was in the memory you weren't supposed to read
- Now CPU finds that branch wasn't taken after all
  - And so nothing happens, neither the illegal load nor the "load not taken"

# But Something *Did* Happen!

- The final "load not taken" got taken!
  - So it will be cached
  - And that load was dependent on the illegal load
  - So we can discover which "load not taken" got actually taken!
    - Allowing us to read memory we aren't supposed to!
    - Fix involves the OS flushing the TLB and presenting a dummy OS page-table when returning to a user process
    - Greatly increasing the cost of a context switch or interrupt

# Countering Meltdown and Spectre...

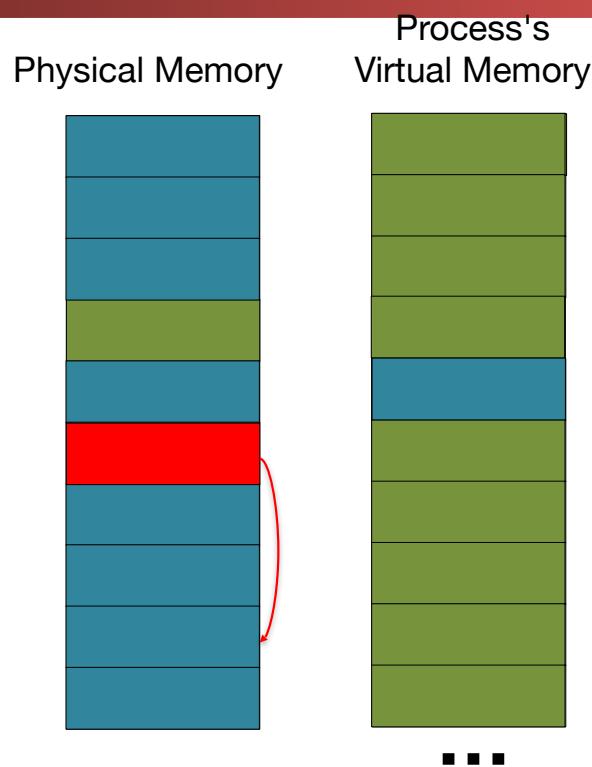
- Meltdown was really a bug...
  - TLB check not acted on right away
- Spectre and variants are really features of caches
  - You could train a branch-prediction buffer that you won't do it..  
Then you did it anyway
- Counteracting Spectre requires flushing ***all*** caches on ***every*** context switch
  - No such thing as a lightweight isolation barrier
  - This is why chrome & firefox eat ram with abandon:  
Every web origin runs in a different OS process

# The Ultimate Page-Table Trick: Rowhammer

- An ***unspeakably cool*** security vulnerability...
- DRAM (unless you pay for error correcting (ECC) memory) is actually unreliable
  - Can repeatedly read/write the same location ("hammer the row" and eventually cause an error in ***some physically distinct memory location***)
- Can tell the OS "I want to map this same block of memory at multiple addresses in my process..."
  - Which creates additional page table entries, lots of them. Lots and lots of them. Lots and lots and lots and lots of them...
- Enter ***Rowhammer***
  - It seems all vulnerabilities get named now, but this one is cool enough to deserve a name!
  - Touches on virtual memory, hardware failures, and breaks security

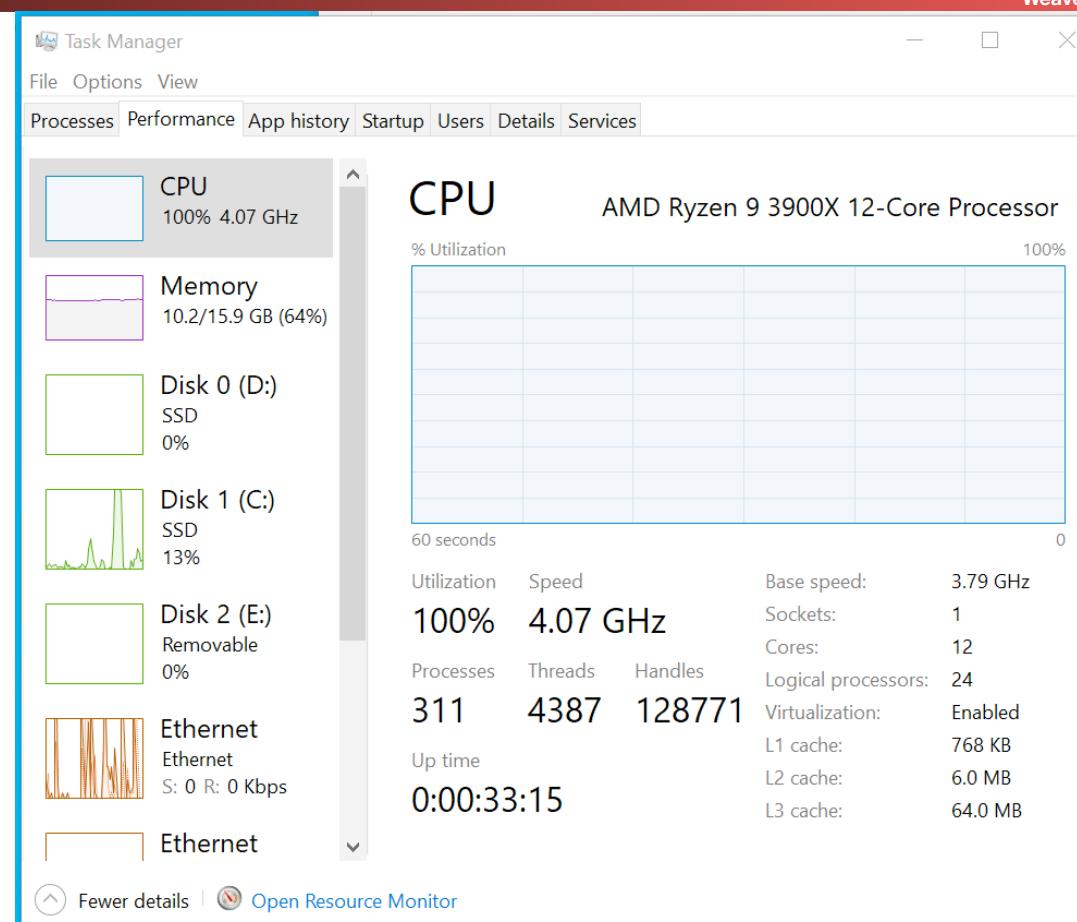
# How RowHammer Works

- Step 1: Allocate a single page of memory
- Step 2: Make the OS make a gazillion page-table entries pointing to the same page
- Step 3: Hammer the DRAM until one of those entries gets corrupted
  - Now causes that memory page to point to a set of page table entries instead
- **Step 4: Profit**
  - Well, the ability to read and write to any physical address in the system, same difference



# Analyzing the Stolen UCOP Data

- Some bad actor stole >4 GB of data from UCOP
  - Basically everything that was on the "secure" file transfer server in December
- The bad actor released at least some of this data **publicly**
  - As a 4GB compressed archive that anyone can download
  - What information about **me** was in the archive?



# What I need to know...

- I already know my social security # got breached
  - They told us that...
  - But I've got fraud alerts & freezes in place already
- But what other information?
  - Address? Phone #? Things I don't know about?
  - Tax information?
  - ***Banking information?***
  - The numbers on the bottom of a check are all an attacker needs to make fake checks

# The Nature of the Dump...

- A ***lot*** of pdf files
  - PDFs are a pain to search, need to convert to text
- A ***lot*** of data tables
  - Some as comma-delimited text, some as excel spreadsheets, some in stada format
- Need to convert it to something reasonable
- Google around...
  - Nice linux OCR pipeline cobbled together:  
PDF -> images -> OCR text
  - pandas can read both xlsx and stada files

# Step 1: File Conversion

- Want to convert everything into text files
- Obscenely parallel problem:
  - For every PDF do X
- But with some gotchas...
  - I can't just spawn 700 PDF->txt conversion programs
    - That would grind my machine to a halt
  - And different invocations take a different amount of time
    - So I can't just spawn off tasks at a reasonable interval
- Two approaches
  - Dynamically tune based on load...
  - Or just say "F-it, and keep X jobs live"

# Keeping X jobs live: Fork/join with a limiter

- Used a simple golang hack
- `capacity := make(chan bool, 10)`  
`done := make(chan bool)`
- `func run(txt string) {`  
    `c <- True;`  
    `....`  
    `<- c;`  
    `done <- True;`  
}
- main just calls "go func()" on every line of stdin...  
and then an equal number of lines of "<- done"

# A bit of tuning...

- ~10 jobs pegs all CPU cores on the OCR pipeline...
  - And it took a couple hours to PDF->txt the lot:  
Driver used multiple threads for single documents
- The .xlsx and stata conversion was a lot faster...
  - Set to ~25 (since python doesn't thread, especially on this task)
  - Took ~10 minutes or so
- Not fully efficient...
  - At the end of the PDF run there was no longer pegged CPU
  - 100% CPU utilization means efficiency loss due to context switching:  
Optimum would be ~95%
  - Also, really stressing the Windows virtualization...
    - I do all my work in "linux" under WSL

# And Now To Search

- Just use the same pipeline with grep...
- But...
  - Ends up **not** pegging the CPUs...  
Instead I'm pegging the "disk"!
- OK for just searching for me, but...
  - Want to be able to do a "for anyone who wants" service
- So to do this, parallelize on an alternate axis:
  - Don't check one person at a time, check **all** people using a single program
  - And then invoke that in parallel across all files
- Gotcha problem: Need to make sure to synchronize writes well
  - Again, golang FTW:  
A channel for each user's results, the search does an atomic write to the channel

# Results So Far...

- Lots of tabular data on everyone
  - For students it isn't just social security # but address, phone, self reported sexual orientation, self reported disability, etc...
- Fortunately no bulk banking or tax records
  - Did see a large number of individual records and documents concerning failed bank transactions
- Tabular data remarkably well organized
- Seeking to create a "Know Your Data" service
  - Request with a google form
  - Reply is a google doc
    - Allows UC/Google to handle all the access control
  - Need UC permission to do this, but I'm Trying!

# Putting CS161 in Context: Nick's Self Defense Strategies...

- **How** and **why** do I protect myself online and in person...
  - **How** I decide what to prepare for (and what not to prepare for)
  - **Why** I've drunk the Apple Kool-Aid™
  - **Why** I use my credit card everywhere but not a debit card
- And my future nightmares:
  - What do I see as the security problems of tomorrow...

# My Personal Threats: The Generic Opportunist

- There are a *lot* of crooks out there
  - And they are rather organized...
- But at the same time, these criminals are generally economically rational
  - So *this* is a bear race: I don't need perfect security, I just need *good enough* security
- I use this to determine security/convenience tradeoffs all the time
  - So no password reuse (use a password manager instead)
  - Full disk encryption & passwords on devices:  
Mitigates the damage from theft
  - Find my iPhone turned on:  
Increases probability of theft recovery

# My Personal Threats: The *Lazy* Nation State

- OK, I'm a high *enough* profile to have to worry about the "Advanced Persistent Threats" ...
  - Trying for a reasonably high profile on computer policy issues
  - A fair amount of stuff studying the NSA's toys and other nation-state tools
  - But only at the Annoying Pestilent Teenager level:  
I'm worth some effort but not an extraordinary amount
- So its only *slightly* more advanced than the everyday attackers...  
With one *huge* exception: Crossing borders
  - Every nation maintains the right to conduct searches of all electronic contents at a border checkpoint

# My Border Crossing Policy: Low Risk Borders

- Not very sensitive borders: Canada, Europe, US, etc...
  - I use full disk encryption with strong passwords on all devices
    - Primary use is to prevent theft from also losing data
  - I have a **very robust** backup strategy
    - Time machine, archived backups in a safe deposit box, working sets under version control backed up to remote systems...
- So, as the plane lands:
  - Power off my devices
    - Device encryption is only **robust** when you aren't logged in
  - Go through the border
- If my devices get seized...
  - "Keep it, we'll let the lawyers sort it out"

# High Risk Borders

- Middle East or, if, god forbid, I visit China or Russia...
  - Need something that doesn't just resist compromise but can also ***tolerate compromise***
- A "burner" iPhone SE with a Bluetooth keyboard
  - The cheapest secure device available
  - Set it up with ***independent*** computer accounts for both Google and Apple
    - Temporarily forward my main email to a temporary gmail account
    - All workflow accessible through Google apps on that device
  - Bluetooth keyboard does leak keystrokes, so don't use it for passwords but its safe for everything else
- Not only is this device very hard to compromise...
  - But there is very low value in ***successfully compromising it***: The attacker would only gain access to dummy accounts that have no additional privileges
- And bonus, I'm not stuck dragging a computer to the ski slopes in Dubai...
  - Since the other unique threat in those environments is the "Evil maid" attack



# My Personal Threats: The Russians... Perhaps

## Click Trajectories: End-to-End Analysis of the Spam Value Chain

Kirill Levchenko\* Andreas Pitsillidis\* Neha Chachra\* Brandon Enright\* Márk Félegyházi† Chris Grier†  
Tristan Halvorson\* Chris Kanich\* Christian Kreibich† He Liu\* Damon McCoy\*  
Nicholas Weaver† Vern Paxson† Geoffrey M. Voelker\* Stefan Savage\*

- This is the paper that killed the Viagra® Spam business
  - A \$100M a year set of organized criminal enterprises in Russia...  
And they put the **organized** in organized crime...
- I've adopted a **detection and response** strategy:
  - The Russians have higher priority targets: The first authors, the last authors, and Brian Krebs
  - If anything suspicious happens to Brian, Kirill, or Stefan, **then** I will start sleeping with a rifle under my bed

# Excluded Threats: Sorta...

- Intimate Partner Threats...
  - But I've had at least one colleague caught up with that.
- Aggressive Nation States...
  - \$50M will buy the latest version of Pegasus malcode
- The US government...
  - The surveillance powers of the US government are awesome and terrifying to behold...

# Passwords and 2-Factor....

- I **love** security keys:
  - I have one in each of my main computers... and one on the keychain
- ANY site that supports multiple security keys has that as the primary 2-factor method
  - Both more convenient **and** more secure than the alternatives...
- I also religiously use a password manager
  - "Credential stuffing" is the biggest threat individuals face
- I personally use 1password, but others are equally good
  - In particular you can get LastPass premium through software@berkeley

# The Apple Kool-Aid...

- The iPhone is perhaps the most secure commodity device available...
  - Not only does it receive patches but since the 5S it gained a dedicated cryptographic coprocessor
- The **Secure Enclave Processor** is the trusted base for the phone
  - Even the main operating system isn't fully trusted by the phone!
- A dedicated ARM v7 coprocessor
  - Small amount of memory, a true RNG, cryptographic engine, etc...
  - Important: A collection of **randomly** set fuses
    - Should not be able to extract these bits without taking the CPU apart:  
Even the Secure Enclave can only use them as keys to the AES engine, not read them directly!
  - But bulk of the memory is shared with the main CPU
- GOOD documentation:
  - The iOS security guide is something you should at least skim....  
I find that the design decisions behind how iOS does things make **great** final exam questions
- But it isn't perfect: Nation-state actors will pay big \$ for exploits
  - So keep it patched
  - And iOS 14.5: New Emoji and **turning on PAC all over the place!**

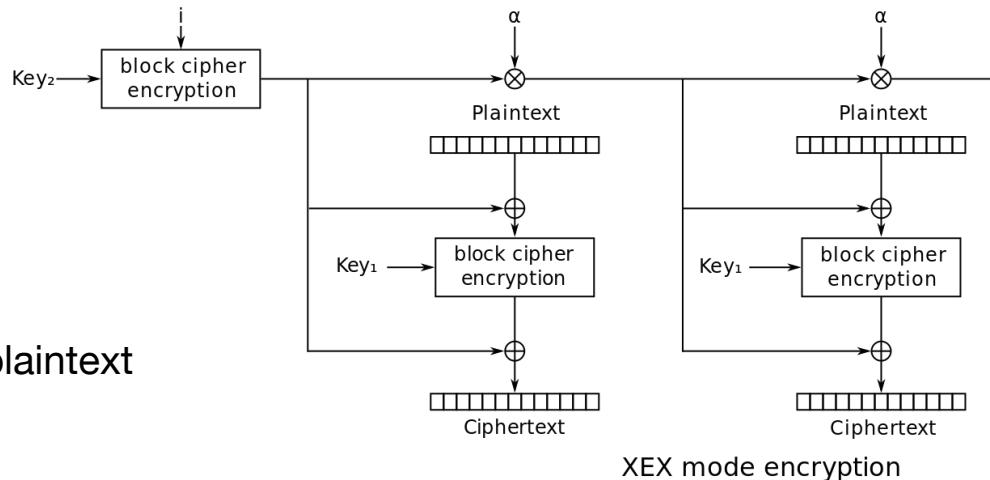
# The Roll of the SEP...

## Things *too important* to allow the OS to handle

- Key management for the encrypted data store
  - The CPU has to ask for access to data!
- Managing the user's passphrase and related information
- User authentication:
  - **Encrypted** channel to the fingerprint reader/face recognition camera
- Storing credit cards
  - ApplePay is cheap for merchants **because it is secure**:  
Designed to have very low probability of fraud!

# AES-256-XEX mode

- A ***confidentiality-only*** mode developed by Phil Rogaway...
  - Designed for encrypting data within a filesystem block  $i$ 
    - Known plaintext, when encrypted, can't be replaced to produce known output, only "random" output
    - Within a block: Same cypher text implies different plaintext
    - Between blocks: Same cypher text implies nothing!
    - $\alpha$  is a galios multiplication and is very quick:  
In practice this enables parallel encryption/decryption
  - Used by the SEP to encrypt its own memory...
    - Since it has to share main memory with the main processor
  - Opens a limited attack surface from the main processor:
    - Main processor can replace 128b blocks with ***random*** corruption



# User Passwords...

- Data is encrypted with the user's password
  - When you power on the phone, most data is completely encrypted
- The master key is PBKDF2(password || on-chip-secret)
  - So you need **both** to generate the master key
  - Some other data has the key as F(on-chip-secret) for stuff that is always available from boot
- The master keys encrypt a block in the flash that holds all the other keys
  - So if the system can erase this block effectively it can erase the phone by erasing just one block of information
- Apple implemented **effaceable storage**:
  - After x failures, OS command, whatever...  
Overwrite that master block in the flash securely
  - Destroy the keys == erase everything!

# Background: FBI v Apple

- A "terrorist" went on a rampage with a rifle in San Bernardino...
  - Killed several people before being killed in a battle with police
- He left behind a work-owned, passcode-locked iPhone 5 in his other car...
- The FBI **knew** there was no valuable information on this phone
  - But never one to refuse a good test case, they tried to compel Apple in court to force Apple to unlock the phone...
- Apple has serious security on the phone
  - Effectively everything is encrypted with PBKDF2(PW||on-chip-secret):  
>128b of randomly set microscopic fuses
    - Requires that **any** brute force attack either be done on the phone or take apart the CPU
  - Multiple timeouts:
    - 5 incorrect passwords -> starts to slow down
    - 10 incorrect passwords -> optional (opt-in) erase-the-phone

# What the FBI wanted...

- Apple provides a ***modified*** version of the operating system for the Secure Enclave which...
  - Removes the timeout on all password attempts
  - Enables password attempts through the USB connection
  - Enables an ***on-line*** brute force attack..  
but with a 4-digit PIN and 10 tries/second, you do the math...
- Apple ***cryptographically signs the rogue OS version!***
  - A horrific precedent:  
This is ***requiring*** that Apple both create a malicious version of the OS and sign it
  - If the FBI could compel Apple to do this, the NSA could too...  
It would make it ***impossible*** to trust software updates!

# Updating the SEP To Prevent This Possibility...

- The SEP will only accept updates ***signed by Apple***
- The FBI previously asked for this capability against a non-SEP equipped phone
  - "Hey Apple, cryptographically sign a corrupted version of the OS so that we can brute-force a password"
- How to prevent the FBI from asking again?
- Now, an OS update (either to the base OS and/or the SEP) requires the user to be logged in ***and input the password***
  - "To rekey the lock, you must first unlock the lock"
  - The FBI can only even ***attempt*** to ask before they have possession of the phone since once they have the phone they must also have the passcode
  - So when offered the chance to try again with a "Lone Wolf's" iPhone in the Texas church shooting, they haven't bothered
- At this point, Apple has now gone back and allows auto-updates for the base OS
  - (but probably not the SEP)

# The Limits of the SEP...

## The host O/S

- The SEP can keep the host OS from accessing things it shouldn't...
  - Credit cards stored for ApplePay, your fingerprint, etc...
- The SEP can **use** the random secret but not read it...
  - Can encrypt with it but can't read it
- But it can't keep the host OS from things it is supposed to access
  - All the user data when the user is logged in...
- So do have to rely on the host OS as part of **my TCB**
  - Fortunately it is updated continuously when vulnerabilities are found
    - Apple has responded to the discovery of very targeted zero-days in <30 days
  - And Apple has both good sandboxing of user applications and a history of decent vetting
    - So the random apps are **not** in the Trusted Base.

# The SEP and Apple Pay

- The SEP is what makes ApplePay possible
  - It handles the authentication to the user with the fingerprint reader/face reader
    - Verifies that it is the user not somebody random
  - It handles the emulation of the credit card
    - A "tokenized" Near Field Communication (NFC) wireless protocol
    - And a tokenized public key protocol for payments through the app
- ***Very hard*** to conduct a fraudulent transaction
  - Designed to enforce user consent at the SEP
- ***Disadvantage:*** The fingerprint reader is part of the trust domain
  - Which means you need special permission from Apple to replace the fingerprint reader when replacing a broken screen

# I *love* ApplePay...

- It is a **faster** protocol than the chip-and-signature
  - NFC protocol is designed to do the same operation in less time because the protocol is newer
- It is a **more secure** protocol than NFC on the credit card
  - Since it actually enforces user-consent
- It is more **privacy sensitive** than standard credit card payments
  - Generates a unique token for each transaction:  
Merchant is not supposed to link your transactions
- Result is its low cost:
  - Very hard to commit fraud -> less cost to transact
- I use it on my watch all the time

# Transitive Trust in the Apple Ecosystem...

- The most trusted item is the iPhone SEP
  - Assumed to be rock-solid
  - Fingerprint reader/face reader allows it to be convenient
- The watch trusts the phone
  - The pairing process includes a cryptographic key exchange mediated by close proximity and the camera
  - So Unlock the phone -> Unlock the watch
- My computer trusts my watch
  - Distance-bounded cryptographic protocol
  - So my watch unlocks my computer
- Result? I don't have to keep retying my password
  - Allows the use of ***strong passwords everywhere*** without driving myself crazy!



# Credit Card Fraud

- Under US law we have very good protections against fraud
  - Theoretical \$50 limit if we catch it quickly
  - \$0 limit in practice
- So cost of credit card fraud for me is the cost of recovery from fraud
  - Because fraud ***will happen***:
  - The mag stripe is all that is needed to duplicate a swipe-card
    - And you can still use swipe-only at gas pumps and other such locations
  - The numbers front and back is all that is needed for card-not-present fraud
    - And how many systems
- What are the recovery costs?
  - Being without the card for a couple of days...
    - Have a second back-up card
  - Having to change all my autopay items...
    - Grrrr....

# But What About "Debit" Cards?

- Theoretically the fraud protection is the same...
- But two caveats...
  - It is easier to not pay your credit card company than to claw money back from your bank...
  - Until the situation is resolved:
    - Credit card? It is the credit card company's money that is missing
    - Debit card? It is *your* money that is missing
- Result is debit card fraud is more transient disruptions...

# So Two Different Policies...

- Credit card: Hakunna Matata!
  - I use it without reservation, just with a spare in case something happens
  - Probably 2-3 compromise events have happened, and its annoying but ah well
  - The most interesting was \$1 to Tsunami relief in 2004...  
was a way for the attacker to test that the stolen card was valid
- Debit card: Paranoia-city...
  - It is an ATM-ONLY card (no Visa/Mastercard logo!)
  - It is used ONLY in ATMs belonging to my bank
  - Reduce the risk of "skimmers": rogue ATMs that record cards and keystrokes

# And Banking Information...

- ***Watch*** your bank account transactions
  - In case of fraud, you have protection but you need to notice
  - Bank accounts are particularly vulnerable:
    - The information on a cheque is all the data needed to transfer to/from an account!

# Putting Everything Together In the Real World: The "Sad DNS" Attack...

- Over a decade after the Kaminski attacks, DNS cache poisoning returned to the news
- Reminder: Kaminski strategy...
  - You send glue records to actually poison the target:  
So to poison `www.google.com`, you create a query for `a.google.com`...  
And in the additional include `www.google.com A 66.66.66.66`
  - Still have to guess TXID ( $2^{16}$  work factor), but can keep trying!
- Defense was randomize the UDP source port as well...
  - So attacker has to guess the port and TXID at the same time (so  $2^{32}$  work give-or-take)

# Work by Keyu Man et al..

## saddns.net (UC Riverside & Tsinghua University)

- Observation #1, can we detect what UDP port(s) are in use for a particular query?
  - If so, it turns the problem from expected  $2^{32}$  work to  $2^{16} + 2^{16}$  work!
  - You search for the open port, and if you get lucky, do the random TXIDs...
- Observation #2, can we cause the DNS authority for a domain to ***not respond?***
  - If so, enables us to have a lot more time for an attack
  - Which can make it far easier to be successful
- Answer to both is ***yes!***

# Answer to 1: Just Ask the DNS Resolver!

- By default you get a response if there is no open UDP port
  - "ICMP port unreachable"
- And UDP ports are not host specific by default...
  - So if you call `sendto()` and then `recvfrom()` ...  
you won't send an ICMP back for that port
  - Behavior is not the same for `connect()` semantics:  
Connect will only not send back an ICMP if the UDP packet is from the remote IP
- So just scan all  $2^{16}$  ports to see if you get a response!
- But there are gotchas...
  - ICMP packet sending has both a per-IP and global rate limit

# So Enter Side Channels....

- Spoof a bunch of packets that will *just* trigger the global rate limit
  - Then send a packet from your IP to a port you *know* will trigger an ICMP response
- If you get a response...
  - One of the ports you checked was open!
  - So divide and conquer
- If you don't... Wait the short 20ms timeout and go onto the next block of ports to check
- Oh, and if they use `connect()` for UDP...
  - You only don't get an ICMP back if you are the IP that was connected to...  
So just spoof the real server with the side channel check!

# And Now To Buy Time...

## Another Rate Limit...

- DNS servers can be used for reflected DOS attacks
  - Spoof the IP address of the target and send a packet to the DNS server
  - DNS server then replies...  
Making the attack look like its coming from the DNS server
  - And since DNS replies are bigger, this is an amplifier for DOS attacks
- So DNS authority servers have their own rate limit
  - Too many requests from a single IP and they will start ignoring some request
- So use ***that*** to buy time...
  - Send just enough requests spoofing the target resolver's IP address for nonsense requests
  - Target resolver ignores the replies (after all, they were never made)
  - But the DNS authority server will now ignore the target resolver's DNS request!

# Solution #1: DNSSEC

- If the resolver (or better yet client) validates DNSSEC...
  - Now it doesn't matter!
  - Fortunately DNSSEC serving is getting easier
    - Most people are using a few outsourced DNS services
    - So they can easily add in DNSSEC if they aren't already
    - **Any** managed DNS service should use DNSSEC these days

# Solution #2: Detection & Response

- Still relies on Kaminsky-style glue records for poisoning
  - Otherwise you can only race once on failure until the record's TTL expires
- This is **VERY NOISY**
  - Hundreds or thousands of non-matching responses
  - This is even noisier than standard Kaminsky:  
Lots of bogus replies from the real server to suppress the legitimate reply
- So detect and respond
  - Don't query once, query multiple times and accept majority
  - Don't promote glue into the cache
  - Or just don't resolve the targeted name(s)
  - Nobody does this however