

CS162  
Operating Systems and  
Systems Programming  
Lecture 7

Synchronization 2: Semaphores (Con't)  
Lock Implementation, Atomic Instructions

February 9<sup>th</sup>, 2021

Profs. Natacha Crooks and Anthony D. Joseph

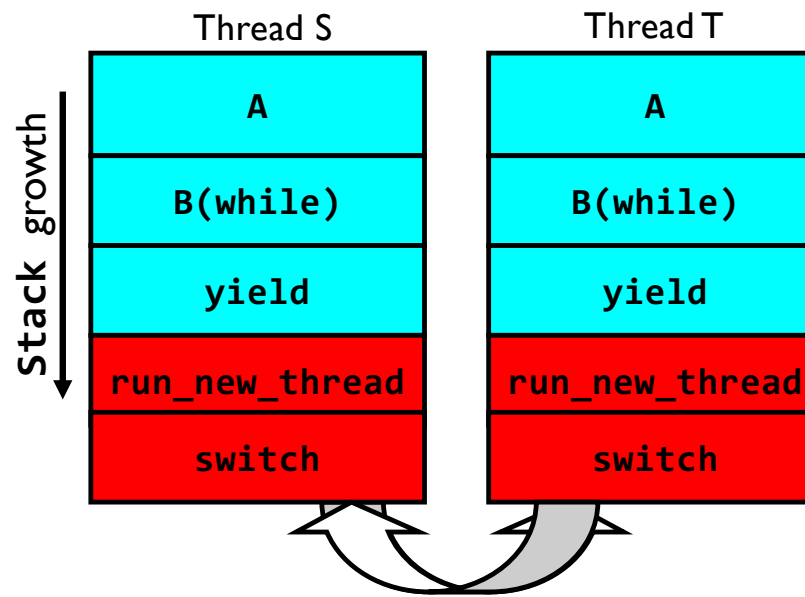
<http://cs162.eecs.Berkeley.edu>

## Recall: Multithreaded Stack Example

- Consider the following code blocks:

```
proc A() {  
    B();  
}  
proc B() {  
    while(TRUE) {  
        yield();  
    }  
}
```

- Suppose we have 2 threads:
  - Threads S and T




Thread S's switch returns to  
Thread T's (and vice versa)

## Recall: Use of Timer Interrupt to Return Control

---

- Solution to our dispatcher problem
  - Use the timer interrupt to force scheduling decisions



- Timer Interrupt routine:

```
TimerInterrupt() {  
    DoPeriodicHouseKeeping();  
    run_new_thread();  
}
```

# Hardware Context Switch Support in x86

- Syscall/Intr (U  $\rightarrow$  K)
  - PL 3  $\rightarrow$  0;
  - TSS  $\leftarrow$  EFLAGS, CS:EIP;
  - SS:ESP  $\leftarrow$  k-thread stack (TSS PL 0);
  - push (old) SS:ESP onto (new) k-stack
  - push (old) eflags, cs:eip, <err>
  - CS:EIP  $\leftarrow$  <k target handler>

- Then

- *Handler then saves other regs, etc*
- *Does all its works, possibly choosing other threads, changing PTBR (CR3)*
- kernel thread has set up user GPRs

- iret (K  $\rightarrow$  U)

- PL 0  $\rightarrow$  3;
- Eflags, CS:EIP  $\leftarrow$  popped off k-stack
- SS:ESP  $\leftarrow$  popped off k-stack

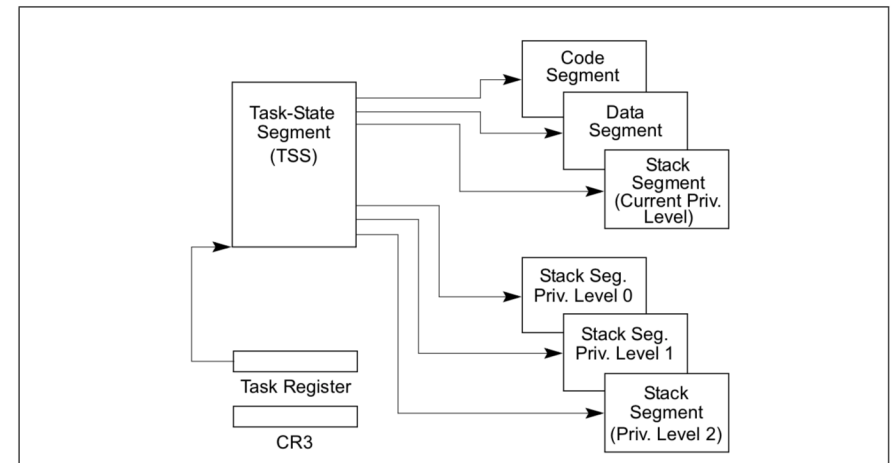



Figure 7-1. Structure of a Task

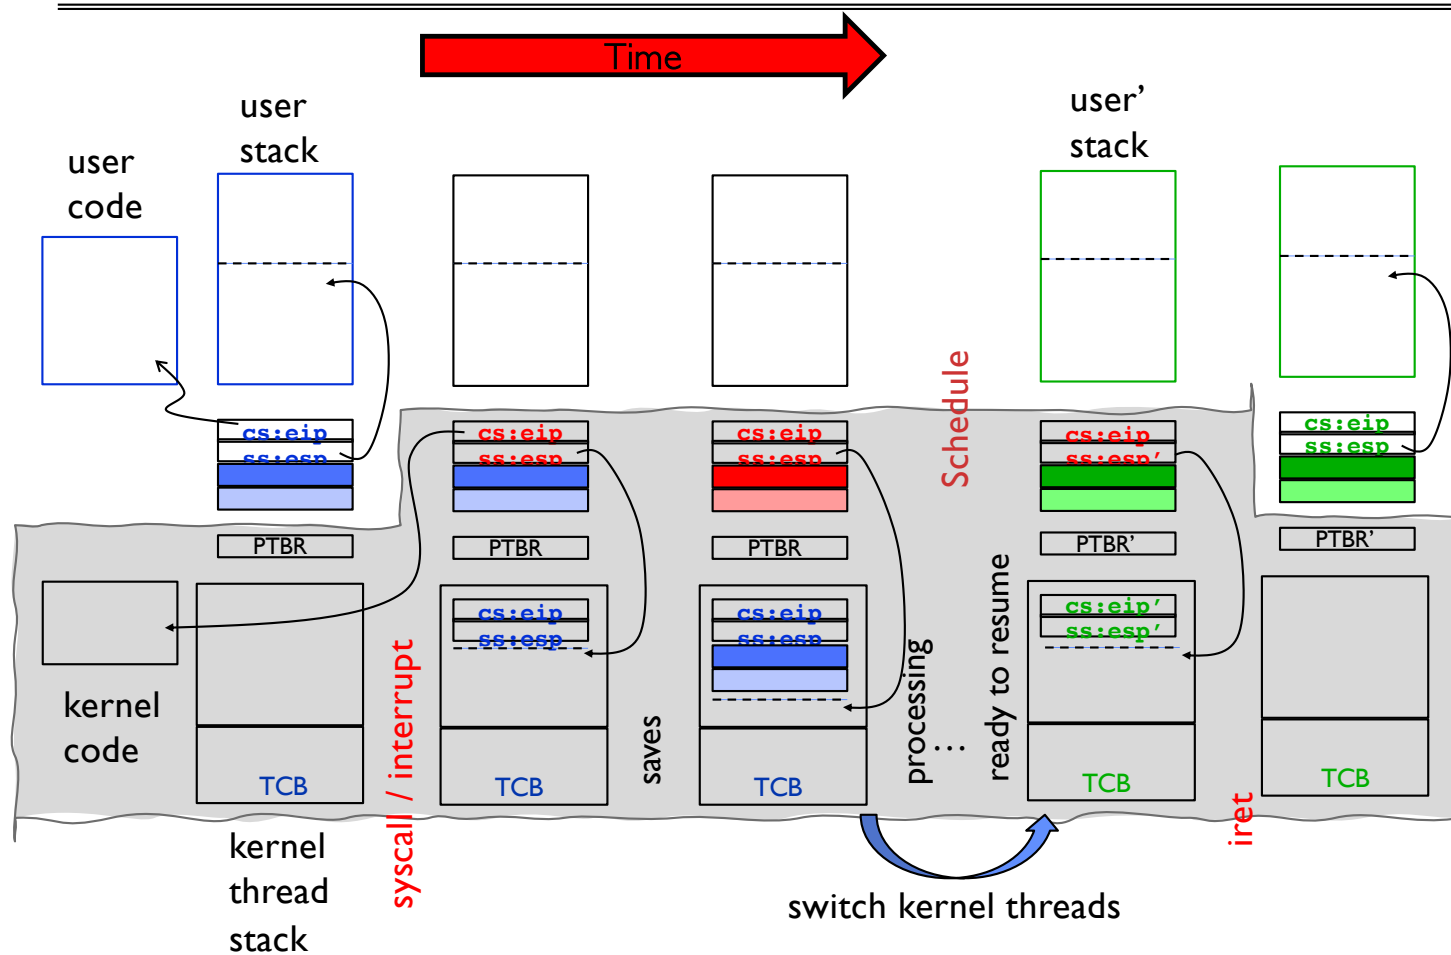
pg 2,942 of 4,922 of x86 reference manual

**Pintos: tss.c, intr-stubs.S**

# Pintos: Kernel Crossing on Syscall or Interrupt



# Pintos: Context Switch – Scheduling




Pintos: switch.S

## Recall: Fix banking problem with Locks!

- Identify critical sections (atomic instruction sequences) and add locking:

```
Deposit(acctId, amount) {  
    Lock.acquire()    // Wait if someone else in critical section!  
    acct = GetAccount(actId);  
    acct->balance += amount;  
    StoreAccount(acct);  
    Lock.release()    // Release someone into critical section  
}
```




- Must use SAME lock with all of the methods (Withdraw, etc...)

## Recall: Red-Black tree example

Thread A

```
Insert(3) {  
  acquire(&treelock)  
  Tree.Insert(3)  
  release(&treelock)  
}
```



Tree-Based Set Data Structure

Thread B


```
Insert(4) {  
  acquire(&treelock)  
  Tree.insert(4)  
  release(&treelock)  
}  
Get(6) {  
  acquire(&treelock)  
  Tree.search(6)  
  release(&treelock)  
}
```

- Here, the Lock is associated with the root of the tree
  - Restricts parallelism but makes sure that tree *always* consistent
  - No races at the operation level
- Threads are exchange information through a consistent data structure
- Could you make it faster with one lock per node? Perhaps, but must be careful!
  - Need to define invariants that are always true despite many simultaneous threads...




## Producer-Consumer with a Bounded Buffer

- Problem Definition
  - Producer(s) put things into a shared buffer
  - Consumer(s) take them out
  - Need synchronization to coordinate producer/consumer
- Don't want producer and consumer to have to work in lockstep, so put a fixed-size buffer between them
  - Need to synchronize access to this buffer
  - Producer needs to wait if buffer is full
  - Consumer needs to wait if buffer is empty
- Example 1: GCC compiler
  - `cpp | cc1 | cc2 | as | ld`
- Example 2: Coke machine
  - Producer can put limited number of Cokes in machine
  - Consumer can't take Cokes out if machine is empty
- Others: Web servers, Routers, ....



## Circular Buffer Data Structure (sequential case)

```
typedef struct buf {  
    int write_index;  
    int read_index;  
    <type> *entries[BUFSIZE];  
} buf_t;
```



- Insert: write & bump write ptr (enqueue)
- Remove: read & bump read ptr (dequeue)
- *How to tell if Full (on insert) Empty (on remove)?*
- *And what do you do if it is?*
- *What needs to be atomic?*

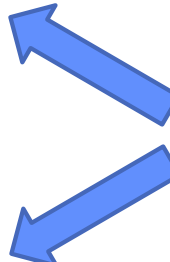
## Circular Buffer – first cut

---

mutex buf\_lock = <initially unlocked>

```
Producer(item) {  
    acquire(&buf_lock);  
    while (buffer full) {}; // Wait for a free slot  
    enqueue(item);  
    release(&buf_lock);  
}
```

```
Consumer() {  
    acquire(&buf_lock);  
    while (buffer empty) {}; // Wait for arrival  
    item = dequeue();  
    release(&buf_lock);  
    return item  
}
```



Will we ever come out of the wait loop?

## Circular Buffer – 2<sup>nd</sup> cut

mutex buf\_lock = <initially unlocked>


```
Producer(item) {  
    acquire(&buf_lock);  
    while (buffer full) {release(&buf_lock); acquire(&buf_lock);}  
    enqueue(item);  
    release(&buf_lock);  
}
```

```
Consumer() {  
    acquire(&buf_lock);  
    while (buffer empty) {release(&buf_lock); acquire(&buf_lock);}  
    item = dequeue();  
    release(&buf_lock);  
    return item  
}
```

What happens when one is waiting for the other?

- Multiple cores ?
- Single core ?





1

CHAPTER 1

1

CHAPTER 1


## Higher-level Primitives than Locks

---

- What is right abstraction for synchronizing threads that share memory?
  - Want as high a level primitive as possible
- Good primitives and practices important!
  - Since execution is not entirely sequential, really hard to find bugs, since they happen rarely
  - UNIX is pretty stable now, but up until about mid-80s (10 years after started), systems running UNIX would crash every week or so – concurrency bugs
- Synchronization is a way of coordinating multiple concurrent activities that are using shared state
  - This lecture and the next presents some ways of structuring sharing

## Recall: Semaphores

---




- Semaphores are a kind of generalized lock
  - First defined by Dijkstra in late 60s
  - Main synchronization primitive used in original UNIX
- Definition: a Semaphore has a non-negative integer value and supports the following two operations:
  - **Down()** or **P()**: an atomic operation that waits for semaphore to become positive, then decrements it by 1
    - » Think of this as the wait() operation
  - **Up()** or **V()**: an atomic operation that increments the semaphore by 1, waking up a waiting P, if any
    - » Think of this as the signal() operation
  - Note that **P()** stands for “*proberen*” (to test) and **V()** stands for “*verhogen*” (to increment) in Dutch

## Semaphores Like Integers Except...

---

- Semaphores are like integers, except:
  - No negative values
  - Only operations allowed are P and V – can't read or write value, except initially
  - Operations must be atomic
    - » Two P's together can't decrement value below zero
    - » Thread going to sleep in P won't miss wakeup from V – even if both happen at same time
- POSIX adds ability to read value, but technically not part of proper interface!
- Semaphore from railway analogy
  - Here is a semaphore initialized to 2 for resource control:





## Two Uses of Semaphores

---

Mutual Exclusion (initial value = 1)

- Also called “Binary Semaphore” or “mutex”.
- Can be used for mutual exclusion, just like a lock:

```
semaP(&mysem);  
    // Critical section goes here  
semaV(&mysem);
```

Scheduling Constraints (initial value = 0)

- Allow thread 1 to wait for a signal from thread 2
  - thread 2 **schedules** thread 1 when a given **event** occurs
- Example: suppose you had to implement ThreadJoin which must wait for thread to terminate:

**Initial value of semaphore = 0**

```
ThreadJoin {  
    semaP(&mysem);  
}  
ThreadFinish {  
    semaV(&mysem);  
}
```



## Revisit Bounded Buffer: Correctness constraints for solution

---

- Correctness Constraints:
  - Consumer must wait for producer to fill buffers, if none full (scheduling constraint)
  - Producer must wait for consumer to empty buffers, if all full (scheduling constraint)
  - Only one thread can manipulate buffer queue at a time (mutual exclusion)
- Remember why we need mutual exclusion
  - Because computers are stupid
  - Imagine if in real life: the delivery person is filling the machine and somebody comes up and tries to stick their money into the machine
- General rule of thumb: Use a separate semaphore for each constraint
  - Semaphore fullBuffers; // consumer's constraint
  - Semaphore emptyBuffers; // producer's constraint
  - Semaphore mutex; // mutual exclusion

# Full Solution to Bounded Buffer (coke machine)

```
Semaphore fullSlots = 0;    // Initially, no coke
Semaphore emptySlots = bufSize; // Initially, num empty slots
Semaphore mutex = 1;        // No one using machine
```



```
Producer(item) {
    semaP(&emptySlots); // Wait until space
    semaP(&mutex);       // Wait until machine free
    Enqueue(item);
    semaV(&mutex);
    semaV(&fullSlots);  // Tell consumers there is
                        // more coke
}

Consumer() {
    semaP(&fullSlots); // Check if there's a coke
    semaP(&mutex);     // Wait until machine free
    item = Dequeue();
    semaV(&mutex);
    semaV(&emptySlots); // tell producer need more
    return item;
}
```

emptySlots  
signals space

fullSlots signals coke

Critical sections  
using mutex  
protect integrity of  
the queue

## Discussion about Solution

- Why asymmetry?

- Producer does: **semaP(&emptyBuffer), semaV(&fullBuffer)**
- Consumer does: **semaP(&fullBuffer), semaV(&emptyBuffer)**

Decrease # of  
empty slots

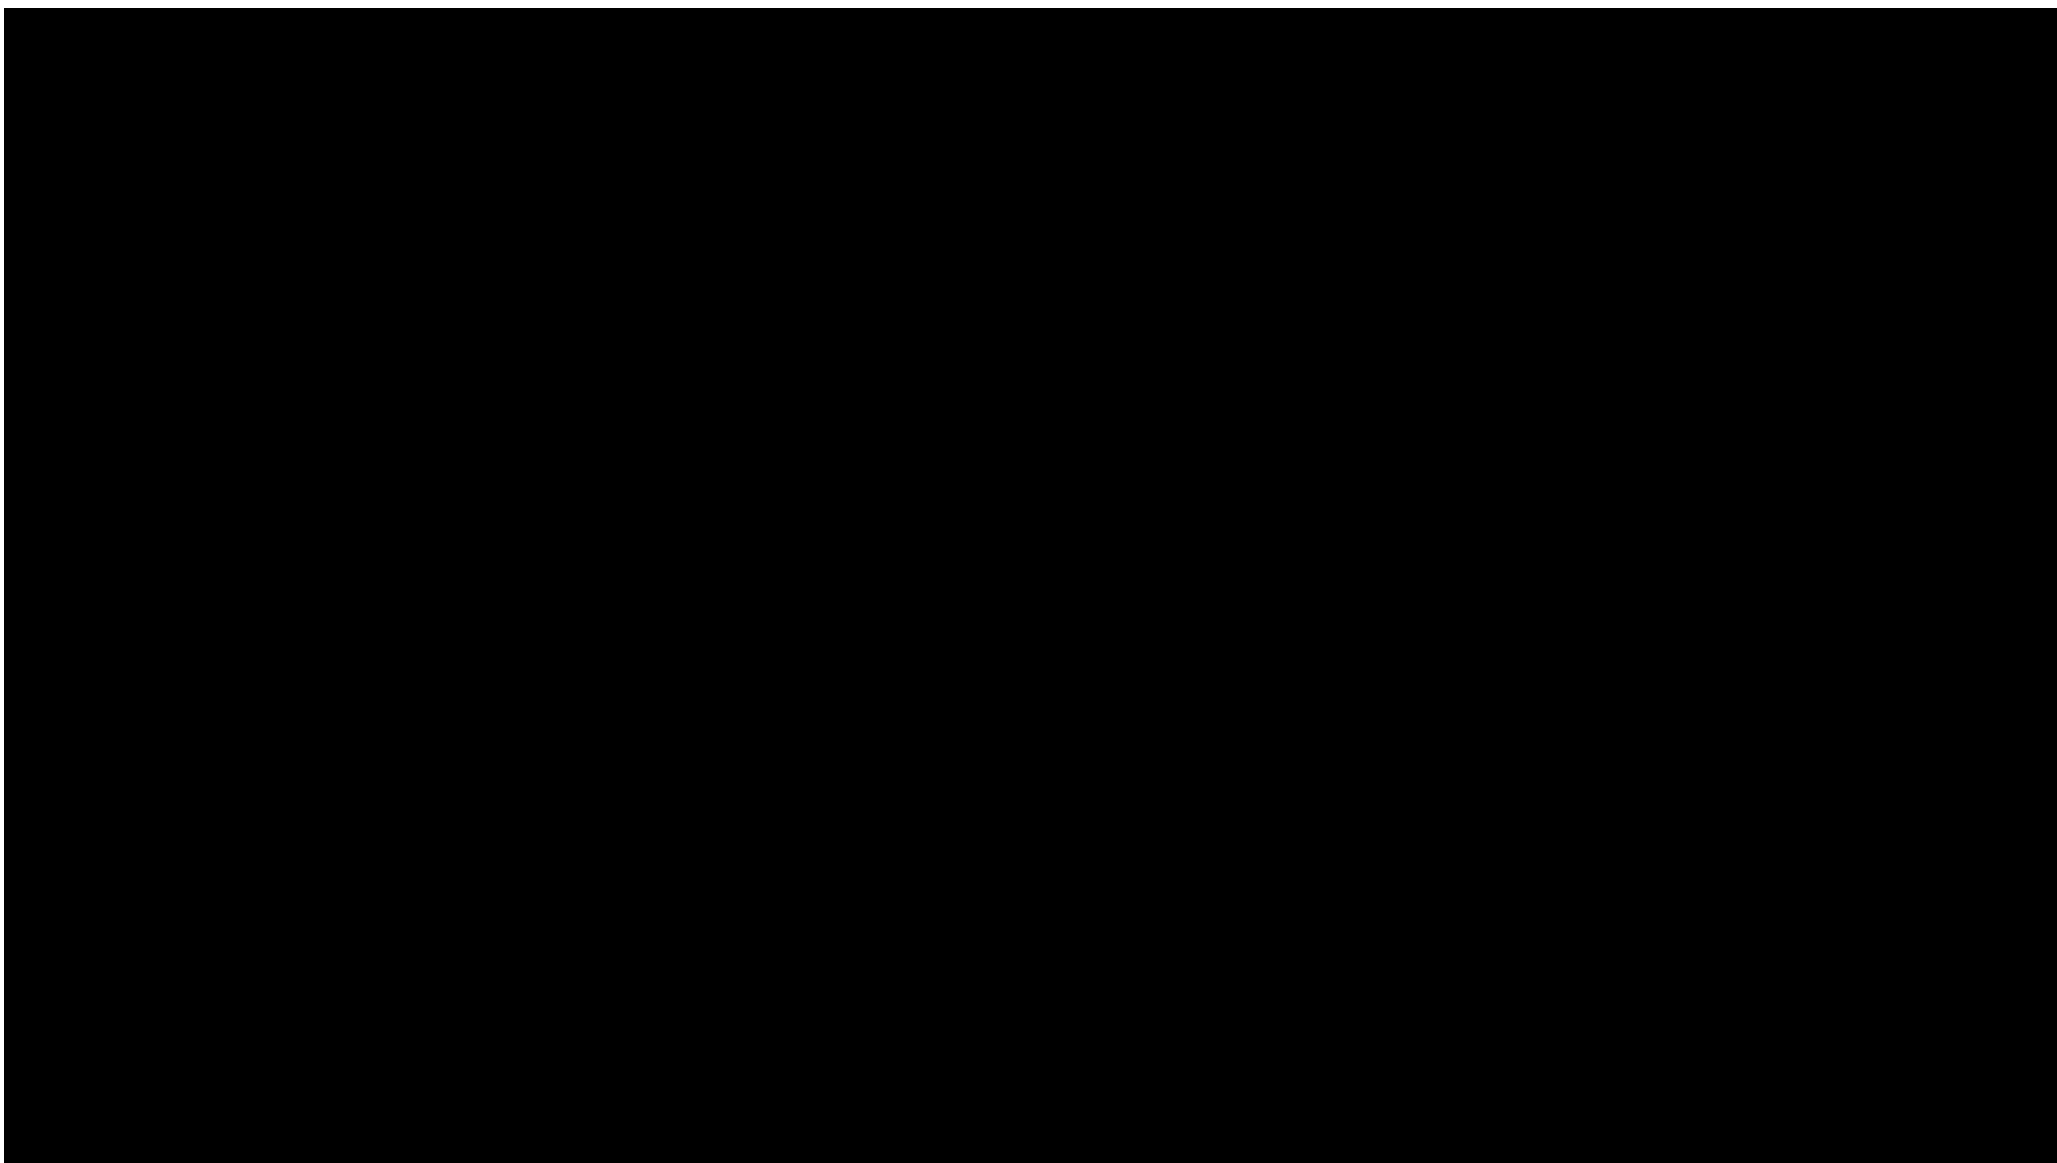
Increase # of  
occupied slots

Decrease # of  
occupied slots

Increase # of  
empty slots

- Is order of P's important?
- Is order of V's important?
- What if we have 2 producers  
or 2 consumers?

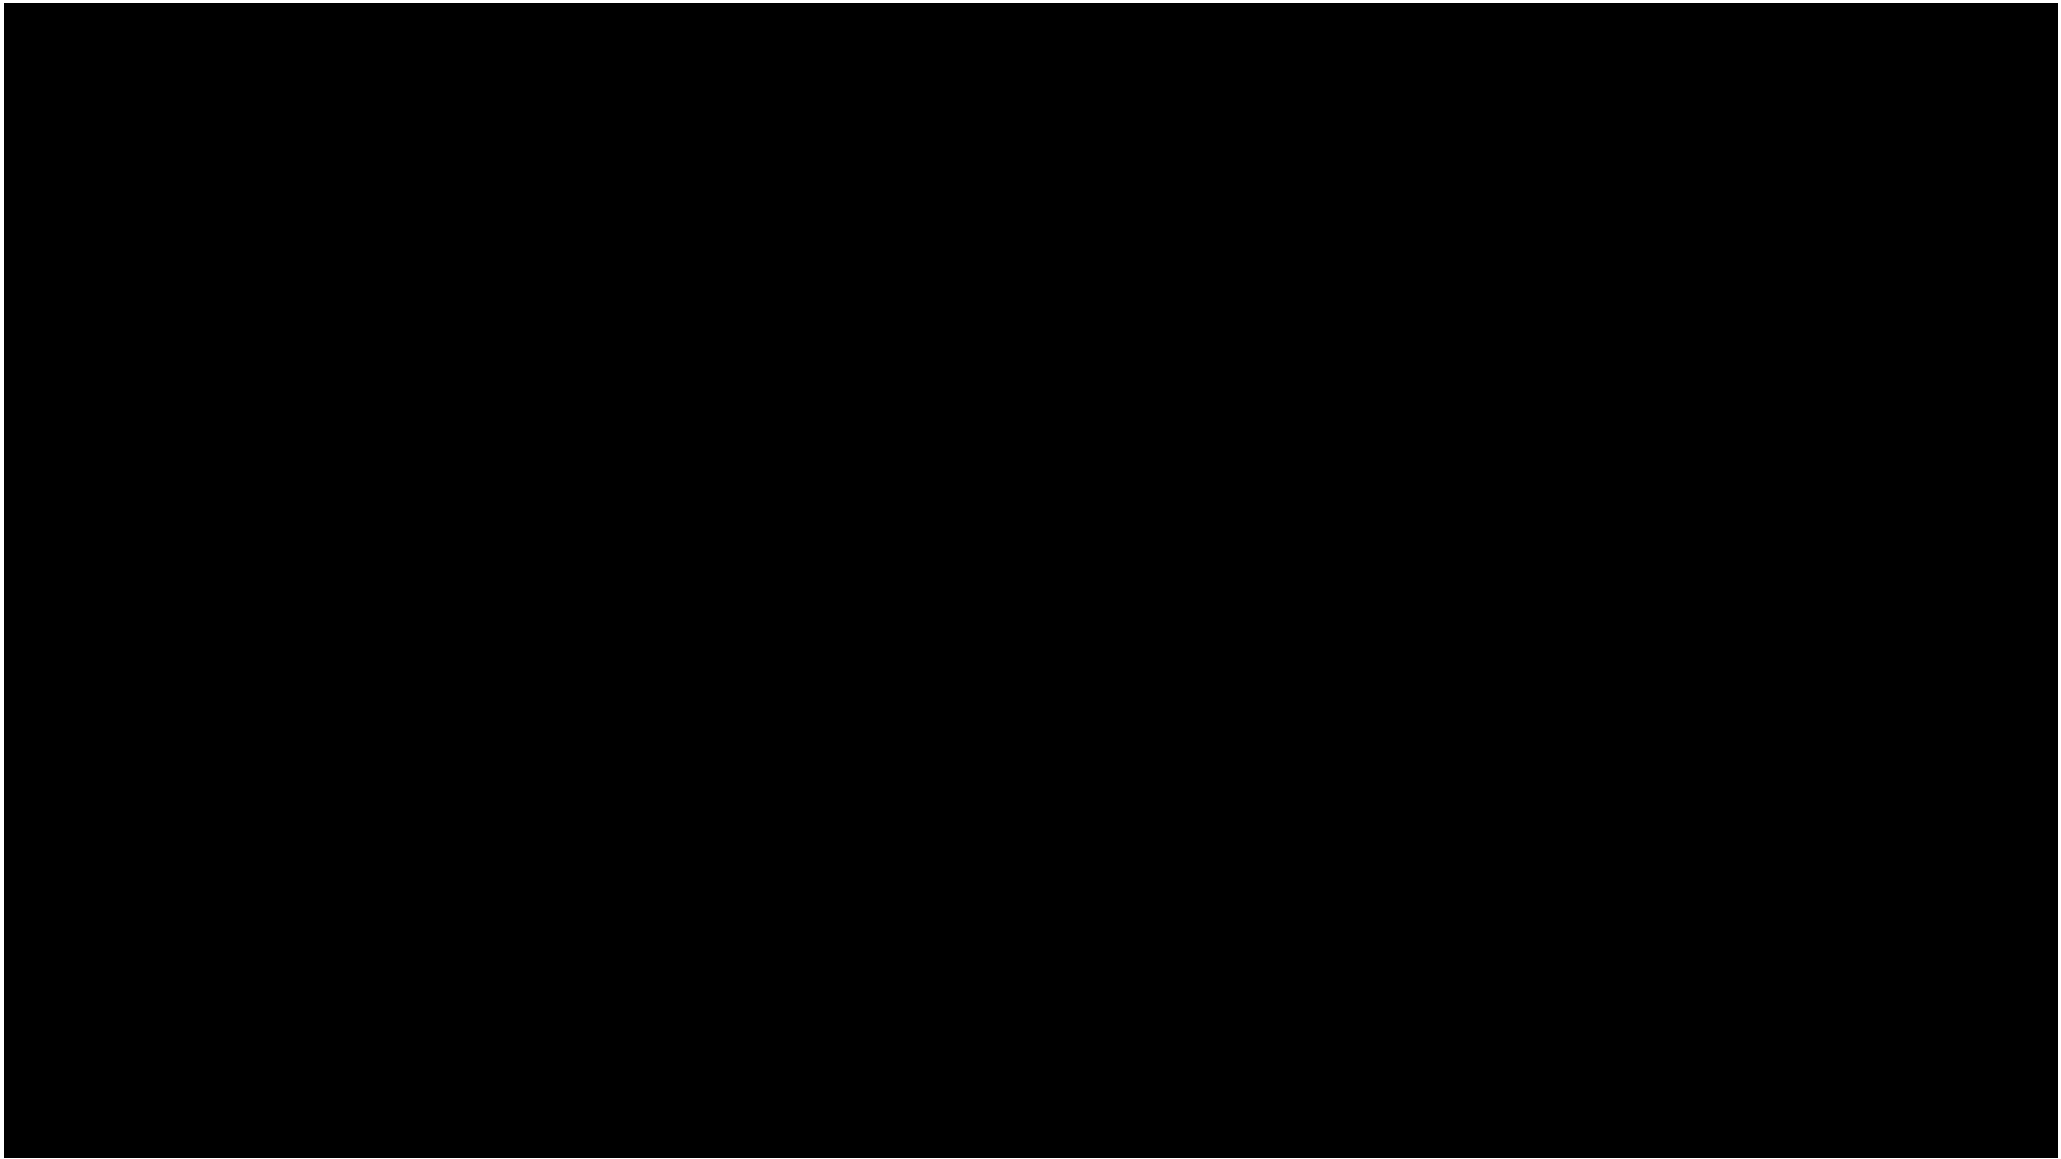
```
Producer(item) {  
    semaP(&mutex);  
    semaP(&emptySlots);  
    Enqueue(item);  
    semaV(&mutex);  
    semaV(&fullSlots);  
}  
Consumer() {  
    semaP(&fullSlots);  
    semaP(&mutex);  
    item = Dequeue();  
    semaV(&mutex);  
    semaV(&emptySlots);  
    return item;  
}
```



## Administrivia

---

- Midterm I: Thu February 18<sup>th</sup>, 5-6:30PM (9 days from today!)
  - Video Proctored, Use of computer to answer questions
  - More details as we get closer to exam
- Midterm topics:
  - Everything up to lecture 9 – lecture will be released early
  - Homework I and Project I (high-level design) are fair game
- Midterm Review: Tuesday February 16<sup>th</sup>, 5-7pm
  - Details TBA



## Where are we going with synchronization?

Programs	Shared Programs
Higher-level API	Locks   Semaphores   Monitors   Send/Receive
Hardware	Load/Store   Disable Ints   Test&Set   Compare&Swap

- We are going to implement various higher-level synchronization primitives using atomic operations
  - Everything is pretty painful if only atomic primitives are load and store
  - Need to provide primitives useful at user-level



## Motivating Example: “Too Much Milk”

---

- Great thing about OS's – analogy between problems in OS and problems in real life
  - Help you understand real life problems better
  - But, computers are much stupider than people
- Example: People need to coordinate:



Time	Person A	Person B
3:00	Look in Fridge. Out of milk	
3:05	Leave for store	
3:10	Arrive at store	Look in Fridge. Out of milk
3:15	Buy milk	Leave for store
3:20	Arrive home, put milk away	Arrive at store
3:25		Buy milk
3:30		Arrive home, put milk away

## Recall: What is a lock?

---

- **Lock**: prevents someone from doing something
  - **Lock** before entering critical section and before accessing shared data
  - **Unlock** when leaving, after accessing shared data
  - **Wait** if locked
    - » Important idea: all synchronization involves waiting
- For example: fix the milk problem by putting a key on the refrigerator
  - Lock it and take key if you are going to go buy milk
  - Fixes too much: roommate angry if only wants OJ



- Of Course – We don't know how to make a lock yet
  - Let's see if we can answer this question!

## Too Much Milk: Correctness Properties

---


- Need to be careful about correctness of concurrent programs, since non-deterministic
  - Impulse is to start coding first, then when it doesn't work, pull hair out
  - Instead, think first, then code
  - Always write down behavior first
- What are the correctness properties for the “Too much milk” problem???
- Never more than one person buys
- Someone buys if needed
- First attempt: Restrict ourselves to use only atomic load and store operations as building blocks

## Too Much Milk: Solution #1

---

- Use a note to avoid buying too much milk:
  - Leave a note before buying (kind of “lock”)
  - Remove note after buying (kind of “unlock”)
  - Don’t buy if note (wait)
- Suppose a computer tries this (remember, only memory read/write are atomic):

```
if (noMilk) {  
    if (noNote) {  
        leave Note;  
        buy milk;  
        remove note;  
    }  
}
```



## Too Much Milk: Solution #1

---

- Use a note to avoid buying too much milk:
  - Leave a note before buying (kind of “lock”)
  - Remove note after buying (kind of “unlock”)
  - Don’t buy if note (wait)
- Suppose a computer tries this (remember, only memory read/write are atomic):

```
Thread A  
if (noMilk) {  
  
    if (noNote) {  
        leave Note;  
        buy Milk;  
        remove Note;  
    }  
}
```


```
Thread B  
  
if (noMilk) {  
    if (noNote) {  
  
  
  
        leave Note;  
        buy Milk;  
        remove Note;  
    }  
}
```

## Too Much Milk: Solution #1

---

- Use a note to avoid buying too much milk:
  - Leave a note before buying (kind of “lock”)
  - Remove note after buying (kind of “unlock”)
  - Don’t buy if note (wait)
- Suppose a computer tries this (remember, only memory read/write are atomic):

```
if (noMilk) {  
    if (noNote) {  
        leave Note;  
        buy milk;  
        remove note;  
    }  
}
```



- Result?
  - Still too much milk **but only occasionally!**
  - Thread can get context switched after checking milk and note but before buying milk!
- Solution makes problem worse since fails **intermittently**
  - Makes it really hard to debug...
  - Must work despite what the dispatcher does!

## Too Much Milk: Solution #1 ½

---

- Clearly the Note is not quite blocking enough
  - Let's try to fix this by placing note first
- Another try at previous solution:

```
leave Note;  
if (noMilk) {  
    if (noNote) {  
        buy milk;  
    }  
}  
remove Note;
```

- What happens here?
  - Well, with human, probably nothing bad
  - With computer: no one ever buys milk



## Too Much Milk Solution #2

---

- How about labeled notes?
  - Now we can leave note before checking
- Algorithm looks like this:

Thread A  
leave note A;  
if (noNote B) {  
    if (noMilk) {  
        buy Milk;  
    }  
}  
remove note A;


Thread B  
leave note B;  
if (noNoteA) {  
    if (noMilk) {  
        buy Milk;  
    }  
}  
remove note B;

- Does this work?
- Possible for neither thread to buy milk
  - Context switches at exactly the wrong times can lead each to think that the other is going to buy
- Really insidious:
  - **Extremely unlikely** this would happen, but will at worse possible time
  - Probably something like this in UNIX



## Too Much Milk Solution #2: problem!

---



- *I'm not getting milk, You're getting milk*
- This kind of lockup is called “starvation!”

## Too Much Milk Solution #3

---

- Here is a possible two-note solution:

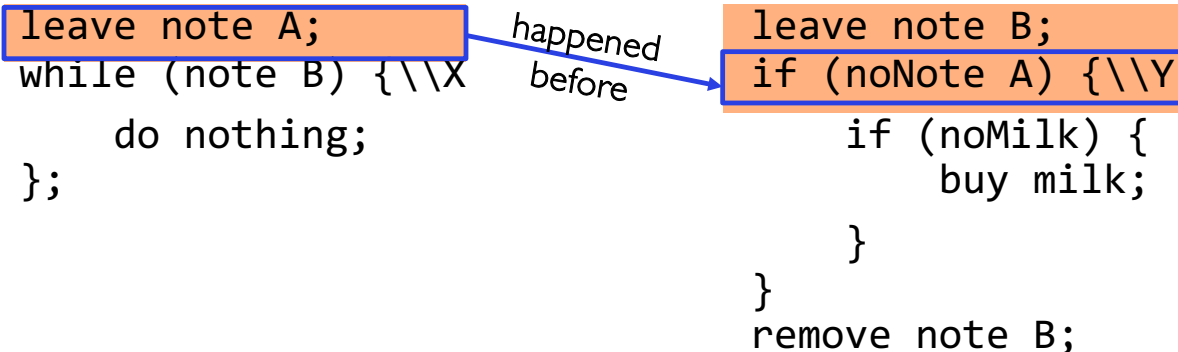
<u>Thread A</u>	<u>Thread B</u>
leave note A;	leave note B;
while (note B) {\\X	if (noNote A) {\\Y
do nothing;	if (noMilk) {
}	buy milk;
if (noMilk) {	}
buy milk;	}
}	remove note B;
remove note A;	

- Does this work? **Yes**. Both can guarantee that:
  - It is safe to buy, or
  - Other will buy, ok to quit
- At **X**:
  - If no note B, safe for A to buy,
  - Otherwise wait to find out what will happen
- At **Y**:
  - If no note A, safe for B to buy
  - Otherwise, A is either buying or waiting for B to quit

## Case I

---

- “leave note A” happens before “if (noNote A)”



```
leave note A;
while (note B) {\X
    do nothing;
};

if (noMilk) {
    buy milk;}
}
remove note A;
```

```
leave note B;
if (noNote A) {\Y
    if (noMilk) {
        buy milk;
    }
}
remove note B;
```

## Case I

---

- “leave note A” happens before “if (noNote A)”

```
leave note A;  
while (note B) {\X  
    do nothing;  
};
```

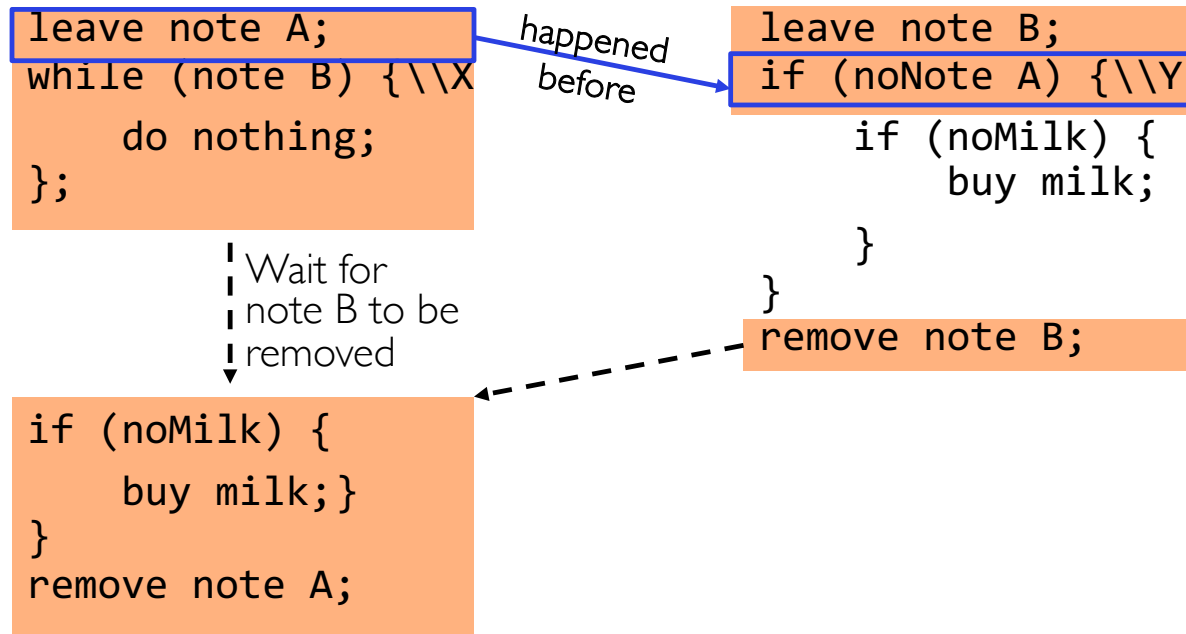
*happened  
before*

```
leave note B;  
if (noNote A) {\Y  
    if (noMilk) {  
        buy milk;  
    }  
}  
remove note B;
```

```
if (noMilk) {  
    buy milk;}  
}  
remove note A;
```

## Case I

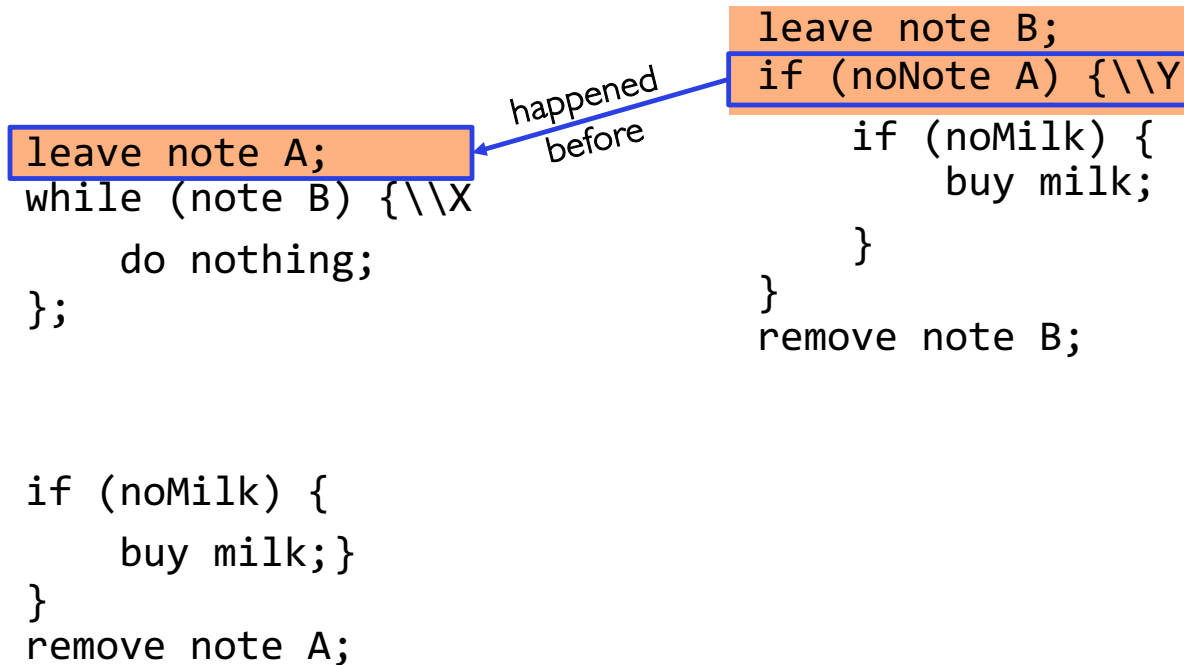
- “leave note A” happens before “if (noNote A)”



## Case 2

---


- “if (noNote A)” happens before “leave note A”



## Case 2

---

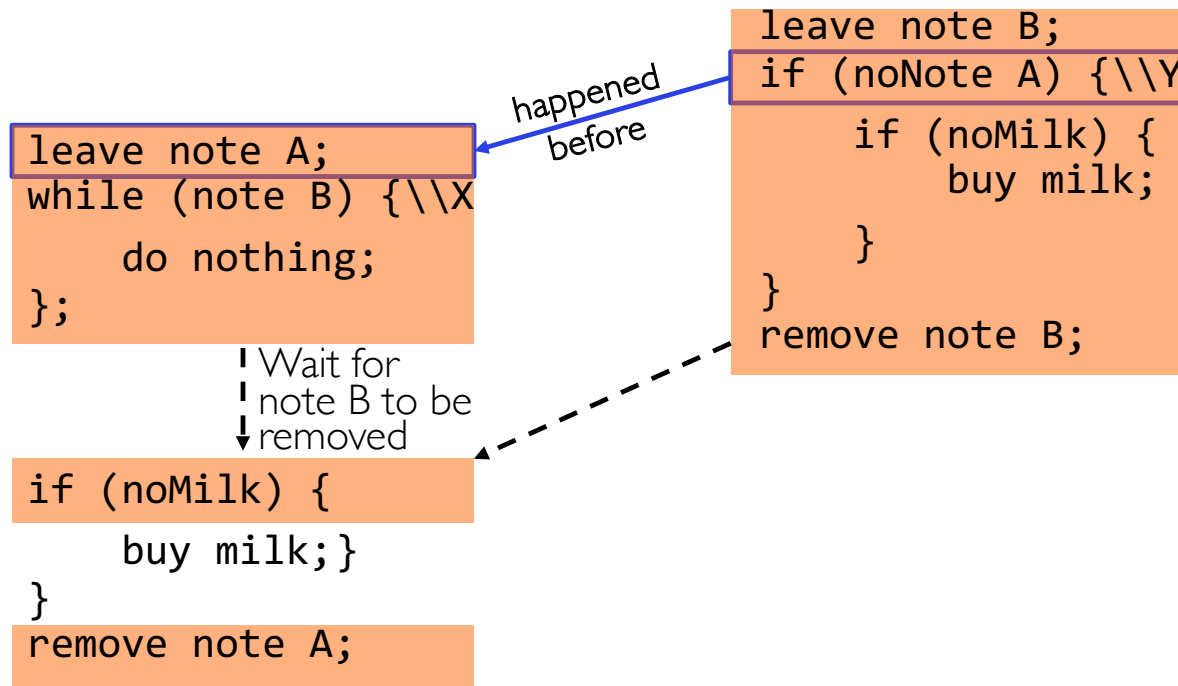
- “if (noNote A)” happens before “leave note A”



```
if (noMilk) {  
    buy milk;  
}  
remove note A;
```

## Case 2

- “if (noNote A)” happens before “leave note A”





---

## This Generalizes to $n$ Threads...

- Leslie Lamport's "Bakery Algorithm" (1974)

Computer  
Systems

G. Bell, D. Siewiorek,  
and S.H. Fuller, Editors

---

### A New Solution of Dijkstra's Concurrent Programming Problem

Leslie Lamport  
Massachusetts Computer Associates, Inc.

---

**A simple solution to the mutual exclusion problem is  
presented which allows the system to continue to operate**

## Solution #3 discussion

---

- Our solution protects a single “Critical-Section” piece of code for each thread:

```
if (noMilk) {  
    buy milk;  
}
```


- Solution #3 works, but it is really unsatisfactory
  - Really complex – even for this simple an example
    - » Hard to convince yourself that this really works
  - A’s code is different from B’s – what if lots of threads?
    - » Code would have to be slightly different for each thread
  - While A is waiting, it is consuming CPU time
    - » This is called “busy-waiting”
- There must be a better way!
  - Have hardware provide higher-level primitives than atomic load & store
  - Build even higher-level programming abstractions on this hardware support

## Too Much Milk: Solution #4?

---

- Recall our target lock interface:
  - `acquire(&milklock)` – wait until lock is free, then grab
  - `release(&milklock)` – Unlock, waking up anyone waiting
  - These must be atomic operations – if two threads are waiting for the lock and both see it is free, only one succeeds to grab the lock
- Then, our milk problem is easy:

```
acquire(&milklock);  
if (nomilk)  
    buy milk;  
release(&milklock);
```



## Back to: How to Implement Locks?

---

- **Lock**: prevents someone from doing something
  - Lock before entering critical section and before accessing shared data
  - Unlock when leaving, after accessing shared data
  - Wait if locked
    - » Important idea: all synchronization involves waiting
    - » Should *sleep* if waiting for a long time
- Atomic Load/Store: yields a solution like Milk #3
  - Pretty complex and error prone
- Hardware Lock instruction
  - Is this a good idea?
  - What about putting a task to sleep?
    - » What is the interface between the hardware and scheduler?
  - Complexity?
    - » Done in the Intel 432
    - » Each feature makes HW more complex and slow



## Naïve use of Interrupt Enable/Disable

---

- How can we build multi-instruction atomic operations?
  - Recall: dispatcher gets control in two ways.
    - » Internal: Thread does something to relinquish the CPU
    - » External: Interrupts cause dispatcher to take CPU
  - On a uniprocessor, can avoid context-switching by:
    - » Avoiding internal events (although virtual memory tricky)
    - » Preventing external events by disabling interrupts

- Consequently, naïve Implementation of locks:

```
LockAcquire { disable Ints; }  
LockRelease { enable Ints; }
```

- Problems with this approach:
  - Can't let user do this! Consider following:

```
LockAcquire();  
While(TRUE) {;;}
```

- Real-Time system—no guarantees on timing!
  - » Critical Sections might be arbitrarily long
- What happens with I/O or other important events?
  - » “Reactor about to meltdown. Help?”



## Better Implementation of Locks by Disabling Interrupts

---

- Key idea: maintain a lock variable and impose mutual exclusion only during operations on that variable

`int value = FREE;`



```
Acquire() {  
    disable interrupts;  
    if (value == BUSY) {  
        put thread on wait queue;  
        Go to sleep();  
        // Enable interrupts?  
    } else {  
        value = BUSY;  
    }  
    enable interrupts;  
}
```

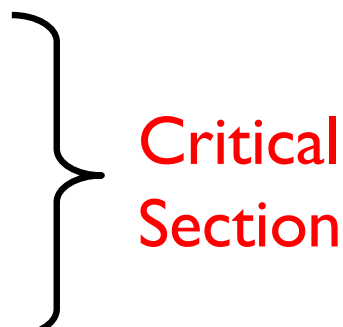
```
Release() {  
    disable interrupts;  
    if (anyone on wait queue) {  
        take thread off wait queue  
        Place on ready queue;  
    } else {  
        value = FREE;  
    }  
    enable interrupts;  
}
```

## New Lock Implementation: Discussion

---

- Why do we need to disable interrupts at all?
  - Avoid interruption between checking and setting lock value
  - Otherwise two threads could think that they both have lock

```
Acquire() {  
    disable interrupts;  
    if (value == BUSY) {  
        put thread on wait queue;  
        Go to sleep();  
        // Enable interrupts?  
    } else {  
        value = BUSY;  
    }  
    enable interrupts;  
}
```



Critical  
Section

- Note: unlike previous solution, the critical section (inside **Acquire()**) is very short
  - User of lock can take as long as they like in their own critical section: doesn't impact global machine behavior
  - Critical interrupts taken in time!



## Interrupt Re-enable in Going to Sleep

---

- What about re-enabling ints when going to sleep?


```
Acquire() {  
    disable interrupts;  
    if (value == BUSY) {  
        put thread on wait queue;  
        Go to sleep();  
    } else {  
        value = BUSY;  
    }  
    enable interrupts;  
}
```

## Interrupt Re-enable in Going to Sleep

---

- What about re-enabling ints when going to sleep?

```
Acquire() {  
    disable interrupts;  
    if (value == BUSY) {  
        put thread on wait queue;  
        Go to sleep();  
    } else {  
        value = BUSY;  
    }  
    enable interrupts;  
}
```


Enable Position 

- Before Putting thread on the wait queue?

## Interrupt Re-enable in Going to Sleep

---

- What about re-enabling ints when going to sleep?

Enable Position 

```
Acquire() {  
    disable interrupts;  
    if (value == BUSY) {  
        put thread on wait queue;  
        Go to sleep();  
    } else {  
        value = BUSY;  
    }  
    enable interrupts;  
}
```


- Before Putting thread on the wait queue?
  - Release can check the queue and not wake up thread

## Interrupt Re-enable in Going to Sleep

---

- What about re-enabling ints when going to sleep?

```
Acquire() {  
    disable interrupts;  
    if (value == BUSY) {  
        put thread on wait queue;  
        Go to sleep();  
    } else {  
        value = BUSY;  
    }  
    enable interrupts;  
}
```

Enable Position 

- Before Putting thread on the wait queue?
  - Release can check the queue and not wake up thread
- After putting the thread on the wait queue

## Interrupt Re-enable in Going to Sleep

---

- What about re-enabling ints when going to sleep?

```
Acquire() {  
    disable interrupts;  
    if (value == BUSY) {  
        put thread on wait queue;  
        Go to sleep();  
    } else {  
        value = BUSY;  
    }  
    enable interrupts;  
}
```

Enable Position →


- Before Putting thread on the wait queue?
  - Release can check the queue and not wake up thread
- After putting the thread on the wait queue
  - Release puts the thread on the ready queue, but the thread still thinks it needs to go to sleep
  - Misses wakeup and still holds lock (deadlock!)

## Interrupt Re-enable in Going to Sleep

---

- What about re-enabling ints when going to sleep?

```
Acquire() {  
    disable interrupts;  
    if (value == BUSY) {  
        put thread on wait queue;  
        Go to sleep();  
    } else {  
        value = BUSY;  
    }  
    enable interrupts;  
}
```

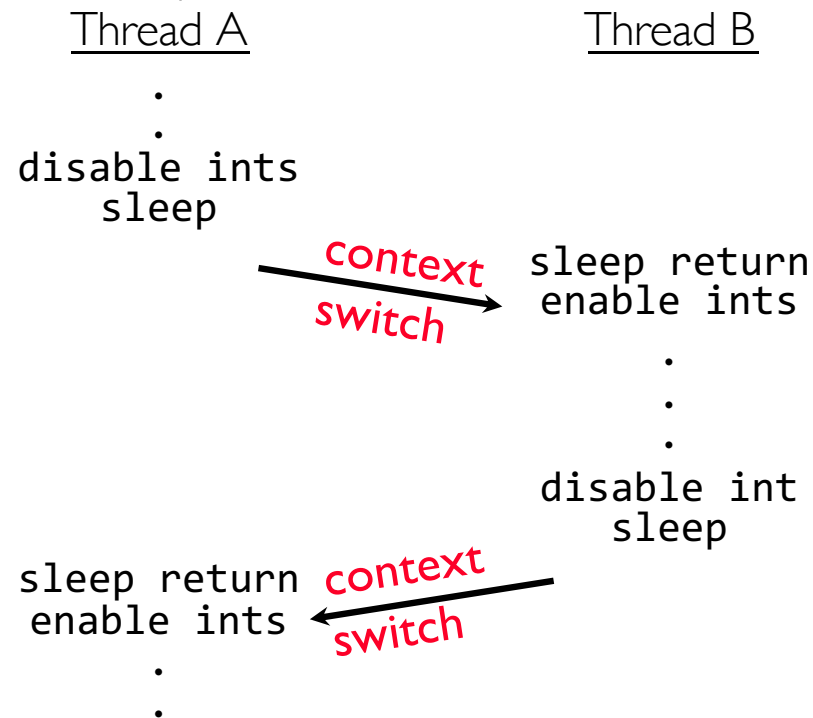
Enable Position 

- Before Putting thread on the wait queue?
  - Release can check the queue and not wake up thread
- After putting the thread on the wait queue
  - Release puts the thread on the ready queue, but the thread still thinks it needs to go to sleep
  - Misses wakeup and still holds lock (deadlock!)
- Want to put it after **sleep()**. But – how?

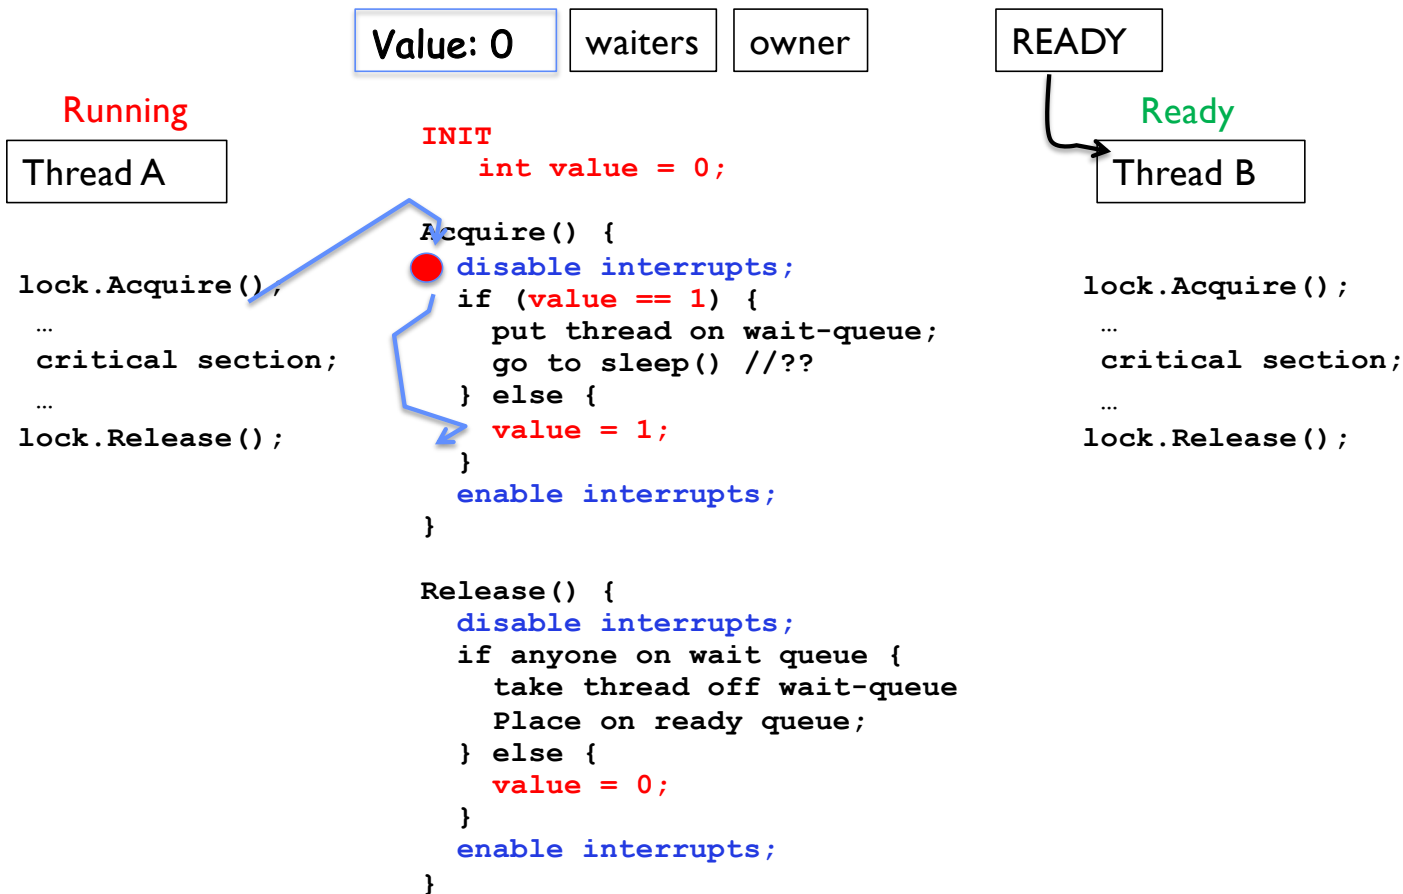
## How to Re-enable After Sleep()?

---

- In scheduler, since interrupts are disabled when you call sleep:
  - Responsibility of the next thread to re-enable ints
  - When the sleeping thread wakes up, returns to acquire and re-enables interrupts

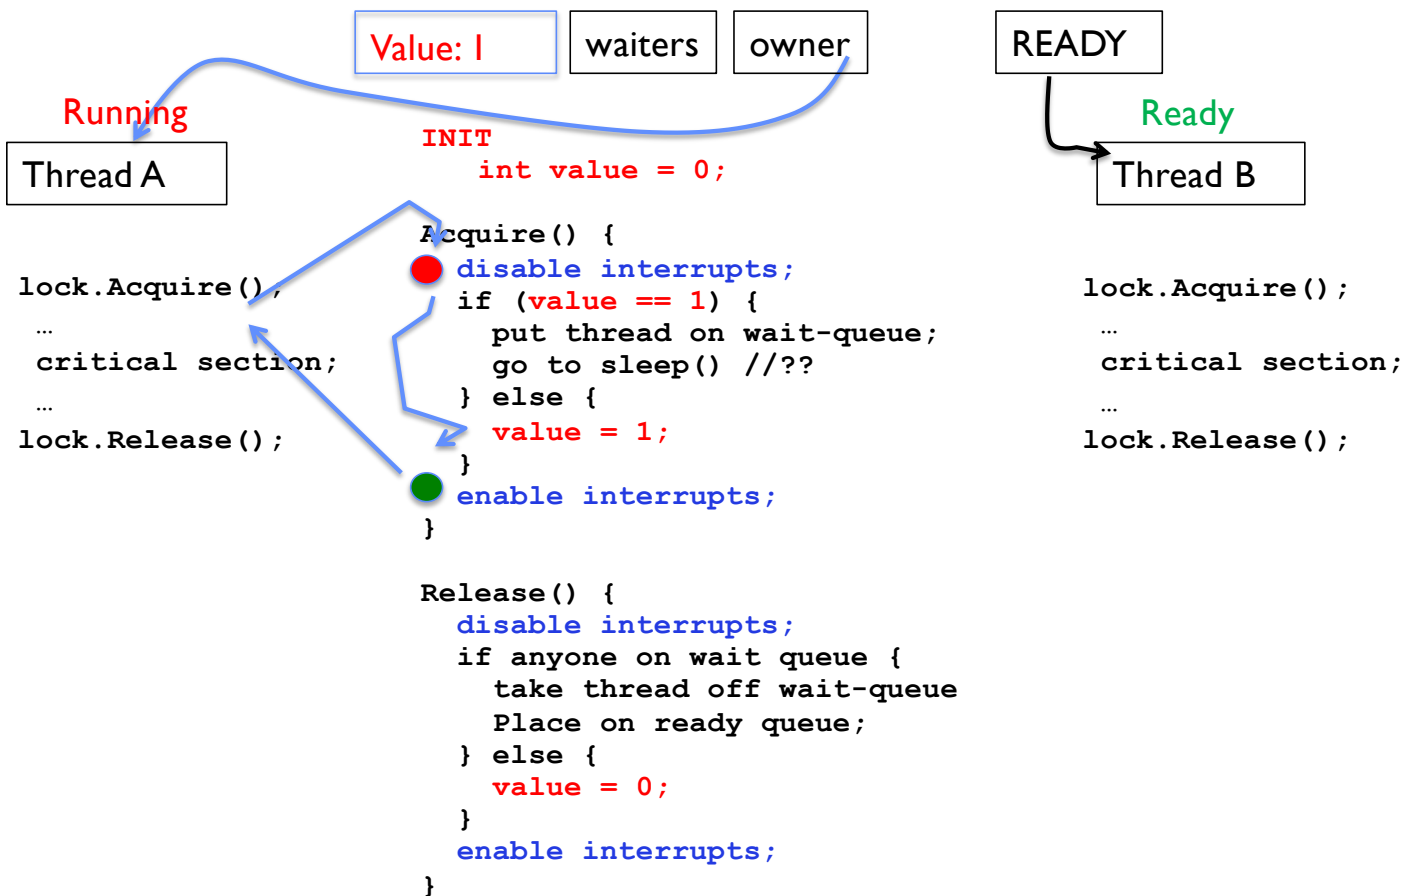


# In-Kernel Lock: Simulation

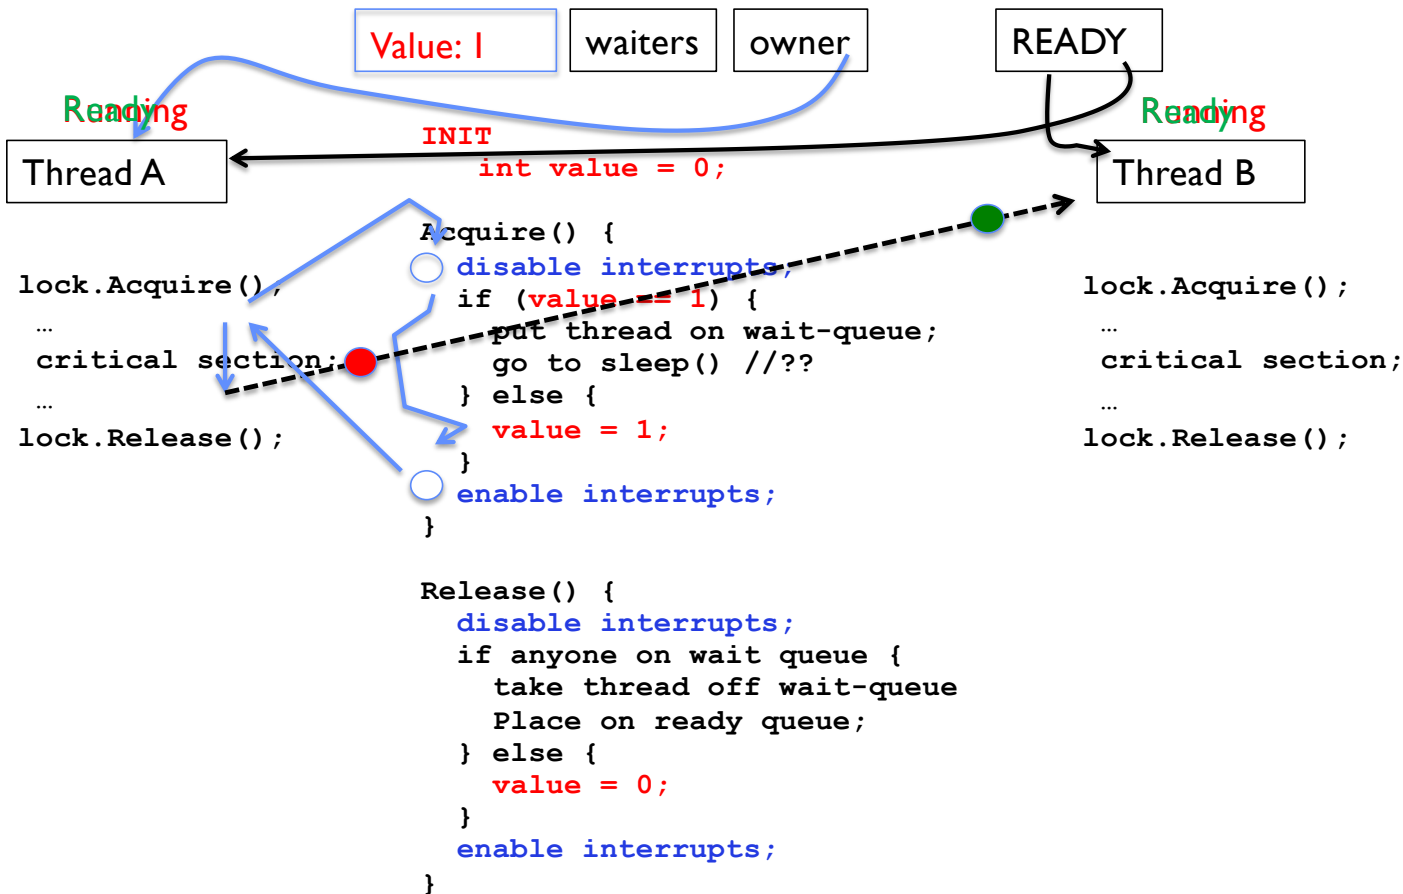




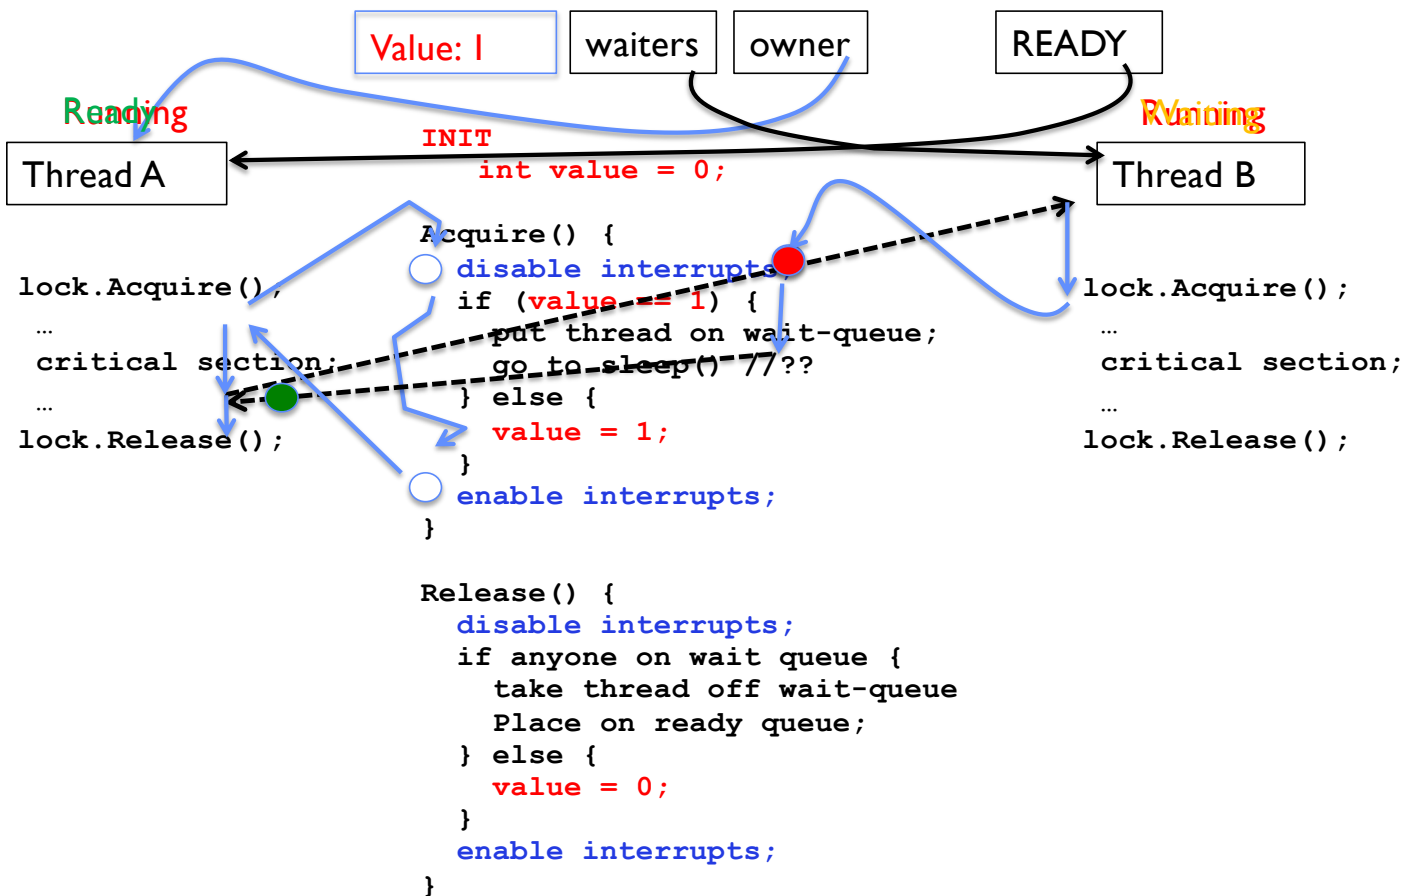
# In-Kernel Lock: Simulation



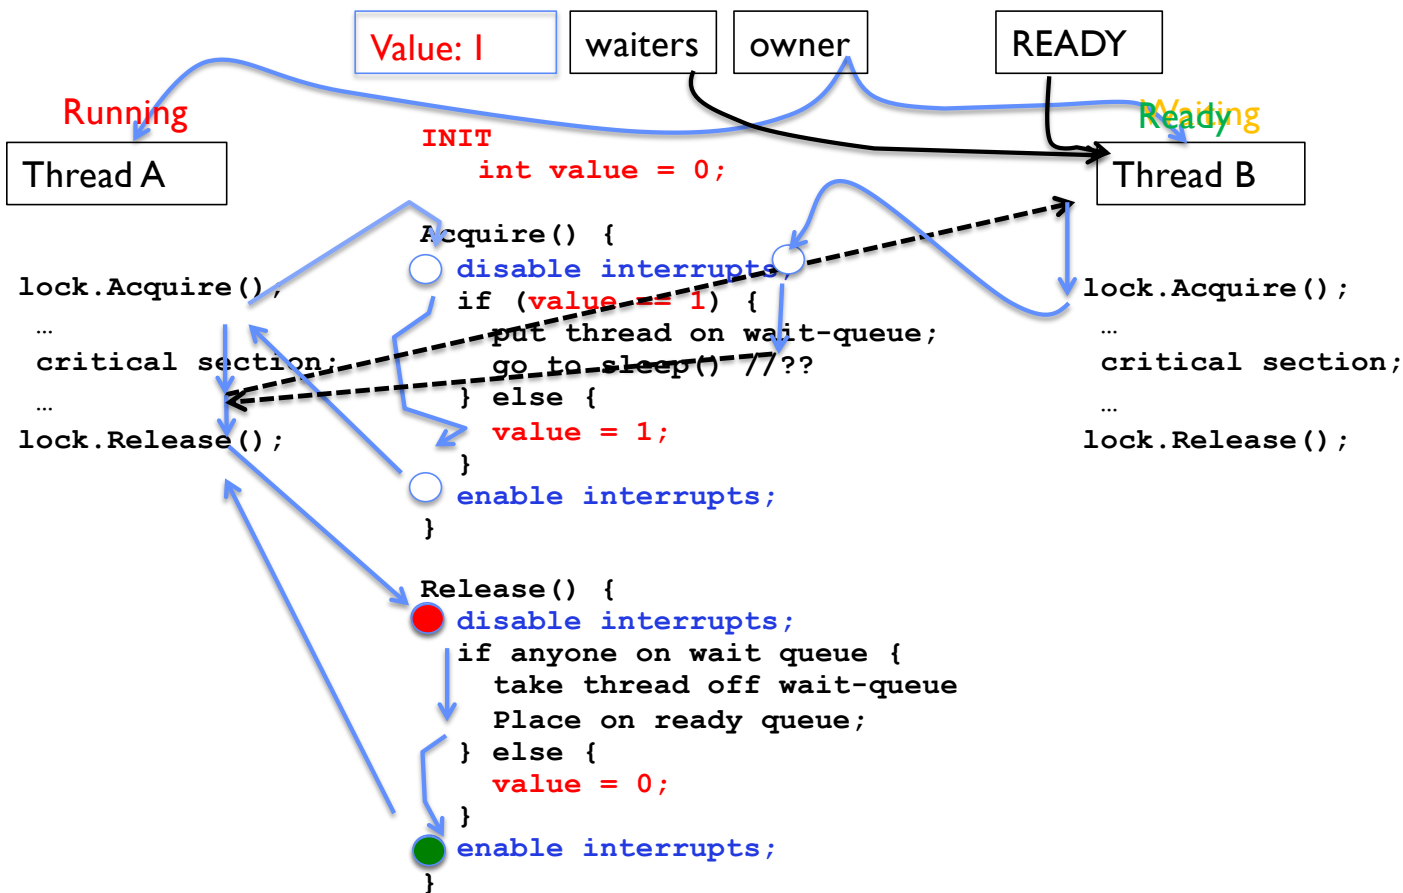
# In-Kernel Lock: Simulation




# In-Kernel Lock: Simulation

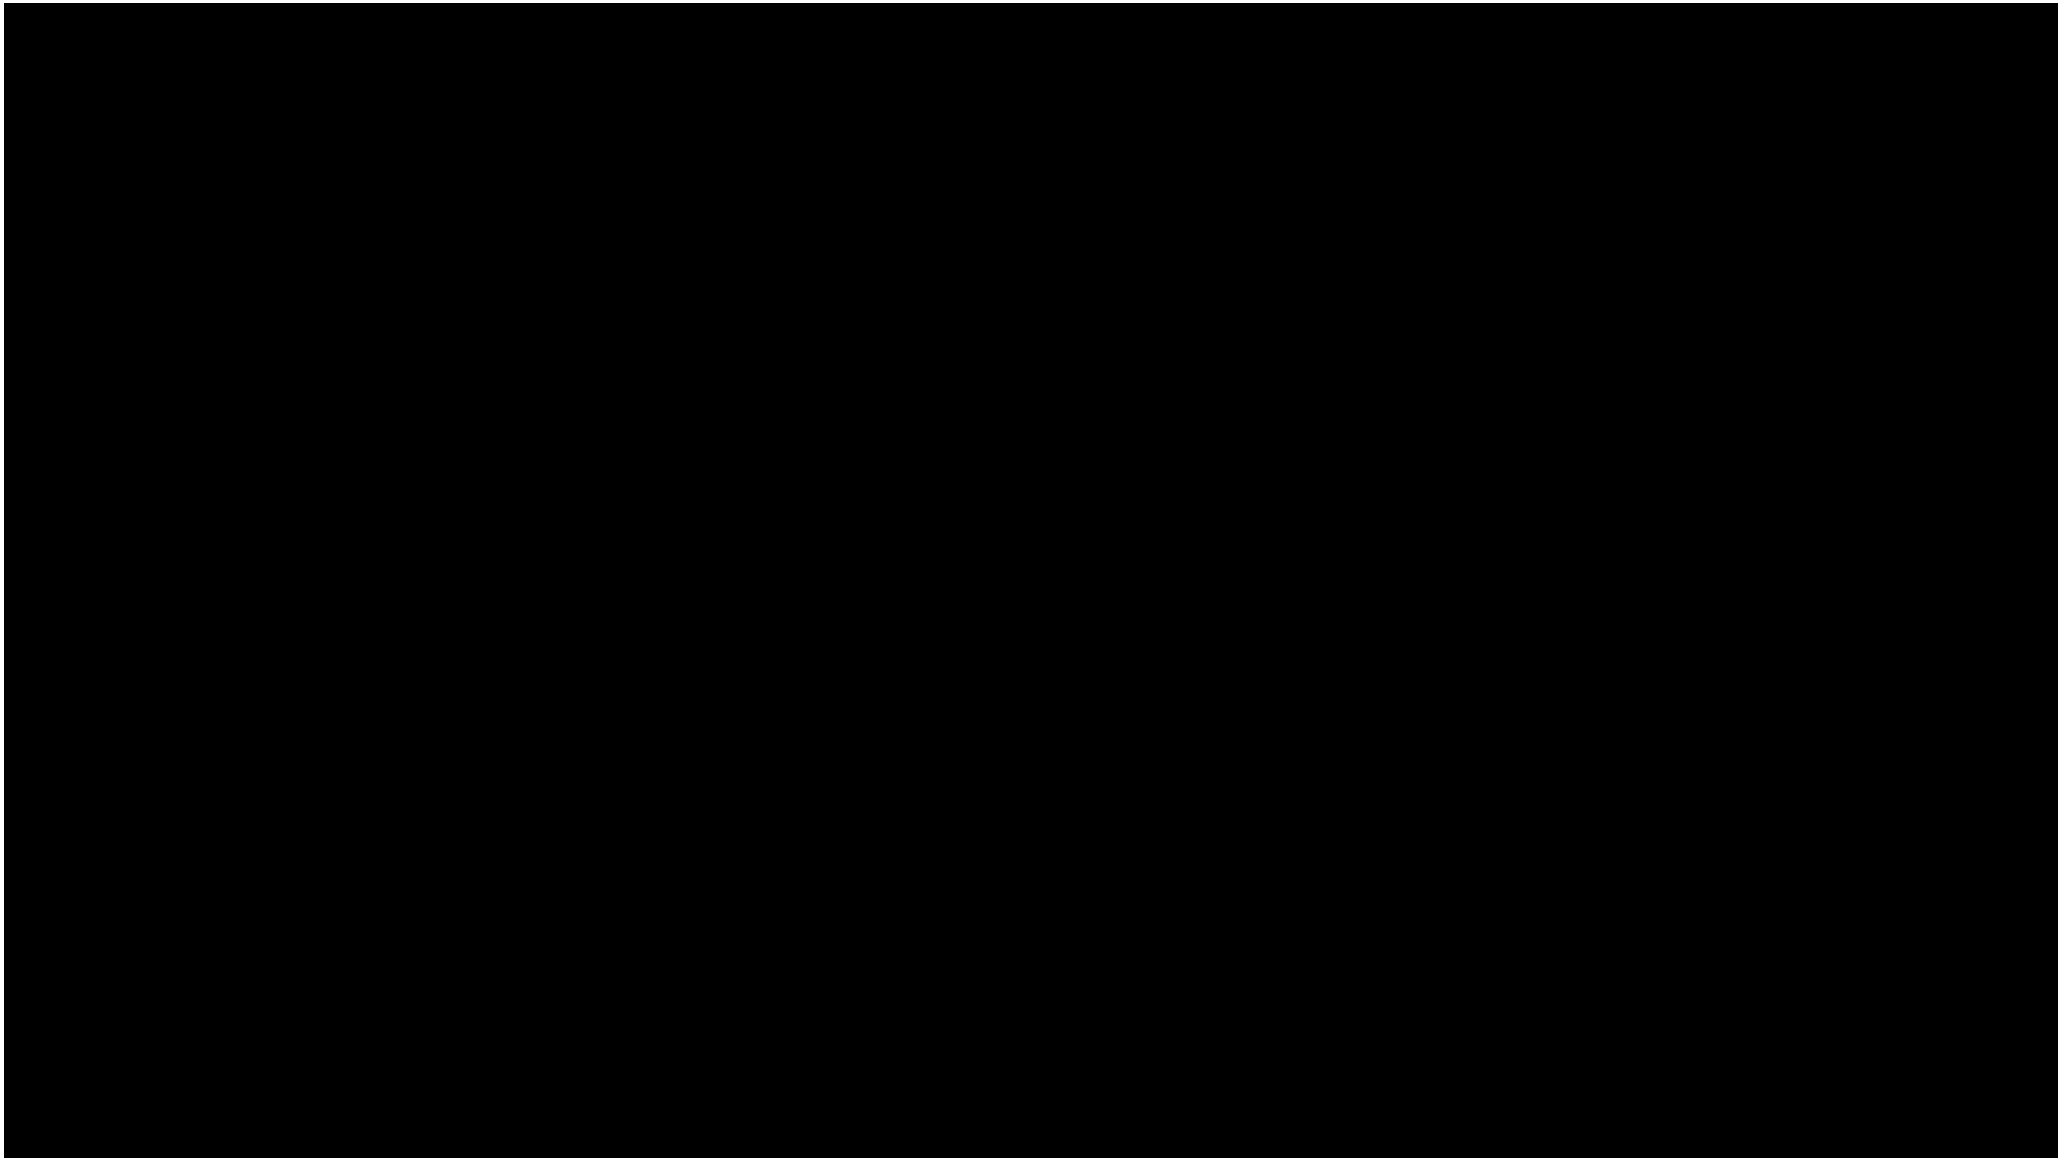


# In-Kernel Lock: Simulation



# In-Kernel Lock: Simulation






## Recall: Multithreaded Server

---

- Bounded pool of worker threads
  - Allocated in **advance**: no thread creation overhead
  - Queue of pending requests




## Simple Performance Model

---

- Given that the overhead of a critical section is  $X$ 
  - User->Kernel Context Switch
  - Acquire Lock
  - Kernel->User Context Switch
  - <perform exclusive work>
  - User->Kernel Context Switch
  - Release Lock
  - Kernel->User Context Switch
- Even if everything else is infinitely fast, with any number of threads and cores
- What is the maximum rate of operations that involve this overhead?



## Highly Contended Case – in a picture



# Back to system performance

## More Practical Motivation

### Back to Jeff Dean's "Numbers everyone should know"

Handle I/O in  
separate thread,  
avoid blocking  
other progress


L1 cache reference	0.5 ns
Branch mispredict	5 ns
L2 cache reference	7 ns
Mutex lock/unlock	25 ns
Main memory reference	100 ns
Compress 1K bytes with Zippy	3,000 ns
Send 2K bytes over 1 Gbps network	20,000 ns
Read 1 MB sequentially from memory	250,000 ns
Round trip within same datacenter	500,000 ns
Disk seek	10,000,000 ns
Read 1 MB sequentially from disk	20,000,000 ns
Send packet CA->Netherlands->CA	150,000,000 ns




- $X = 1\text{ms} \Rightarrow 1,000 \text{ ops/sec}$

## Uncontended Many-Lock Case


---



## Recall: Basic cost of a system call



- Min System call  $\sim 25\times$  cost of function call
- Scheduling could be many times more
- Streamline system processing as much as possible
- Other optimizations seek to process as much of the call in user space as possible (eg, Linux vDSO)



# Atomic Read-Modify-Write Instructions

---

- Problems with previous solution:
  - Can't give lock implementation to users
  - Doesn't work well on multiprocessor
    - » Disabling interrupts on all processors requires messages and would be very time consuming
- Alternative: [atomic instruction sequences](#)
  - These instructions read a value and write a new value atomically
  - Hardware is responsible for implementing this correctly
    - » on both uniprocessors (not too hard)
    - » and multiprocessors (requires help from cache coherence protocol)
  - Unlike disabling interrupts, can be used on both uniprocessors and multiprocessors

# Examples of Read-Modify-Write

---

- `test&set (&address) {` `/* most architectures */`  
    `result = M[address];` `// return result from "address" and`  
    `M[address] = 1;` `// set value at "address" to 1`  
    `return result;`  
}
- `swap (&address, register) {` `/* x86 */`  
    `temp = M[address];` `// swap register's value to`  
    `M[address] = register;` `// value at "address"`  
    `register = temp;`  
}
- `compare&swap (&address, reg1, reg2) {` `/* 68000 */`  
    `if (reg1 == M[address]) {` `// If memory still == reg1,`  
        `M[address] = reg2;` `// then put reg2 => memory`  
        `return success;`  
    `} else {` `// Otherwise do not change memory`  
        `return failure;`  
    `}`  
}
- `load-linked&store-conditional(&address) {` `/* R4000, alpha, ARM, RISC-V */`  
    `loop:`  
        `ll r1, M[address];`  
        `movi r2, 1;` `// Can do arbitrary computation`  
        `sc r2, M[address];`  
        `beqz r2, loop;`  
    `}`

## Conclusion

---

- Important concept: Atomic Operations
  - An operation that runs to completion or not at all
  - These are the primitives on which to construct various synchronization primitives
- Talked about hardware atomicity primitives:
  - Disabling of Interrupts, test&set, swap, compare&swap, load-linked & store-conditional
- Showed several constructions of Locks
  - Must be very careful not to waste/tie up machine resources
    - » Shouldn't disable interrupts for long
    - » Shouldn't spin wait for long
  - Key idea: Separate lock variable, use hardware mechanisms to protect modifications of that variable



