

CS162  
Operating Systems and  
Systems Programming  
Lecture 5

Sockets and IPC (Finished)  
Concurrency: Processes and Threads

February 1<sup>st</sup>, 2022  
Prof. Anthony Joseph and John Kubiatowicz  
<http://cs162.eecs.Berkeley.edu>

## Recall: Key Unix I/O Design Concepts

- **Uniformity – Everything Is a File!**
  - file operations, device I/O, and interprocess communication through open, read/write, close
  - Allows simple composition of programs
    - » find | grep | wc ...
- **Open before use**
  - Provides opportunity for access control and arbitration
  - Sets up the underlying machinery, i.e., data structures
- **Byte-oriented**
  - Even if blocks are transferred, addressing is in bytes
- **Kernel buffered reads**
  - Streaming and block devices looks the same, read blocks yielding processor to other task
- **Kernel buffered writes**
  - Completion of out-going transfer decoupled from the application, allowing it to continue
- **Explicit close**

2/1/2022

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 5.2

## Recall: Low-Level vs High-Level file API

- |   |  |
|---|--|
| <ul style="list-style-type: none"><li>• Low-level direct use of syscall interface:<br/><code>open()</code>, <code>read()</code>, <code>write()</code>, <code>close()</code></li><li>• Opening of file returns file descriptor:<br/><code>int myfile = open(...);</code></li><li>• File descriptor only meaningful to kernel<ul style="list-style-type: none"><li>– Index into process (PDB) which holds pointers to kernel-level structure ("file description") describing file.</li></ul></li><li>• Every <code>read()</code> or <code>write()</code> causes syscall no matter how small (could read a single byte)</li><li>• Consider loop to get 4 bytes at a time using <code>read()</code>:<ul style="list-style-type: none"><li>– Each iteration enters kernel for 4 bytes.</li></ul></li></ul> | <ul style="list-style-type: none"><li>• High-level buffered access:<br/><code>fopen()</code>, <code>fread()</code>, <code>fwrite()</code>, <code>fclose()</code></li><li>• Opening of file returns ptr to FILE:<br/><code>FILE *myfile = fopen(...);</code></li><li>• FILE structure is user space contains:<ul style="list-style-type: none"><li>– a chunk of memory for a buffer</li><li>– the file descriptor for the file (<code>fopen()</code> will call <code>open()</code> automatically)</li></ul></li><li>• Every <code>fread()</code> or <code>fwrite()</code> filters through buffer and may not call <code>read()</code> or <code>write()</code> on every call.</li><li>• Consider loop to get 4 bytes at a time using <code>fread()</code>:<ul style="list-style-type: none"><li>– First call to <code>fread()</code> calls <code>read()</code> for block of bytes (say 1024). Puts in buffer and returns first 4 to user.</li><li>– Subsequent <code>fread()</code> grab bytes from buffer</li></ul></li></ul> |
|---|--|

2/1/2022

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 5.3

## Recall: Low-Level vs. High-Level File API

### Low-Level Operation:

```
ssize_t read(...) {
```

asm code ... syscall # into %eax  
put args into registers %ebx, ...  
*special trap instruction*

Kernel:  
get args from regs  
dispatch to system func  
Do the work to read from the file  
Store return value in %eax

get return values from regs

Return data to caller

```
};
```

### High-Level Operation:

```
ssize_t fread(...) {
```

Check buffer for contents  
Return data to caller if available

asm code ... syscall # into %eax  
put args into registers %ebx, ...  
*special trap instruction*

Kernel:  
get args from regs  
dispatch to system func  
Do the work to read from the file  
Store return value in %eax

get return values from regs

Update buffer with excess data  
Return data to caller

```
};
```

2/1/2022

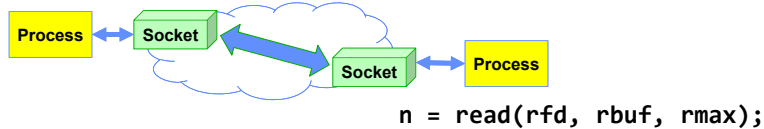
Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 5.4

## Recall: Sockets: An Endpoint for Communication

- **Key Idea:** Communication across the world looks like File I/O

```
write(wfd, wbuf, wlen);
```



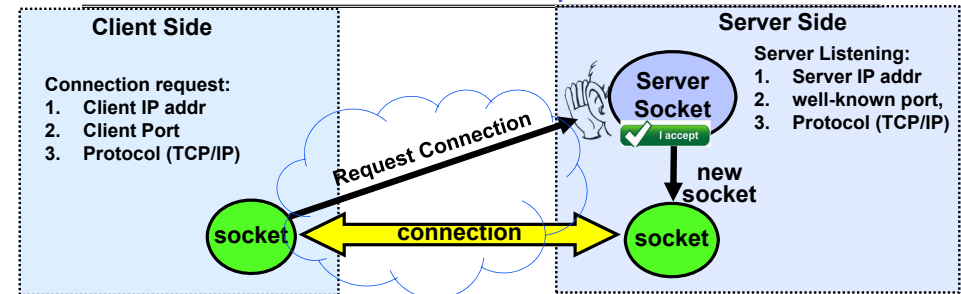
- Sockets: Endpoint for Communication
  - Queues to temporarily hold results
- Connection: Two Sockets Connected Over the network  $\Rightarrow$  IPC over network!
  - How to **open()**?
  - What is the namespace?
  - How are they connected in time?

2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.5

## Recall: Connection Setup over TCP/IP



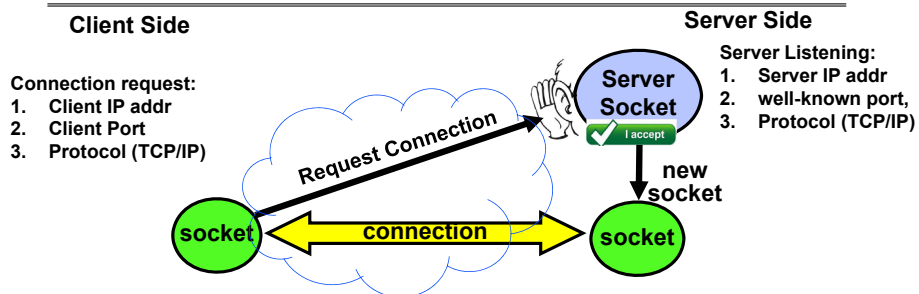
- Special kind of socket: **server socket**
  - Has file descriptor
  - Can't read or write
- Two operations:
  1. **listen():** Start allowing clients to connect
  2. **accept():** Create a *new socket* for a *particular* client

2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.6

## Recall: Connection Setup over TCP/IP



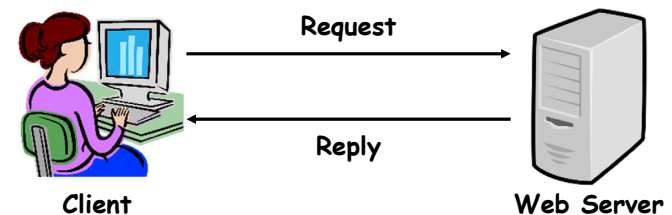
- 5-Tuple identifies each connection:
  1. Source IP Address
  2. Destination IP Address
  3. Source Port Number
  4. Destination Port Number
  5. Protocol (always TCP here)
- Often, Client Port “randomly” assigned
  - Done by OS during client socket setup
- Server Port often “well known”
  - 80 (web), 443 (secure web), 25 (sendmail), etc
  - Well-known ports from 0—1023

2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.7

## Web Server

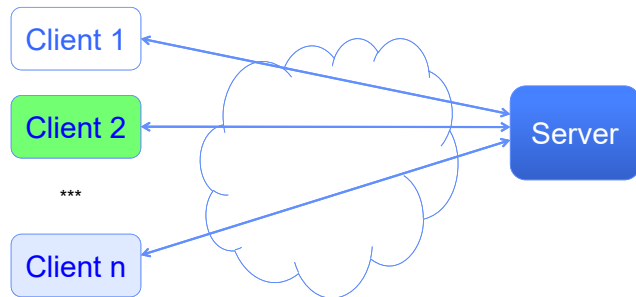


2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.8

## Client-Server Models



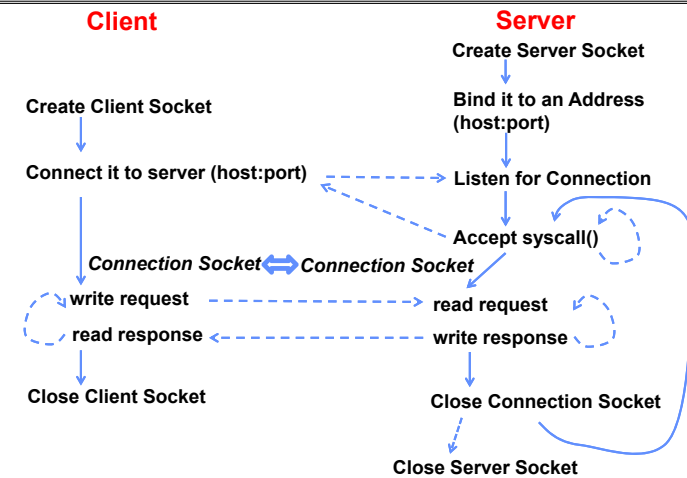
- File servers, web, FTP, Databases, ...
- Many clients accessing a common server

2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.9

## Sockets in concept



2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.10

## Client Protocol

```
char *host_name, *port_name;

// Create a socket
struct addrinfo *server = lookup_host(host_name, port_name);
int sock_fd = socket(server->ai_family, server->ai_socktype,
                    server->ai_protocol);

// Connect to specified host and port
connect(sock_fd, server->ai_addr, server->ai_addrlen);

// Carry out Client-Server protocol
run_client(sock_fd);

/* Clean up on termination */
close(sock_fd);
```

2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.11

## Server Protocol (v1)

```
// Create socket to listen for client connections
char *port_name;
struct addrinfo *server = setup_address(port_name);
int server_socket = socket(server->ai_family,
                          server->ai_socktype, server->ai_protocol);

// Bind socket to specific port
bind(server_socket, server->ai_addr, server->ai_addrlen);
// Start listening for new client connections
listen(server_socket, MAX_QUEUE);

while (1) {
    // Accept a new client connection, obtaining a new socket
    int conn_socket = accept(server_socket, NULL, NULL);
    serve_client(conn_socket);
    close(conn_socket);
}
close(server_socket);
```

2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.12

## How Could the Server Protect Itself?

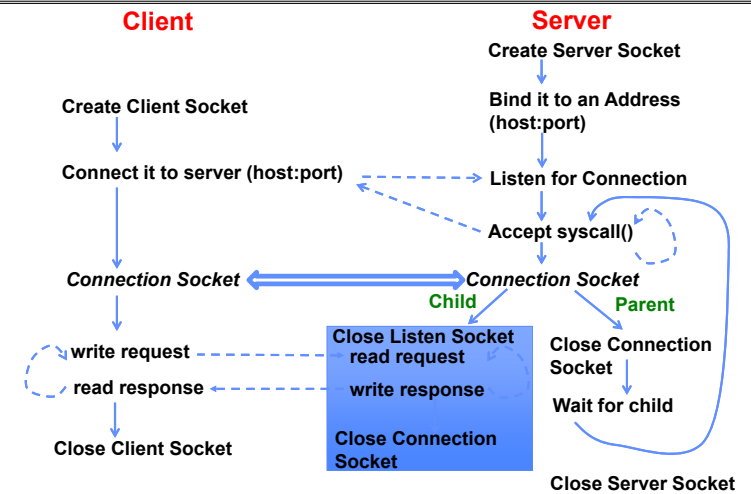
- Handle each connection in a separate process
  - This will mean that the logic serving each request will be “sandboxed” away from the main server process

2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.13

## Sockets With Protection (each connection has own process)



2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.14

## Server Protocol (v2)

```

// Socket setup code elided...
listen(server_socket, MAX_QUEUE);
while (1) {
    // Accept a new client connection, obtaining a new socket
    int conn_socket = accept(server_socket, NULL, NULL);
    pid_t pid = fork();
    if (pid == 0) {
        close(server_socket);
        serve_client(conn_socket);
        close(conn_socket);
        exit(0);
    } else {
        close(conn_socket);
        wait(NULL);
    }
}
close(server_socket);
    
```

2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.15

## Concurrent Server

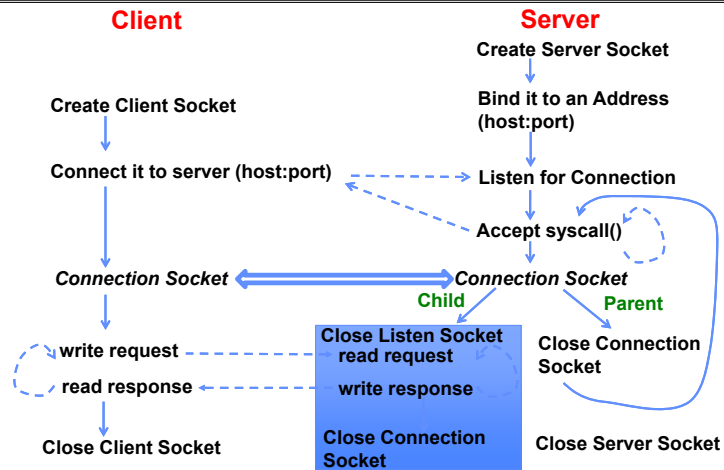
- So far, in the server:
  - Listen will queue requests
  - Buffering present elsewhere
  - But server waits for each connection to terminate before servicing the next
- A concurrent server can handle and service a new connection before the previous client disconnects

2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.16

## Sockets With Protection and Concurrency



2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.17

## Server Protocol (v3)

```

// Socket setup code elided...
listen(server_socket, MAX_QUEUE);
while (1) {
    // Accept a new client connection, obtaining a new socket
    int conn_socket = accept(server_socket, NULL, NULL);
    pid_t pid = fork();
    if (pid == 0) {
        close(server_socket);
        serve_client(conn_socket);
        close(conn_socket);
        exit(0);
    } else {
        close(conn_socket);
        //wait(NULL);
    }
}
close(server_socket);

```

2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.18

## Server Address: Itself

```

struct addrinfo *setup_address(char *port) {
    struct addrinfo *server;
    struct addrinfo hints;
    memset(&hints, 0, sizeof(hints));
    hints.ai_family = AF_UNSPEC;
    hints.ai_socktype = SOCK_STREAM;
    hints.ai_flags = AI_PASSIVE;
    getaddrinfo(NULL, port, &hints, &server);
    return server;
}

```

- Accepts any connections on the specified port

2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.19

## Client: Getting the Server Address

```

struct addrinfo *lookup_host(char *host_name, char *port) {
    struct addrinfo *server;
    struct addrinfo hints;
    memset(&hints, 0, sizeof(hints));
    hints.ai_family = AF_UNSPEC;
    hints.ai_socktype = SOCK_STREAM;

    int rv = getaddrinfo(host_name, port_name,
                        &hints, &server);

    if (rv != 0) {
        printf("getaddrinfo failed: %s\n", gai_strerror(rv));
        return NULL;
    }
    return server;
}

```

2/1/2022

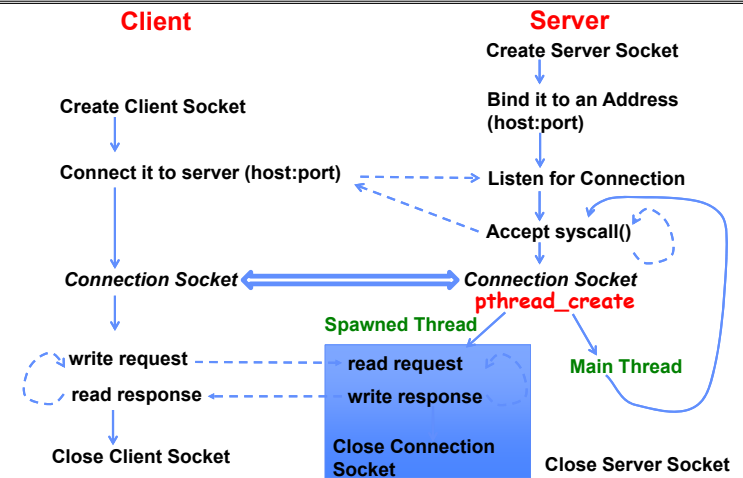
Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.20

## Concurrent Server without Protection

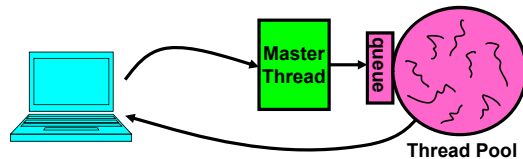
- Spawn a new thread to handle each connection
- Main thread initiates new client connections without waiting for previously spawned threads
- Why give up the protection of separate processes?
  - More efficient to create new threads
  - More efficient to switch between threads

## Sockets with Concurrency, without Protection



## Thread Pools: More Later!

- Problem with previous version: Unbounded Threads
  - When web-site becomes too popular – throughput sinks
- Instead, allocate a bounded “pool” of worker threads, representing the maximum level of multiprogramming



```

master() {
    allocThreads(worker,queue);
    while(TRUE) {
        con=AcceptCon();
        Enqueue(queue,con);
        wakeUp(queue);
    }
}

worker(queue) {
    while(TRUE) {
        con=Dequeue(queue);
        if (con==null)
            sleepOn(queue);
        else
            ServiceWebPage(con);
    }
}

```

## Administrivia

- Kubiatowicz Office Hours:
  - 1-2pm, Tuesday & Wednesday
- Friday was drop deadline. If you forgot to drop, we can't help you!
  - You need to speak with advisor services in your department about how to drop
- Recommendation: Read assigned readings *before* lecture
- Group sign up should have happened already
  - If you don't have 4 members in your group, we will try to find you other partners
  - Want everyone in your group to have the same TA
  - Go to your assigned section on Friday, starting this week!
- Midterm 1 conflicts
  - We will handle these conflicts next week

## Administrivia (Con't)

- Back in person this week!
  - Please be up-to-date with vaccinations and wear masks!
    - » You should have a green pass if you come to class
  - I will be in VLSB 2050 on Tuesday/Thursday 3:30-5:00
    - » Will be trying to get synchronous zoom working. May take a couple of tries to get right
    - » Screen Cast for sure. If I can project it, it will be recorded...
  - We will be trying to make virtual options available for people who are sick
- Start Planning on how your group will collaborate on projects!
  - Meet regularly, in person as regularly as possible
    - » We will have more suggestions on collaborating as term goes on
  - Virtual Interactions: Plan ways of *also* collaborating remotely
    - » Virtual Coffee Hours with your group (with camera)
    - » Regular Brainstorming meetings?
  - Try to meet multiple times a week



2/1/2022

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 5.25

## Computers (Cars/other things) in the news

- Y2K22? January 2022 saw a whole new class of bugs:
  - Well, welcome to Y2K22 bugs. If you write a date/time in YYMMDDHHMM format (which is year, month, day, hour, and minute), it now exceeds 31 bits!
  - Meaning – if they use unsigned instead of signed 32-bit numbers it breaks!
  - So, a bunch of systems are now broken:



Exchange (Email)

SONICWALL™  
Email Security/Firewall



Honda Car Clocks/Navigation Systems

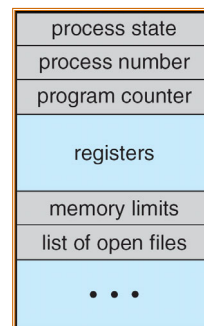
2/1/2022

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 5.26

## Recall: The Process Control Block

- Kernel represents each process as a process control block (PCB)
  - Status (running, ready, blocked, ...)
  - Register state (when not ready)
  - Process ID (PID), User, Executable, Priority, ...
  - Execution time, ...
  - Memory space, translation, ...
- Kernel *Scheduler* maintains a data structure containing the PCBs
  - Give out CPU to different processes
  - This is a Policy Decision
- Give out non-CPU resources
  - Memory/I/O
  - Another policy decision



Process  
Control  
Block

2/1/2022

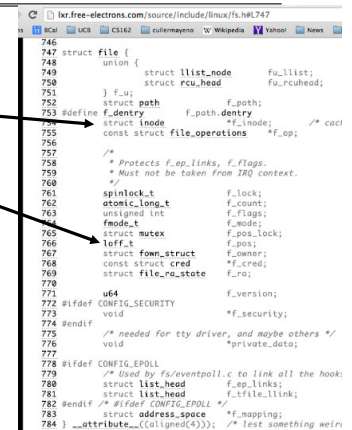
Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 5.27

## Recall: What's in an Open File Description?

For our purposes, the two most important things are:

- Where to find the file data on disk
- The current position within the file

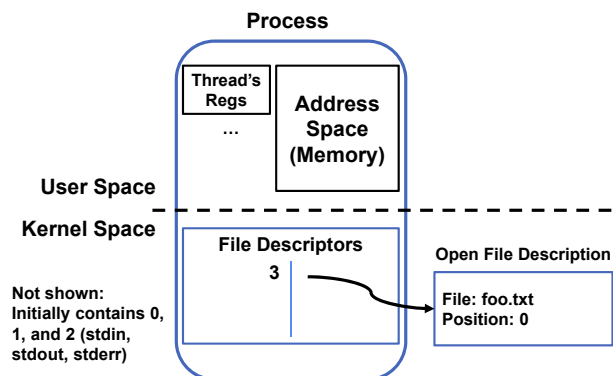


2/1/2022

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 5.28

## Abstract Representation of a Process



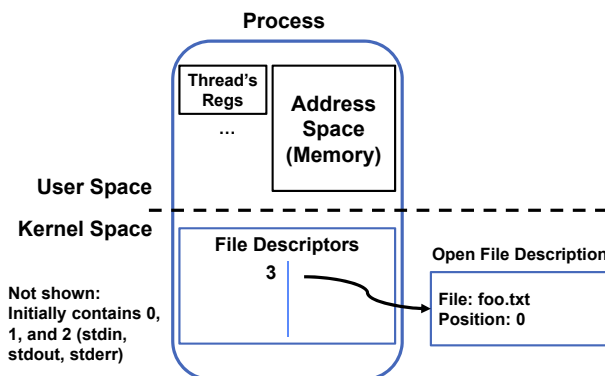
Suppose that we execute  
open("foo.txt")  
and that the result is 3

2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.29

## Abstract Representation of a Process



Suppose that we execute  
open("foo.txt")  
and that the result is 3

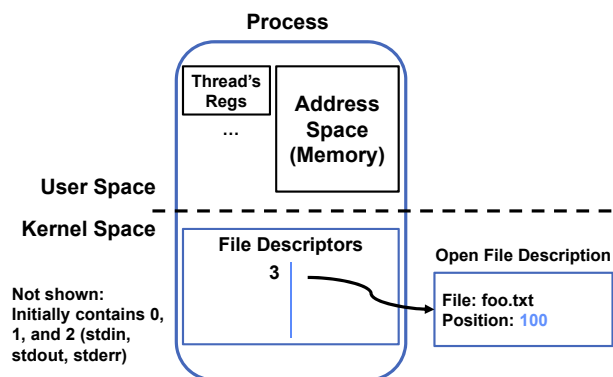
Next, suppose that we execute  
read(3, buf, 100)  
and that the result is 100

2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.30

## Abstract Representation of a Process



Suppose that we execute  
open("foo.txt")  
and that the result is 3

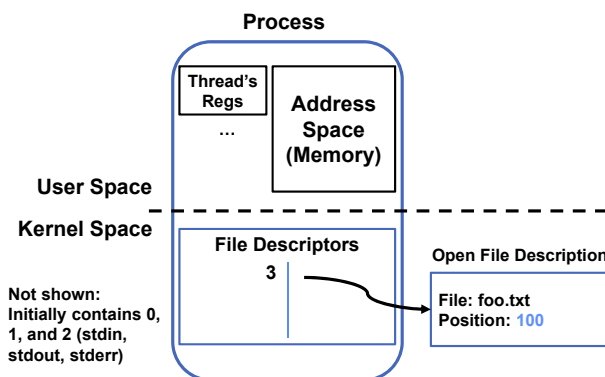
Next, suppose that we execute  
read(3, buf, 100)  
and that the result is 100

2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.31

## Abstract Representation of a Process



Suppose that we execute  
open("foo.txt")  
and that the result is 3

Next, suppose that we execute  
read(3, buf, 100)  
and that the result is 100

Finally, suppose that we  
execute  
close(3)

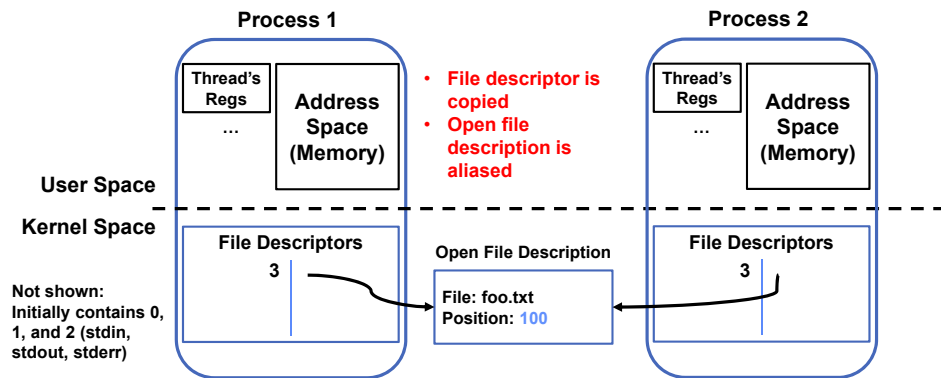
2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.32



## Instead of Closing, let's fork()!

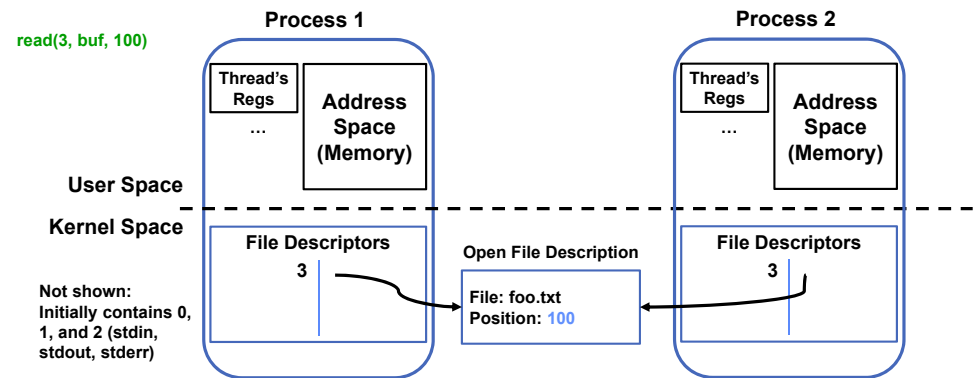


2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.33

## Open File Description is *Aliased*

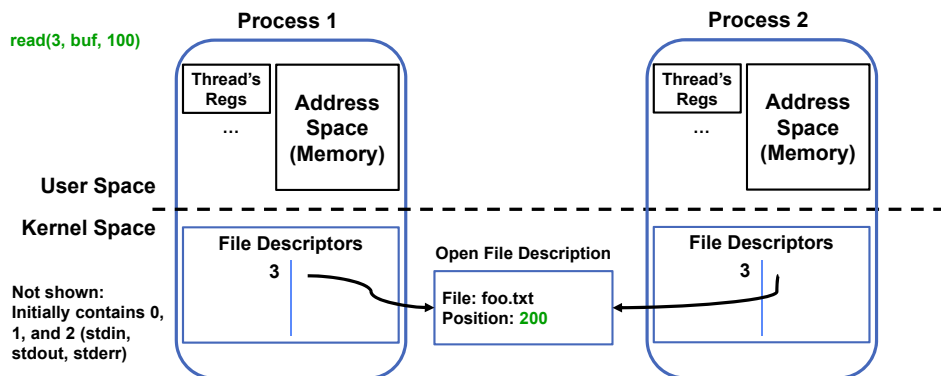


2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.34

## Open File Description is *Aliased*

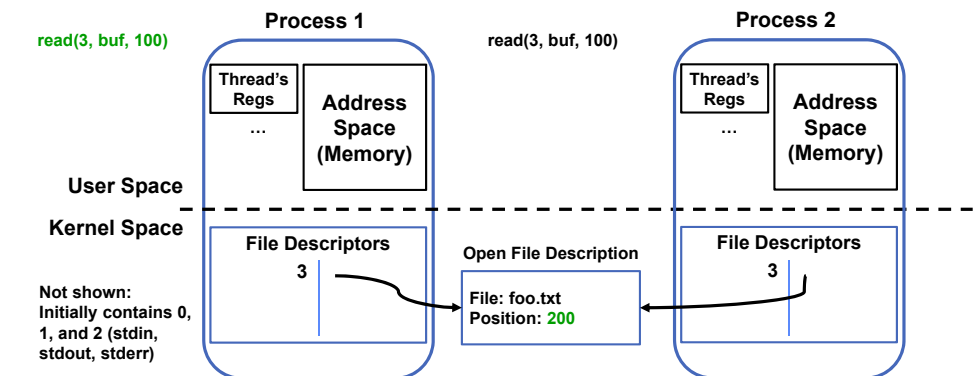


2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.35

## Open File Description is *Aliased*

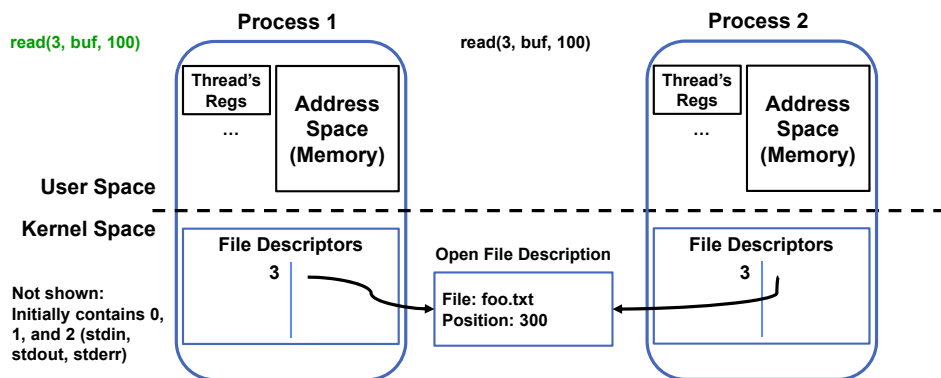


2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.36

## Open File Description is *Aliased*

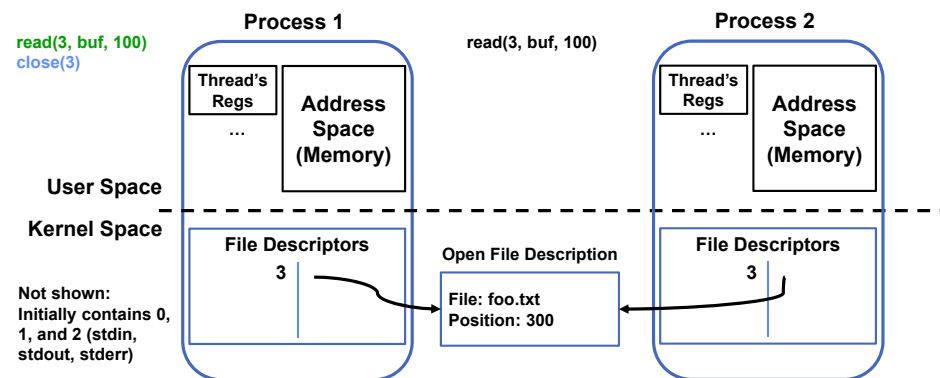


2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.37

## File Descriptor is *Copied*

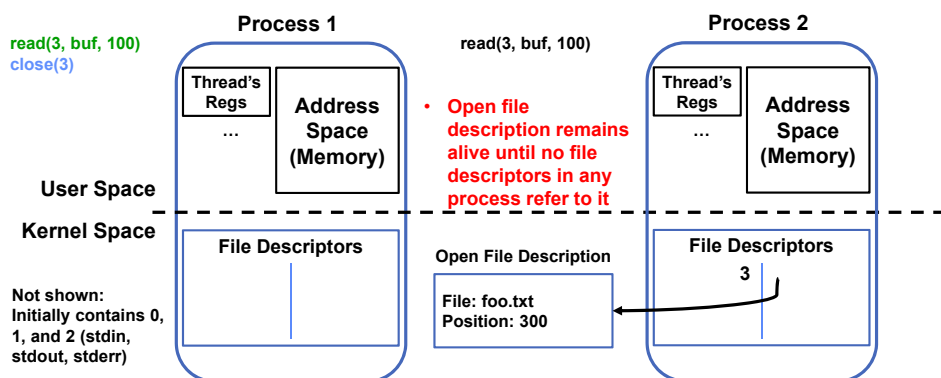


2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.38

## File Descriptor is *Copied*



2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.39

## Why is Aliasing the Open File Description a Good Idea?

- It allows for *shared resources* between processes

2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.40

## Recall: In POSIX, Everything is a “File”

- Identical interface for:
  - Files on disk
  - Devices (terminals, printers, etc.)
  - Regular files on disk
  - Networking (sockets)
  - Local interprocess communication (pipes, sockets)
- Based on the system calls **open()**, **read()**, **write()**, and **close()**

2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.41

## Example: Shared Terminal Emulator

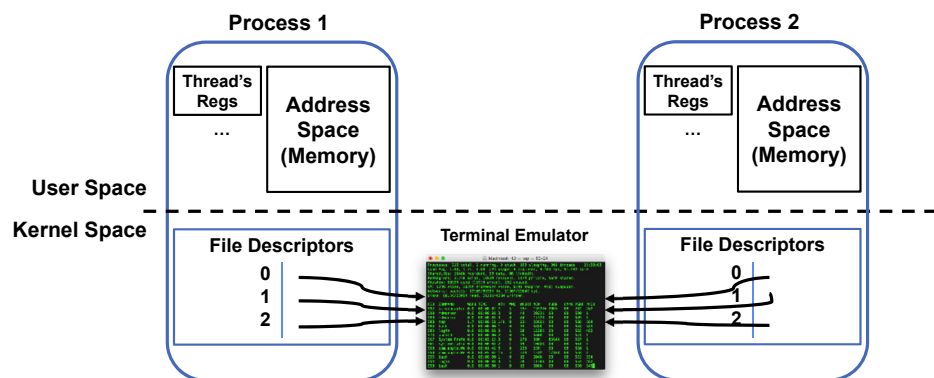
- When you `fork()` a process, the parent's and child's `printf` outputs go to the same terminal

2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.42

## Example: Shared Terminal Emulator

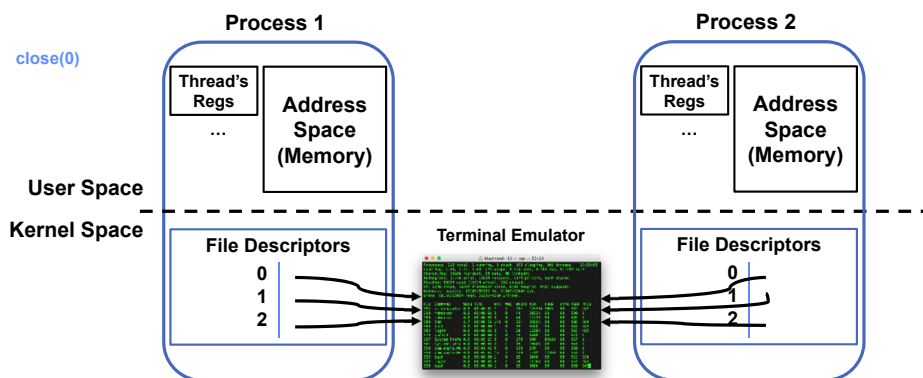


2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.43

## Example: Shared Terminal Emulator

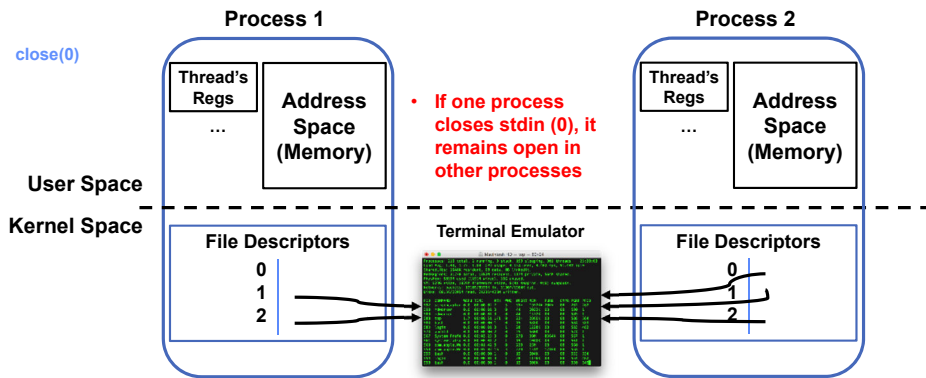


2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.44

## Example: Shared Terminal Emulator



2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.45

## Other Examples

- Shared network connections after `fork()`
  - Allows handling each connection in a separate process
  - We'll explore this next time
- Shared access to pipes
  - Useful for interprocess communication
  - And in writing a shell (Homework 2)

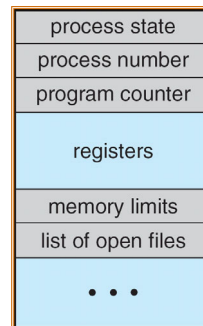
2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.46

## Recall: How do we Multiplex Processes?

- The current state of process held in a process control block (PCB):
  - This is a "snapshot" of the execution and protection environment
  - Only one PCB active at a time
- Give out CPU time to different processes (Scheduling):
  - Only one process "running" at a time
  - Give more time to important processes
- Give pieces of resources to different processes (Protection):
  - Controlled access to non-CPU resources
  - Example mechanisms:
    - » Memory Translation: Give each process their own address space
    - » Kernel/User duality: Arbitrary multiplexing of I/O through system calls



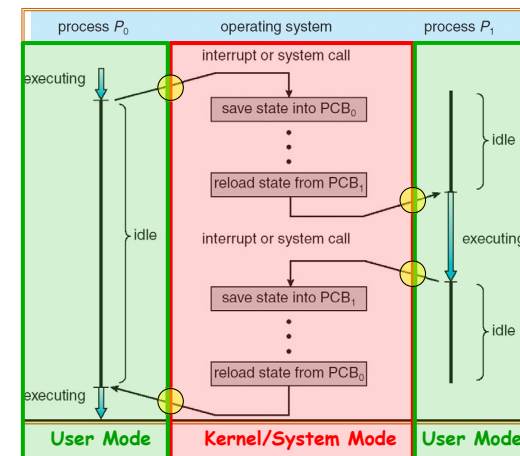
Process Control Block

2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.47

## Recall: CPU Switch From Process A to Process B

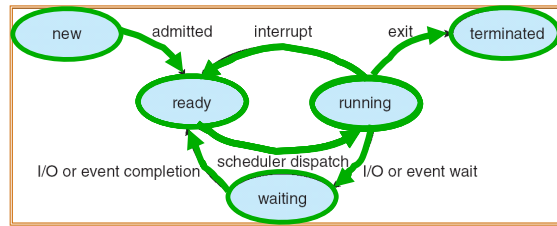


2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.48

## Lifecycle of a Process



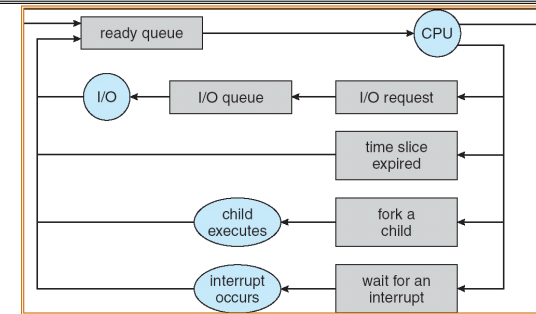
- As a process executes, it changes state:
  - new**: The process is being created
  - ready**: The process is waiting to run
  - running**: Instructions are being executed
  - waiting**: Process waiting for some event to occur
  - terminated**: The process has finished execution

2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.49

## Process Scheduling



- PCBs move from queue to queue as they change state
  - Decisions about which order to remove from queues are **Scheduling** decisions
  - Many algorithms possible (few weeks from now)

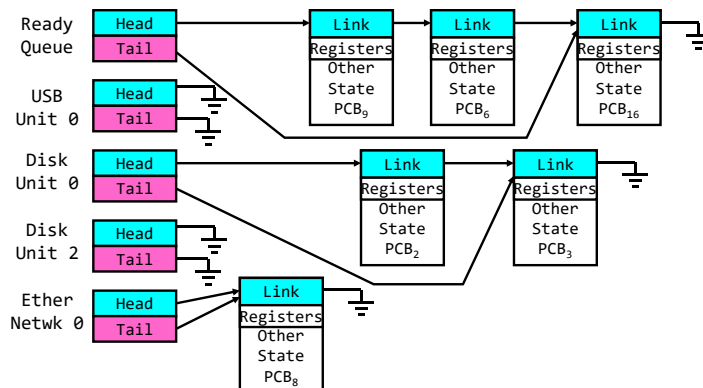
2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.50

## Ready Queue And Various I/O Device Queues

- Process not running  $\Rightarrow$  PCB is in some scheduler queue
  - Separate queue for each device/signal/condition
  - Each queue can have a different scheduler policy



2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.51

## Modern Process with Threads

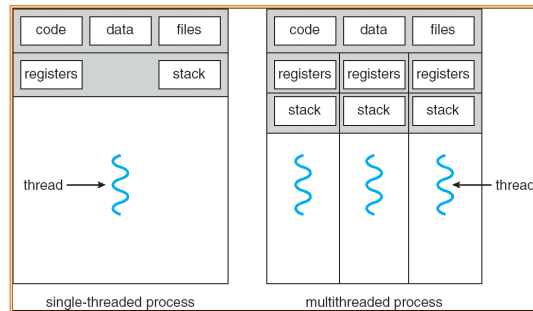
- Thread: *a sequential execution stream within process* (Sometimes called a **“Lightweight process”**)
  - Process still contains a single Address Space
  - No protection between threads
- Multithreading: *a single program made up of a number of different concurrent activities*
  - Sometimes called multitasking, as in Ada ...
- Why separate the concept of a thread from that of a process?
  - Discuss the “thread” part of a process (concurrency)
  - Separate from the “address space” (protection)
  - Heavyweight Process  $\equiv$  Process with one thread

2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.52

## Single and Multithreaded Processes



- Threads encapsulate concurrency: “Active” component
- Address spaces encapsulate protection: “Passive” part
  - Keeps buggy program from trashing the system
- Why have multiple threads per address space?

2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.53

## Thread State

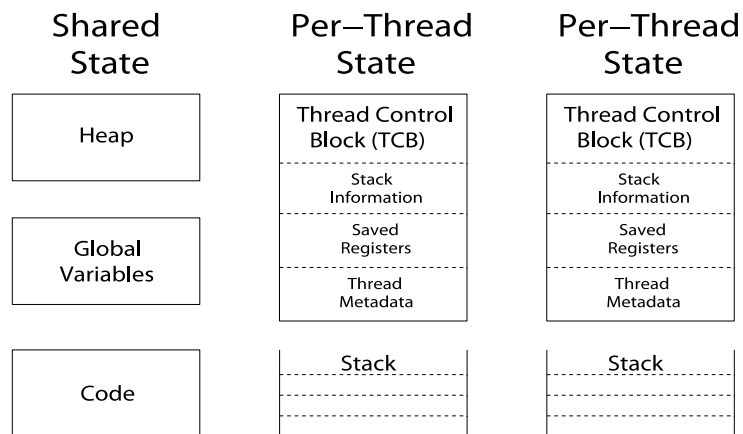
- State shared by all threads in process/address space
  - Content of memory (global variables, heap)
  - I/O state (file descriptors, network connections, etc)
- State “private” to each thread
  - Kept in **TCB** = **Thread Control Block**
  - CPU registers (including, program counter)
  - Execution stack – what is this?
- Execution Stack
  - Parameters, temporary variables
  - Return PCs are kept while called procedures are executing

2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.54

## Shared vs. Per-Thread State



2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.55

## Execution Stack Example

```

A(int tmp) {
A:   if (tmp<2)
A+1:   B();
A+2:   printf(tmp);
      }
      B() {
B:     C();
B+1:   }
      C() {
C:     A(2);
C+1:   }
      A(1);
exit:
    
```

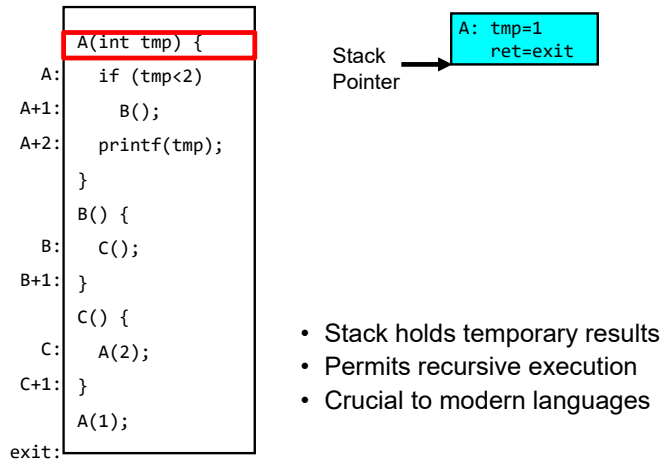
- Stack holds temporary results
- Permits recursive execution
- Crucial to modern languages

2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.56

## Execution Stack Example

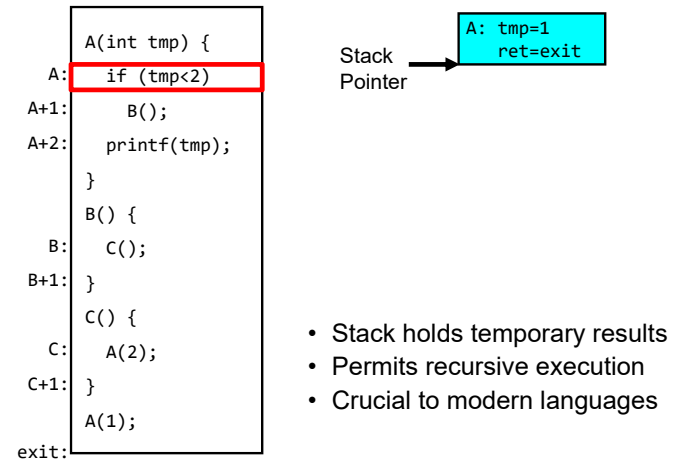


2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.57

## Execution Stack Example

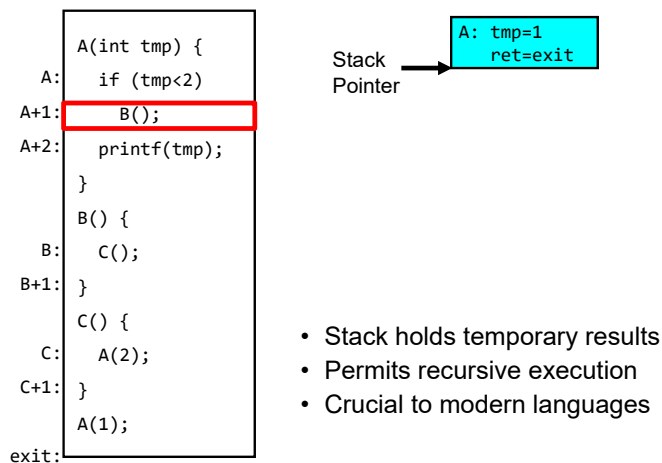


2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.58

## Execution Stack Example

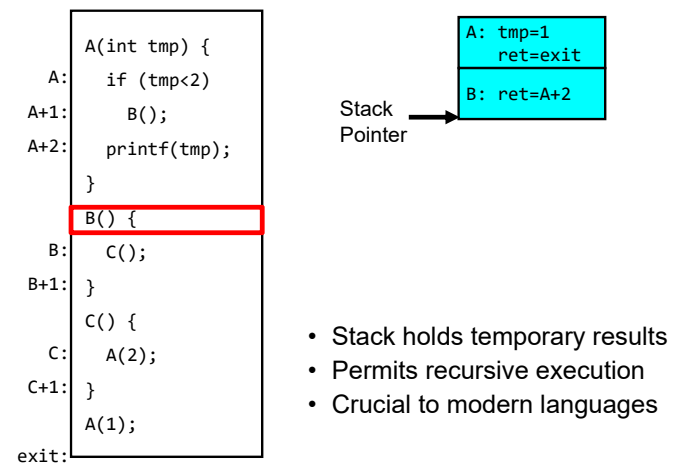


2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.59

## Execution Stack Example

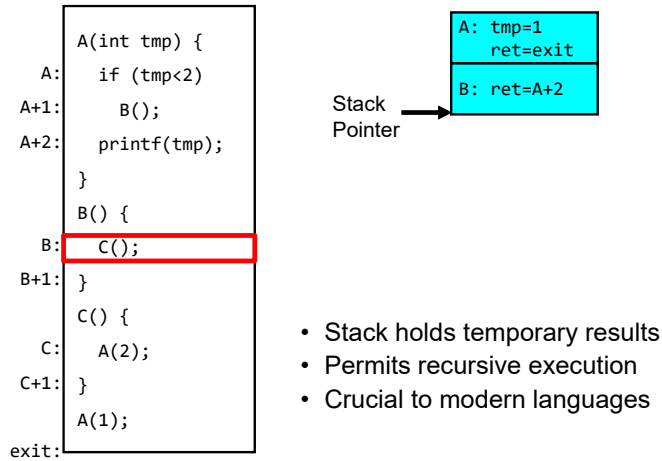


2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.60

## Execution Stack Example

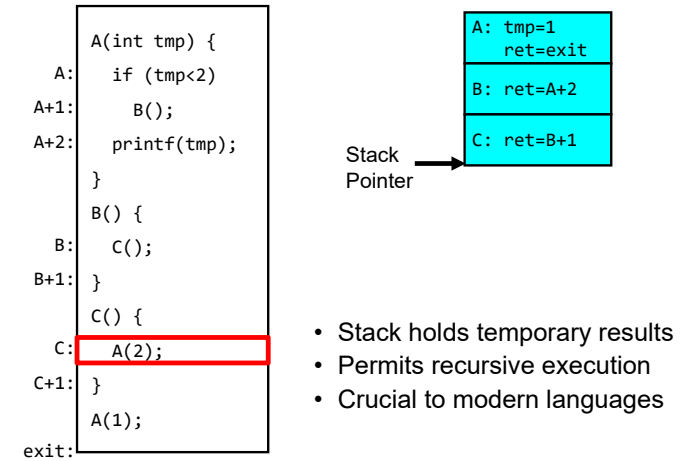


2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.61

## Execution Stack Example

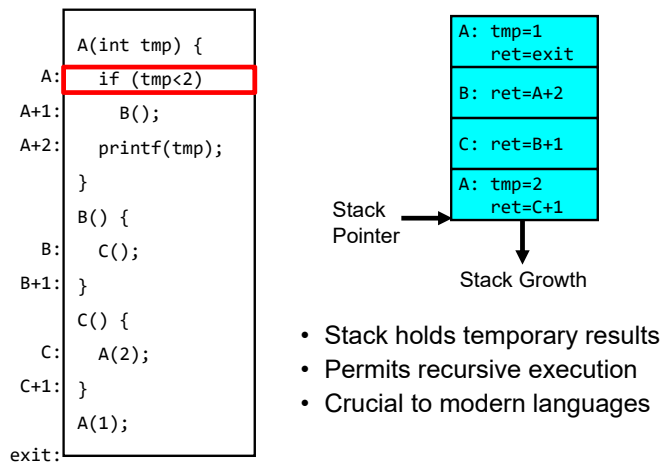


2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.62

## Execution Stack Example

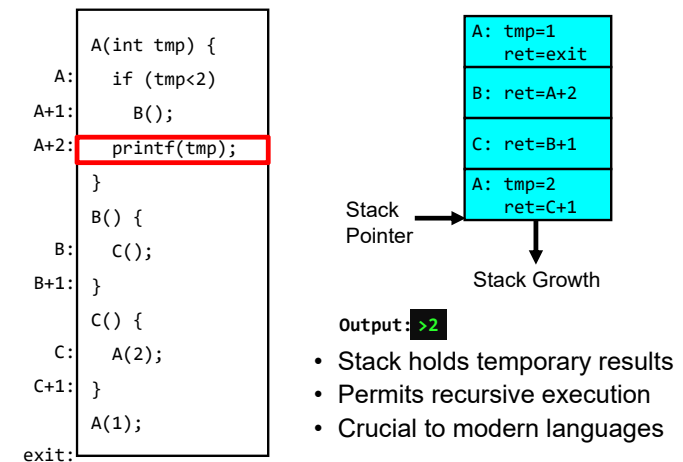


2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.63

## Execution Stack Example



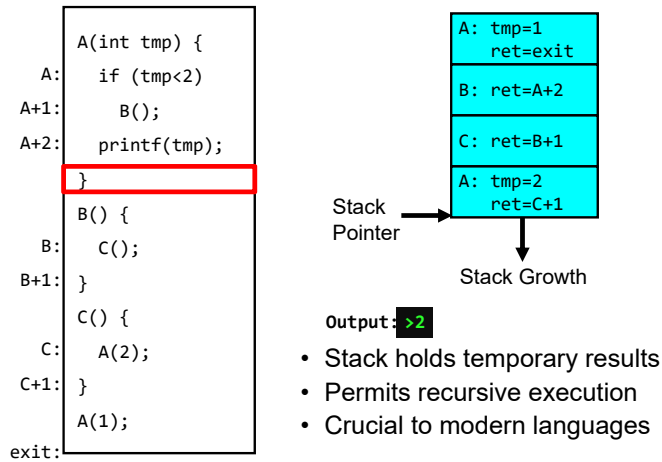
2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.64



## Execution Stack Example

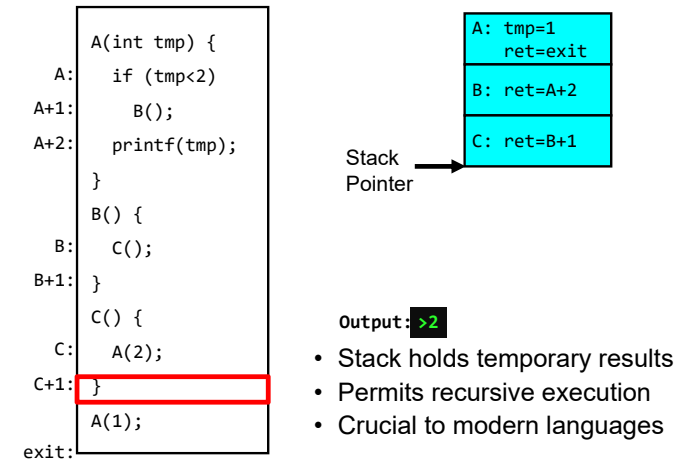


2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.65

## Execution Stack Example

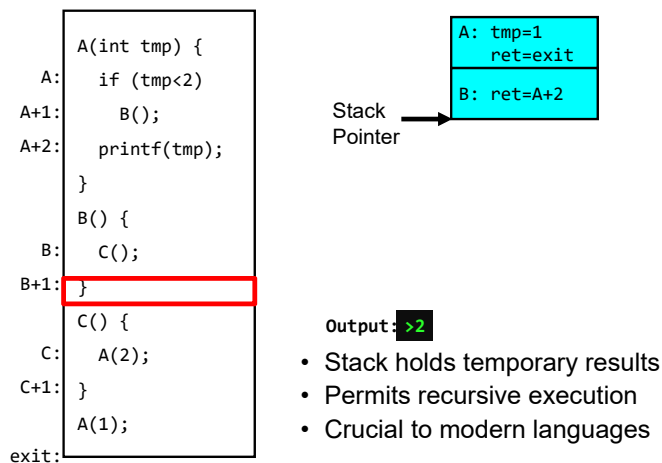


2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.66

## Execution Stack Example

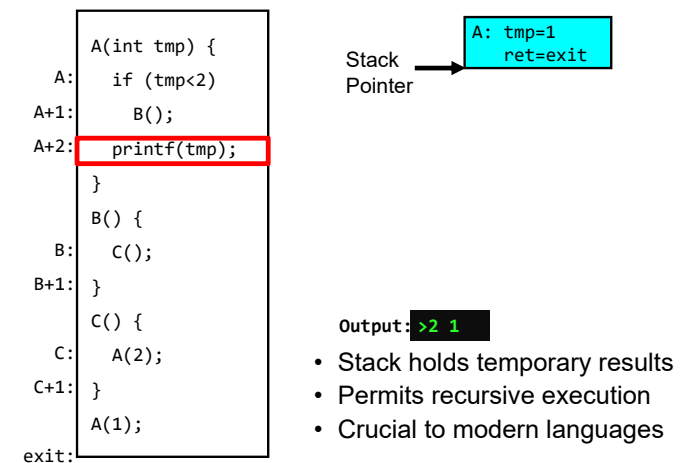


2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.67

## Execution Stack Example

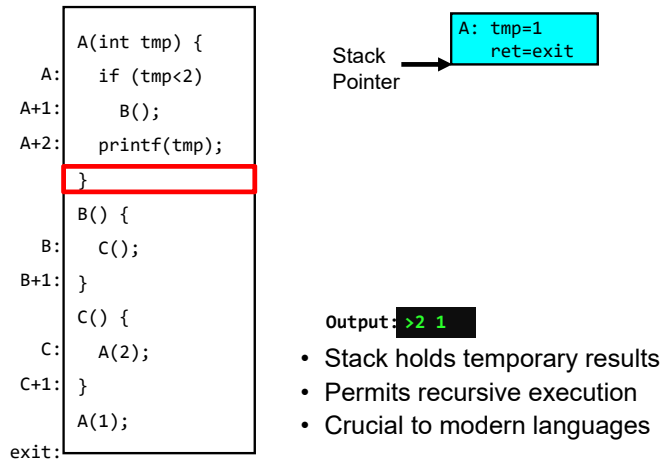


2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.68

## Execution Stack Example

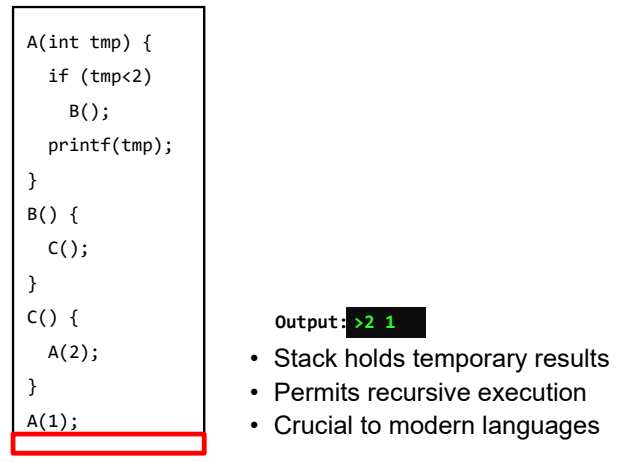


2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.69

## Execution Stack Example

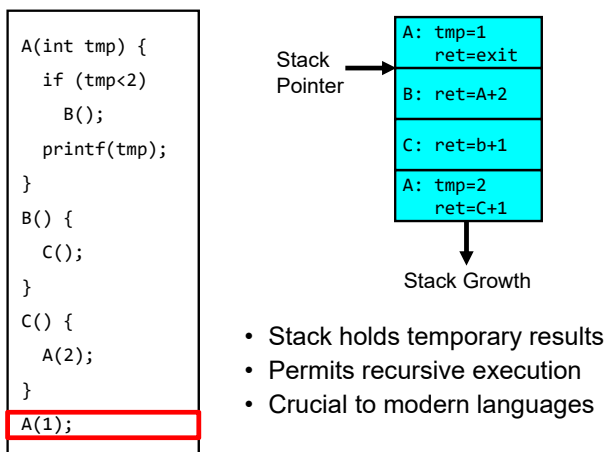


2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.70

## Execution Stack Example



2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.71

## Motivational Example for Threads

- Imagine the following C program:

```
main() {  
    ComputePI("pi.txt");  
    PrintClassList("classlist.txt");  
}
```

- What is the behavior here?
  - Program would never print out class list
  - Why? ComputePI would never finish

2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

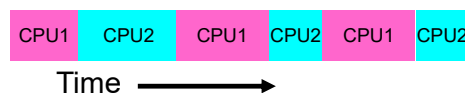
Lec 5.72

## Use of Threads

- Version of program with Threads (loose syntax):

```
main() {  
    ThreadFork(ComputePI, "pi.txt" );  
    ThreadFork(PrintClassList, "classlist.txt");  
}
```

- What does ThreadFork() do?
  - Start independent thread running given procedure
- What is the behavior here?
  - Now, you would actually see the class list
  - This *should* behave as if there are two separate CPUs



2/1/2022

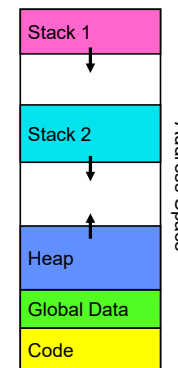
Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.73

## Memory Footprint: Two-Threads

- If we stopped this program and examined it with a debugger, we would see
  - Two sets of CPU registers
  - Two sets of Stacks

- Questions:
  - How do we position stacks relative to each other?
  - What maximum size should we choose for the stacks?
  - What happens if threads violate this?
  - How might you catch violations?



2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.74

## OS Library API for Threads: *pthread*s

**pthread**s: POSIX standard for thread programming  
[POSIX.1c, Threads extensions (IEEE Std 1003.1c-1995)]

```
int pthread_create(pthread_t *thread, const pthread_attr_t *attr,  
                  void *(*start_routine)(void*), void *arg);  
    - thread is created executing start_routine with arg as its sole argument.  
    - return is implicit call to pthread_exit  
  
void pthread_exit(void *value_ptr);  
    - terminates the thread and makes value_ptr available to any successful join  
  
int pthread_yield();  
    - causes the calling thread to yield the CPU to other threads  
  
int pthread_join(pthread_t thread, void **value_ptr);  
    - suspends execution of the calling thread until the target thread terminates.  
    - On return with a non-NULL value_ptr the value passed to pthread_exit() by the  
      terminating thread is made available in the location referenced by value_ptr.
```

prompt% man pthread  
<https://pubs.opengroup.org/onlinepubs/7908799/xsh/pthread.h.html>

2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.75

## Dispatch Loop

- Conceptually, the dispatching loop of the operating system looks as follows:

```
Loop {  
    RunThread();  
    ChooseNextThread();  
    SaveStateOfCPU(curTCB);  
    LoadStateOfCPU(newTCB);  
}
```

- This is an *infinite* loop
  - One could argue that this is all that the OS does
- Should we ever exit this loop???
  - When would that be?

2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.76

## Running a thread

Consider first portion: `RunThread()`

- How do I run a thread?
  - Load its state (registers, PC, stack pointer) into CPU
  - Load environment (virtual memory space, etc)
  - Jump to the PC
- How does the dispatcher get control back?
  - Internal events: thread returns control voluntarily
  - External events: thread gets *preempted*

2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.77

## Internal Events

- Blocking on I/O
  - The act of requesting I/O implicitly yields the CPU
- Waiting on a “signal” from other thread
  - Thread asks to wait and thus yields the CPU
- Thread executes a `yield()`
  - Thread volunteers to give up CPU

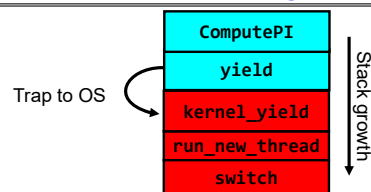
```
computePI() {  
    while(TRUE) {  
        ComputeNextDigit();  
        yield();  
    }  
}
```

2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.78

## Stack for Yielding Thread



- How do we run a new thread?

```
run_new_thread() {  
    newThread = PickNewThread();  
    switch(curThread, newThread);  
    ThreadHouseKeeping(); /* Do any cleanup */  
}
```
- How does dispatcher switch to a new thread?
  - Save anything next thread may trash: PC, regs, stack pointer
  - Maintain isolation for each thread

2/1/2022

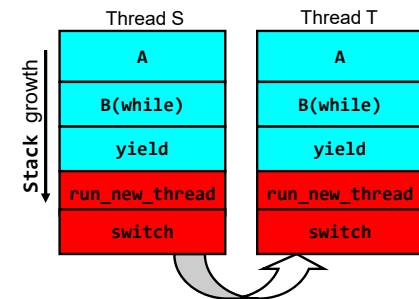
Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.79

## What Do the Stacks Look Like?

- Consider the following code blocks:

```
proc A() {  
    B();  
}  
  
proc B() {  
    while(TRUE) {  
        yield();  
    }  
}
```



- Suppose we have 2 threads:
  - Threads S and T

2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.80

## Saving/Restoring state (often called “Context Switch”)

```
Switch(tCur,tNew) {
    /* Unload old thread */
    TCB[tCur].regs.r7 = CPU.r7;
    ...
    TCB[tCur].regs.r0 = CPU.r0;
    TCB[tCur].regs.sp = CPU.sp;
    TCB[tCur].regs.retpc = CPU.retpc; /*return addr*/

    /* Load and execute new thread */
    CPU.r7 = TCB[tNew].regs.r7;
    ...
    CPU.r0 = TCB[tNew].regs.r0;
    CPU.sp = TCB[tNew].regs.sp;
    CPU.retpc = TCB[tNew].regs.retpc;
    return; /* Return to CPU.retpc */
}
```

2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.81

## Switch Details (continued)

- What if you make a mistake in implementing switch?
  - Suppose you forget to save/restore register 32
  - Get intermittent failures depending on when context switch occurred and whether new thread uses register 32
  - System will give wrong result without warning
- Can you devise an exhaustive test to test switch code?
  - No! Too many combinations and inter-leavings
- Cautionary tale:
  - For speed, Topaz kernel saved one instruction in switch()
  - Carefully documented! Only works as long as kernel size < 1MB
  - What happened?
    - » Time passed, People forgot
    - » Later, they added features to kernel (no one removes features!)
    - » Very weird behavior started happening
  - Moral of story: Design for simplicity

2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.82

## Conclusion

- Socket: an abstraction of a network I/O queue (IPC mechanism)
- Processes have two parts
  - One or more Threads (Concurrency)
  - Address Spaces (Protection)
- Concurrency accomplished by multiplexing CPU Time:
  - Unloading current thread (PC, registers)
  - Loading new thread (PC, registers)
  - Such context switching may be voluntary (yield(), I/O operations) or involuntary (timer, other interrupts)

2/1/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 5.83