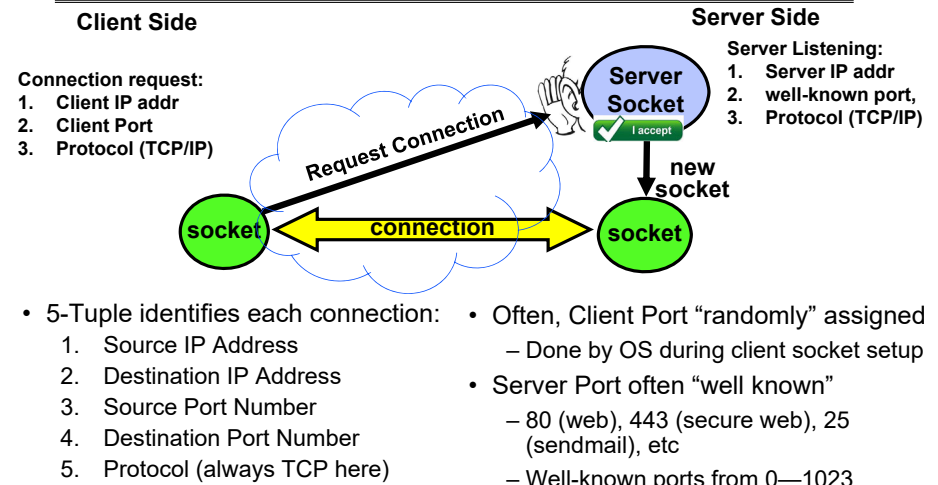


CS162
Operating Systems and
Systems Programming
Lecture 6

Synchronization 1: Concurrency
and Mutual Exclusion

February 3rd, 2022
Prof. Anthony Joseph and John Kubiawicz
<http://cs162.eecs.Berkeley.edu>

Recall: Connection Setup over TCP/IP



2/3/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 6.2

Recall: Server Protocol (v3)

```
// Socket setup code elided...
listen(server_socket, MAX_QUEUE);
while (1) {
    // Accept a new client connection, obtaining a new socket
    int conn_socket = accept(server_socket, NULL, NULL);
    pid_t pid = fork();
    if (pid == 0) {
        close(server_socket);
        serve_client(conn_socket);
        close(conn_socket);
        exit(0);
    } else {
        close(conn_socket);
        // wait(NULL);
    }
}
close(server_socket);
```


2/3/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 6.3

Recall: Multiplexing Processes: The Process Control Block

- Kernel represents each process as a process control block (PCB)
 - Status (running, ready, blocked, ...)
 - Register state (when not ready)
 - Process ID (PID), User, Executable, Priority, ...
 - Execution time, ...
 - Memory space, translation, ...
- Kernel *Scheduler* maintains a data structure containing the PCBs
 - Give out CPU to different processes
 - This is a Policy Decision
- Give out non-CPU resources
 - Memory/I/O
 - Another policy decision




Process
Control
Block

2/3/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 6.4


Recall: CPU Switch From Process A to Process B



Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 6.5

Recall: Lifecycle of a Process or Thread

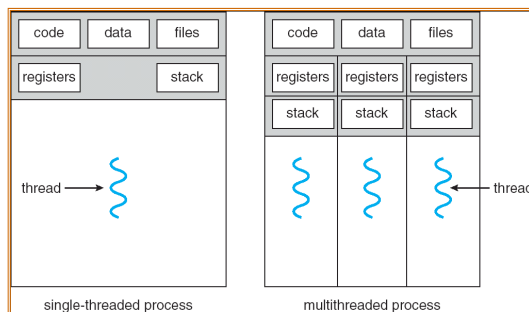


- As a process executes, it changes state:
 - new**: The process/thread is being created
 - ready**: The process is waiting to run
 - running**: Instructions are being executed
 - waiting**: Process waiting for some event to occur
 - terminated**: The process has finished execution

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 6.6

Recall: Single and Multithreaded Processes

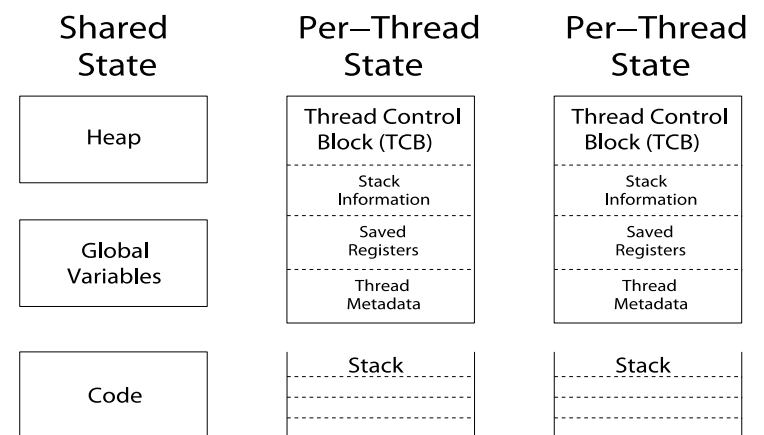


- Threads encapsulate concurrency: "Active" component
- Address spaces encapsulate protection: "Passive" part
 - Keeps buggy program from trashing the system
- Why have multiple threads per address space?

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 6.7

Recall: Shared vs. Per-Thread State



Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 6.8

The Core of Concurrency: the Dispatch Loop

- Conceptually, the scheduling loop of the operating system looks as follows:

```
Loop {
    RunThread();
    ChooseNextThread();
    SaveStateOfCPU(curTCB);
    LoadStateOfCPU(newTCB);
}
```

- This is an *infinite* loop
 - One could argue that this is all that the OS does
- Should we ever exit this loop???
 - When would that be?

2/3/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 6.9

Running a thread

Consider first portion: RunThread()

- How do I run a thread?
 - Load its state (registers, PC, stack pointer) into CPU
 - Load environment (virtual memory space, etc)
 - Jump to the PC
- How does the dispatcher get control back?
 - Internal events: thread returns control voluntarily
 - External events: thread gets *preempted*

2/3/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 6.10

Internal Events

- Blocking on I/O
 - The act of requesting I/O implicitly yields the CPU
- Waiting on a “signal” from other thread
 - Thread asks to wait and thus yields the CPU
- Thread executes a yield()
 - Thread volunteers to give up CPU


```
computePI() {
    while(TRUE) {
        ComputeNextDigit();
        yield();
    }
}
```

2/3/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 6.11

Stack for Yielding Thread



- How do we run a new thread?

```
run_new_thread() {
    newThread = PickNewThread();
    switch(curThread, newThread);
    ThreadHouseKeeping(); /* Do any cleanup */
}
```

- How does dispatcher switch to a new thread?
 - Save anything next thread may trash: PC, regs, stack pointer
 - Maintain isolation for each thread

2/3/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022


Lec 6.12

What Do the Stacks Look Like?

- Consider the following code blocks:

```
proc A() {  
    B();  
}  
  
proc B() {  
    while(TRUE) {  
        yield();  
    }  
}
```

- Suppose we have 2 threads:
 - Threads S and T



Saving/Restoring state (often called "Context Switch")

```
Switch(tCur, tNew) {  
    /* Unload old thread */  
    TCB[tCur].regs.r7 = CPU.r7;  
    ...  
    TCB[tCur].regs.r0 = CPU.r0;  
    TCB[tCur].regs.sp = CPU.sp;  
    TCB[tCur].regs.retpc = CPU.retpc; /*return addr*/  
  
    /* Load and execute new thread */  
    CPU.r7 = TCB[tNew].regs.r7;  
    ...  
    CPU.r0 = TCB[tNew].regs.r0;  
    CPU.sp = TCB[tNew].regs.sp;  
    CPU.retpc = TCB[tNew].regs.retpc;  
    return; /* Return to CPU.retpc */  
}
```

Switch Details (continued)

- What if you make a mistake in implementing switch?
 - Suppose you forget to save/restore register 32
 - Get intermittent failures depending on when context switch occurred and whether new thread uses register 32
 - System will give wrong result without warning
- Can you devise an exhaustive test to test switch code?
 - No! Too many combinations and inter-leavings
- Cautionary tale:
 - For speed, Topaz kernel saved one instruction in switch()
 - Carefully documented! Only works as long as kernel size < 1MB
 - What happened?
 - Time passed, People forgot
 - Later, they added features to kernel (no one removes features!)
 - Very weird behavior started happening
 - Moral of story: Design for simplicity


Administrivia

- Project 1 in full swing! Released Yesterday!
 - We expect that your design document will give intuitions behind your designs, not just a dump of pseudo-code
 - Think of this you are in a company and your TA is your manager
- Paradox: need code for design document?
 - Not full code, just enough prove you have thought through complexities of design
- Should be attending your permanent discussion section!
 - Discussion section attendance is mandatory, but don't come if sick!!
 - We have given a mechanism to make up for missed sections—see piazza
 - We will have a rotating recording of sections for later viewing as well
- Midterm 1: February 17th, 7-9PM (Two weeks from today!)
 - Fill out conflict request by tomorrow!

Are we still switching contexts with previous examples?

- Yes, but **much cheaper** than switching processes
 - No need to change address space
- Some numbers from Linux:
 - Frequency of context switch: 10-100ms
 - Switching between processes: 3-4 μ sec.
 - Switching between threads: 100 ns
- Even cheaper: switch threads (using “yield”) in user-space!


What we are talking about in Today's lecture




Simple One-to-One Threading Model

2/3/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022




Many-to-One



Many-to-Many

Lec 6.17

What happens when thread blocks on I/O?



- What happens when a thread requests a block of data from the file system?
 - User code invokes a system call
 - Read operation is initiated
 - Run new thread/switch
- Thread communication similar
 - Wait for Signal/Join
 - Networking

2/3/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 6.18

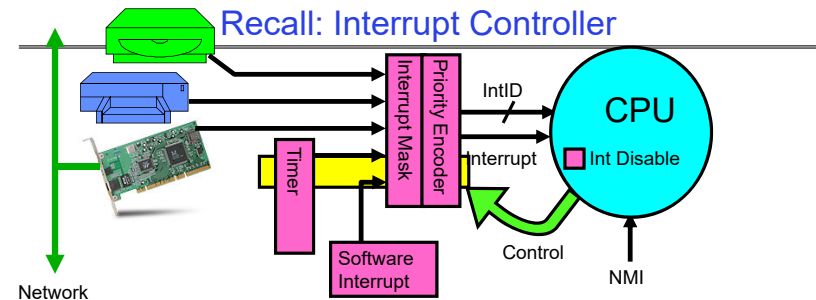
External Events

- What happens if thread never does any I/O, never waits, and never yields control?
 - Could the ComputePI program grab all resources and never release the processor?
 - » What if it didn't print to console?
 - Must find way that dispatcher can regain control!
- Answer: utilize external events
 - Interrupts: signals from hardware or software that stop the running code and jump to kernel
 - Timer: like an alarm clock that goes off every some milliseconds
- If we make sure that external events occur frequently enough, can ensure dispatcher runs

2/3/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 6.19



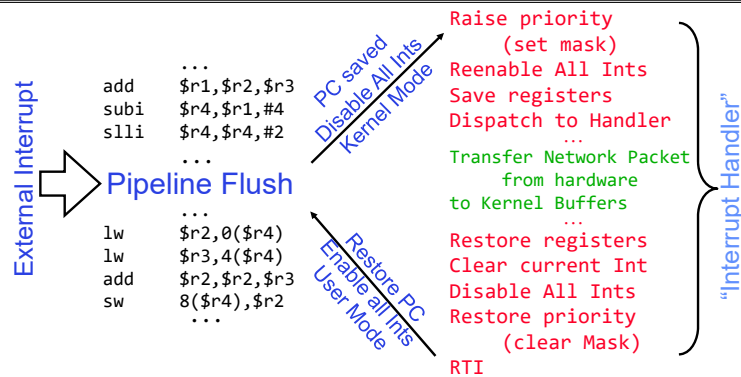
- Interrupts invoked with interrupt lines from devices
- Interrupt controller chooses interrupt request to honor
 - Interrupt identity specified with ID line
 - Mask enables/disables interrupts
 - Priority encoder picks highest enabled interrupt
 - Software Interrupt Set/Cleared by Software
- CPU can disable all interrupts with internal flag
- Non-Maskable Interrupt line (NMI) can't be disabled

2/3/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 6.20

Example: Network Interrupt



- An interrupt is a hardware-invoked context switch
 - No separate step to choose what to run next
 - Always run the interrupt handler immediately


2/3/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 6.21

Use of Timer Interrupt to Return Control

- Solution to our dispatcher problem
 - Use the timer interrupt to force scheduling decisions



- Timer Interrupt routine:

```

TimerInterrupt() {
    DoPeriodicHouseKeeping();
    run_new_thread();
}

```

2/3/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 6.22

ThreadFork(): Create a New Thread

- ThreadFork() is a user-level procedure that creates a new thread and places it on ready queue
- Arguments to ThreadFork()
 - Pointer to application routine (fcnPtr)
 - Pointer to array of arguments (fcnArgPtr)
 - Size of stack to allocate
- Implementation
 - Sanity check arguments
 - Enter Kernel-mode and Sanity Check arguments again
 - Allocate new Stack and TCB
 - Initialize TCB and place on ready list (Runnable)

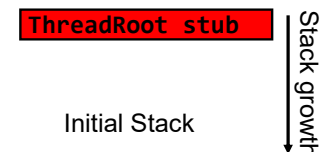
2/3/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 6.23

How do we initialize TCB and Stack?

- Initialize Register fields of TCB
 - Stack pointer made to point at stack
 - PC return address \Rightarrow OS (asm) routine ThreadRoot()
 - Two arg registers (a0 and a1) initialized to fcnPtr and fcnArgPtr, respectively
- Initialize stack data?
 - Minimal initialization \Rightarrow setup return to go to beginning of ThreadRoot()
 - Important part of stack frame is in registers for RISC-V (ra)
 - X86: need to push a return address on stack
 - Think of stack frame as just before body of ThreadRoot() really gets started




2/3/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 6.24

How does Thread get started?



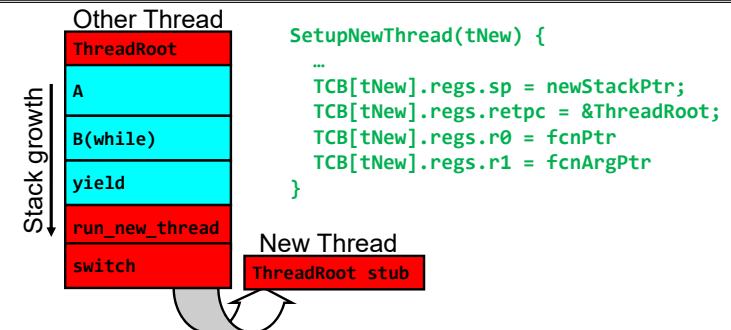
- Eventually, `run_new_thread()` will select this TCB and return into beginning of `ThreadRoot()`
 - This really starts the new thread

2/3/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 6.25

How does a thread get started?



- How do we make a **new** thread?
 - Setup TCB/kernel thread to point at new user stack and `ThreadRoot` code
 - Put pointers to start function and args in registers or top of stack
 - » This depends heavily on the calling convention (i.e. RISC-V vs x86)
- Eventually, `run_new_thread()` will select this TCB and return into beginning of `ThreadRoot()`
 - This really starts the new thread

2/3/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022


Lec 6.26

What does `ThreadRoot()` look like?

- `ThreadRoot()` is the root for the thread routine:

```

ThreadRoot(fcnPTR, fcnArgPtr) {
    DoStartupHousekeeping();
    UserModeSwitch(); /* enter user mode */
    Call fcnPtr(fcnArgPtr);
    ThreadFinish();
}
    
```




- Startup Housekeeping
 - Includes things like recording start time of thread
 - Other statistics
- Stack will grow and shrink with execution of thread
- Final return from thread returns into `ThreadRoot()` which calls `ThreadFinish()`
 - `ThreadFinish()` wake up sleeping threads

2/3/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 6.27

Processes vs. Threads: One Core



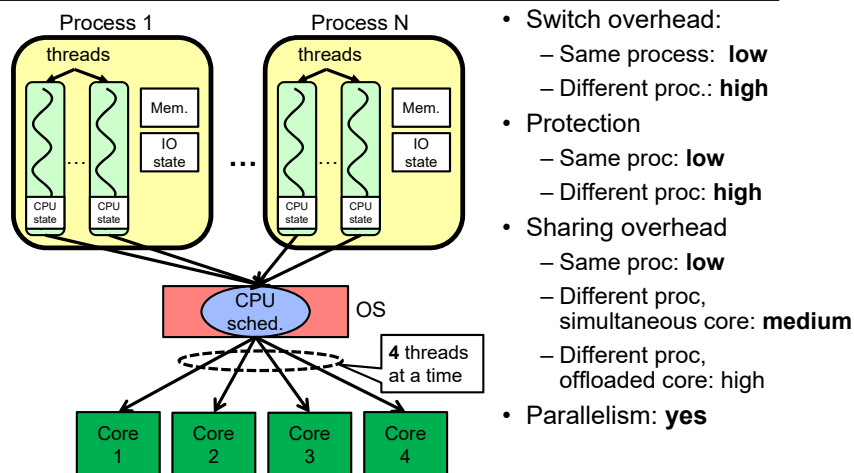
- Switch overhead:
 - Same process: **low**
 - Different proc.: **high**
- Protection
 - Same proc: **low**
 - Different proc: **high**
- Sharing overhead
 - Same proc: **low**
 - Different proc: **high**
- Parallelism: **no**

2/3/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 6.28

Processes vs. Threads: MultiCore



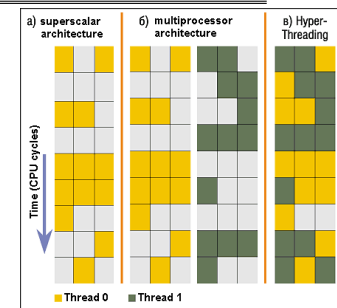
2/3/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 6.29

Recall: Simultaneous MultiThreading/Hyperthreading

- Hardware scheduling technique
 - Superscalar processors can execute multiple instructions that are independent.
 - Hyperthreading duplicates register state to make a second “thread,” allowing more instructions to run.
- Can schedule each thread as if were separate CPU
 - But, sub-linear speedup!
- Original technique called “Simultaneous Multithreading”
 - <http://www.cs.washington.edu/research/smt/index.html>
 - SPARC, Pentium 4/Xeon (“Hyperthreading”), Power 5



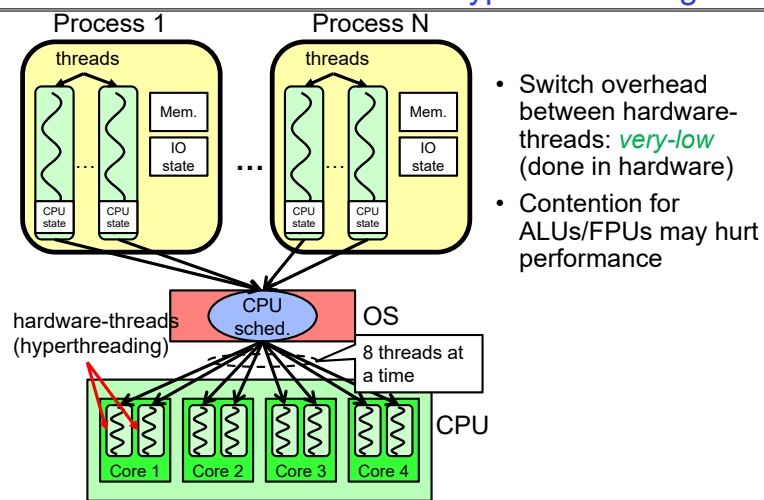
Colored blocks show instructions executed

2/3/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 6.30

Processes vs. Threads: Hyper-Threading



2/3/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 6.31

Threads vs Address Spaces: Options

# threads Per AS:	# of addr spaces:	One	Many
One		MS/DOS, early Macintosh	Traditional UNIX
Many		Embedded systems (Geoworks, VxWorks, JavaOS, etc) JavaOS, Pilot(PC)	Mach, OS/2, Linux Windows 10 Win NT to XP, Solaris, HP-UX, OS X

- Most operating systems have either
 - One or many address spaces
 - One or many threads per address space


2/3/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 6.32

Multiprocessing vs Multiprogramming

- Some Definitions:
 - Multiprocessing \equiv Multiple CPUs
 - Multiprogramming \equiv Multiple Jobs or Processes
 - Multithreading \equiv Multiple threads per Process
- What does it mean to run two threads “concurrently”?
 - Scheduler is free to run threads in any order and interleaving: FIFO, Random, ...
 - Dispatcher can choose to run each thread to completion or time-slice in big chunks or small chunks



Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 6.33

Correctness for systems with concurrent threads

- If dispatcher can schedule threads in any way, programs must work under all circumstances
 - Can you test for this?
 - How can you know if your program works?
- Independent Threads:**
 - No state shared with other threads
 - Deterministic \Rightarrow Input state determines results
 - Reproducible \Rightarrow Can recreate Starting Conditions, I/O
 - Scheduling order doesn't matter (if `switch()` works!!!)
- Cooperating Threads:**
 - Shared State between multiple threads
 - Non-deterministic
 - Non-reproducible
- Non-deterministic and Non-reproducible means that bugs can be intermittent
 - Sometimes called “Heisenbugs”

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 6.34

Interactions Complicate Debugging

- Is any program truly independent?
 - Every process shares the file system, OS resources, network, etc
 - Extreme example: buggy device driver causes thread A to crash “independent thread” B
- You probably don't realize how much you depend on reproducibility:
 - Example: Evil C compiler
 - Modifies files behind your back by inserting errors into C program unless you insert debugging code
 - Example: Debugging statements can overrun stack
- Non-deterministic errors are really difficult to find
 - Example: Memory layout of kernel+user programs
 - depends on scheduling, which depends on timer/other things
 - Original UNIX had a bunch of non-deterministic errors
 - Example: Something which does interesting I/O
 - User typing of letters used to help generate secure keys

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 6.35


Why allow cooperating threads?

- People cooperate; computers help/enhance people's lives, so computers must cooperate
 - By analogy, the non-reproducibility/non-determinism of people is a notable problem for “carefully laid plans”
- Advantage 1: Share resources
 - One computer, many users
 - One bank balance, many ATMs
 - What if ATMs were only updated at night?
 - Embedded systems (robot control: coordinate arm & hand)
- Advantage 2: Speedup
 - Overlap I/O and computation
 - Many different file systems do read-ahead
 - Multiprocessors – chop up program into parallel pieces
- Advantage 3: Modularity
 - More important than you might think
 - Chop large problem up into simpler pieces
 - To compile, for instance, gcc calls `cpp` | `cc1` | `cc2` | `as` | `ld`
 - Makes system easier to extend

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 6.36

Recall: High-level Example: Web Server



- Server must handle many requests
- Non-cooperating version:


```
serverLoop() {
    con = AcceptCon();
    ProcessFork (ServiceWebPage (), con);
}
```
- What are some disadvantages of this technique?

2/3/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 6.37

Recall: Threaded Web Server

- Now, use a single process
- Multithreaded (cooperating) version:


```
serverLoop() {
    connection = AcceptCon();
    ThreadFork (ServiceWebPage (), connection);
}
```
- Looks almost the same, but has many advantages:
 - Can share file caches kept in memory, results of CGI scripts, other things
 - Threads are *much* cheaper to create than processes, so this has a lower per-request overhead
- Question: would a user-level (say one-to-many) thread package make sense here?
 - When one request blocks on disk, all block...
- What about Denial of Service attacks or digg / Slash-dot effects?




2/3/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 6.38

Thread Pools: Bounded Concurrency

- Problem with previous version: Unbounded Threads
 - When web-site becomes too popular – throughput sinks
- Instead, allocate a bounded “pool” of worker threads, representing the maximum level of multiprogramming



```
master() {
    allocThreads(worker, queue);
    while(TRUE) {
        con=AcceptCon();
        Enqueue(queue, con);
        wakeUp(queue);
    }
}

worker(queue) {
    while(TRUE) {
        con=Dequeue(queue);
        if (con==null)
            sleepOn(queue);
        else
            ServiceWebPage(con);
    }
}
```

2/3/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 6.39

Correctness with Concurrent Threads?


- Non-determinism:
 - Scheduler can run threads in **any order**
 - Scheduler can switch threads **at any time**
 - This can make testing very difficult
- *Independent Threads*
 - No state shared with other threads
 - Deterministic, reproducible conditions
- *Cooperating Threads*
 - Shared state between multiple threads
- **Goal: Correctness by Design**

2/3/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 6.40

ATM Bank Server



- ATM server problem:
 - Service a set of requests
 - Do so without corrupting database
 - Don't hand out too much money

2/3/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 6.41

ATM bank server example

- Suppose we wanted to implement a server process to handle requests from an ATM network:

```
BankServer() {
    while (TRUE) {
        ReceiveRequest(&op, &acctId, &amount);
        ProcessRequest(op, acctId, amount);
    }
}

ProcessRequest(op, acctId, amount) {
    if (op == deposit) Deposit(acctId, amount);
    else if ...
}

Deposit(acctId, amount) {
    acct = GetAccount(acctId); /* may use disk I/O */
    acct->balance += amount;
    StoreAccount(acct); /* Involves disk I/O */
}
```

- How could we speed this up?
 - More than one request being processed at once
 - Event driven (overlap computation and I/O)
 - Multiple threads (multi-proc, or overlap comp and I/O)

2/3/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 6.42

Event Driven Version of ATM server

- Suppose we only had one CPU
 - Still like to overlap I/O with computation
 - Without threads, we would have to rewrite in event-driven style
- Example

```
BankServer() {
    while (TRUE) {
        event = WaitForNextEvent();
        if (event == ATMRequest)
            StartOnRequest();
        else if (event == AcctAvail)
            ContinueRequest();
        else if (event == AcctStored)
            FinishRequest();
    }
}
```

– This technique is used for graphical programming

- Complication:
 - What if we missed a blocking I/O step?
 - What if we have to split code into hundreds of pieces which could be blocking?

2/3/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 6.43

Can Threads Make This Easier?

- Threads yield overlapped I/O and computation without “deconstructing” code into non-blocking fragments
 - One thread per request
- Requests proceeds to completion, blocking as required:

```
Deposit(acctId, amount) {
    acct = GetAccount(acctId); /* May use disk I/O */
    acct->balance += amount;
    StoreAccount(acct); /* Involves disk I/O */
}
```

- Unfortunately, shared state can get corrupted:


<u>Thread 1</u>	<u>Thread 2</u>
load r1, acct->balance	
	load r1, acct->balance
	add r1, amount2
	store r1, acct->balance
add r1, amount1	
store r1, acct->balance	

2/3/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 6.44

Recall: Possible Executions



2/3/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 6.45

Problem is at the Lowest Level

- Most of the time, threads are working on separate data, so scheduling doesn't matter:

Thread A

$x = 1;$

Thread B

$y = 2;$

- However, what about (Initially, $y = 12$):

Thread A

$x = 1;$

$x = y + 1;$

Thread B

$y = 2;$

$y = y * 2;$

- What are the possible values of x ?

- Or, what are the possible values of x below?

Thread A

$x = 1;$

Thread B

$x = 2;$

- x could be 1 or 2 (non-deterministic!)

- Could even be 3 for serial processors:

» Thread A writes 0001, B writes 0010 → scheduling order ABABABBA yields 3!

2/3/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 6.46

Atomic Operations

- To understand a concurrent program, we need to know what the underlying indivisible operations are!
- Atomic Operation:** an operation that always runs to completion or not at all
 - It is *indivisible*: it cannot be stopped in the middle and state cannot be modified by someone else in the middle
 - Fundamental building block – if no atomic operations, then have no way for threads to work together
- On most machines, memory references and assignments (i.e. loads and stores) of words are atomic
 - Consequently – weird example that produces “3” on previous slide can't happen
- Many instructions are not atomic
 - Double-precision floating point store often not atomic
 - VAX and IBM 360 had an instruction to copy a whole array

2/3/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 6.47

Locks

- Lock:** prevents someone from doing something
 - Lock** before entering critical section and before accessing shared data
 - Unlock** when leaving, after accessing shared data
 - Wait** if locked
 - » Important idea: all synchronization involves waiting
- Locks need to be allocated and initialized:
 - structure Lock mylock or pthread_mutex_t mylock;
 - lock_init(&mylock) or mylock = PTHREAD_MUTEX_INITIALIZER;
- Locks provide two **atomic** operations:
 - acquire(&mylock)** – wait until lock is free; then mark it as busy
 - » After this returns, we say the calling thread *holds* the lock
 - release(&mylock)** – mark lock as free
 - » Should only be called by a thread that currently holds the lock
 - » After this returns, the calling thread no longer holds the lock



2/3/2022

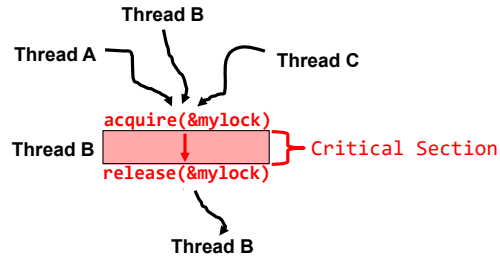
Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 6.48

Fix banking problem with Locks!

- Identify critical sections (atomic instruction sequences) and add locking:

```
Deposit(acctId, amount) {  
    acquire(&mylock)      // Wait if someone else in critical section!  
    acct = GetAccount(acctId);  
    acct->balance += amount;  
    StoreAccount(acct);  
    release(&mylock)      // Release someone into critical section  
}
```



Threads serialized by lock
through critical section.
Only one thread at a time

- Must use SAME lock (mylock) with all of the methods (Withdraw, etc...)
 - Shared with all threads!

Conclusion

- Processes have two parts
 - Threads (Concurrency)
 - Address Spaces (Protection)
- Various textbooks talk about *processes*
 - When this concerns concurrency, really talking about thread portion of a process
 - When this concerns protection, talking about address space portion of a process
- Concurrent threads are a very useful abstraction
 - Allow transparent overlapping of computation and I/O
 - Allow use of parallel processing when available
- Concurrent threads introduce problems when accessing shared data
 - Programs must be insensitive to arbitrary interleavings
 - Without careful design, shared variables can become completely inconsistent