

CS162
Operating Systems and
Systems Programming
Lecture 7

Synchronization 2: Concurrency (Con't),
Lock Implementation, Atomic Instructions

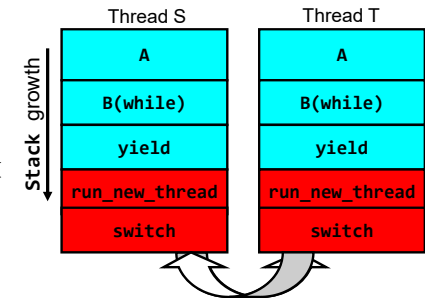
February 8th, 2022
Prof. Anthony Joseph and John Kubiawicz
<http://cs162.eecs.Berkeley.edu>

Recall: Multithreaded Stack Example

- Consider the following code blocks:

```
proc A() {
    B();
}

proc B() {
    while(TRUE) {
        yield();
    }
}
```



- Suppose we have 2 threads:
 - Threads S and T

Thread S's switch returns to Thread T's (and vice versa)

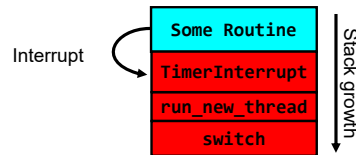
2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.2

Recall: Use of Timer Interrupt to Return Control

- Solution to our dispatcher problem
 - Use the timer interrupt to force scheduling decisions



- Timer Interrupt routine:

```
TimerInterrupt() {
    DoPeriodicHouseKeeping();
    run_new_thread();
}
```

Hardware context switch support in x86

- Syscall/Intr (U → K)
 - PL 3 → 0;
 - TSS ← EFLAGS, CS:EIP;
 - SS:ESP ← k-thread stack (TSS PL 0);
 - push (old) SS:ESP onto (new) k-stack
 - push (old) eflags, cs:eip, <err>
 - CS:EIP ← <k target handler>
- Then
 - Handler saves other regs, etc
 - Does all its work, possibly choosing other threads, changing PTBR (CR3)
 - kernel thread has set up user GPRs
- iret (K → U)
 - PL 0 → 3;
 - Eflags, CS:EIP ← popped off k-stack
 - SS:ESP ← popped off k-stack

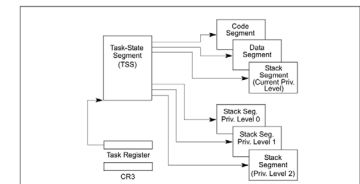


Figure 7-1. Structure of a Task

pg 2,942 of 4,922 of x86 reference manual

Pintos: `tss.c`, `intr-stubs.S`

2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

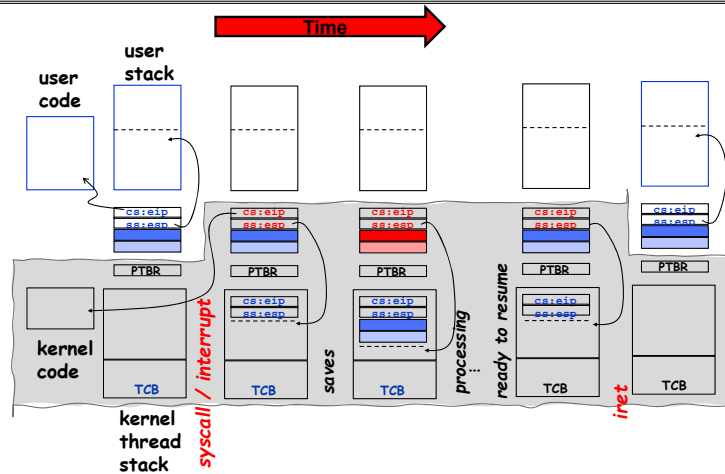
Lec 7.3

2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.4

Pintos: Kernel Crossing on Syscall or Interrupt

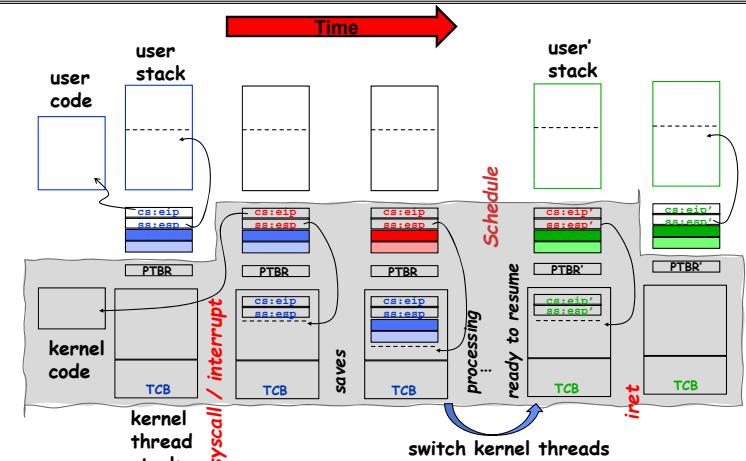


2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.5

Pintos: Context Switch – Scheduling



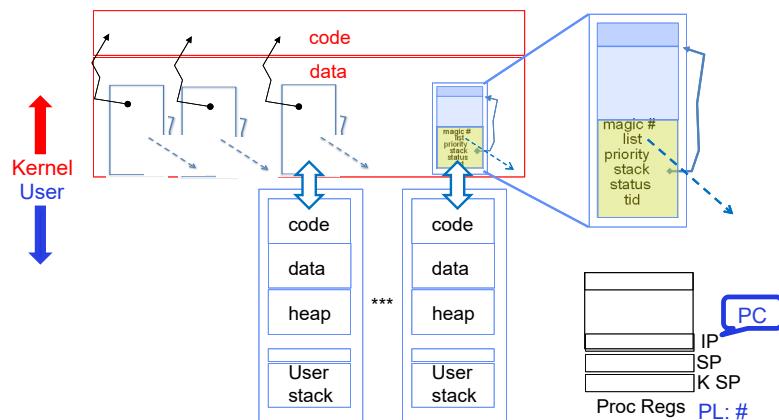
Pintos: switch.S

2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.6

MT Kernel 1T Process ala Pintos/x86



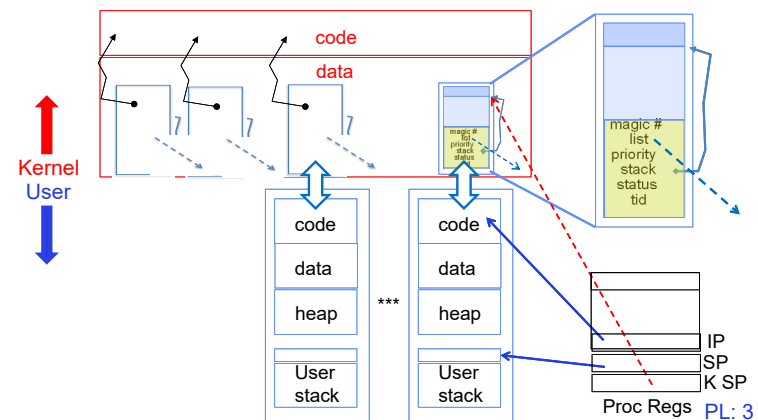
- Each user process/thread associated with a kernel thread, described by a 4KB page object containing TCB and kernel stack for the kernel thread

2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.7

In User thread, w/ Kernel thread waiting



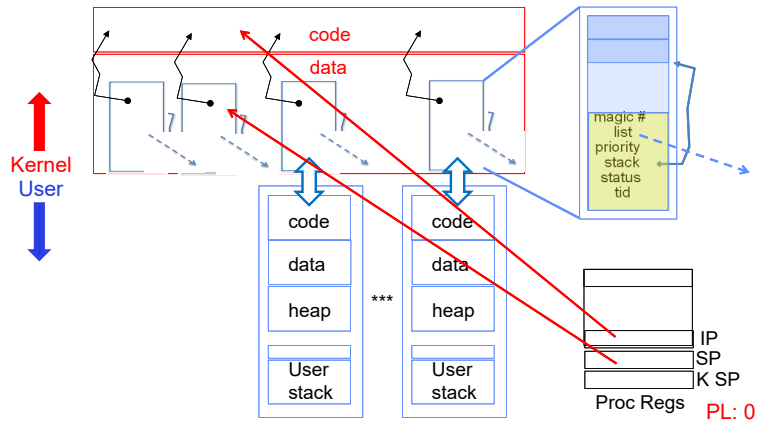
- x86 CPU holds interrupt SP in register
- During user thread execution, associated kernel thread is "standing by"

2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.8

In Kernel Thread: No User Component



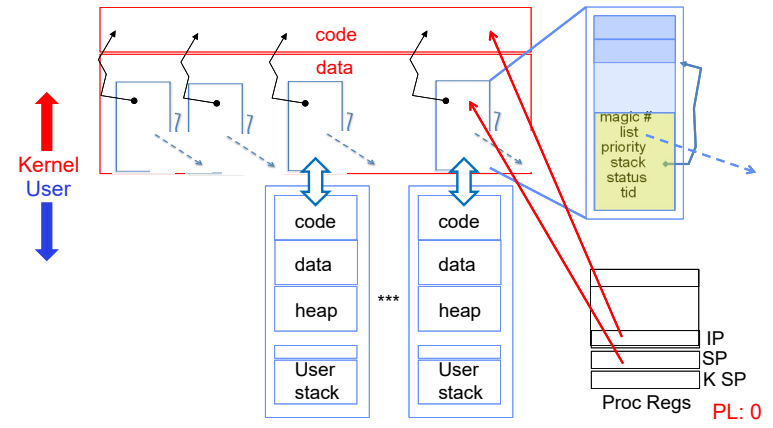
- Kernel threads execute with small stack in thread structure
- Pure kernel threads have no corresponding user-mode thread

2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.9

User → Kernel (exceptions, syscalls)



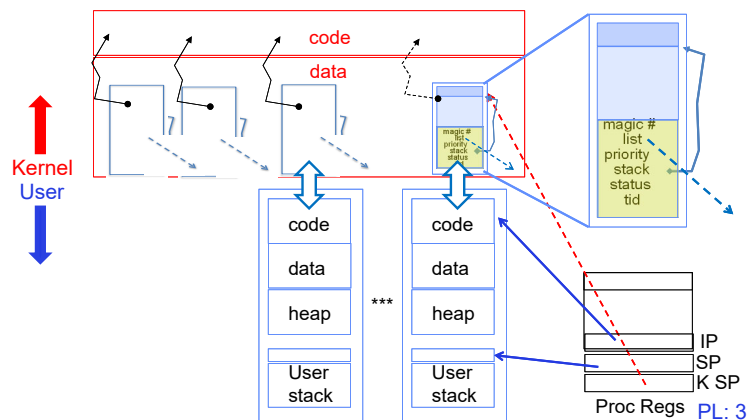
- Mechanism to resume k-thread goes through interrupt vector

2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.10

Kernel → User



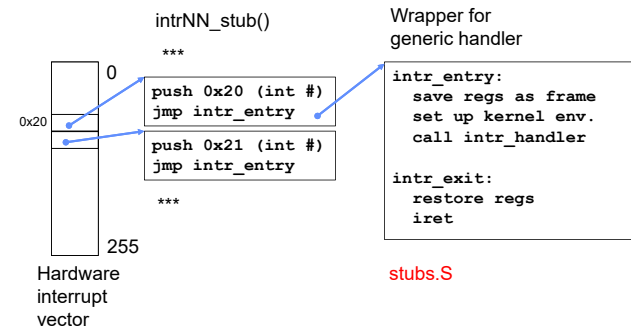
- Interrupt return (iret) restores user stack, IP, and PL

2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.11

Pintos Interrupt Processing



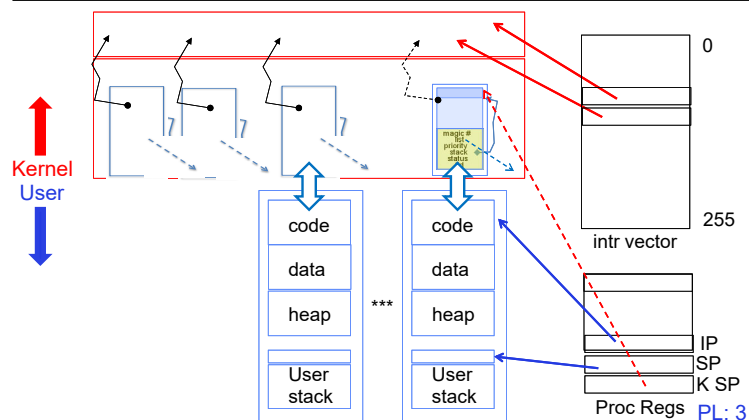
stubs.S

2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.12

User → Kernel via interrupt vector



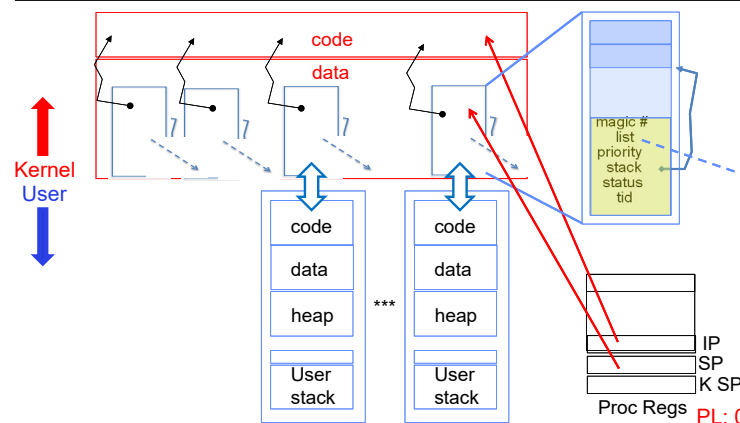
- Interrupt transfers control through the Interrupt Vector (IDT in x86)
- `iret` restores user stack and priority level (PL)

2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.13

Switch to Kernel Thread for Process

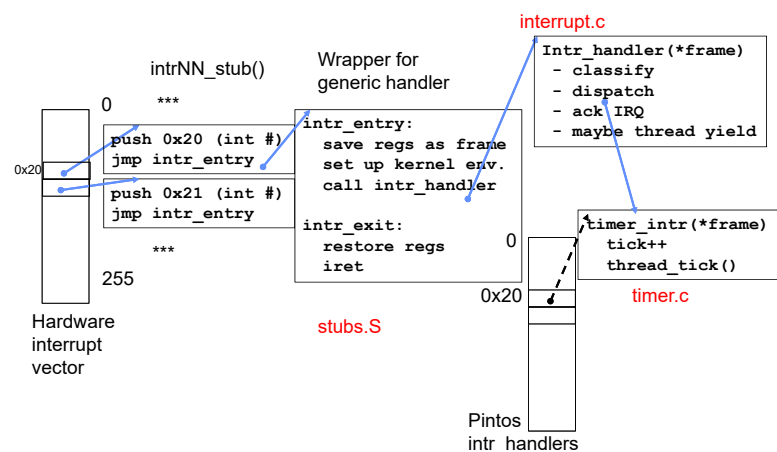


2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.14

Pintos Interrupt Processing



2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.15

Timer may trigger thread switch

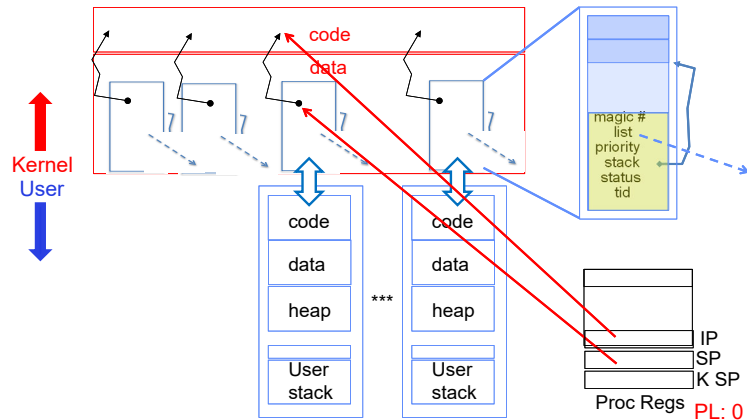
- `thread_tick`
 - Updates thread counters
 - If quanta exhausted, sets yield flag
- `thread_yield`
 - On path to `rtn` from interrupt
 - Sets current thread back to `READY`
 - Pushes it back on `ready_list`
 - Calls `schedule` to select next thread to run upon `iret`
- `Schedule`
 - Selects next thread to run
 - Calls `switch_threads` to change regs to point to stack for thread to resume
 - Sets its status to `RUNNING`
 - If user thread, activates the process
 - Returns back to `intr_handler`

2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.16

Thread Switch (switch.S)



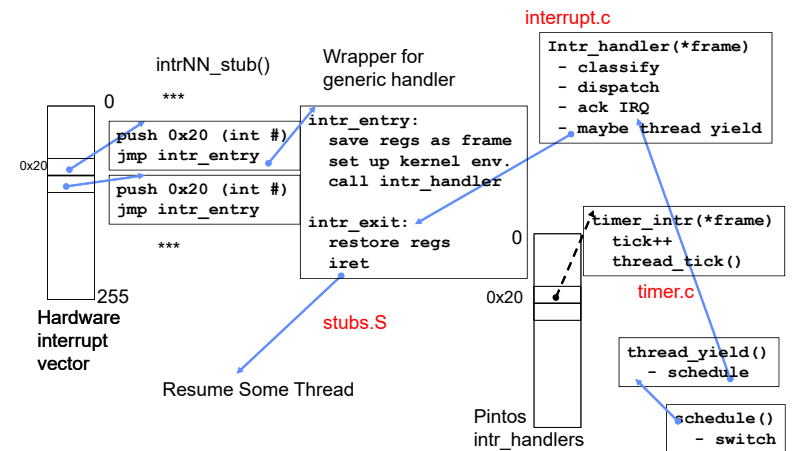
- switch_threads: save regs on current small stack, change SP, return from destination threads call to switch_threads

2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.17

Pintos Return from Processing

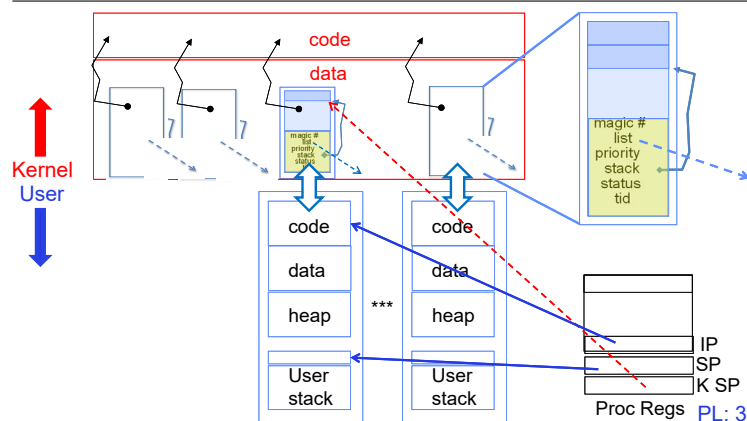


2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.18

Kernel → Different User Thread



- iret restores user stack and priority level (PL)

2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.19

Famous Quote WRT Scheduling: Dennis Richie

Dennis Richie,
Unix V6, slp.c:

```

2230 /*
2231  * If the new process paused because it was
2232  * swapped out, set the stack level to the last call
2233  * to savu(u_ssav). This means that the return
2234  * which is executed immediately after the call to aretu
2235  * actually returns from the last routine which did
2236  * the savu.
2237  *
2238  * You are not expected to understand this.
2239  */
    
```

"If the new process paused because it was swapped out, set the stack level to the last call to savu(u_ssav). This means that the return which is executed immediately after the call to aretu actually returns from the last routine which did the savu."

"You are not expected to understand this."

Source: Dennis Ritchie, Unix V6 slp.c (context-switching code) as per The Unix Heritage Society(tuhs.org); gif by Eddie Koehler.

Included by Ali R. Butt in CS3204 from Virginia Tech

2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.20

Administrivia

- Midterm Thursday 2/17
 - No class on day of midterm
 - 7-9PM
- Project 1 Design Document due next Friday 2/11
- Project 1 Design reviews upcoming
 - High-level discussion of your approach
 - » What will you modify?
 - » What algorithm will you use?
 - » How will things be linked together, etc.
 - » Do not need final design (complete with all semicolons!)
 - You will be asked about testing
 - » Understand testing framework
 - » Are there things you are doing that are not tested by tests we give you?
- Do your own work!
 - Please do not try to find solutions from previous terms
 - We will be on the look out for anyone doing this...today

2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.21

Goals for Rest of Today

- Challenges and Pitfalls of Concurrency
- Synchronization Operations/Critical Sections
- How to build a lock?
- Atomic Instructions



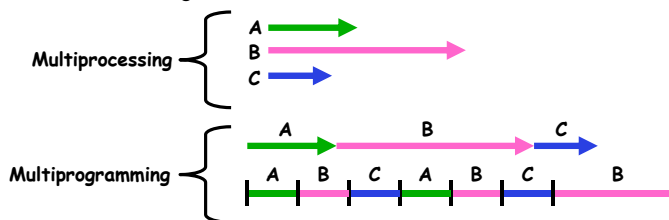
2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.22

Recall: Multiprocessing vs Multiprogramming

- Some Definitions:
 - Multiprocessing \equiv Multiple CPUs
 - Multiprogramming \equiv Multiple Jobs or Processes
 - Multithreading \equiv Multiple threads per Process
- What does it mean to run two threads "concurrently"?
 - Scheduler is free to run threads in any order and interleaving: FIFO, Random, ...
 - Dispatcher can choose to run each thread to completion or time-slice in big chunks or small chunks

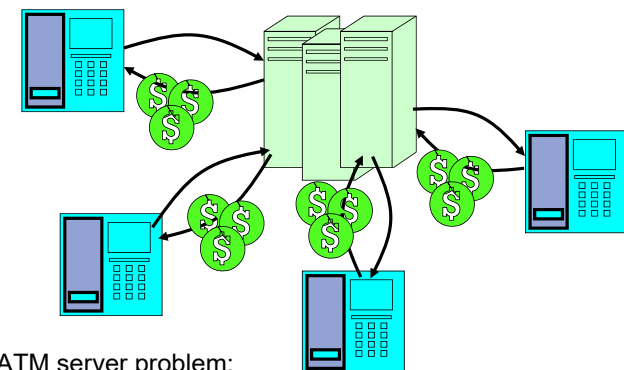


2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.23

ATM Bank Server



- ATM server problem:
 - Service a set of requests
 - Do so without corrupting database
 - Don't hand out too much money

2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.24

ATM bank server example

- Suppose we wanted to implement a server process to handle requests from an ATM network:

```
BankServer() {
    while (TRUE) {
        ReceiveRequest(&op, &acctId, &amount);
        ProcessRequest(op, acctId, amount);
    }
}

ProcessRequest(op, acctId, amount) {
    if (op == deposit) Deposit(acctId, amount);
    else if ...
}

Deposit(acctId, amount) {
    acct = GetAccount(acctId); /* may use disk I/O */
    acct->balance += amount;
    StoreAccount(acct); /* Involves disk I/O */
}
```

- How could we speed this up?
 - More than one request being processed at once
 - Event driven (overlap computation and I/O)
 - Multiple threads (multi-proc, or overlap comp and I/O)

2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.25

Event Driven Version of ATM server

- Suppose we only had one CPU
 - Still like to overlap I/O with computation
 - Without threads, we would have to rewrite in event-driven style

- Example

```
BankServer() {
    while(TRUE) {
        event = WaitForNextEvent();
        if (event == ATMRequest)
            StartOnRequest();
        else if (event == AcctAvail)
            ContinueRequest();
        else if (event == AcctStored)
            FinishRequest();
    }
}
```

- This technique is used for graphical programming

- Complication:

- What if we missed a blocking I/O step?
- What if we have to split code into hundreds of pieces which could be blocking?

2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.26

Can Threads Make This Easier?

- Threads yield overlapped I/O and computation without “deconstructing” code into non-blocking fragments

- One thread per request

- Requests proceeds to completion, blocking as required:

```
Deposit(acctId, amount) {
    acct = GetAccount(acctId); /* May use disk I/O */
    acct->balance += amount;
    StoreAccount(acct); /* Involves disk I/O */
}
```

- Unfortunately, shared state can get corrupted:

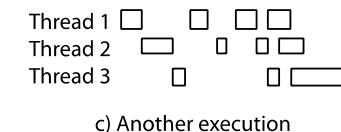
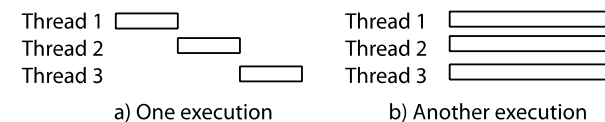
<u>Thread 1</u>	<u>Thread 2</u>
load r1, acct->balance	
	load r1, acct->balance
	add r1, amount2
	store r1, acct->balance
add r1, amount1	
store r1, acct->balance	

2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.27

Recall: Possible Executions



2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.28

Problem is at the Lowest Level

- Most of the time, threads are working on separate data, so scheduling doesn't matter:

Thread A

x = 1;

Thread B

y = 2;

- However, what about (Initially, y = 12):

Thread A

x = 1;

x = y+1;

Thread B

y = 2;

y = y*2;

– What are the possible values of x?

- Or, what are the possible values of x below?

Thread A

x = 1;

Thread B

x = 2;

– X could be 1 or 2 (non-deterministic!)

– Could even be 3 for serial processors:

» Thread A writes 0001, B writes 0010 → scheduling order ABABABBA yields 3!

2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.29

Atomic Operations

- To understand a concurrent program, we need to know what the underlying indivisible operations are!
- Atomic Operation:** an operation that always runs to completion or not at all
 - It is *indivisible*: it cannot be stopped in the middle and state cannot be modified by someone else in the middle
 - Fundamental building block – if no atomic operations, then have no way for threads to work together
- On most machines, memory references and assignments (i.e. loads and stores) of words are atomic
 - Consequently – weird example that produces “3” on previous slide can't happen
- Many instructions are not atomic
 - Double-precision floating point store often not atomic
 - VAX and IBM 360 had an instruction to copy a whole array

2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.30

Another Concurrent Program Example

- Two threads, A and B, compete with each other
 - One tries to increment a shared counter
 - The other tries to decrement the counter

Thread A

i = 0;

while (i < 10)

i = i + 1;

printf("A wins!");

Thread B

i = 0;

while (i > -10)

i = i - 1;

printf("B wins!");

- Assume that memory loads and stores are atomic, but incrementing and decrementing are *not* atomic
- Who wins? Could be either
- Is it guaranteed that someone wins? Why or why not?
- What if both threads have their own CPU running at same speed? Is it guaranteed that it goes on forever?

2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.31

Hand Simulation Multiprocessor Example

- Inner loop looks like this:

Thread A

r1=0 load r1, M[i]

r1=1 add r1, r1, 1

M[i]=1 store r1, M[i]

Thread B

r1=0 load r1, M[i]

r1=-1 sub r1, r1, 1

M[i]=-1 store r1, M[i]

- Hand Simulation:**
 - And we're off. A gets off to an early start
 - B says “hmph, better go fast” and tries really hard
 - A goes ahead and writes “1”
 - B goes and writes “-1”
 - A says “HUH??? I could have sworn I put a 1 there”
- Could this happen on a uniprocessor? With Hyperthreads?
 - Yes! Unlikely, but if you are depending on it not happening, it will and your system will break...

2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.32

Definitions

- **Synchronization**: using atomic operations to ensure cooperation between threads
 - For now, only loads and stores are atomic
 - We are going to show that its hard to build anything useful with only reads and writes
- **Mutual Exclusion**: ensuring that only one thread does a particular thing at a time
 - One thread *excludes* the other while doing its task
- **Critical Section**: piece of code that only one thread can execute at once. Only one thread at a time will get into this section of code
 - Critical section is the result of mutual exclusion
 - Critical section and mutual exclusion are two ways of describing the same thing

2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.33

Locks

- **Lock**: prevents someone from doing something
 - **Lock()** before entering critical section and before accessing shared data
 - **Unlock()** when leaving, after accessing shared data
 - **Wait** if locked
 - » Important idea: all synchronization involves waiting
- Locks need to be allocated and initialized:
 - structure Lock mylock or pthread_mutex_t mylock;
 - lock_init(&mylock) or mylock = PTHREAD_MUTEX_INITIALIZER;
- Locks provide two **atomic** operations:
 - **acquire(&mylock)** – wait until lock is free; then mark it as busy
 - » After this returns, we say the calling thread *holds* the lock
 - **release(&mylock)** – mark lock as free
 - » Should only be called by a thread that currently holds the lock
 - » After this returns, the calling thread no longer holds the lock



2/8/22

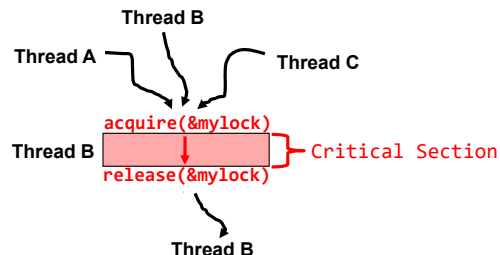
Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.34

Fix banking problem with Locks!

- Identify critical sections (atomic instruction sequences) and add locking:

```
Deposit(acctId, amount) {
    acquire(&mylock) // Wait if someone else in critical section!
    acct = GetAccount(acctId);
    acct->balance += amount;
    StoreAccount(acct);
    release(&mylock) // Release someone into critical section
}
```



- Must use **SAME** lock (mylock) with all of the methods (Withdraw, etc...)
 - Shared with all threads!

2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.35

Correctness Requirements

- Threaded programs must work for all interleavings of thread instruction sequences
 - Cooperating threads inherently non-deterministic and non-reproducible
 - Really hard to debug unless carefully designed!
- Example: Therac-25
 - Machine for radiation therapy
 - » Software control of electron accelerator and electron beam/Xray production
 - » Software control of dosage
 - Software errors caused the death of several patients
 - » A series of race conditions on shared variables and poor software design
 - » "They determined that data entry speed during editing was the key factor in producing the error condition: If the prescription data was edited at a fast pace, the overdose occurred."

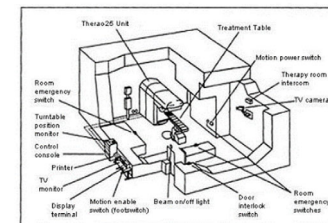


Figure 1. Typical Therac-25 facility

2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.36

Motivating Example: “Too Much Milk”

- Great thing about OS’s – analogy between problems in OS and problems in real life
 - Help you understand real life problems better
 - But, computers are much stupider than people
- Example: People need to coordinate:



Time	Person A	Person B
3:00	Look in Fridge. Out of milk	
3:05	Leave for store	
3:10	Arrive at store	Look in Fridge. Out of milk
3:15	Buy milk	Leave for store
3:20	Arrive home, put milk away	Arrive at store
3:25		Buy milk
3:30		Arrive home, put milk away

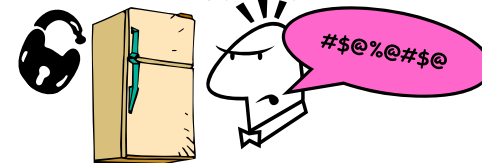
2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.37

Solve with a lock?

- **Recall:** Lock prevents someone from doing something
 - Lock before entering critical section
 - Unlock when leaving
 - Wait if locked
- » Important idea: all synchronization involves waiting
- For example: fix the milk problem by putting a key on the refrigerator
 - Lock it and take key if you are going to go buy milk
 - Fixes too much: roommate angry if only wants OJ



- Of Course – We don’t know how to make a lock yet
 - Let’s see if we can answer this question!

2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.38

Too Much Milk: Correctness Properties

- Need to be careful about correctness of concurrent programs, since non-deterministic
 - Impulse is to start coding first, then when it doesn’t work, pull hair out
 - Instead, think first, then code
 - Always write down behavior first
- What are the correctness properties for the “Too much milk” problem???
- Never more than one person buys
- Someone buys if needed
- **First attempt: Restrict ourselves to use only atomic load and store operations as building blocks**

2/8/22

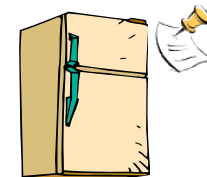
Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.39

Too Much Milk: Solution #1

- Use a note to avoid buying too much milk:
 - Leave a note before buying (kind of “lock”)
 - Remove note after buying (kind of “unlock”)
 - Don’t buy if note (wait)
- Suppose a computer tries this (remember, only memory read/write are atomic):

```
if (noMilk) {  
    if (noNote) {  
        leave Note;  
        buy milk;  
        remove note;  
    }  
}
```



2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.40

Too Much Milk: Solution #1

- Use a note to avoid buying too much milk:
 - Leave a note before buying (kind of "lock")
 - Remove note after buying (kind of "unlock")
 - Don't buy if note (wait)
- Suppose a computer tries this (remember, only memory read/write are atomic):

```

Thread A
if (noMilk) {
    if (noNote) {
        leave Note;
        buy Milk;
        remove Note;
    }
}

Thread B
if (noMilk) {
    if (noNote) {
        leave Note;
        buy Milk;
        remove Note;
    }
}
    
```

2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

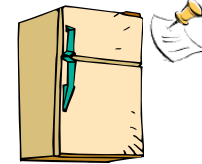
Lec 7.41

Too Much Milk: Solution #1

- Use a note to avoid buying too much milk:
 - Leave a note before buying (kind of "lock")
 - Remove note after buying (kind of "unlock")
 - Don't buy if note (wait)
- Suppose a computer tries this (remember, only memory read/write are atomic):

```

if (noMilk) {
    if (noNote) {
        leave Note;
        buy milk;
        remove note;
    }
}
    
```



- Result?
 - Still too much milk **but only occasionally!**
 - Thread can get context switched after checking milk and note but before buying milk!
- Solution makes problem worse since fails **intermittently**
 - Makes it really hard to debug...
 - Must work despite what the dispatcher does!

2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.42

Too Much Milk: Solution #1½

- Clearly the Note is not quite blocking enough
 - Let's try to fix this by placing note first
- Another try at previous solution:

```

leave Note;
if (noMilk) {
    if (noNote) {
        buy milk;
    }
}
remove Note;
    
```



- What happens here?
 - Well, with human, probably nothing bad
 - With computer: no one ever buys milk

2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.43

Too Much Milk Solution #2

- How about labeled notes?
 - Now we can leave note before checking
- Algorithm looks like this:

```

Thread A
leave note A;
if (noNote B) {
    if (noMilk) {
        buy Milk;
    }
}
remove note A;

Thread B
leave note B;
if (noNote A) {
    if (noMilk) {
        buy Milk;
    }
}
remove note B;
    
```

- Does this work?
- Possible for neither thread to buy milk
 - Context switches at exactly the wrong times can lead each to think that the other is going to buy
- Really insidious:
 - Extremely unlikely** this would happen, but will at worse possible time
 - Probably something like this in UNIX

2/8/22

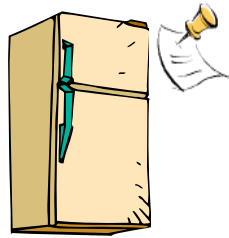
Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.44

Too Much Milk Solution #2: problem!



- I'm not getting milk, You're getting milk
- This kind of lockup is called "starvation!"



2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.45

Too Much Milk Solution #3

- Here is a possible two-note solution:

Thread A

```
leave note A;
while (note B) {\X
  do nothing;
}
if (noMilk) {
  buy milk;
}
remove note A;
```

Thread B

```
leave note B;
if (noNote A) {\Y
  if (noMilk) {
    buy milk;
  }
}
remove note B;
```

- Does this work? **Yes**. Both can guarantee that:
 - It is safe to buy, or
 - Other will buy, ok to quit
- At X:
 - If no note B, safe for A to buy,
 - Otherwise wait to find out what will happen
- At Y:
 - If no note A, safe for B to buy
 - Otherwise, A is either buying or waiting for B to quit

2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.46

Case 1

- "leave note A" happens before "if (noNote A)"

```

leave note A;
while (note B) {\X
  do nothing;
};

if (noMilk) {
  buy milk;
}
remove note A;

      happened
      before
leave note B;
if (noNote A) {\Y
  if (noMilk) {
    buy milk;
  }
}
remove note B;
```

2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.47

Case 1

- "leave note A" happens before "if (noNote A)"

```

leave note A;
while (note B) {\X
  do nothing;
};

if (noMilk) {
  buy milk;
}
remove note A;

      happened
      before
leave note B;
if (noNote A) {\Y
  if (noMilk) {
    buy milk;
  }
}
remove note B;
```

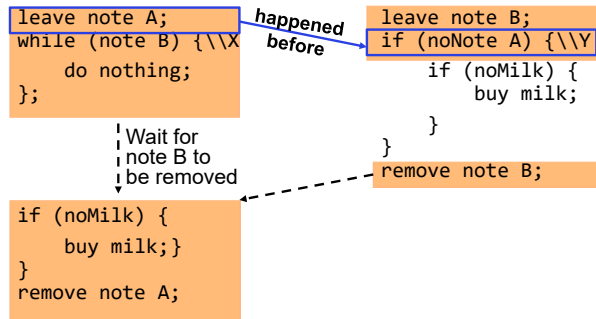
2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.48

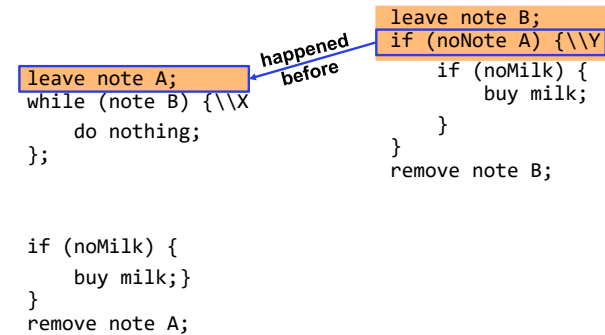
Case 1

- “leave note A” happens before “if (noNote A)”



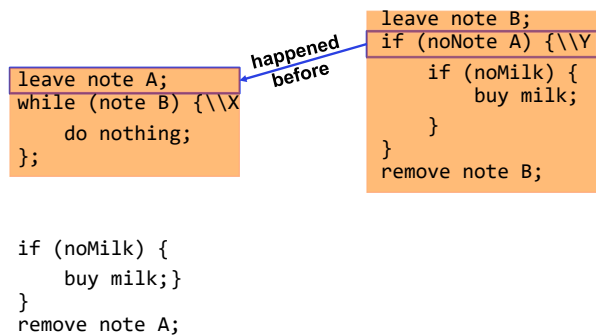
Case 2

- “if (noNote A)” happens before “leave note A”



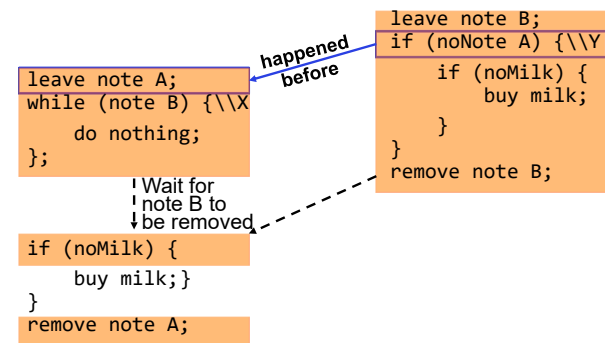
Case 2

- “if (noNote A)” happens before “leave note A”



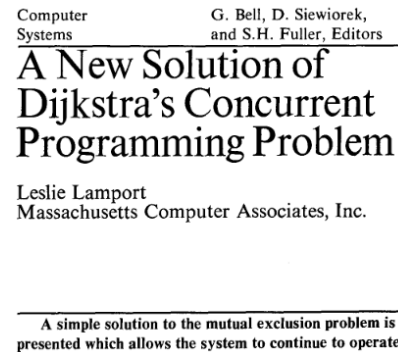
Case 2

- “if (noNote A)” happens before “leave note A”



This Generalizes to n Threads...

- Leslie Lamport's "Bakery Algorithm" (1974)



2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.53

Solution #3 discussion

- Our solution protects a single "Critical-Section" piece of code for each thread:

```
if (noMilk) {  
    buy milk;  
}
```

- Solution #3 works, but it's really unsatisfactory
 - Really complex – even for this simple an example
 - » Hard to convince yourself that this really works
 - A's code is different from B's – what if lots of threads?
 - » Code would have to be slightly different for each thread
 - While A is waiting, it is consuming CPU time
 - » This is called "busy-waiting"
- There's got to be a better way!
 - Have hardware provide higher-level primitives than atomic load & store
 - Build even higher-level programming abstractions on this hardware support

2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.54

Too Much Milk: Solution #4?

- Recall our target lock interface:
 - `acquire(&milklock)` – wait until lock is free, then grab
 - `release(&milklock)` – Unlock, waking up anyone waiting
 - These must be atomic operations – if two threads are waiting for the lock and both see it's free, only one succeeds to grab the lock
- Then, our milk problem is easy:

```
acquire(&milklock);  
if (nomilk)  
    buy milk;  
release(&milklock);
```

2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.55

Where are we going with synchronization?

Programs	Shared Programs
Higher-level API	Locks Semaphores Monitors Send/Receive
Hardware	Load/Store Disable Ints Test&Set Compare&Swap

- We are going to implement various higher-level synchronization primitives using atomic operations
 - Everything is pretty painful if only atomic primitives are load and store
 - Need to provide primitives useful at user-level

2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.56

Back to: How to Implement Locks?

- **Lock**: prevents someone from doing something
 - Lock before entering critical section and before accessing shared data
 - Unlock when leaving, after accessing shared data
 - Wait if locked
 - » Important idea: all synchronization involves waiting
 - » Should *sleep* if waiting for a long time
- Atomic Load/Store: get solution like Milk #3
 - Pretty complex and error prone
- Hardware Lock instruction
 - Is this a good idea?
 - What about putting a task to sleep?
 - » What is the interface between the hardware and scheduler?
 - Complexity?
 - » Done in the Intel 432
 - » Each feature makes HW more complex and slow



2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.57

Naïve use of Interrupt Enable/Disable

- How can we build multi-instruction atomic operations?
 - Recall: dispatcher gets control in two ways.
 - » Internal: Thread does something to relinquish the CPU
 - » External: Interrupts cause dispatcher to take CPU
 - On a uniprocessor, can avoid context-switching by:
 - » Avoiding internal events (although virtual memory tricky)
 - » Preventing external events by disabling interrupts
- Consequently, naïve Implementation of locks:

```
LockAcquire { disable Ints; }
LockRelease { enable Ints; }
```
- Problems with this approach:
 - **Can't let user do this!** Consider following:

```
LockAcquire();
While(TRUE) {}
```
 - Real-Time system—no guarantees on timing!
 - » Critical Sections might be arbitrarily long
 - What happens with I/O or other important events?
 - » “Reactor about to meltdown. Help?”



2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.58

Better Implementation of Locks by Disabling Interrupts

- Key idea: maintain a lock variable and impose mutual exclusion only during operations on that variable

```
int value = FREE;

Acquire() {
    disable interrupts;
    if (value == BUSY) {
        put thread on wait queue;
        Go to sleep();
        // Enable interrupts?
    } else {
        value = BUSY;
        enable interrupts;
    }
}

Release() {
    disable interrupts;
    if (anyone on wait queue) {
        take thread off wait queue
        Place on ready queue;
    } else {
        value = FREE;
    }
    enable interrupts;
}
```



2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.59

New Lock Implementation: Discussion

- Why do we need to disable interrupts at all?
 - Avoid interruption between checking and setting lock value
 - Otherwise two threads could think that they both have lock
- ```
Acquire() {
 disable interrupts;
 if (value == BUSY) {
 put thread on wait queue;
 Go to sleep();
 // Enable interrupts?
 } else {
 value = BUSY;
 }
 enable interrupts;
}
```
- } **Critical Section**
- Note: unlike previous solution, the critical section (inside Acquire()) is very short
    - User of lock can take as long as they like in their own critical section: doesn't impact global machine behavior
    - Critical interrupts taken in time!

2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.60

## Interrupt Re-enable in Going to Sleep

- What about re-enabling ints when going to sleep?

```
Acquire() {
 disable interrupts;
 if (value == BUSY) {
 put thread on wait queue;
 Go to sleep();
 } else {
 value = BUSY;
 }
 enable interrupts;
}
```

## Interrupt Re-enable in Going to Sleep

- What about re-enabling ints when going to sleep?

```
Acquire() {
 disable interrupts;
 if (value == BUSY) {
 put thread on wait queue;
 Go to sleep();
 } else {
 value = BUSY;
 }
 enable interrupts;
}
```

Enable Position →

- Before Putting thread on the wait queue?

## Interrupt Re-enable in Going to Sleep

- What about re-enabling ints when going to sleep?

```
Acquire() {
 disable interrupts;
 if (value == BUSY) {
 put thread on wait queue;
 Go to sleep();
 } else {
 value = BUSY;
 }
 enable interrupts;
}
```

Enable Position →

- Before Putting thread on the wait queue?
  - Release can check the queue and not wake up thread

## Interrupt Re-enable in Going to Sleep

- What about re-enabling ints when going to sleep?

```
Acquire() {
 disable interrupts;
 if (value == BUSY) {
 put thread on wait queue;
 Go to sleep();
 } else {
 value = BUSY;
 }
 enable interrupts;
}
```

Enable Position →

- Before Putting thread on the wait queue?
  - Release can check the queue and not wake up thread
- After putting the thread on the wait queue



## Interrupt Re-enable in Going to Sleep

- What about re-enabling ints when going to sleep?

```

Acquire() {
 disable interrupts;
 if (value == BUSY) {
 put thread on wait queue;
 Go to sleep();
 } else {
 value = BUSY;
 }
 enable interrupts;
}

```

Enable Position →

- Before Putting thread on the wait queue?
  - Release can check the queue and not wake up thread
- After putting the thread on the wait queue
  - Release puts the thread on the ready queue, but the thread still thinks it needs to go to sleep
  - Misses wakeup and still holds lock (deadlock!)

2/8/22

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 7.65

## Interrupt Re-enable in Going to Sleep

- What about re-enabling ints when going to sleep?

```

Acquire() {
 disable interrupts;
 if (value == BUSY) {
 put thread on wait queue;
 Go to sleep();
 } else {
 value = BUSY;
 }
 enable interrupts;
}

```

Enable Position →

- Before Putting thread on the wait queue?
  - Release can check the queue and not wake up thread
- After putting the thread on the wait queue
  - Release puts the thread on the ready queue, but the thread still thinks it needs to go to sleep
  - Misses wakeup and still holds lock (deadlock!)
- Want to put it after sleep(). But – how?

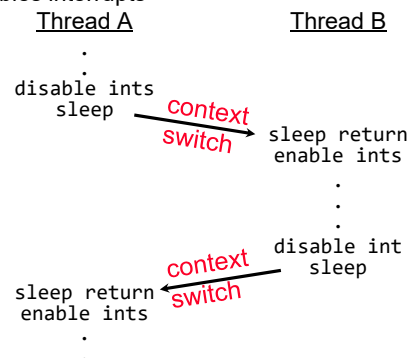
2/8/22

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 7.66

## How to Re-enable After Sleep()

- In scheduler, since interrupts are disabled when you call sleep:
  - Responsibility of the next thread to re-enable ints
  - When the sleeping thread wakes up, returns to acquire and re-enables interrupts

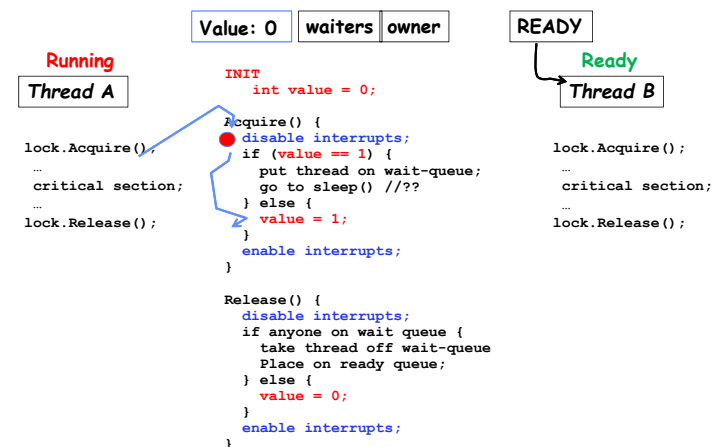


2/8/22

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 7.67

## In-Kernel Lock: Simulation

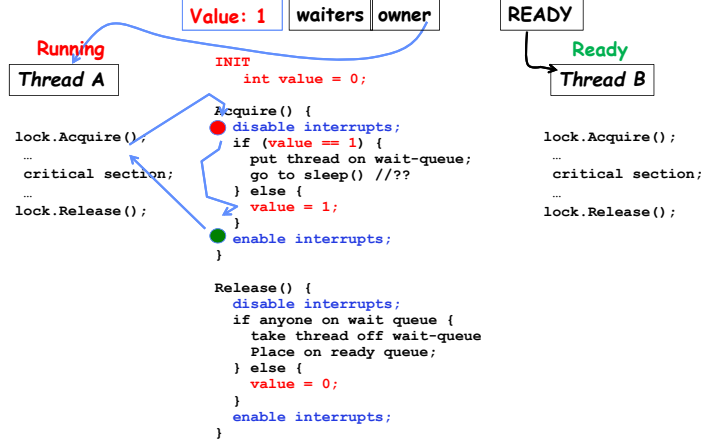


2/8/22

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 7.68

## In-Kernel Lock: Simulation

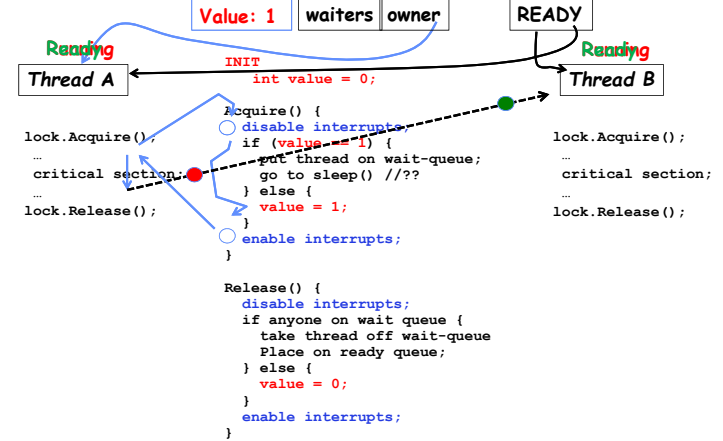


2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.69

## In-Kernel Lock: Simulation

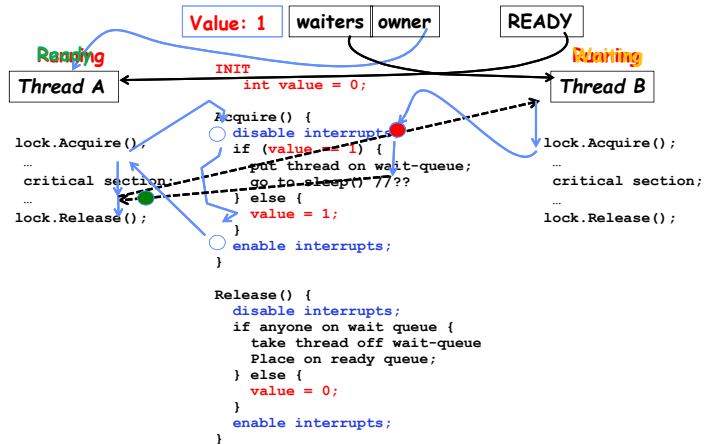


2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.70

## In-Kernel Lock: Simulation

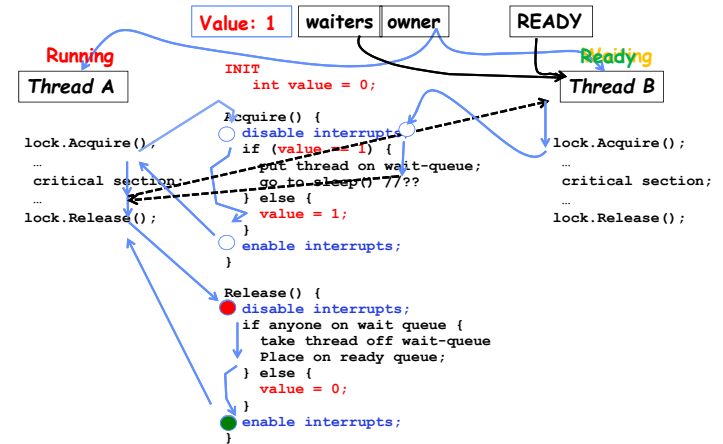


2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.71

## In-Kernel Lock: Simulation

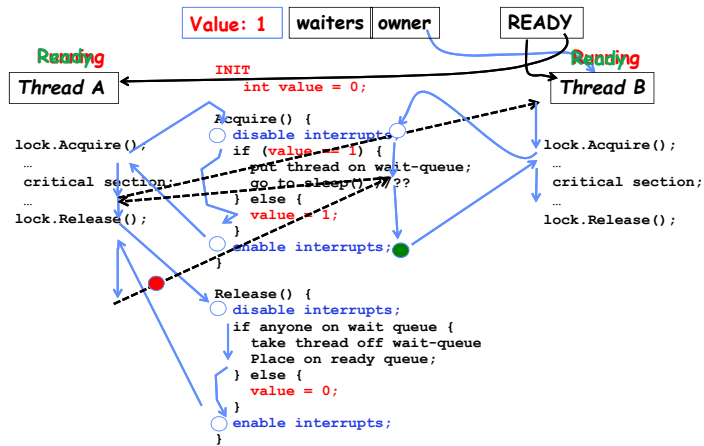


2/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 7.72

## In-Kernel Lock: Simulation



## Conclusion

- Concurrent threads introduce problems when accessing shared data
  - Programs must be insensitive to arbitrary interleavings
  - Without careful design, shared variables can become completely inconsistent
- Important concept: **Atomic Operations**
  - An operation that runs to completion or not at all
  - These are the primitives on which to construct various synchronization primitives
- Talked about hardware atomicity primitives:
  - Disabling of Interrupts, test&set, swap, compare&swap, load-locked & store-conditional
- Showed several constructions of Locks
  - Must be very careful not to waste/tie up machine resources
    - Shouldn't disable interrupts for long
    - Shouldn't spin wait for long
  - Key idea: Separate lock variable, use hardware mechanisms to protect modifications of that variable