


CS162  
Operating Systems and  
Systems Programming  
Lecture 14

Memory 2: Virtual Memory (Con't), Caching and TLBs

March 8<sup>th</sup>, 2022

Prof. Anthony Joseph and John Kubiawicz  
<http://cs162.eecs.Berkeley.edu>

Recall: General Address translation



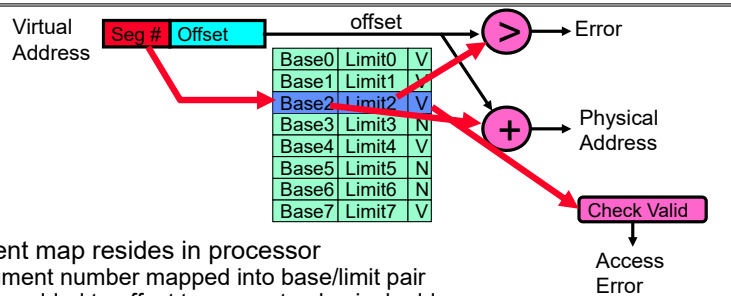
- Consequently, two views of memory:
  - View from the CPU (what program sees, virtual memory)
  - View from memory (physical memory)
- Translation box** (Memory Management Unit or MMU) converts between the two views
- Translation  $\Rightarrow$  much easier to implement protection!**
  - If task A cannot even gain access to task B's data, no way for A to adversely affect B
  - Extra benefit: every program can be linked/loaded into same region of user address space

3/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 14.2

Recall: Multi-Segment Model




- Segment map resides in processor
  - Segment number mapped into base/limit pair
  - Base added to offset to generate physical address
  - Error check catches offset out of range
- As many chunks of physical memory as entries
  - Segment addressed by portion of virtual address
  - However, could be included in instruction instead:
    - x86 Example: `mov [es:bx], ax`
- What is "V/N" (valid / not valid)?
  - Can mark segments as invalid; requires check as well

3/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 14.3

What if not all segments fit in memory?



- Extreme form of Context Switch: **Swapping**
  - To make room for next process, some or all of the previous process is moved to disk
    - Likely need to send out complete segments
  - This greatly increases the cost of context-switching
- What might be a desirable alternative?
  - Some way to keep only active portions of a process in memory at any one time**
  - Need finer granularity control over physical memory

3/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 14.4

## Problems with Segmentation

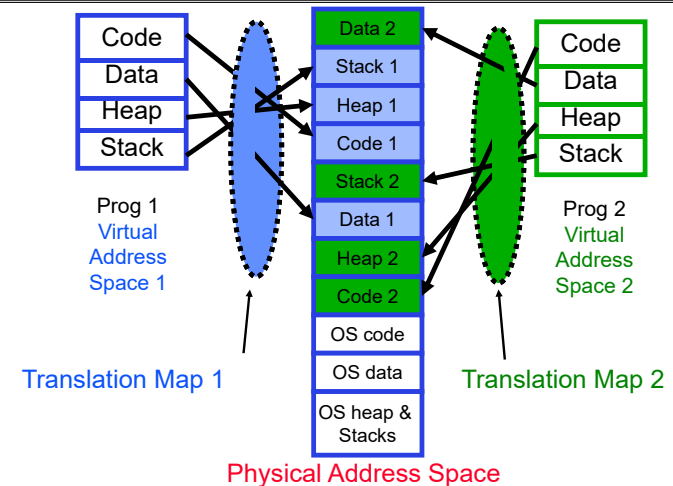
- Must fit variable-sized chunks into physical memory
- May move processes multiple times to fit everything
- Limited options for swapping to disk
- **Fragmentation**: wasted space
  - **External**: free gaps between allocated chunks
  - **Internal**: don't need all memory within allocated chunks

3/8/22

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 14.5

## Recall: General Address Translation



3/8/22

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 14.6

## Paging: Physical Memory in Fixed Size Chunks


- Solution to fragmentation from segments?
  - Allocate physical memory in **fixed size** chunks (“pages”)
  - Every chunk of physical memory is equivalent
    - » Can use simple vector of bits to handle allocation: 00110001110001101 ... 110010
    - » Each bit represents page of physical memory  
1 ⇒ allocated, 0 ⇒ free
- Should pages be as big as our previous segments?
  - No: Can lead to lots of internal fragmentation
    - » Typically have small pages (1K-16K)
  - Consequently: need multiple pages/segment

3/8/22

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 14.7

## How to Implement Simple Paging?



- Page Table (One per process)
  - Resides in physical memory
  - Contains physical page and permission for each virtual page (e.g. Valid bits, Read, Write, etc)
- Virtual address mapping
  - Offset from Virtual address copied to Physical Address
    - » Example: 10 bit offset ⇒ 1024-byte pages
  - Virtual page # is all remaining bits
    - » Example for 32-bits: 32-10 = 22 bits, i.e. 4 million entries
    - » Physical page # copied from table into physical address
  - Check Page Table bounds and permissions

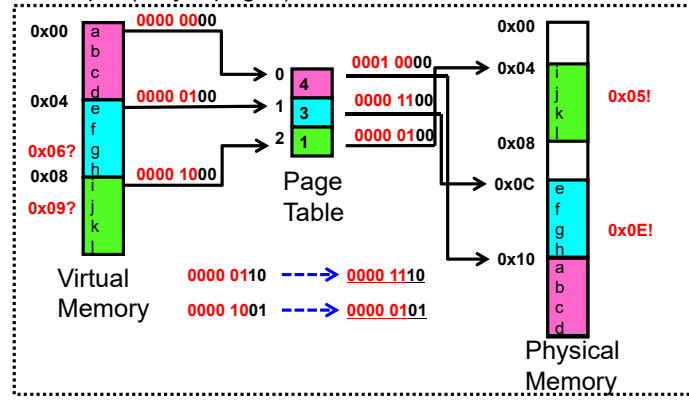
3/8/22

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 14.8

## Simple Page Table Example

Example (4 byte pages)

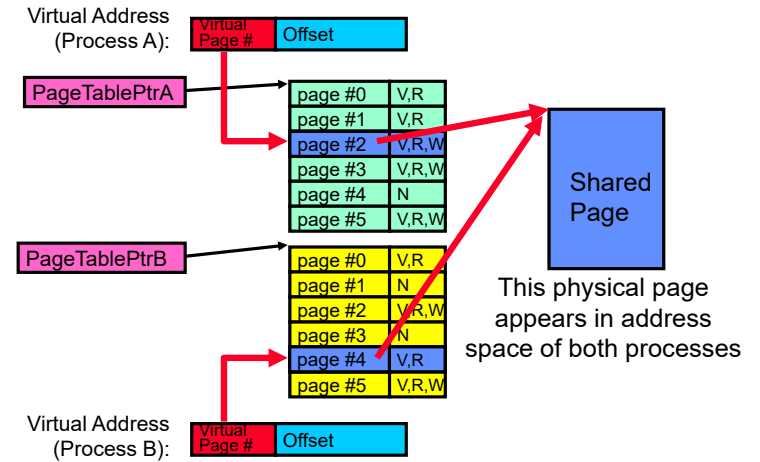


3/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 14.9

## What about Sharing?



3/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 14.10

## Where is page sharing used ?

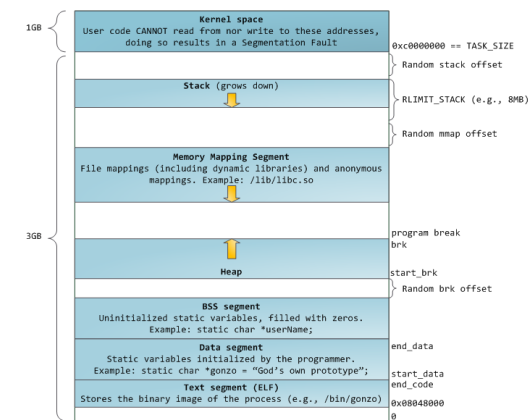
- The “kernel region” of every process has the same page table entries
  - The process cannot access it at user level
  - But on U->K switch, kernel code can access it AS WELL AS the region for THIS user
    - » What does the kernel need to do to access other user processes?
- Different processes running same binary!
  - Execute-only, but do not need to duplicate code segments
- User-level system libraries (execute only)
- Shared-memory segments between different processes
  - Can actually share objects directly between processes
    - » Must map page into same place in address space!
  - This is a limited form of the sharing that threads have within a single process

3/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 14.11

## Recall: Memory Layout for Linux 32-bit (Pre-Meltdown patch!)



<http://static.duartes.org/img/blogPosts/linuxFlexibleAddressSpaceLayout.png>

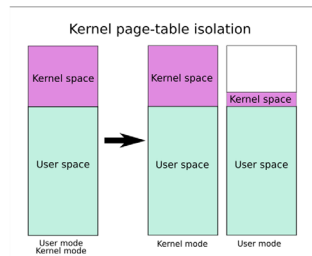
3/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 14.12

## Some simple security measures

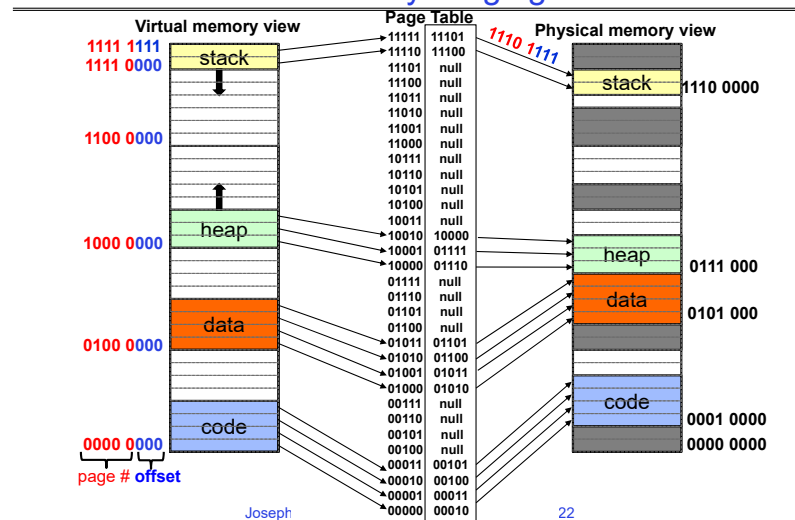
- Address Space Randomization
  - Position-Independent Code  $\Rightarrow$  can place user code anywhere in address space
    - » Random start address makes much harder for attacker to cause jump to code that it seeks to take over
  - Stack & Heap can start anywhere, so randomize placement
- Kernel address space isolation
  - Don't map whole kernel space into each process, switch to kernel page table
  - Meltdown  $\Rightarrow$  map none of kernel into user mode!



Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 14.13

## Summary: Paging

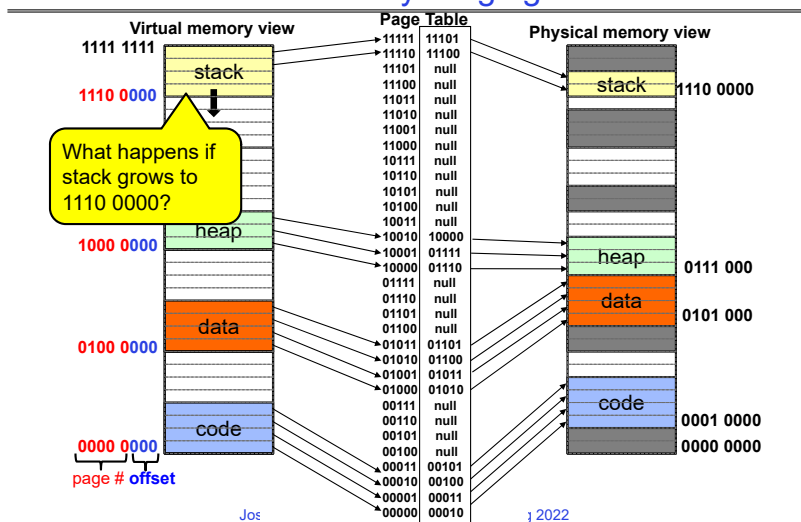


Joseph

22

Lec 14.14

## Summary: Paging

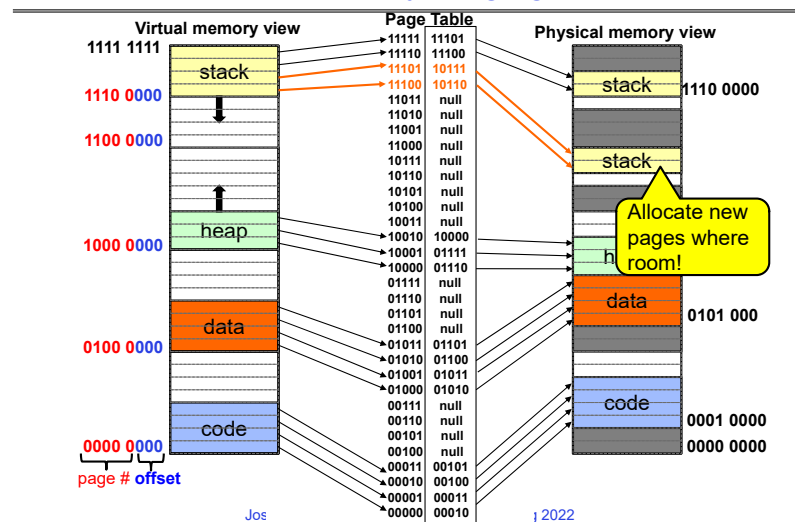


Jos

j 2022

Lec 14.15

## Summary: Paging



Jos

j 2022

Lec 14.16

## How big do things get?

- 32-bit address space  $\Rightarrow 2^{32}$  bytes (**4 GB**)
  - Note: "b" = bit, and "B" = byte
  - And *for memory*:
    - "K"(kilo) =  $2^{10} = 1024 \approx 10^3$  (But not quite!): Sometimes called "Ki" (Kibi)
    - "M"(mega) =  $2^{20} = (1024)^2 = 1,048,576 \approx 10^6$  (But not quite!): Sometimes called "Mi" (Mibi)
    - "G"(giga) =  $2^{30} = (1024)^3 = 1,073,741,824 \approx 10^9$  (But not quite!): Sometimes called "Gi" (Gibi)
- Typical page size: 4 KB
  - how many bits of the address is that? (remember  $2^{10} = 1024$ )
  - Ans –  $4KB = 4 \times 2^{10} = 2^{12} \Rightarrow 12$  bits of the address
- So how big is the simple page table for *each* process?
  - $2^{32}/2^{12} = 2^{20}$  (that's about a million entries)  $\times 4$  bytes each  $\Rightarrow$  **4 MB**
  - When 32-bit machines got started (vax 11/780, intel 80386), 16 MB was a LOT of memory
- How big is a simple page table on a 64-bit processor (x86\_64)?
  - $2^{64}/2^{12} = 2^{52}$  (that's  $4.5 \times 10^{15}$  or 4.5 exa-entries)  $\times 8$  bytes each =  **$36 \times 10^{15}$  bytes or 36 exa-bytes!!!!** This is a ridiculous amount of memory!
  - This is really a lot of space – for only the page table!!!
- The address space is *sparse*, i.e. has holes that are not mapped to physical memory
  - So, most of this space is taken up by page tables mapped to nothing

3/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 14.17

## Page Table Discussion

- What needs to be switched on a context switch?
  - Page table pointer and limit
- What provides protection here?
  - Translation (per process) *and* dual-mode!
  - Can't let process alter its own page table!
- Analysis
  - Pros
    - Simple memory allocation
    - Easy to share
  - Con: What if address space is sparse?
    - E.g., on UNIX, code starts at 0, stack starts at  $(2^{31}-1)$
    - With 1K pages, need 2 million page table entries!
  - Con: What if table really big?
    - Not all pages used all the time  $\Rightarrow$  would be nice to have working set of page table in memory
- Simple Page table is way too big!
  - Does it all need to be in memory?
  - How about multi-level paging?
  - or combining paging and segmentation


3/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 14.18

## How to Structure a Page Table

- Page Table is a *map* (function) from VPN to PPN



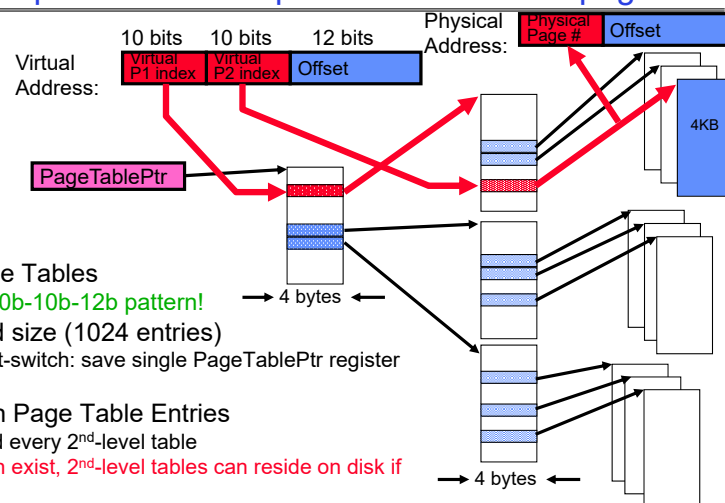
- Simple page table corresponds to a *very large* lookup table
  - VPN is index into table, each entry contains PPN
- What other map structures can you think of?
  - Trees?
  - Hash Tables?

3/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 14.19

## Fix for sparse address space: The two-level page table



- Tree of Page Tables
  - "Magic" 10b-10b-12b pattern!
- Tables fixed size (1024 entries)
  - On context-switch: save single PageTablePtr register (i.e. CR3)
- Valid bits on Page Table Entries
  - Don't need every 2<sup>nd</sup>-level table
  - Even when exist, 2<sup>nd</sup>-level tables can reside on disk if not in use

3/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 14.20

## Example: x86 classic 32-bit address translation

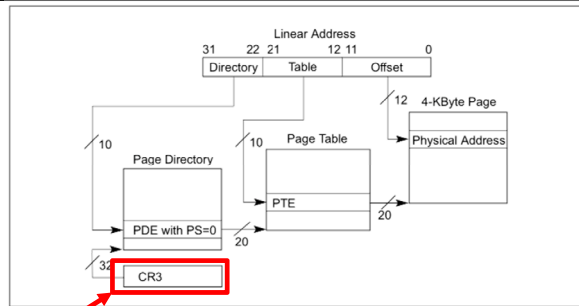


Figure 4-2. Linear-Address Translation to a 4-KByte Page using 32-Bit Paging

- Intel terminology: Top-level page-table called a "Page Directory"
  - With "Page Directory Entries"
- CR3 provides physical address of the page directory
  - This is what we have called the "PageTablePtr" in previous slides
  - Change in CR3 changes the whole translation table!


3/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 14.21

## What is in a Page Table Entry (PTE)?

- What is in a Page Table Entry (or PTE)?
  - Pointer to next-level page table or to actual page
  - Permission bits: valid, read-only, read-write, write-only
- Example: Intel x86 architecture PTE:
  - Address same format previous slide (10, 10, 12-bit offset)
  - Intermediate page tables called "Directories"



- P: Present (same as "valid" bit in other architectures)
- W: Writeable
- U: User accessible
- PWT: Page write transparent: external cache write-through
- PCD: Page cache disabled (page cannot be cached)
- A: Accessed: page has been accessed recently
- D: Dirty (PTE only): page has been modified recently
- PS: Page Size: PS=1 ⇒ 4MB page (directory only). Bottom 22 bits of virtual address serve as offset

3/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 14.22

## Examples of how to use a PTE

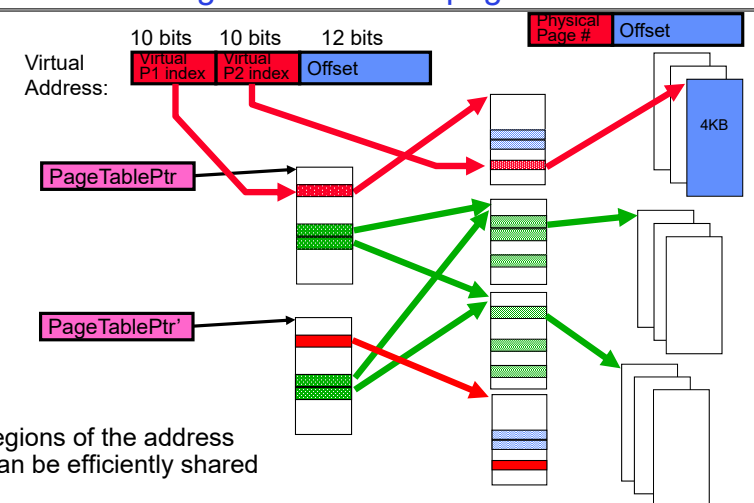
- How do we use the PTE?
  - Invalid PTE can imply different things:
    - Region of address space is actually invalid or
    - Page/directory is just somewhere else than memory
  - Validity checked first
    - OS can use other (say) 31 bits for location info
- Usage Example: Demand Paging
  - Keep only active pages in memory
  - Place others on disk and mark their PTEs invalid
- Usage Example: Copy on Write
  - UNIX fork gives copy of parent address space to child
    - Address spaces disconnected after child created
  - How to do this cheaply?
    - Make copy of parent's page tables (point at same memory)
    - Mark entries in both sets of page tables as read-only
    - Page fault on write creates two copies
- Usage Example: Zero Fill On Demand
  - New data pages must carry no information (say be zeroed)
  - Mark PTEs as invalid; page fault on use gets zeroed page
  - Often, OS creates zeroed pages in background

3/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 14.23

## Sharing with multilevel page tables



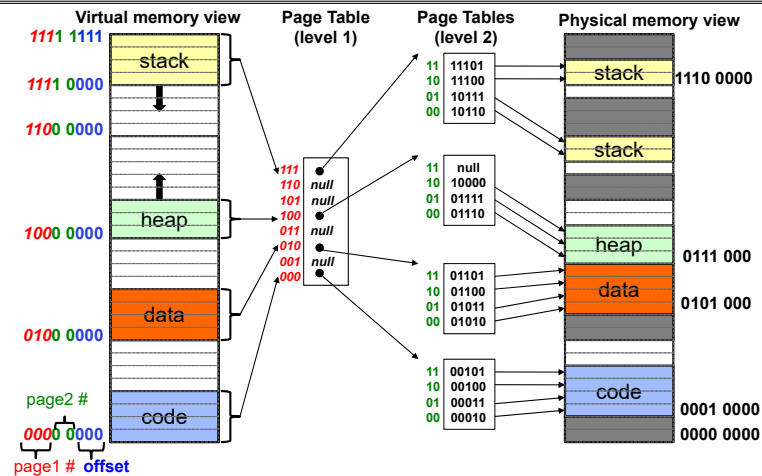
- Entire regions of the address space can be efficiently shared

3/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 14.24

## Summary: Two-Level Paging

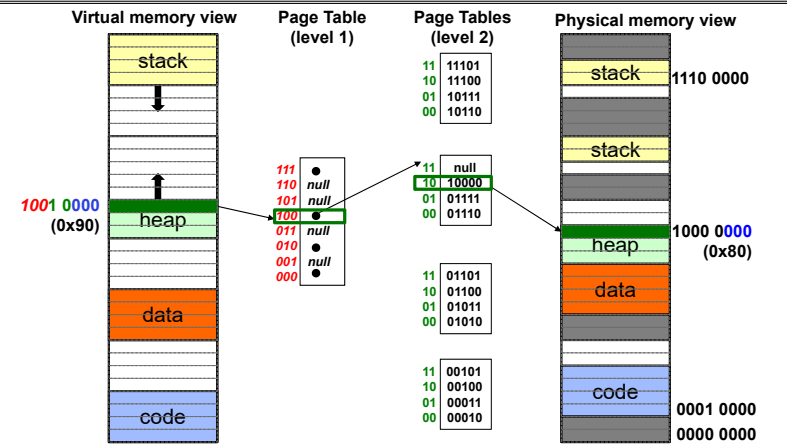


3/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 14.25

## Summary: Two-Level Paging



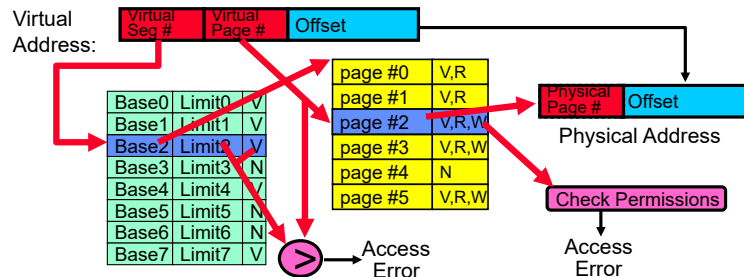
3/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 14.26

## Multi-level Translation: Segments + Pages

- What about a tree of tables?
  - Lowest level page table  $\Rightarrow$  memory still allocated with bitmap
  - Higher levels often segmented
- Could have any number of levels. Example (top segment):



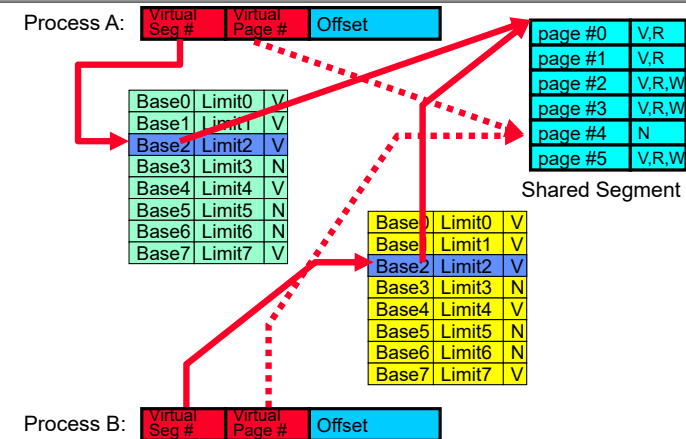
- What must be saved/restored on context switch?
  - Contents of top-level segment registers (for this example)
  - Pointer to top-level table (page table)

3/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 14.27

## What about Sharing (Complete Segment)?



3/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 14.28



## Multi-level Translation Analysis

- Pros:
  - Only need to allocate as many page table entries as we need for application
    - » In other words, sparse address spaces are easy
  - Easy memory allocation
  - Easy Sharing
    - » Share at segment or page level (need additional reference counting)
- Cons:
  - One pointer per page (typically 4K – 16K pages today)
  - Page tables need to be contiguous
    - » However, the 10b-10b-12b configuration keeps tables to exactly one page in size
  - Two (or more, if >2 levels) lookups per reference
    - » Seems very expensive!

3/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 14.29

## Recall: Dual-Mode Operation

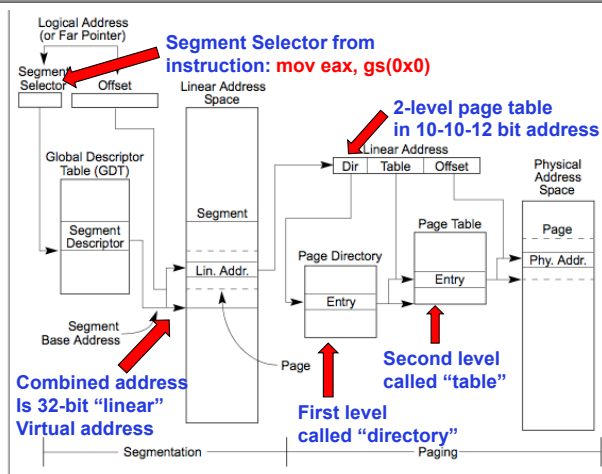
- Can a process modify its own translation tables? **NO!**
  - If it could, could get access to all of physical memory (no protection!)
- To Assist with Protection, **Hardware** provides at least two modes (Dual-Mode Operation):
  - “Kernel” mode (or “supervisor” or “protected”)
  - “User” mode (Normal program mode)
  - Mode set with bit(s) in control register only accessible in Kernel mode
  - Kernel can easily switch to user mode; User program must invoke an exception of some sort to get back to kernel mode (more in moment)
- Note that x86 model actually has more modes:
  - Traditionally, four “rings” representing priority; most OSes use only two:
    - » Ring 0 ⇒ Kernel mode, Ring 3 ⇒ User mode
    - » Called “Current Privilege Level” or CPL
  - Newer processors have additional mode for hypervisor (“Ring -1”)
- **Certain operations restricted to Kernel mode:**
  - Modifying page table base (CR3 in x86), and segment descriptor tables
    - » Have to transition into Kernel mode before you can change them!
  - Also, all page-table pages must be mapped only in kernel mode

3/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 14.30

## Making it real: X86 Memory model with segmentation (16/32-bit)




3/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 14.31

## X86 Segment Descriptors (32-bit Protected Mode)

- Segments are implicit in the instruction (e.g. code segments) or part of the instruction
  - There are 6 registers: SS, CS, DS, ES, FS, GS
- What is in a segment register?
  - A *pointer* to the actual segment description:
  - G/L selects between GDT and LDT tables (global vs local descriptor tables)
  - RPL: Requestor's Privilege Level (**RPL of CS ⇒ Current Privilege Level**)
- Two registers: GDTR/LDTR hold pointers to global/local descriptor tables in memory
  - Descriptor format (64 bits)



G: Granularity of segment [ Limit Size ] (0: 16bit, 1: 4KiB unit)

DB: Default operand size (0: 16bit, 1: 32bit)

A: Freely available for use by software

P: Segment present

DPL: Descriptor Privilege Level: Access requires  $\text{Max}(\text{CPL}, \text{RPL}) \leq \text{DPL}$

S: System Segment (0: System, 1: code or data)

3/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 14.32



## How are segments used?

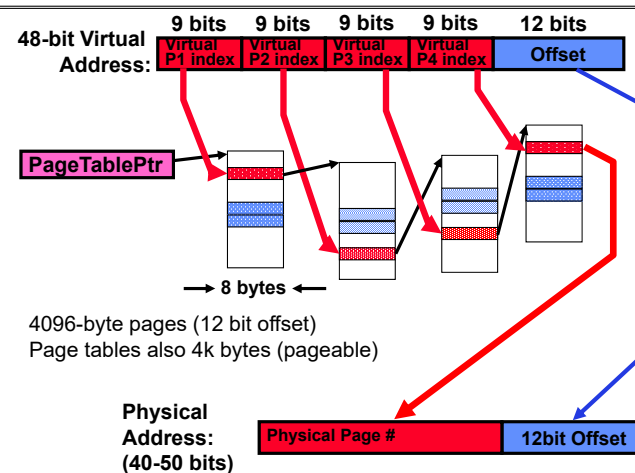
- One set of global segments (GDT) for everyone, different set of local segments (LDT) for every process
- In legacy applications (16-bit mode):
  - Segments provide protection for different components of user programs
  - Separate segments for chunks of code, data, stacks
    - » RPL of Code Segment  $\Rightarrow$  CPL (Current Privilege Level)
  - Limited to 64K segments
- Modern use in 32-bit Mode:
  - Even though there is full segment functionality, segments are set up as "flattened", i.e. every segment is 4GB in size
  - One exception: Use of GS (or FS) as a pointer to "Thread Local Storage" (TLS)
    - » A thread can make accesses to TLS like this:  
`mov eax, gs(0x0)`
- Modern use in 64-bit ("long") mode
  - Most segments (SS, CS, DS, ES) have zero base and no length limits
  - Only FS and GS retain their functionality TLS

3/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 14.33

## X86\_64: Four-level page table!



3/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 14.34

## From x86\_64 architecture specification

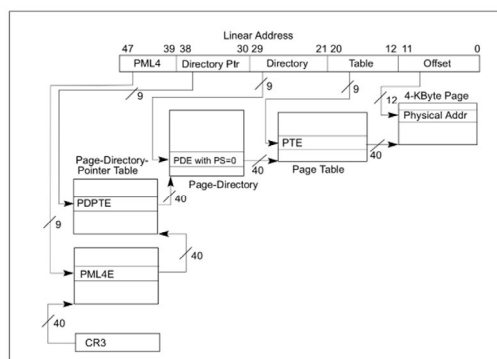


Figure 4-8. Linear-Address Translation to a 4-KByte Page using 4-Level Paging

- All current x86 processor support a 64 bit operation
- 64-bit words (so ints are 8 bytes) but 48-bit addresses

3/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 14.35

## Larger page sizes supported as well

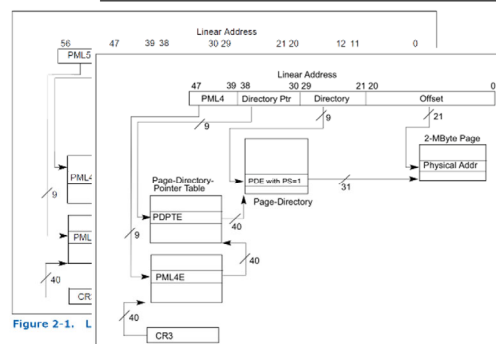


Figure 4-9. Linear-Address Translation to a 2-MByte Page using 4-Level Paging

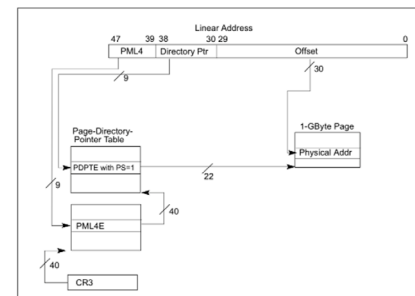


Figure 4-10. Linear-Address Translation to a 1-GByte Page using 4-Level Paging

- Larger page sizes (2MB, 1GB) make sense since memory is now cheap
  - Great for kernel, large libraries, etc
  - Use limited primarily by internal fragmentation...

3/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 14.36

## IA64: 64bit addresses: Six-level page table?!

64bit Virtual Address:	7 bits	9 bits	9 bits	9 bits	9 bits	9 bits	12 bits
	Virtual P1 index	Virtual P2 index	Virtual P3 index	Virtual P4 index	Virtual P5 index	Virtual P6 index	Offset

No!

Too slow  
Too many almost-empty tables


3/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 14.37

## Alternative: Inverted Page Table

- With all previous examples ("Forward Page Tables")
  - Size of page table is at least as large as amount of virtual memory allocated to processes
  - Physical memory may be much less
    - » Much of process space may be out on disk or not in use



- Answer: use a hash table
  - Called an "Inverted Page Table"
  - Size is independent of virtual address space
  - Directly related to amount of physical memory
  - Very attractive option for 64-bit address spaces
    - » PowerPC, UltraSPARC, IA64
- Cons:
  - Complexity of managing hash chains: Often in hardware!
  - Poor cache locality of page table

3/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 14.38

## Address Translation Comparison

	Advantages	Disadvantages
Simple Segmentation	Fast context switching (segment map maintained by CPU)	External fragmentation
Paging (Single-Level)	No external fragmentation Fast and easy allocation	Large table size (~ virtual memory) Internal fragmentation
Paged Segmentation	Table size ~ # of pages in virtual memory	Multiple memory references per page access
Multi-Level Paging	Fast and easy allocation	
Inverted Page Table	Table size ~ # of pages in physical memory	Hash function more complex No cache locality of page table

3/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 14.39

## Administrivia

- Prof Joseph's office hours: Tuesdays 1-2pm and Thursdays 12-1 (Soda 447A)
- Project 2 design docs are due Friday 3/11
- Midterm 2: Coming up on Thursday 3/17 7-9pm
  - Topics: up until Lecture 16: Scheduling, Deadlock, Address Translation, Virtual Memory, Caching, TLBs, Demand Paging
- Review Session: Wednesday 3/16 (Details TBA)

3/8/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 14.40