


CS162  
Operating Systems and  
Systems Programming  
Lecture 15

Memory 3: Caching and TLBs (Con't), Demand Paging

March 10<sup>th</sup>, 2022

Prof. Anthony Joseph and John Kubiatowicz  
<http://cs162.eecs.Berkeley.edu>

Recall: The two-level page table




- Tree of Page Tables
  - “Magic” 10b-10b-12b pattern!
- Tables fixed size (1024 entries)
  - On context-switch: save single PageTablePtr register (i.e. CR3)
- Valid bits on Page Table Entries
  - Don’t need every 2<sup>nd</sup>-level table
  - Even when exist, 2<sup>nd</sup>-level tables can reside on disk if not in use

3/10/22

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 15.2

How is the Translation Accomplished?




- The MMU must translate virtual address to physical address on:
  - Every instruction fetch
  - Every load
  - Every store
- What does the MMU need to do to translate an address?
  - 1-level Page Table
    - » Read PTE from memory, check valid, merge address
    - » Set “accessed” bit in PTE, Set “dirty bit” on write
  - 2-level Page Table
    - » Read and check first level
    - » Read, check, and update PTE
  - N-level Page Table ...
- MMU does **page table Tree Traversal** to translate each address

3/10/22

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 15.3

Where and What is the MMU ?




- The processor requests READ Virtual-Address to memory system
  - Through the MMU to the cache (to the memory)
- Some time later, the memory system responds with the data stored at the physical address (resulting from virtual → physical) translation
  - Fast on a cache hit, slow on a miss
- So what is the MMU doing?
- On every reference (I-fetch, Load, Store) read (multiple levels of) page table entries to get physical frame or FAULT
  - Through the caches to the memory
  - Then read/write the physical location

3/10/22

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 15.4

## Recall: CS61c Caching Concept



- **Cache**: a repository for copies that can be accessed more quickly than the original
  - Make frequent case fast and infrequent case less dominant
- Caching underlies many techniques used today to make computers fast
  - Can cache: memory locations, address translations, pages, file blocks, file names, network routes, etc...
- Only good if:
  - Frequent case frequent enough and
  - Infrequent case not too expensive
- Important measure: Average Access time =  
 $(\text{Hit Rate} \times \text{Hit Time}) + (\text{Miss Rate} \times \text{Miss Time})$


3/10/22

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 15.5

## Recall: In Machine Structures (eg. 61C) ...

- Caching is the key to memory system performance



Average Memory Access Time (AMAT)  
 $= (\text{Hit Rate} \times \text{HitTime}) + (\text{Miss Rate} \times \text{MissTime})$   
 Where  $\text{HitRate} + \text{MissRate} = 1$

HitRate = 90%  $\Rightarrow$  AMAT =  $(0.9 \times 1) + (0.1 \times 101) = 11$  ns

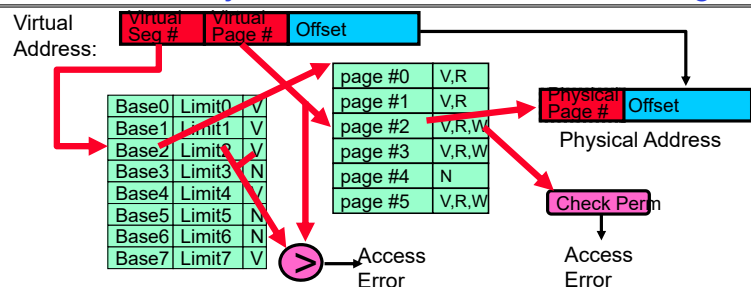
HitRate = 99%  $\Rightarrow$  AMAT =  $(0.99 \times 1) + (0.01 \times 101) = 2.01$  ns

$\text{MissTime}_{L_1}$  includes  $\text{HitTime}_{L_1} + \text{MissPenalty}_{L_1} \equiv \text{HitTime}_{L_1} + \text{AMAT}_{L_2}$

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 15.6

## Another Major Reason to Deal with Caching




- Cannot afford to translate on every access
  - At least three DRAM accesses per actual DRAM access
  - Or: perhaps I/O if page table partially on disk!
- Even worse: What if we are using caching to make memory access faster than DRAM access?
- Solution? Cache translations!
  - Translation Cache: TLB ("Translation Lookaside Buffer")

3/10/22


Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 15.7

## Why Does Caching Help? Locality!



- **Temporal Locality** (Locality in Time):
  - Keep recently accessed data items closer to processor
- **Spatial Locality** (Locality in Space):
  - Move contiguous blocks to the upper levels




Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 15.8

## Recall: Memory Hierarchy

- Caching: Take advantage of the principle of locality to:
  - Present the illusion of having as much memory as in the cheapest technology
  - Provide average speed similar to that offered by the fastest technology



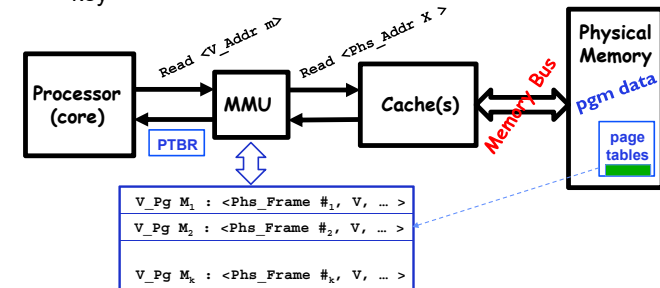
3/10/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 15.9

## How do we make Address Translation Fast?

- Cache results of recent translations !
  - Different from a traditional cache
  - Cache Page Table Entries using Virtual Page # as the key



3/10/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 15.10

## Translation Look-Aside Buffer

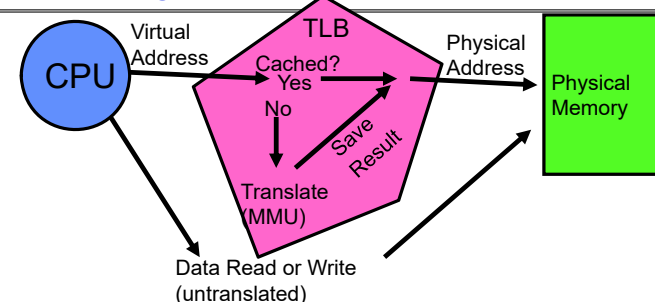
- Record recent Virtual Page # to Physical Frame # translation
- If present, have the physical address without reading any of the page tables !!!
  - Even if the translation involved multiple levels
  - Caches the end-to-end result
- Was invented by Sir Maurice Wilkes – *prior to caches*
  - When you come up with a new concept, you get to name it!
  - People realized “if it’s good for page tables, why not the rest of the data in memory?”
- On a *TLB miss*, the page tables may be cached, so only go to memory when both miss

3/10/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 15.11

## Caching Applied to Address Translation




- Question is one of page locality: does it exist?
  - Instruction accesses spend a lot of time on the same page (since accesses sequential)
  - Stack accesses have definite locality of reference
  - Data accesses have less page locality, but still some...
- Can we have a TLB hierarchy?
  - Sure: multiple levels at different sizes/speeds

3/10/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 15.12

## What kind of Cache for TLB?



- Remember all those cache design parameters and trade-offs?
  - Amount of Data =  $N * L * K$
  - Tag is portion of address that identifies line (w/o line offset)
  - Write Policy (write-thru, write-back), Eviction Policy (LRU, ...)

3/10/22

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 15.13

## How might organization of TLB differ from that of a conventional instruction or data cache?

- Let's do some review ...

3/10/22

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 15.14

## A Summary on Sources of Cache Misses


- Compulsory** (cold start or process migration, first reference): first access to a block
  - “Cold” fact of life: not a whole lot you can do about it
  - Note: If you are going to run “billions” of instruction, Compulsory Misses are insignificant
- Capacity:**
  - Cache cannot contain all blocks access by the program
  - Solution: increase cache size
- Conflict** (collision):
  - Multiple memory locations mapped to the same cache location
  - Solution 1: increase cache size
  - Solution 2: increase associativity
- Coherence** (Invalidation): other process (e.g., I/O) updates memory

3/10/22

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 15.15

## How is a Block found in a Cache?



- Block** is minimum quantum of caching
  - Data select field used to select data within block
  - Many caching applications don't have data select field
- Index** Used to Lookup Candidates in Cache
  - Index identifies the set
- Tag** used to identify actual copy
  - If no candidates match, then declare cache miss

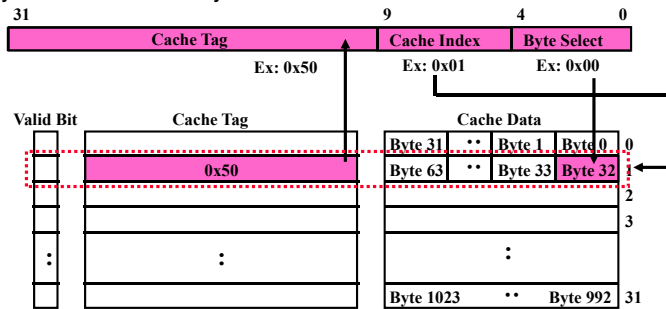
3/10/22

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 15.16

## Review: Direct Mapped Cache

- **Direct Mapped  $2^N$  byte cache:**
  - The uppermost (32 - N) bits are always the Cache Tag
  - The lowest M bits are the Byte Select (Block Size =  $2^M$ )
- Example: 1 KB Direct Mapped Cache with 32 B Blocks
  - Index chooses potential block
  - Tag checked to verify block
  - Byte select chooses byte within block




3/10/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 15.17

## Review: Set Associative Cache

- **N-way set associative:** N entries per Cache Index
  - N direct mapped caches operates in parallel
- Example: Two-way set associative cache
  - Cache Index selects a "set" from the cache
  - Two tags in the set are compared to input in parallel
  - Data is selected based on the tag result




3/10/22

Lec 15.18

## Review: Fully Associative Cache

- **Fully Associative:** Every block can hold any line
  - Address does not include a cache index
  - Compare Cache Tags of all Cache Entries in Parallel
- Example: Block Size=32B blocks
  - We need N 27-bit comparators
  - Still have byte select to choose from within block




3/10/22

Joseph & Kubiawicz CS162 © UCB Spring 2022


Lec 15.19

## Where does a Block Get Placed in a Cache?


- Example: Block 12 placed in 8 block cache




**Direct mapped:**  
block 12 can go  
only into block 4  
(12 mod 8)



**Set associative:**  
block 12 can go  
anywhere in set 0  
(12 mod 4)



**Fully associative:**  
block 12 can go  
anywhere



3/10/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 15.20

## Which block should be replaced on a miss?

- Easy for Direct Mapped: Only one possibility
- Set Associative or Fully Associative:
  - Random
  - LRU (Least Recently Used)

- Miss rates for a workload:

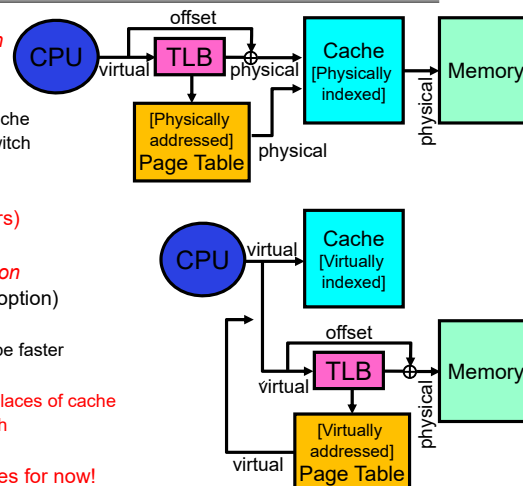
Size	2-way		4-way		8-way	
	LRU	Random	LRU	Random	LRU	Random
16 KB	5.2%	5.7%	4.7%	5.3%	4.4%	5.0%
64 KB	1.9%	2.0%	1.5%	1.7%	1.4%	1.5%
256 KB	1.15%	1.17%	1.13%	1.13%	1.12%	1.12%

## Review: What happens on a write?

- **Write through:** The information is written to both the block in the cache and to the block in the lower-level memory
- **Write back:** The information is written only to the block in the cache
  - Modified cache block is written to main memory only when it is replaced
  - Question is block clean or dirty?
- Pros and Cons of each?
  - WT:
    - » PRO: read misses cannot result in writes
    - » CON: Processor held up on writes unless writes buffered
  - WB:
    - » PRO: repeated writes not sent to DRAM  
processor not held up on writes
    - » CON: More complex  
Read miss may require writeback of dirty data

## Physically-Indexed vs Virtually-Indexed Caches


- Physically-Indexed Caches
  - Address handed to cache **after translation**
  - Page Table holds *physical* addresses
  - Benefits:
    - » Every piece of data has single place in cache
    - » Cache can stay unchanged on context switch
  - **Challenges:**
    - » TLB is in critical path of lookup!
    - Pretty Common today (e.g. x86 processors)
- Virtually-Indexed Caches
  - Address handed to cache **before translation**
  - Page Table holds *virtual* addresses (one option)
  - Benefits:
    - » TLB not in critical path of lookup, so can be faster
  - **Challenges:**
    - » Same data could be mapped in multiple places of cache
    - » May need to flush cache on context switch
- We will stick with Physically Addressed Caches for now!



## Administrivia

- Prof Joseph's office hours: Tuesdays 1-2pm and Thursdays 12-1 (Soda 447A)
- Project 2 design docs are due **TOMORROW** Friday 3/11
- Midterm 2: Coming up on next Thursday 3/17 7-9pm
  - Topics: up until Lecture 16: Scheduling, Deadlock, Address Translation, Virtual Memory, Caching, TLBs, Demand Paging
- Review Session: Wednesday 3/16 (Details TBA)

## What TLB Organization Makes Sense?



- Needs to be really fast
  - Critical path of memory access
    - In simplest view: before the cache
    - Thus, this adds to access time (reducing cache speed)
  - Seems to argue for Direct Mapped or Low Associativity
- However, needs to have very few conflicts!
  - With TLB, the Miss Time extremely high! (PT traversal)
    - Cost of Conflict (Miss Time) is high**
    - Hit Time – dictated by clock cycle**
- Thrashing**: continuous conflicts between accesses
  - What if use low order bits of virtual page number as index into TLB?
    - First page of code, data, stack may map to same entry
    - Need 3-way associativity at least?
  - What if use high order bits as index?
    - TLB mostly unused for small programs

3/10/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 15.25

## TLB organization: include protection

- How big does TLB actually have to be?
  - Usually small: 128-512 entries (larger now)
  - Not very big, can support higher associativity
- Small TLBs usually organized as fully-associative cache**
  - Lookup is by Virtual Address
  - Returns Physical Address + other info
- What happens when fully-associative is too slow?
  - Put a small (4-16 entry) direct-mapped cache in front
  - Called a “TLB Slice”
- Example for MIPS R3000:

Virtual Address	Physical Address	Dirty	Ref	Valid	Access	ASID
0xFA00	0x0003	Y	N	Y	R/W	34
0x0040	0x0010	N	Y	Y	R	0
0x0041	0x0011	N	Y	Y	R	0

3/10/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 15.26

## Example: R3000 pipeline includes TLB “stages”

MIPS R3000 Pipeline

Inst Fetch	Dcd/ Reg	ALU / E.A	Memory	Write Reg
TLB	I-Cache	RF	Operation	WB
		E.A.	TLB	D-Cache

TLB

64 entry, on-chip, fully associative, software TLB fault handler

Virtual Address Space

ASID	V. Page Number	Offset
6	20	12

0xx User segment (caching based on PT/TLB entry)  
 100 Kernel physical space, cached  
 101 Kernel physical space, uncached  
 11x Kernel virtual space

Allows context switching among  
 64 user processes without TLB flush

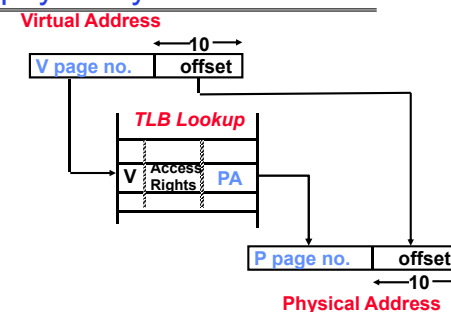
3/10/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 15.27

## Reducing translation time for physically-indexed caches

- As described, TLB lookup is in serial with cache lookup
  - Consequently, speed of TLB can impact speed of access to cache
- Machines with TLBs go one step further: overlap TLB lookup with cache access
  - Works because offset available early
  - Offset in virtual address exactly covers the “cache index” and “byte select”
  - Thus can select the cached byte(s) in parallel to perform address translation



virtual address: **Virtual Page #** **Offset**

physical address: **tag / page #** **index** **byte**

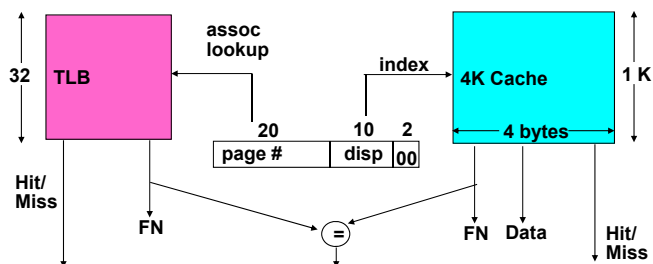
3/10/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 15.28

## Overlapping TLB & Cache Access

- Here is how this might work with a 4K cache:



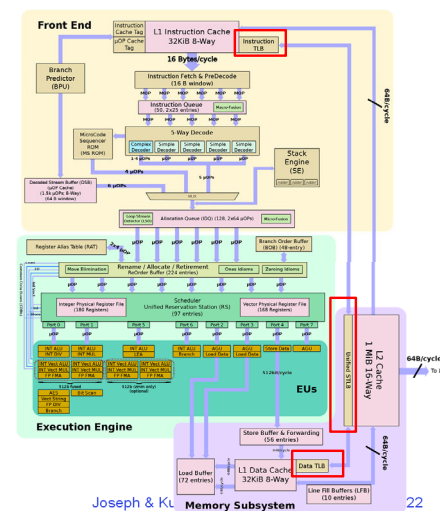
- What if cache size is increased to 8KB?
  - Overlap not complete
  - Need to do something else. See CS152/252
- Another option: Virtual Caches would make this faster
  - Tags in cache are virtual addresses
  - Translation only happens on cache misses

3/10/22

Joseph & Kubiatowicz CS 162 © UCB Spring 2022

Lec 15.29

## Current Intel x86 (Skylake, Cascade Lake)



3/10/22

Joseph & Kubiatowicz CS 162 © UCB Spring 2022

Lec 15.30

## Current Example: Memory Hierarchy

- Caches (all 64 B line size)
  - L1 I-Cache: 32 KiB/core, 8-way set assoc.
  - L1 D Cache: 32 KiB/core, 8-way set assoc., 4-5 cycles load-to-use, Write-back policy
  - L2 Cache: 1 MiB/core, 16-way set assoc., Inclusive, Write-back policy, 14 cycles latency
  - L3 Cache: 1.375 MiB/core, 11-way set assoc., shared across cores, Non-inclusive victim cache, Write-back policy, 50-70 cycles latency
- TLB
  - L1 ITLB, 128 entries; 8-way set assoc. for 4 KB pages
    - 8 entries per thread; fully associative, for 2 MiB / 4 MiB page
  - L1 DTLB 64 entries; 4-way set associative for 4 KB pages
    - 32 entries; 4-way set associative, 2 MiB / 4 MiB page translations:
    - 4 entries; 4-way associative, 1G page translations:
  - L2 STLB: 1536 entries; 12-way set assoc. 4 KiB + 2 MiB pages
    - 16 entries; 4-way set associative, 1 GiB page translations:

3/10/22

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 15.31

## What happens on a Context Switch?

- Need to do something, since TLBs map virtual addresses to physical addresses
  - Address Space just changed, so TLB entries no longer valid!
- Options?
  - Invalidate TLB: simple but might be expensive
    - What if switching frequently between processes?
  - Include ProcessID in TLB
    - This is an architectural solution: needs hardware
- What if translation tables change?
  - For example, to move page from memory to disk or vice versa...
  - Must invalidate TLB entry!
    - Otherwise, might think that page is still in memory!
  - Called "TLB Consistency"
- Aside: with Virtually-Indexed cache, need to flush cache!
  - Remember, everyone has their own version of the address "0"!

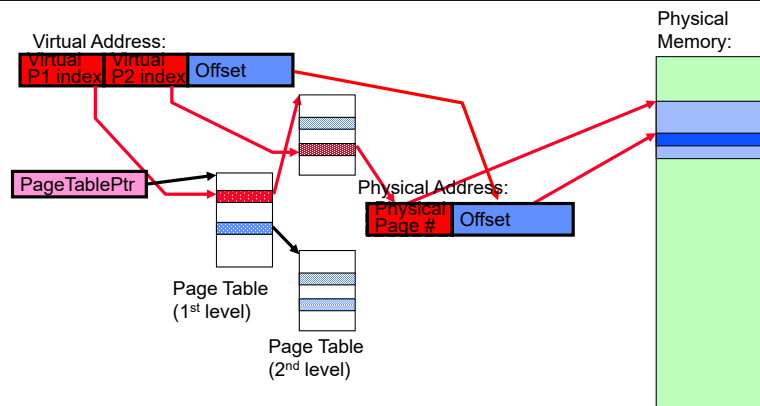
3/10/22

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 15.32



## Putting Everything Together: Address Translation

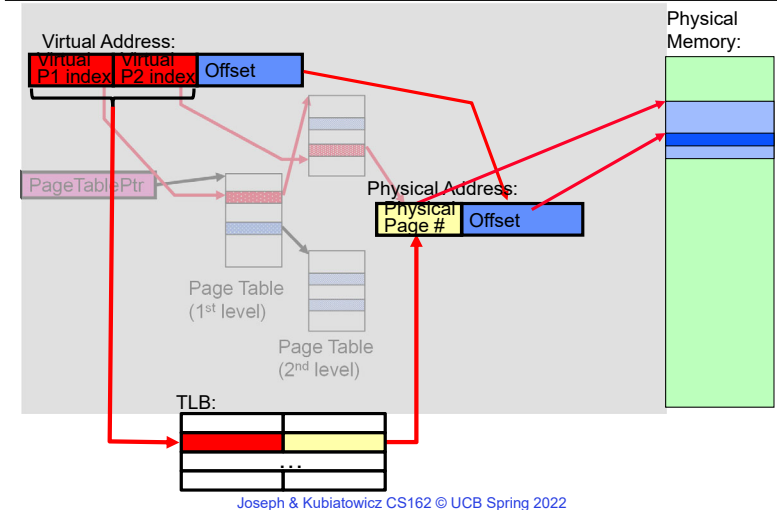


3/10/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 15.33

## Putting Everything Together: TLB

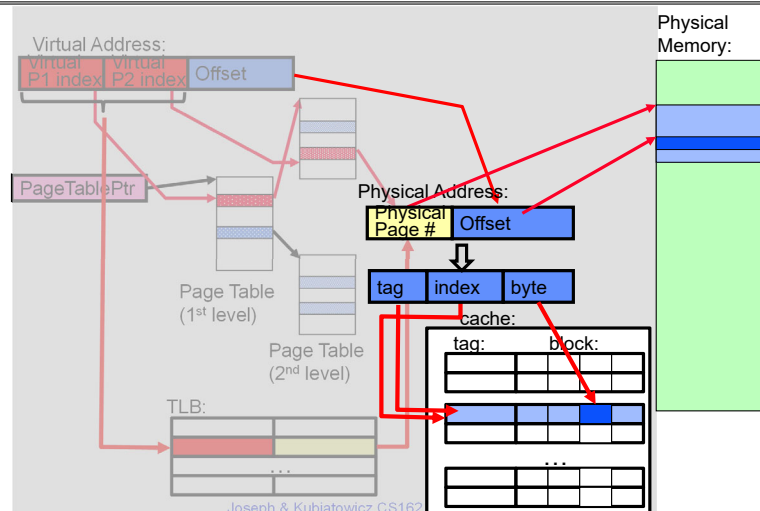


3/10/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 15.34

## Putting Everything Together: Cache



3/10/22

Joseph & Kubiawicz CS162

Lec 15.35

## Page Fault

- The Virtual-to-Physical Translation fails
  - PTE marked invalid, Priv. Level Violation, Access violation, or does not exist
  - Causes an Fault / Trap
    - » Not an interrupt because synchronous to instruction execution
  - May occur on instruction fetch or data access
  - Protection violations typically terminate the instruction
- Other Page Faults engage operating system to fix the situation and retry the instruction
  - Allocate an additional stack page, or
  - Make the page accessible - Copy on Write,
  - Bring page in from secondary storage to memory – demand paging
- Fundamental inversion of the hardware / software boundary


3/10/22

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 15.36

## Demand Paging

- Modern programs require a lot of physical memory
  - Memory per system growing faster than 25%-30%/year
- But they don't use all their memory all of the time
  - 90-10 rule: programs spend 90% of their time in 10% of their code
  - Wasteful to require all of user's code to be in memory
- Solution: use main memory as "cache" for disk




3/10/22

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 15.37

## Page Fault ⇒ Demand Paging



3/10/22

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 15.38

## Summary (1/2)

- The Principle of Locality:
  - Program likely to access a relatively small portion of the address space at any instant of time.
    - » **Temporal Locality**: Locality in Time
    - » **Spatial Locality**: Locality in Space
- Three (+1) Major Categories of Cache Misses:
  - **Compulsory Misses**: sad facts of life. Example: cold start misses.
  - **Conflict Misses**: increase cache size and/or associativity
  - **Capacity Misses**: increase cache size
  - **Coherence Misses**: Caused by external processors or I/O devices
- Cache Organizations:
  - Direct Mapped: single block per set
  - Set associative: more than one block per set
  - Fully associative: all entries equivalent

3/10/22

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 15.39

## Summary (2/2)

- "Translation Lookaside Buffer" (TLB)
  - Small number of PTEs and optional process IDs (< 512)
  - Often Fully Associative (Since conflict misses expensive)
  - On TLB miss, page table must be traversed and if located PTE is invalid, cause Page Fault
  - On change in page table, TLB entries must be invalidated
- Demand Paging: Treating the DRAM as a cache on disk
  - Page table tracks which pages are in memory
  - Any attempt to access a page that is not in memory generates a page fault, which causes OS to bring missing page into memory
- Replacement policies
  - FIFO: Place pages on queue, replace page at end
  - MIN: Replace page that will be used farthest in future
  - LRU: Replace page used farthest in past

3/10/22

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 15.40