

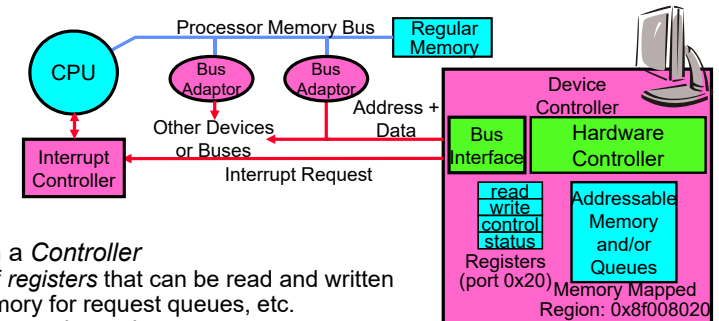
CS162 Operating Systems and Systems Programming Lecture 18

General I/O (Con't), Storage Devices, Performance

March 31st, 2022

Prof. Anthony Joseph and John Kubitowicz
<http://cs162.eecs.Berkeley.edu>

Recall: How does the Processor Talk to the Device?



- CPU interacts with a *Controller*
 - Contains a set of *registers* that can be read and written
 - May contain memory for request queues, etc.
- Processor accesses registers in two ways:
 - **Port-Mapped I/O**: in/out instructions
 - » Example from the Intel architecture: out 0x21, AL
 - **Memory-mapped I/O**: load/store instructions
 - » Registers/memory appear in physical address space
 - » I/O accomplished with load and store instructions

3/31/2022

Joseph & Kubitowicz CS162 © UCB Spring 2022

Lec 18.2

Port-Mapped I/O in Pintos Speaker Driver

Pintos: `devices/speaker.c`


```
13 /* Sets the PC speaker to emit a tone at the given FREQUENCY, in
14 Hz. */
15 void
16 speaker_on (int frequency)
17 {
18     if (frequency >= 20 && frequency <= 20000)
19     {
20         /* Set the timer channel that's connected to the speaker to
21          output a square wave at the given FREQUENCY, then
22          connect the timer channel output to the speaker. */
23         enum intr_level old_level = intr_disable ();
24         pit_configure_channel (2, 3, frequency);
25         outb (SPEAKER_PORT_GATE, inb (SPEAKER_PORT_GATE) | SPEAKER_GATE_ENABLE);
26         intr_set_level (old_level);
27     }
28     else
29     {
30         /* FREQUENCY is outside the range of normal human hearing.
31          Just turn off the speaker. */
32         speaker_off ();
33     }
34 }
35
36 /* Turn off the PC speaker, by disconnecting the timer channel's
37 output from the speaker. */
38 void
39 speaker_off (void)
40 {
41     intr_level old_level = intr_disable ();
42     outb (SPEAKER_PORT_GATE, inb (SPEAKER_PORT_GATE) & ~SPEAKER_GATE_ENABLE);
43     intr_set_level (old_level);
44 }
```

Pintos: `threads/io.h`

```
7 /* Reads and returns a byte from PORT. */
8 static inline uint8_t
9 inb (uint16_t port)
10 {
11     /* See [IA32-v2a] "IN". */
12     uint8_t data;
13     asm volatile ("inb %d, %b0" : "=a" (data) : "Nd" (port));
14     return data;
15 }
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64 /* Writes byte DATA to PORT. */
65 static inline void
66 outb (uint16_t port, uint8_t data)
67 {
68     /* See [IA32-v2b] "OUT". */
69     asm volatile ("outb %b0, %d" : : "a" (data), "Nd" (port));
70 }
```

Example: Memory-Mapped Display Controller

- Memory-Mapped:
 - Hardware maps control registers and display memory into physical address space
 - » Addresses set by HW jumpers or at boot time
 - Simply writing to display memory (also called the “frame buffer”) changes image on screen
 - » Addr: 0x8000F000 — 0x8000FFFF
 - Writing graphics description to cmd queue
 - » Say enter a set of triangles describing some scene
 - » Addr: 0x80010000 — 0x8001FFFF
 - Writing to the command register may cause on-board graphics hardware to do something
 - » Say render the above scene
 - » Addr: 0x0007F004
- Can protect with address translation



3/31/2022

Joseph & Kubitowicz CS162 © UCB Spring 2022

Lec 18.3

3/31/2022

Joseph & Kubitowicz CS162 © UCB Spring 2022

Lec 18.4

There's more than just a CPU in there!




3/31/2022

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 18.5

Chip-scale Features of 2015 x86 (Sky Lake)



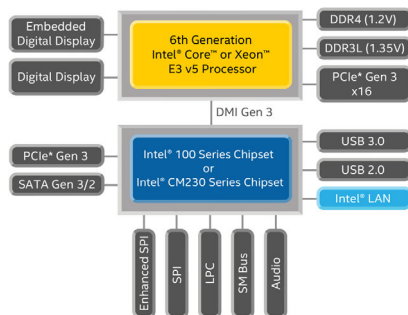
3/31/2022

- Significant pieces:
 - Four OOO cores with deeper buffers
 - » Intel MPX (Memory Protection Extensions)
 - » Intel SGX (Software Guard Extensions)
 - » Issue up to 6 μ -ops/cycle
 - GPU, System Agent (Mem, Fast I/O)
 - Large shared L3 cache with on-chip ring bus
 - » 2 MB/core instead of 1.5 MB/core
 - » High-BW access to L3 Cache
- Integrated I/O
 - Integrated memory controller (IMC)
 - » Two independent channels of DRAM
 - High-speed PCI-Express (for Graphics cards)
 - Direct Media Interface (DMI) Connection to PCH (Platform Control Hub)

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 18.6

Sky Lake I/O: PCH



Sky Lake
System Configuration

- **Platform Controller Hub**
 - Connected to processor with proprietary bus
 - » Direct Media Interface
- Types of I/O on PCH:
 - USB, Ethernet
 - Thunderbolt 3
 - Audio, BIOS support
 - More PCI Express (lower speed than on Processor)
 - SATA (for Disks)

3/31/2022

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 18.7

Operational Parameters for I/O

- Data granularity: Byte vs. Block
 - Some devices provide single byte at a time (e.g., keyboard)
 - Others provide whole blocks (e.g., disks, networks, etc.)
- Access pattern: Sequential vs. Random
 - Some devices must be accessed sequentially (e.g., tape)
 - Others can be accessed "randomly" (e.g., disk, cd, etc.)
 - » Fixed overhead to start transfers
 - Some devices require continual monitoring
 - Others generate interrupts when they need service
- Transfer Mechanism: Programmed IO and DMA

3/31/2022

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 18.8

Transferring Data To/From Controller

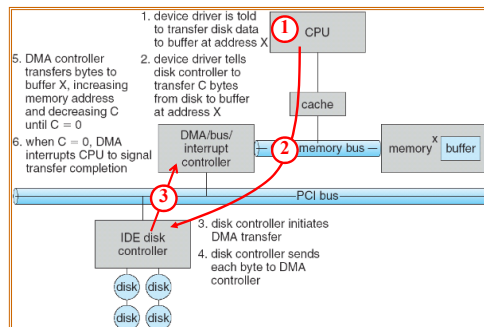
- **Programmed I/O:**

- Each byte transferred via processor in/out or load/store
- Pro: Simple hardware, easy to program
- Con: Consumes processor cycles proportional to data size

- **Direct Memory Access:**

- Give controller access to memory bus
- Ask it to transfer data blocks to/from memory directly

- Sample interaction with DMA controller (from OSC book):



3/31/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 18.9

Transferring Data To/From Controller

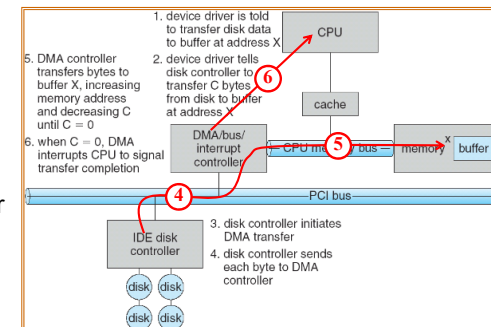
- **Programmed I/O:**

- Each byte transferred via processor in/out or load/store
- Pro: Simple hardware, easy to program
- Con: Consumes processor cycles proportional to data size

- **Direct Memory Access:**

- Give controller access to memory bus
- Ask it to transfer data blocks to/from memory directly

- Sample interaction with DMA controller (from OSC book):



3/31/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 18.10

I/O Device Notifying the OS

- The OS needs to know when:

- The I/O device has completed an operation
- The I/O operation has encountered an error

- **I/O Interrupt:**

- Device generates an interrupt whenever it needs service
- Pro: handles unpredictable events well
- Con: interrupts relatively high overhead

- **Polling:**

- OS periodically checks a device-specific status register
 - » I/O device puts completion information in status register
- Pro: low overhead
- Con: may waste many cycles on polling if infrequent or unpredictable I/O operations

- **Actual devices combine both polling and interrupts**

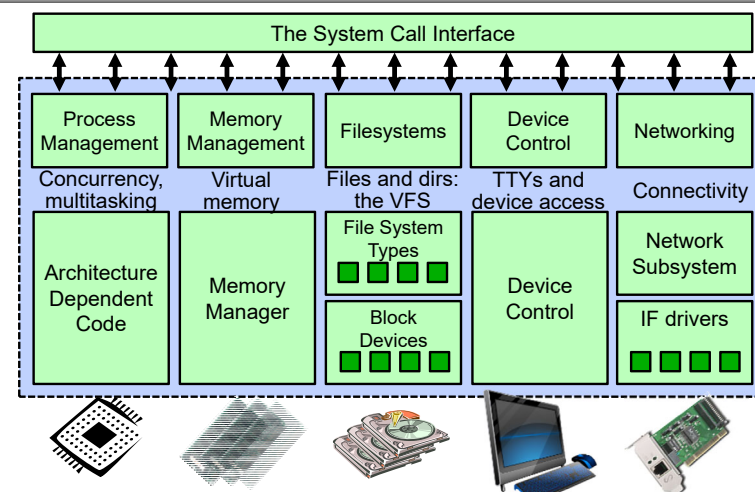
- For instance – High-bandwidth network adapter:
 - » Interrupt for first incoming packet
 - » Poll for following packets until hardware queues are empty

3/31/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 18.11

Kernel Device Structure



3/31/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 18.12

Recall: Device Drivers

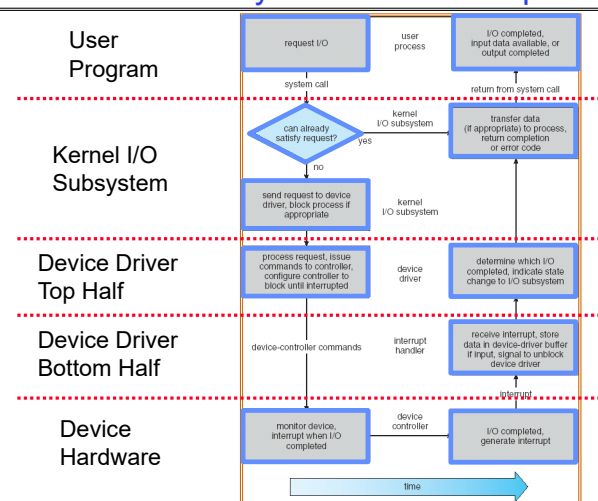
- **Device Driver:** Device-specific code in the kernel that interacts directly with the device hardware
 - Supports a standard, internal interface
 - Same kernel I/O system can interact easily with different device drivers
 - Special device-specific configuration supported with the `ioctl()` system call
- Device Drivers typically divided into two pieces:
 - Top half: accessed in call path from system calls
 - » implements a set of **standard, cross-device calls** like `open()`, `close()`, `read()`, `write()`, `ioctl()`, `strategy()`
 - » This is the kernel's interface to the device driver
 - » Top half will *start* I/O to device, may put thread to sleep until finished
 - Bottom half: run as interrupt routine
 - » Gets input or transfers next block of output
 - » May wake sleeping threads if I/O now complete

3/31/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 18.13

Recall: Life Cycle of An I/O Request



3/31/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 18.14

The Goal of the I/O Subsystem

- Provide Uniform Interfaces, Despite Wide Range of Different Devices
 - This code works on many different devices:


```
FILE fd = fopen("/dev/something", "rw");
for (int i = 0; i < 10; i++) {
    fprintf(fd, "Count %d\n", i);
}
close(fd);
```
 - Why? Because code that controls devices ("device driver") implements standard interface
- We will try to get a flavor for what is involved in actually controlling devices in rest of lecture
 - Can only scratch surface!

3/31/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 18.15

Want Standard Interfaces to Devices

- **Block Devices:** e.g. disk drives, tape drives, DVD-ROM
 - Access blocks of data
 - Commands include `open()`, `read()`, `write()`, `seek()`
 - Raw I/O or file-system access
 - Memory-mapped file access possible
- **Character Devices:** e.g. keyboards, mice, serial ports, some USB devices
 - Single characters at a time
 - Commands include `get()`, `put()`
 - Libraries layered on top allow line editing
- **Network Devices:** e.g. Ethernet, Wireless, Bluetooth
 - Different enough from block/character to have own interface
 - Unix and Windows include **socket** interface
 - » Separates network protocol from network operation
 - » Includes `select()` functionality
 - Usage: pipes, FIFOs, streams, queues, mailboxes

3/31/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 18.16

How Does User Deal with Timing?

- **Blocking Interface:** "Wait"
 - When request data (e.g. `read()` system call), put process to sleep until data is ready
 - When write data (e.g. `write()` system call), put process to sleep until device is ready for data
- **Non-blocking Interface:** "Don't Wait"
 - Returns quickly from read or write request with count of bytes successfully transferred
 - Read may return nothing, write may write nothing
- **Asynchronous Interface:** "Tell Me Later"
 - When request data, take pointer to user's buffer, return immediately; later kernel fills buffer and notifies user
 - When send data, take pointer to user's buffer, return immediately; later kernel takes data and notifies user

3/31/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 18.17

Storage Devices


- **Magnetic disks**
 - Storage that rarely becomes corrupted
 - Large capacity at low cost
 - Block level random access (except for SMR – later!)
 - Slow performance for random access
 - Better performance for sequential access
- **Flash memory**
 - Storage that rarely becomes corrupted
 - Capacity at intermediate cost (5-20x disk)
 - Block level random access
 - Good performance for reads; worse for random writes
 - Erasure requirement in large blocks
 - Wear patterns issue

3/31/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 18.18

Hard Disk Drives (HDDs)



IBM/Hitachi Microdrive



Read/Write Head Side View

IBM Personal Computer/AT (1986)
30 MB hard disk - \$500
30-40ms seek time
0.7-1 MB/s (est.)


3/31/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 18.19

The Amazing Magnetic Disk

- **Unit of Transfer: Sector**
 - Ring of sectors form a track
 - Stack of tracks form a cylinder
 - Heads position on cylinders
- **Disk Tracks ~ 1μm (micron) wide**
 - Wavelength of light is ~ 0.5μm
 - Resolution of human eye: 50μm
 - 100K tracks on a typical 2.5" disk
- **Separated by unused guard regions**
 - Reduces likelihood neighboring tracks are corrupted during writes (still a small non-zero chance)




3/31/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 18.20

The Amazing Magnetic Disk

- Track length varies across disk
 - Outside: More sectors per track, higher bandwidth
 - Disk is organized into regions of tracks with same # of sectors/track
 - Only outer half of radius is used
 - » Most of the disk area in the outer regions of the disk
- Disks so big that some companies (like Google) reportedly only use part of disk for active data
 - Rest is archival data



3/31/2022


Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 18.21


Shingled Magnetic Recording (SMR)

- Overlapping tracks yields greater density, capacity
- Restrictions on writing, complex DSP for reading

Conventional Writes



SMR Writes




3/31/2022

Joseph & Kubiatowicz CS162 © UCB Spring 2022


Lec 18.22

Review: Magnetic Disks

- Cylinders**: all the tracks under the head at a given point on all surfaces
- Read/write data is a three-stage process:
 - **Seek time**: position the head/arm over the proper track
 - **Rotational latency**: wait for desired sector to rotate under r/w head
 - **Transfer time**: transfer a block of bits (sector) under r/w head



$$\text{Disk Latency} = \text{Queueing Time} + \text{Controller time} + \text{Seek Time} + \text{Rotation Time} + \text{Xfer Time}$$



3/31/2022

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 18.23

Typical Numbers for Magnetic Disk

Parameter	Info/Range
Space/Density	Space: 18TB (Seagate), 9 platters, in 3½ inch form factor! Areal Density: ≥ 1 Terabit/square inch! (PMR, Helium, ...)
Average Seek Time	Typically 4-6 milliseconds
Average Rotational Latency	Most laptop/desktop disks rotate at 3600-7200 RPM (16-8 ms/rotation). Server disks up to 15,000 RPM. Average latency is halfway around disk so 4-8 milliseconds
Controller Time	Depends on controller hardware
Transfer Time	Typically 50 to 270 MB/s. Depends on: <ul style="list-style-type: none"> • Transfer size (usually a sector): 512B – 1KB per sector • Rotation speed: 3600 RPM to 15000 RPM • Recording density: bits per inch on a track • Diameter: ranges from 1 in to 5.25 in
Cost	Used to drop by a factor of two every 1.5 years (or faster), now slowing down

3/31/2022

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 18.24

Disk Performance Example

- Assumptions:
 - Ignoring queuing and controller times for now
 - Avg seek time of 5ms,
 - 7200RPM \Rightarrow Time for rotation: $60000 \text{ (ms/min)} / 7200 \text{ (rev/min)} \approx 8\text{ms}$
 - Transfer rate of 50MByte/s, block size of 4Kbyte \Rightarrow
 $4096 \text{ bytes} / 50 \times 10^6 \text{ (bytes/s)} = 81.92 \times 10^{-6} \text{ sec} \approx 0.082 \text{ ms}$ for 1 sector
- Read block from random place on disk:
 - Seek (5ms) + Rot. Delay (4ms) + Transfer (0.082ms) = 9.082ms
 - Approx 9ms to fetch/put data: $4096 \text{ bytes} / 9.082 \times 10^{-3} \text{ s} \approx 451\text{KB/s}$
- Read block from random place in same cylinder:
 - Rot. Delay (4ms) + Transfer (0.082ms) = 4.082ms
 - Approx 4ms to fetch/put data: $4096 \text{ bytes} / 4.082 \times 10^{-3} \text{ s} \approx 1.03\text{MB/s}$
- Read next block on same track:
 - Transfer (0.082ms): $4096 \text{ bytes} / 0.082 \times 10^{-3} \text{ s} \approx 50\text{MB/sec}$
- Key to using disk effectively (especially for file systems) is to minimize seek and rotational delays

3/31/2022

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 18.25

Lots of Intelligence in the Controller


- Sectors contain sophisticated error correcting codes
 - Disk head magnet has a field wider than track
 - Hide corruptions due to neighboring track writes
- Sector sparing
 - Remap bad sectors transparently to spare sectors on the same surface
- Slip sparing
 - Remap all sectors (when there is a bad sector) to preserve sequential behavior
- Track skewing
 - Sector numbers offset from one track to the next, to allow for disk head movement for sequential ops

3/31/2022

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 18.26

Hard Drive Prices over Time



3/31/2022

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 18.27

Example of Current HDDs

- Seagate Exos X18 (2020)
 - 18 TB hard disk
 - 9 platters, 18 heads
 - Helium filled: reduce friction and power
 - 4.16ms average seek time
 - 4096 byte physical sectors
 - 7200 RPMs
 - Dual 6 Gbps SATA /12Gbps SAS interface
 - 270MB/s MAX transfer rate
 - Cache size: 256MB
 - Price: \$ 562 (~ \$0.03/GB)
- IBM Personal Computer/AT (1986)
 - 30 MB hard disk
 - 30-40ms seek time
 - 0.7-1 MB/s (est.)
 - Price: \$500 (\$17K/GB, 340,000x more expensive !!)




3/31/2022

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 18.28

Solid State Disks (SSDs)

- 1995 – Replace rotating magnetic media with non-volatile memory (battery backed DRAM)
- 2009 – Use NAND Multi-Level Cell (2 or 3-bit/cell) flash memory
 - Sector (4 KB page) addressable, but stores 4-64 “pages” per memory block
 - Trapped electrons distinguish between 1 and 0
- No moving parts (no rotate/seek motors)
 - Eliminates seek and rotational delay (0.1-0.2ms access time)
 - Very low power and lightweight
 - Limited “write cycles”
- Rapid advances in capacity and cost ever since!

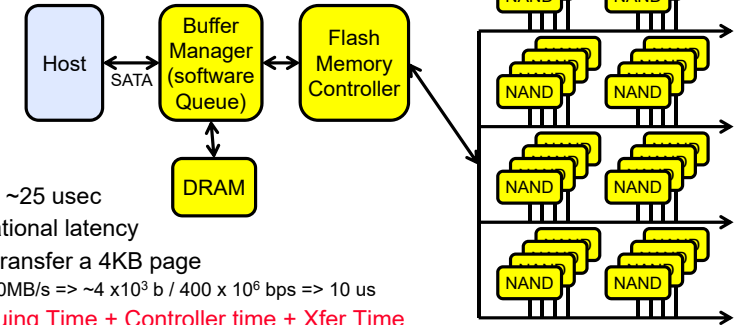


3/31/2022

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 18.29

SSD Architecture – Reads



Read 4 KB Page: ~25 usec

- No seek or rotational latency
- Transfer time: transfer a 4KB page
 - » SATA: 300-600MB/s => $\sim 4 \times 10^3 \text{ b} / 400 \times 10^6 \text{ bps} \Rightarrow 10 \text{ us}$
- **Latency = Queuing Time + Controller time + Xfer Time**
- **Highest Bandwidth:** Sequential OR Random reads


3/31/2022

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 18.30

SSD Architecture – Writes

- Writing data is complex! (~200µs – 1.7ms)
 - Can only write empty pages in a block
 - Erasing a block takes ~1.5ms
 - Controller maintains pool of empty blocks by coalescing used pages (read, erase, write), also reserves some % of capacity
- Rule of thumb: writes 10x reads, erasure 10x writes



https://en.wikipedia.org/wiki/Solid-state_drive

3/31/2022

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 18.31

SSD Architecture – Writes

- SSDs provide same interface as HDDs to OS – read and write chunk (4KB) at a time
- But can only overwrite data 256KB at a time!
- Why not just erase and rewrite new version of entire 256KB block?
 - Erasure is very slow (milliseconds)
 - Each block has a finite lifetime, can only be erased and rewritten about 10K times
 - Heavily used blocks likely to wear out quickly

3/31/2022

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 18.32

Solution – Two Systems Principles

1. Layer of Indirection
 - Maintain a *Flash Translation Layer (FTL)* in SSD
 - Map virtual block numbers (which OS uses) to physical page numbers (which flash mem. controller uses)
 - **Can now freely relocate data w/o OS knowing**
2. Copy on Write
 - Don't overwrite a page when OS updates its data
 - Instead, write new version in a free page
 - Update FTL mapping to point to new location

3/31/2022

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 18.33

Flash Translation Layer

- No need to erase and rewrite entire 256KB block when making small modifications
- SSD controller can assign mappings to spread workload across pages
 - *Wear Levelling*
- What to do with old versions of pages?
 - *Garbage Collection* in background
 - Erase blocks with old pages, add to free list


3/31/2022

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 18.34

Some “Current” (large) 3.5in SSDs

- Seagate Exos SSD: 15.36TB (2017)
 - Dual 12Gb/s interface
 - Seq reads 860MB/s
 - Seq writes 920MB/s
 - Random Reads (IOPS): 102K
 - Random Writes (IOPS): 15K
 - Price (Amazon): \$5495 (\$0.36/GB)
- Nimbus SSD: 100TB (2019)
 - Dual port: 12Gb/s interface
 - Seq reads/writes: 500MB/s
 - Random Read Ops (IOPS): 100K
 - *Unlimited writes for 5 years!*
 - Price: ~ \$40K? (\$0.4/GB)
 - » However, 50TB drive costs \$12500 (\$0.25/GB)



3/31/2022

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 18.35

Amusing calculation: Is a full Kindle heavier than an empty one?

- Actually, “Yes”, but not by much
- Flash works by trapping electrons:
 - So, erased state lower energy than written state
- Assuming that:
 - Kindle has 4GB flash
 - ½ of all bits in full Kindle are in high-energy state
 - High-energy state about 10^{-15} joules higher
 - Then: Full Kindle is 1 attogram (10^{-18} gram) heavier (Using $E = mc^2$)
- Of course, this is less than most sensitive scale can measure (it can measure 10^{-9} grams)
- Of course, this weight difference overwhelmed by battery discharge, weight from getting warm,
- **Source: John Kubiatowicz (New York Times, Oct 24, 2011)**

3/31/2022

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 18.36

SSD Summary

- Pros (vs. hard disk drives):
 - Low latency, high throughput (eliminate seek/rotational delay)
 - No moving parts:
 - » Very light weight, low power, silent, very shock insensitive
 - Read at memory speeds (limited by controller and I/O bus)
- Cons
 - Small storage (0.1-0.5x disk), expensive (3-20x disk)
 - » Hybrid alternative: combine small SSD with large HDD

3/31/2022

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 18.37

SSD Summary

- Pros (vs. hard disk drives):
 - Low latency, high throughput (eliminate seek/rotational delay)
 - No moving parts:
 - » Very light weight, low power, silent, very shock insensitive
 - Read at memory speeds (limited by controller and I/O bus)
- Cons
 - ~~Small storage (0.1-0.5x disk), expensive (3-20x disk)~~
 - » Hybrid alternative: combine small SSD with large HDD
 - Asymmetric block write performance: read pg/erase/write pg
 - » Controller garbage collection (GC) algorithms have major effect on performance
 - Limited drive lifetime
 - » 1-10K writes/page for MLC NAND
 - » Avg failure rate is 6 years, life expectancy is 9-11 years
- These are changing rapidly!

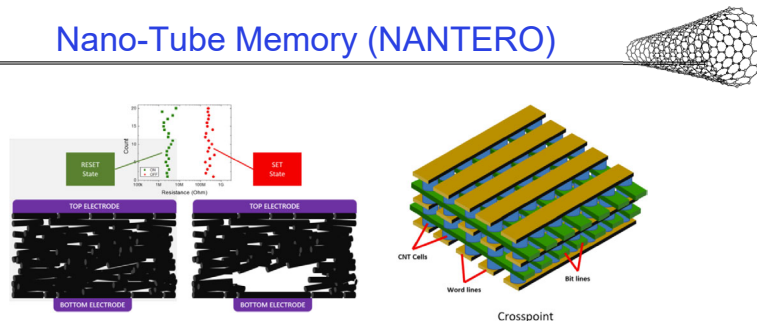
No
longer
true!

3/31/2022

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 18.38

Nano-Tube Memory (NANTERO)



- Yet another possibility: Nanotube memory
 - NanoTubes between two electrodes, slight conductivity difference between ones and zeros
 - No wearout!
- Better than DRAM?
 - Speed of DRAM, no wearout, non-volatile!
 - Nantero promises 512Gb/die for 8Tb/chip! (with 16 die stacking)

3/31/2022

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 18.39

Ways of Measuring Performance: Times (s) and Rates (op/s)

- **Latency** – time to complete a task
 - Measured in units of time (s, ms, us, ..., hours, years)
- **Response Time** - time to initiate and operation and get its response
 - Able to issue one that *depends* on the result
 - Know that it is done (anti-dependence, resource usage)
- **Throughput** or **Bandwidth** – rate at which tasks are performed
 - Measured in units of things per unit time (ops/s, GFLOP/s)
- **Start up or "Overhead"** – time to initiate an operation
- Most I/O operations are roughly linear in b bytes
 - $\text{Latency}(b) = \text{Overhead} + b/\text{TransferCapacity}$
- Performance???
 - Operation time (4 mins to run a mile...)
 - Rate (mph, mpg, ...)

3/31/2022

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 18.40

Example: Overhead in Fast Network


- Consider a 1 Gb/s link ($B = 125 \text{ MB/s}$) with startup cost $S = 1 \text{ ms}$

- Latency: $L(b) = S + \frac{b}{B}$

- Effective Bandwidth:

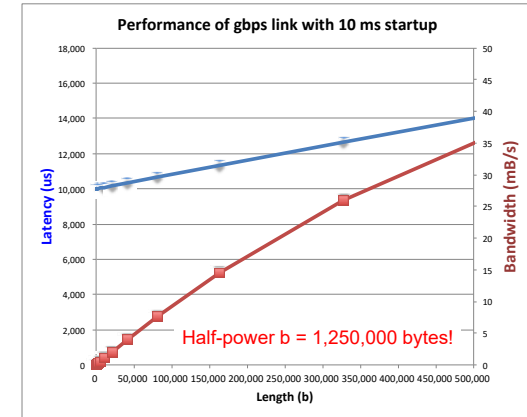
$$E(b) = \frac{b}{S + \frac{b}{B}} = \frac{B \cdot b}{B \cdot S + b} = \frac{B}{\frac{B \cdot S}{b} + 1}$$

- Half-power Bandwidth: $E(b) = \frac{B}{2}$
- For this example, half-power bandwidth occurs at $b = 125 \text{ KB}$



Example: 10 ms Startup Cost (e.g., Disk)

- Half-power bandwidth at $b = 1.25 \text{ MB}$
- Large startup cost can degrade effective bandwidth
- Amortize it by performing I/O in larger blocks

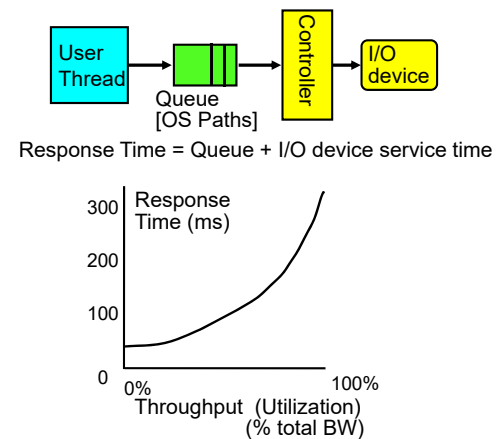


What Determines Peak BW for I/O?

- Bus Speed
 - PCI-X: $1064 \text{ MB/s} = 133 \text{ MHz} \times 64 \text{ bit (per lane)}$
 - ULTRA WIDE SCSI: 40 MB/s
 - Serial Attached SCSI & Serial ATA & IEEE 1394 (firewire): $1.6 \text{ Gb/s full duplex}$ (200 MB/s)
 - USB 3.0 – 5 Gb/s
 - Thunderbolt 3 – 40 Gb/s
- Device Transfer Bandwidth
 - Rotational speed of disk
 - Write / Read rate of NAND flash
 - Signaling rate of network link
- Whatever is the bottleneck in the path...

Overall Performance for I/O Path

- Performance of I/O subsystem
 - Metrics: Response Time, Throughput
 - Effective BW = transfer size / response time
 - Contributing factors to latency:
 - » Software paths (can be loosely modeled by a queue)
 - » Hardware controller
 - » I/O device service time
- Queuing behavior:
 - Can lead to big increases of latency as utilization increases
 - Solutions?



Sequential Server Performance




- Single sequential “server” that can deliver a task in time L operates at rate $\leq \frac{1}{L}$ (on average, in steady state, ...)
- $L = 10 \text{ ms} \rightarrow B = 100 \text{ op/s}$
- $L = 2 \text{ yr} \rightarrow B = 0.5 \text{ op/yr}$
- Applies to a processor, a disk drive, a person, a TA, ...

3/31/2022

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 18.45

Single Pipelined Server




- Single pipelined server of k stages for tasks of length L (i.e., time L/k per stage) delivers at rate $\leq k/L$.
- $L = 10 \text{ ms}, k = 4 \rightarrow B = 400 \text{ op/s}$
- $L = 2 \text{ yr}, k = 2 \rightarrow B = 1 \text{ op/yr}$

3/31/2022

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 18.46

Example Systems “Pipelines”



- Anything with queues between operational process behaves roughly “pipeline like”
- Important difference is that “initiations” are decoupled from processing
 - May have to queue up a burst of operations
 - Not synchronous and deterministic like in 61C

3/31/2022

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 18.47

Multiple Servers



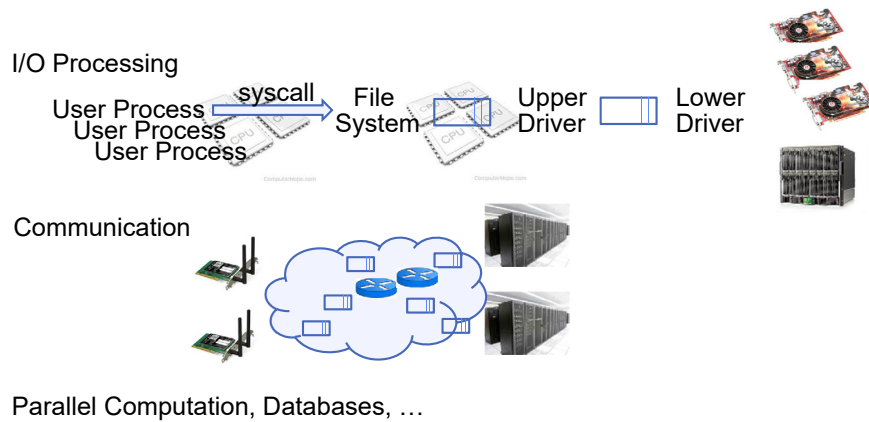
- k servers handling tasks of length L delivers at rate $\leq k/L$.
- $L = 10 \text{ ms}, k = 4 \rightarrow B = 400 \text{ op/s}$
- $L = 2 \text{ yr}, k = 2 \rightarrow B = 1 \text{ op/yr}$
- In 61C you saw multiple processors (cores)
 - Systems present lots of multiple parallel servers
 - Often with lots of queues

3/31/2022

Joseph & Kubiatowicz CS162 © UCB Spring 2022

Lec 18.48

Example Systems “Parallelism”



3/31/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 18.49

Conclusion (1/2)

- Notification mechanisms
 - Interrupts
 - Polling: Report results through status register that processor looks at periodically
- Device drivers interface to I/O devices
 - Provide clean Read/Write interface to OS above
 - Manipulate devices through PIO, DMA & interrupt handling
 - Three types: block, character, and network
- Direct Memory Access (DMA)
 - Permit devices to directly access memory
 - Free up processor from transferring every byte

3/31/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 18.50

Conclusion (2/2)

- Disk Performance:
 - Queuing time + Controller + Seek + Rotational + Transfer
 - Rotational latency: on average $\frac{1}{2}$ rotation
 - Transfer time: spec of disk depends on rotation speed and bit storage density
- Devices have complex interaction and performance characteristics
 - Response time (Latency) = Queue + Overhead + Transfer
 - » Effective BW = $BW \cdot T/(S+T)$
 - HDD: Queuing time + controller + seek + rotation + transfer
 - SSD: Queuing time + controller + transfer (erasure & wear)
- Systems (e.g., file system) designed to optimize performance and reliability
 - Relative to performance characteristics of underlying device
- Next time: Bursts & High Utilization introduce queuing delays
- Next time: Queuing Latency:
 - M/M/1 and M/G/1 queues: simplest to analyze
 - As utilization approaches 100%, latency $\rightarrow \infty$
$$T_q = T_{ser} \times \frac{1}{2}(1+C) \times \rho/(1-\rho)$$

3/31/2022

Joseph & Kubiawicz CS162 © UCB Spring 2022

Lec 18.51