

# EECS 182      Deep Neural Networks

## Fall 2022      Discussion Worksheet

# Discussion 1

This discussion serves as an introduction to neural networks. You will examine a simple example of gradient descent, and build intuition for the ReLU nonlinearity through visualizations in the Jupyter [notebook]. This notebook takes a long time to do the initial network training! You should begin the training process then return to the theory part of this discussion while you wait.

## 1. Gradient Descent Mechanics

Gradient descent is the primary algorithm to search optimal parameters for our models. Typically, we want to solve optimization problems stated as

$$\min_{\theta \in \Theta} \mathcal{L}(f_{\theta}, \mathcal{D})$$

where  $\mathcal{L}$  are differentiable functions. In this example, we look at a simple supervised learning problem where given a dataset  $\mathcal{D} = \{(x_i, y_i)\}_i^N$ , we want to find the optimal parameters  $\theta$  that minimizes some loss. We consider different models for learning the mapping from input to output, and examine the behavior of gradient descent for each model.

- (a) The simplest parametric model entails learning a single-parameter constant function, where we set  $\hat{y}_i = \theta$ . We wish to find

$$\hat{\theta}_{\text{const}} = \min_{\theta \in \mathbb{R}} \mathcal{L}(f_{\theta}, \mathcal{D}) = \min_{\theta \in \mathbb{R}} \frac{1}{N} \sum_{i=1}^N (y_i - \theta)^2$$

- What is the gradient of  $\mathcal{L}$  w.r.t  $\theta$ ?
  - What is the optimal value of  $\theta$ ?
  - Write the gradient descent update.
  - *Stochastic Gradient Descent (SGD)* is an alternative optimization algorithm, where instead of using all  $N$  samples, we use single sample per optimization step to update the model. What is the contribution of each data-point to the full gradient update?
- (b) Instead of constant functions, we now consider a single-parameter linear model  $\hat{y}(x_i) = \theta x_i$ , where we search for  $\theta$  such that

$$\hat{\theta} = \min_{\theta \in \mathbb{R}} \frac{1}{N} \sum_{i=1}^N (y_i - \theta x_i)^2$$

- Repeat questions from part (a).
- Do all points get the same *vote* in the update? Why?

## 2. ReLU SGD Visualization

Work through the notebook to explore how a simple network with ReLU non-linearities adapts to model a function with SGD updates. training the networks takes 5-10 minutes depending on whether you run locally or the server, so you should start the training process (run through the train all layers cell) then return to the theory part of the discussion while training occurs.

As you walk through the notebook, pay attention to how the slopes and elbows of the ReLU functions change during training and how they impact the shape of the final learned function.

### Contributors:

- Kumar Krishna Agrawal.
- Anant Sahai.
- Ashwin Pananjady.
- Josh Sanz.
- Yichao Zhou.