


EECS 182      Deep Neural Networks  
Fall 2022      Anant Sahai

# Midterm Review: Basics

Consider the simple neural network that takes a scalar real input, has 1 hidden layer with  $k$  units in it and a sigmoid nonlinearity for those units, and an output linear (affine) layer to predict a scalar output. We can algebraically write any function that it represents as

$$y_{pred} = C\sigma(Ax + \mathbf{b}) + d$$

The  $\sigma(\cdot)$  represents an arbitrary nonlinearity, with derivative  $\sigma'(\cdot)$ . Where  $x \in \mathbb{R}$ ,  $A \in \mathbb{R}^{k \times 1}$ ,  $\mathbf{b} \in \mathbb{R}^{k \times 1}$ ,  $C \in \mathbb{R}^{1 \times k}$ ,  $d \in \mathbb{R}$ , and  $y_{pred} \in \mathbb{R}$ . We can write it as  $y_{pred} = C\sigma(\mathbf{z}) + d$ , where  $\mathbf{z} = Ax + \mathbf{b}$  and the nonlinearity is applied element-wise. We have the true label  $y_{true}$  for each  $x$ , and we use the L2 Loss  $L(y_{true}, y_{pred}) = (y_{true} - y_{pred})^2$ .



1. (a) Consider the sigmoid nonlinearity function  $\sigma(z) = \frac{1}{1+e^{-z}}$ . Show that  $\frac{d}{dz}\sigma(z) = \sigma(z)(1 - \sigma(z))$

(b) Calculate  $\frac{\partial L}{\partial d}$

(c) Calculate  $\frac{\partial L}{\partial C_i}$

(d) Calculate  $\frac{\partial L}{\partial b_i}$

(e) Calculate  $\frac{\partial L}{\partial A_i}$

(f) Write the gradient-descent update rule for  $\mathbf{b}_{t+1}$  with learning rate  $\alpha$ .

2. Given the Regularized Objective function:

$$\operatorname{argmin}_{\mathbf{x}} \|\mathbf{A}\mathbf{x} - \mathbf{b}\|^2 + \lambda \|\mathbf{x}\|^2$$

Use vector calculus to find the closed form solution for  $\mathbf{x}$ . Interpret what this means in terms of the singular values.

3. Consider a simple neural network that spits out 1-dim values after a nonlinearity. These values for a batch are  $\{1, 7, 7, 9\}$ . What is the output of running batchnorm with this data and  $\gamma = 1$  and  $\beta = 0$ . In other words, standardize the data to have mean 0 and variance 1.

4. Consider a simplified batchnorm layer where we don't actually divide by standard deviation, instead we just de-mean our data before scaling it by  $\gamma$  and shifting it by  $\beta$ , then passing it to the next layer. That is, we calculate our mini-batch mean  $\mu$ , then simply let  $\hat{x}_i = x_i - \mu$ , and  $y_i = \gamma \hat{x}_i + \beta$  is passed onto the next layer. Assume batchsize of  $m$ . If our final loss function is  $L$ , Calculate  $\frac{\partial L}{\partial x_i}$  in terms of  $\frac{\partial L}{\partial y_j}$  for  $j = 1, \dots, m$ ,  $\gamma$ ,  $\beta$ , and  $m$ .