CS 188: Artificial Intelligence

Hidden Markov Models

Instructor: Anca Dragan --- University of California, Berkeley
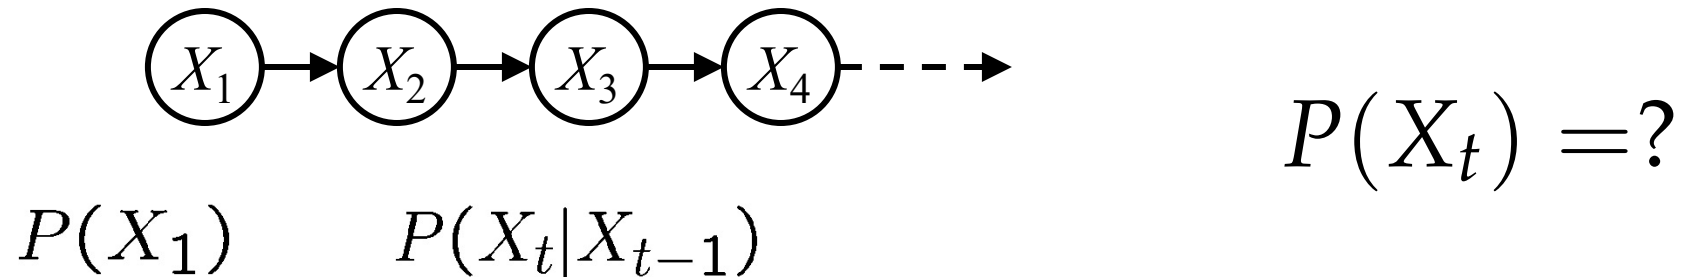
Reasoning over Time or Space

o Often, we want to reason about a sequence of observations

- o Speech recognition
- o Robot localization
- o User attention
- o Medical monitoring

o Need to introduce time (or space) into our models

Markov Models

o Value of X at a given time is called the state

$$P(X_1) \qquad P(X_t | X_{t-1})$$

$$P(X_t) = ?$$

o Parameters: called transition probabilities or dynamics, specify how the state evolves over time (also, initial state probabilities)
o Stationarity assumption: transition probabilities the same at all times
o Same as MDP transition model, but no choice of action
o A (growable) BN: We can always use generic BN reasoning on it if we truncate the chain at a fixed length
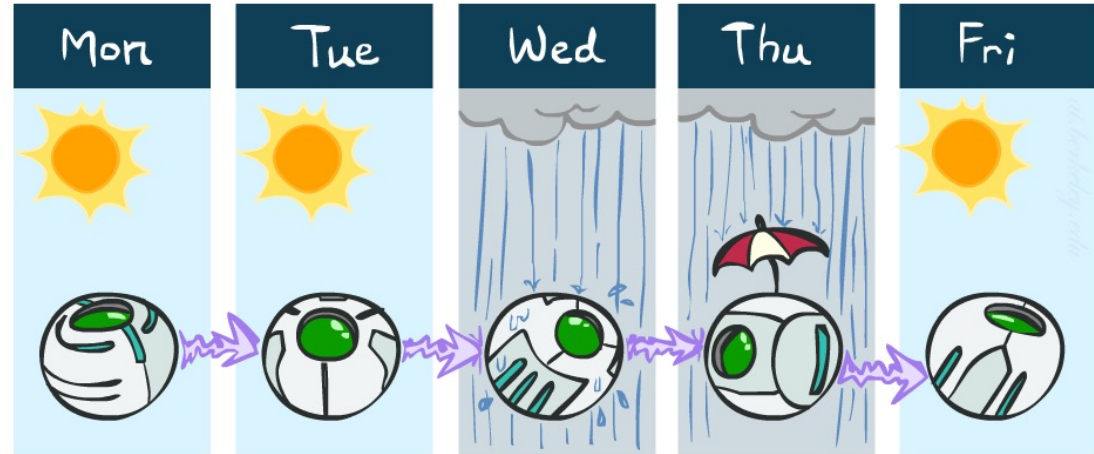
Markov Assumption: Conditional Independence
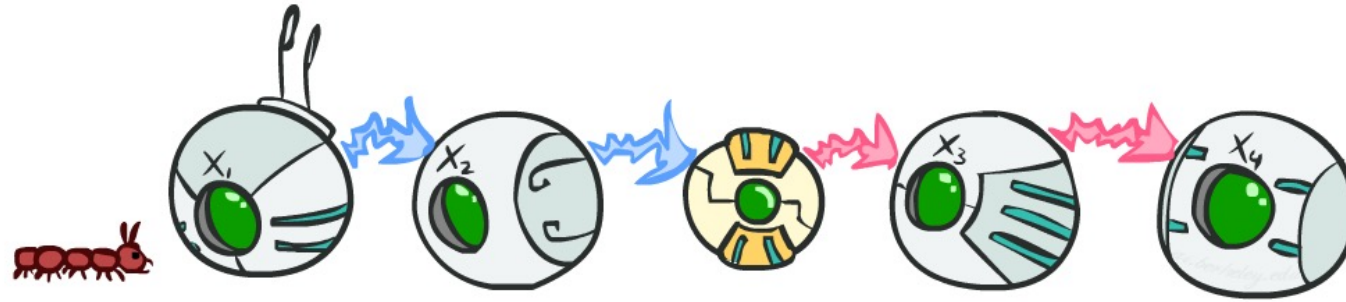
o **Basic conditional independence:**
  - o Past and future independent given the present
  - o Each time step only depends on the previous
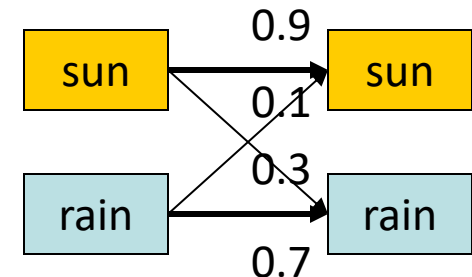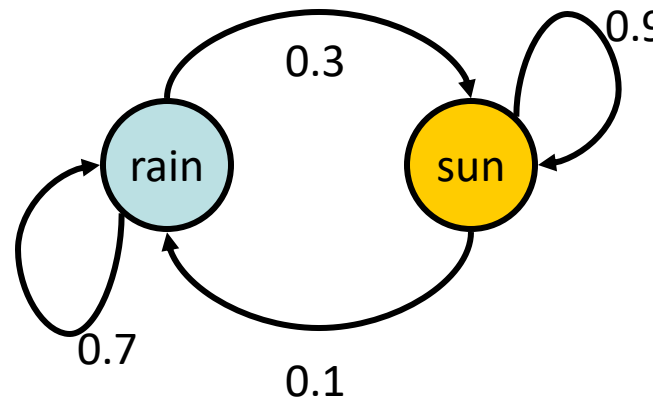 - o This is called the (first order) Markov property

Example Markov Chain: Weather

○ States: X = {rain, sun}

■ Initial distribution: 1.0 sun

■ CPT $P(X_t \mid X_{t-1})$:

Two new ways of representing the same CPT

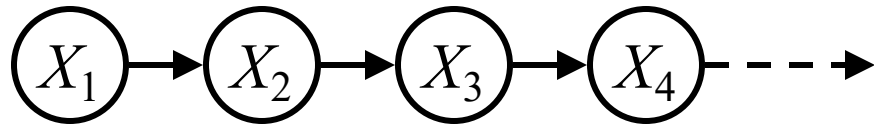| X_{t-1}	X_t	$P(X_t\|X_{t-1})$
sun	sun	0.9
sun	rain	0.1
rain	sun	0.3
rain	rain	0.7

Example Markov Chain: Weather

o Initial distribution: 1.0 sun
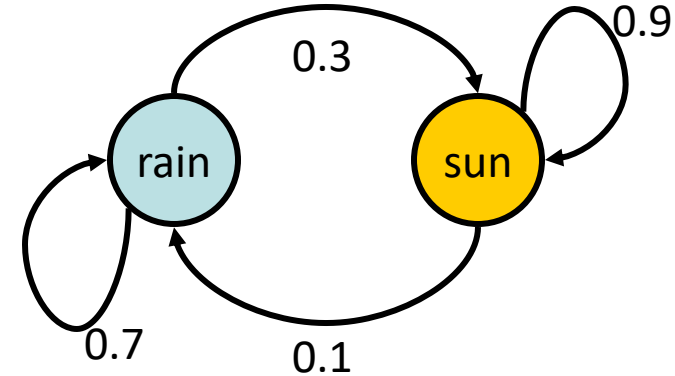
o What is the probability distribution after one step?

$$P(X_2 = sun) = \sum_{x_1} P(x_1, X_2 = sun) = \sum_{x_1} P(X_2 = sun | x_1) P(x_1)$$

$$P(X_2 = \text{sun}) = \quad P(X_2 = sun | X_1 = \text{sun})P(X_1 = \text{sun}) +$$

$$P(X_2 = sun | X_1 = \text{rain})P(X_1 = \text{rain})$$

$$0.9 \cdot 1.0 + 0.3 \cdot 0.0 = 0.9$$

Mini-Forward Algorithm

○ Question: What's P(X) on some day t?

$X_1 \rightarrow X_2 \rightarrow X_3 \rightarrow X_4 \dashrightarrow$

$P(x_1) = $ known

$$P(x_t) = \sum_{x_{t-1}} P(x_{t-1}, x_t)$$

$$= \sum_{x_{t-1}} P(x_t \mid x_{t-1}) P(x_{t-1})$$

Forward simulation
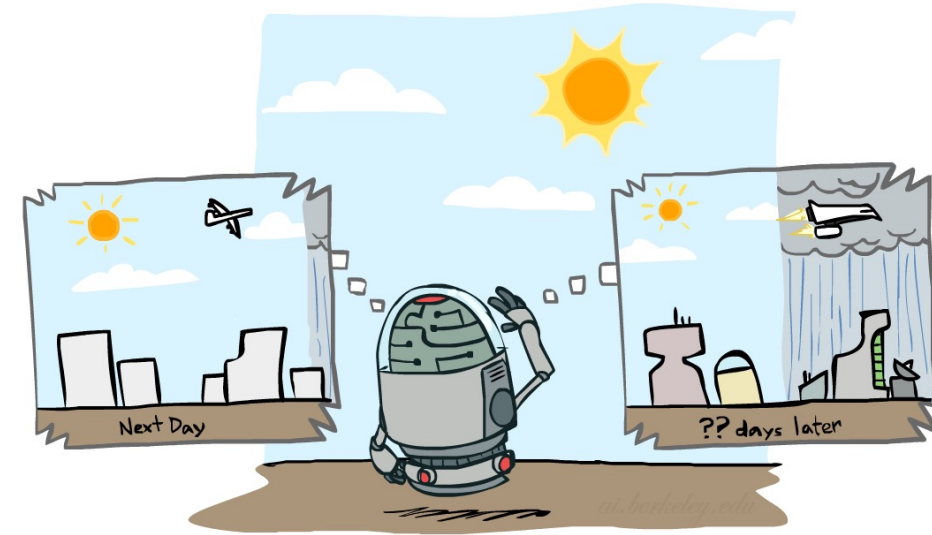
Example Run of Mini-Forward Algorithm

- From initial observation of sun

$$\left\langle \begin{matrix} 1.0 \\ 0.0 \end{matrix} \right\rangle \quad \left\langle \begin{matrix} 0.9 \\ 0.1 \end{matrix} \right\rangle \quad \left\langle \begin{matrix} 0.84 \\ 0.16 \end{matrix} \right\rangle \quad \left\langle \begin{matrix} 0.804 \\ 0.196 \end{matrix} \right\rangle \implies \left\langle \begin{matrix} 0.75 \\ 0.25 \end{matrix} \right\rangle$$

$P(X_1) \qquad P(X_2) \qquad P(X_3) \qquad P(X_4) \qquad\qquad P(X_\infty)$

- From initial observation of rain

$$\left\langle \begin{matrix} 0.0 \\ 1.0 \end{matrix} \right\rangle \quad \left\langle \begin{matrix} 0.3 \\ 0.7 \end{matrix} \right\rangle \quad \left\langle \begin{matrix} 0.48 \\ 0.52 \end{matrix} \right\rangle \quad \left\langle \begin{matrix} 0.588 \\ 0.412 \end{matrix} \right\rangle \implies \left\langle \begin{matrix} 0.75 \\ 0.25 \end{matrix} \right\rangle$$

$P(X_1) \qquad P(X_2) \qquad P(X_3) \qquad P(X_4) \qquad\qquad P(X_\infty)$

- From yet another initial distribution $P(X_1)$:

$$\left\langle \begin{matrix} p \\ 1-p \end{matrix} \right\rangle \qquad \ldots \qquad\qquad\qquad \implies \left\langle \begin{matrix} 0.75 \\ 0.25 \end{matrix} \right\rangle$$

$P(X_1) \qquad\qquad\qquad\qquad\qquad\qquad P(X_\infty)$

[Demo: L13D1,2,3]

Video of Demo Ghostbusters Basic Dynamics

Video of Demo Ghostbusters Circular Dynamics

Video of Demo Ghostbusters Whirlpool Dynamics

Stationary Distributions

o For most chains:

  o Influence of the initial distribution gets less and less over time.

  o The distribution we end up in is independent of the initial distribution
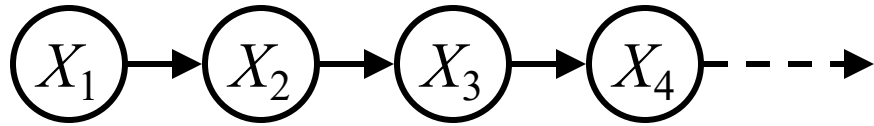
- Stationary distribution:

  - The distribution we end up with is called the stationary distribution P_∞ of the chain

  - It satisfies

$$P_\infty(X) = P_{\infty+1}(X) = \sum_x P(X|x)P_\infty(x)$$

Example: Stationary Distributions

○ Question: What's P(X) at time t = infinity?

$$X_1 \rightarrow X_2 \rightarrow X_3 \rightarrow X_4 \dashrightarrow$$

$$P_\infty(sun) = P(sun|sun)P_\infty(sun) + P(sun|rain)P_\infty(rain)$$
$$P_\infty(rain) = P(rain|sun)P_\infty(sun) + P(rain|rain)P_\infty(rain)$$

$$P_\infty(sun) = 0.9P_\infty(sun) + 0.3P_\infty(rain)$$
$$P_\infty(rain) = 0.1P_\infty(sun) + 0.7P_\infty(rain)$$

$$P_\infty(sun) = 3P_\infty(rain)$$
$$P_\infty(rain) = 1/3P_\infty(sun)$$

Also: $P_\infty(sun) + P_\infty(rain) = 1$

$\Longrightarrow$
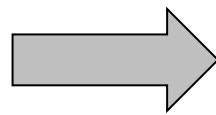
$$P_\infty(sun) = 3/4$$
$$P_\infty(rain) = 1/4$$
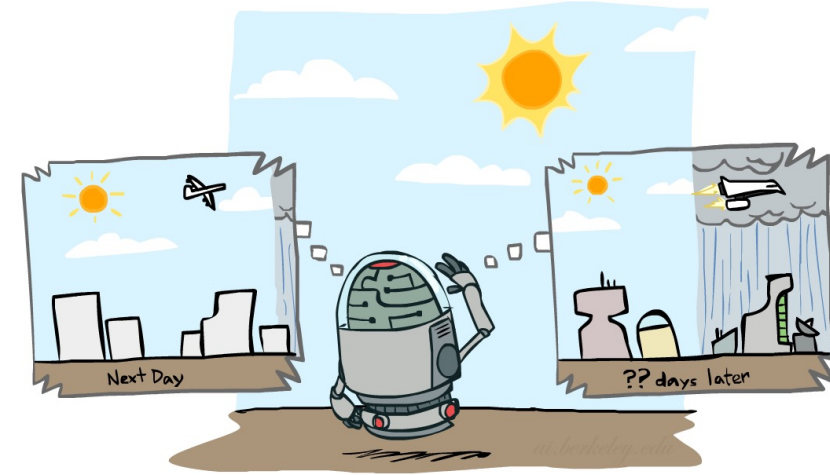
X_{t-1}	X_t	$P(X_t|X_{t-1})$
sun	sun	0.9
sun	rain	0.1
rain	sun	0.3
rain	rain	0.7

Application of Stationary Distribution: Web Link Analysis

o PageRank over a web graph
  o Each web page is a possible value of a state
  o Initial distribution: uniform over pages
  o Transitions:
    o With prob. c, uniform jump to a random page (dotted lines, not all shown)
    o With prob. 1-c, follow a random outlink (solid lines)

o Stationary distribution
  o Will spend more time on highly reachable pages
  o E.g. many ways to get to the Acrobat Reader download page
  o Somewhat robust to link spam.
  o Google 1.0 returned the set of pages containing all your keywords in decreasing rank, now all search engines use link analysis along with many other factors (rank actually getting less important over time)
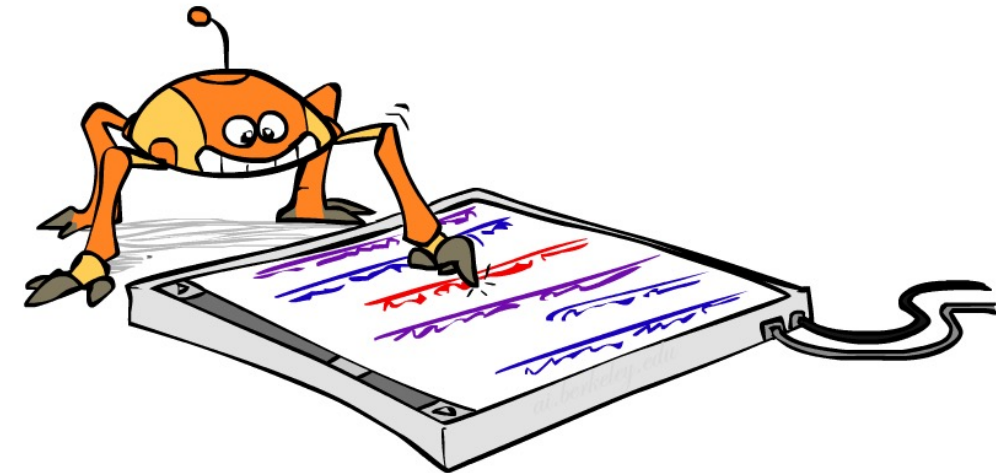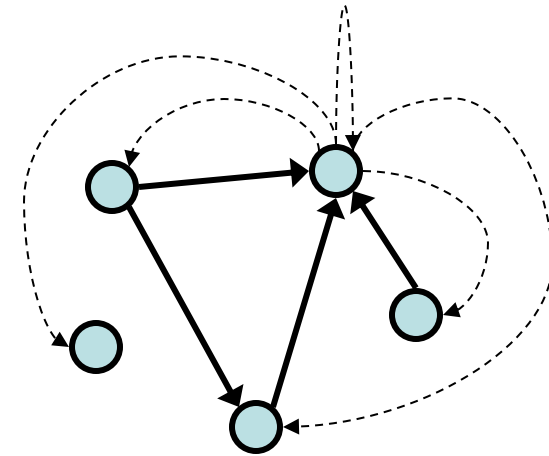
Application of Stationary Distributions: Gibbs Sampling*

o Each joint instantiation over all hidden and query variables is a state: $\{X_1, \ldots, X_n\} = H \cup Q$

o Transitions:

    o With probability $1/n$ resample variable X_j according to

$$P(X_j \mid x_1, x_2, \ldots, x_{j-1}, x_{j+1}, \ldots, x_n, e_1, \ldots, e_m)$$

o Stationary distribution:

    o Conditional distribution $P(X_1, X_2, \ldots, X_n \mid e_1, \ldots, e_m)$

    o Means that when running Gibbs sampling long enough we get a sample from the desired distribution

    o Requires some proof to show this is true!
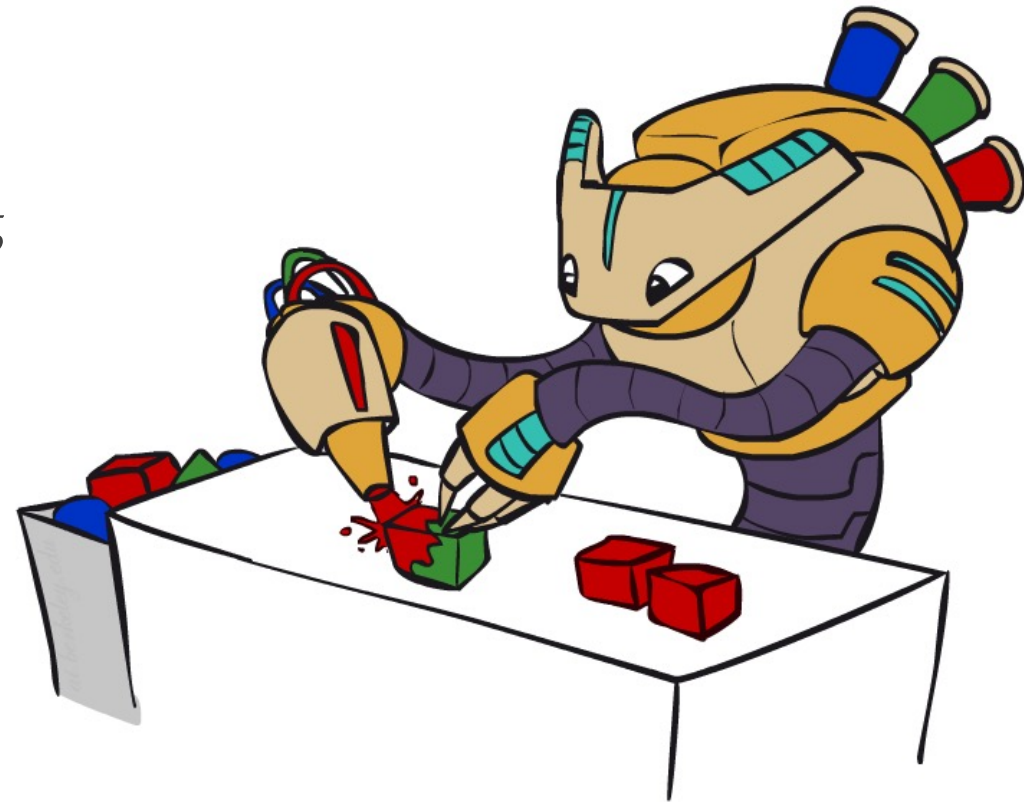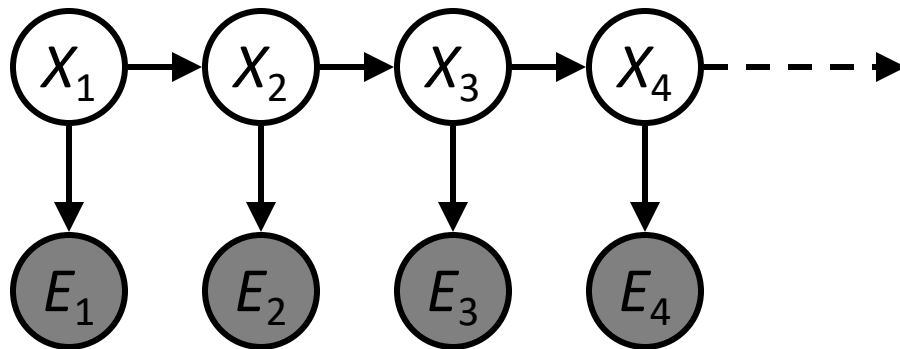
Hidden Markov Models

Pacman – Sonar

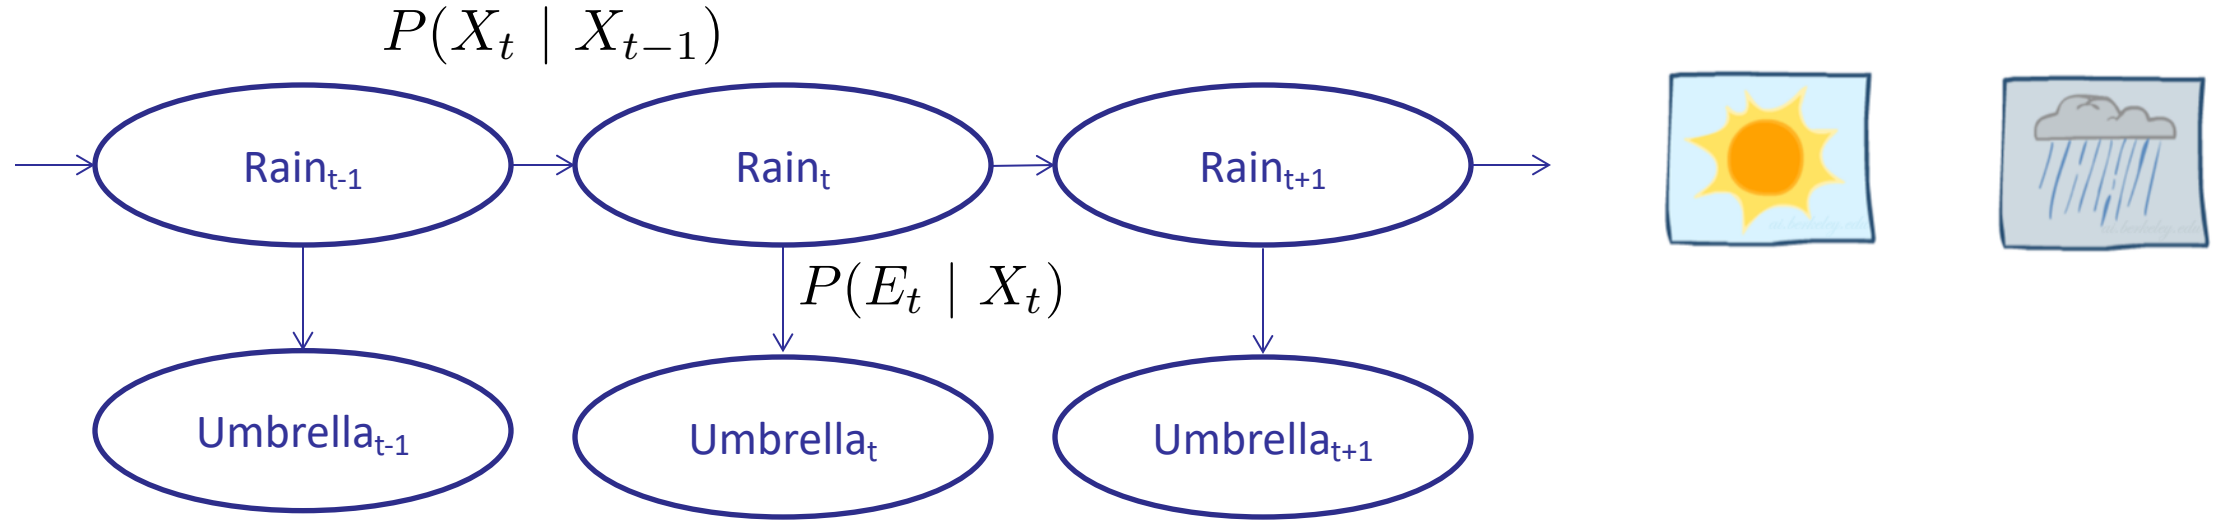# Video of Demo Pacman – Sonar (no beliefs)

Hidden Markov Models

o Markov chains not so useful for most agents
  o Need observations to update your beliefs

o Hidden Markov models (HMMs)
  o Underlying Markov chain over states X
  o You observe outputs (effects) at each time step

Example: Weather HMM

$$P(X_t \mid X_{t-1})$$

$$P(E_t \mid X_t)$$

o An HMM is defined by:
- o Initial distribution: $P(X_1)$
- o Transitions: $P(X_t \mid X_{t-1})$
- o Emissions: $P(E_t \mid X_t)$

R_{t-1}	R_t	$P(R_t|R_{t-1})$
+r	+r	0.7
+r	-r	0.3
-r	+r	0.3
-r	-r	0.7

R_t	U_t	$P(U_t|R_t)$
+r	+u	0.9
+r	-u	0.1
-r	+u	0.2
-r	-u	0.8

Video of Demo Ghostbusters – Circular Dynamics -- HMM

Example: Ghostbusters HMM

- P(X_1) = uniform

- P(X|X') = usually move clockwise, but sometimes move in a random direction or stay in place

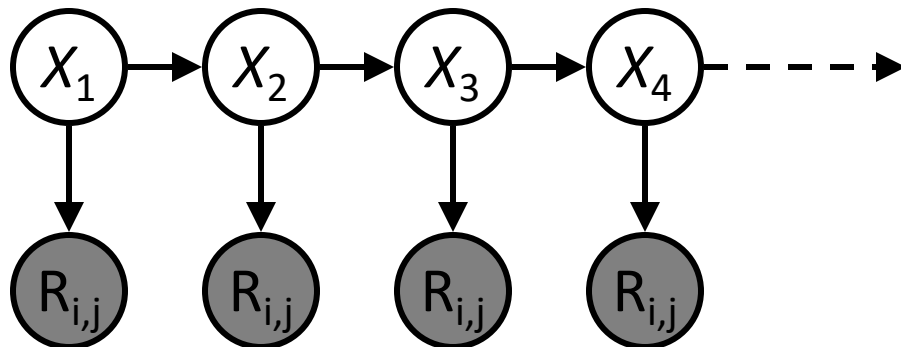- P(R_{ij}|X) = same sensor model as before: red means close, green means far away.
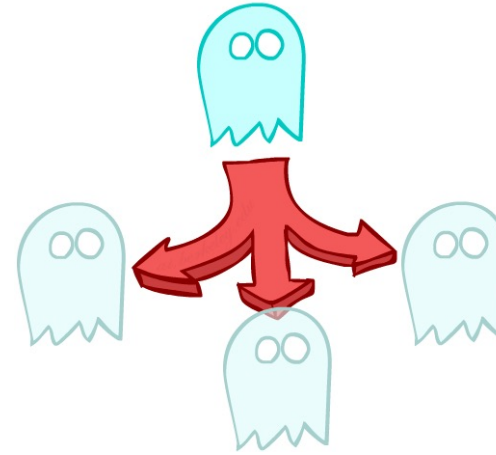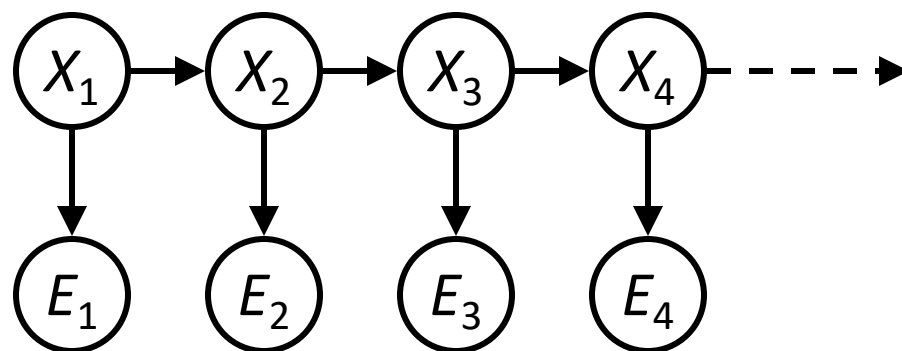
1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

P(X_1)
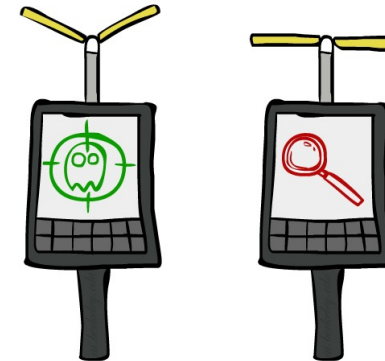
1/6	1/6	1/2
0	1/6	0
0	0	0

P(X|X' =<1,2>)

Conditional Independence

o HMMs have two important independence properties:

  o Markov hidden process: future depends on past via the present

  o Current observation independent of all else given current state

o Does this mean that evidence variables are guaranteed to be independent?

 o [No, they tend to correlated by the hidden state]

Real HMM Examples

o Robot tracking:
  o Observations are range readings (continuous)
  o States are positions on a map (continuous)

o Speech recognition HMMs:
  o Observations are acoustic signals (continuous valued)
  o States are specific positions in specific words (so, tens of thousands)

o Machine translation HMMs:
  o Observations are words (tens of thousands)
  o States are translation options

Filtering / Monitoring

- Filtering, or monitoring, is the task of tracking the distribution $B_t(X) = P_t(X_t \mid e_1, \ldots, e_t)$ (the belief state) over time

- We start with $B_1(X)$ in an initial setting, usually uniform

- As time passes, or we get observations, we update $B(X)$

- The Kalman filter was invented in the 60's and first implemented as a method of trajectory estimation for the Apollo program
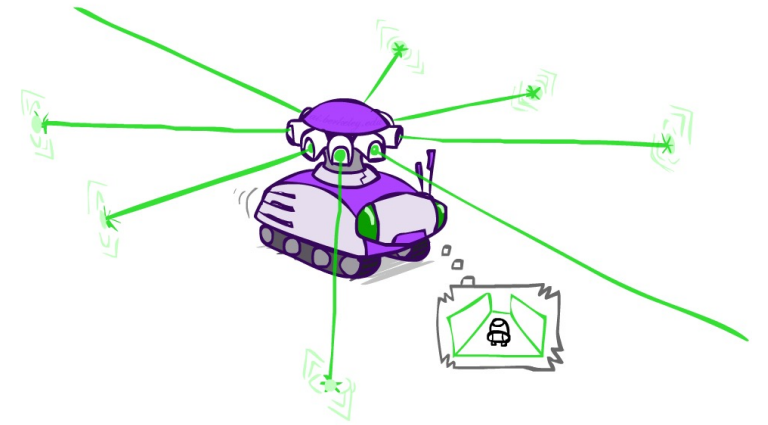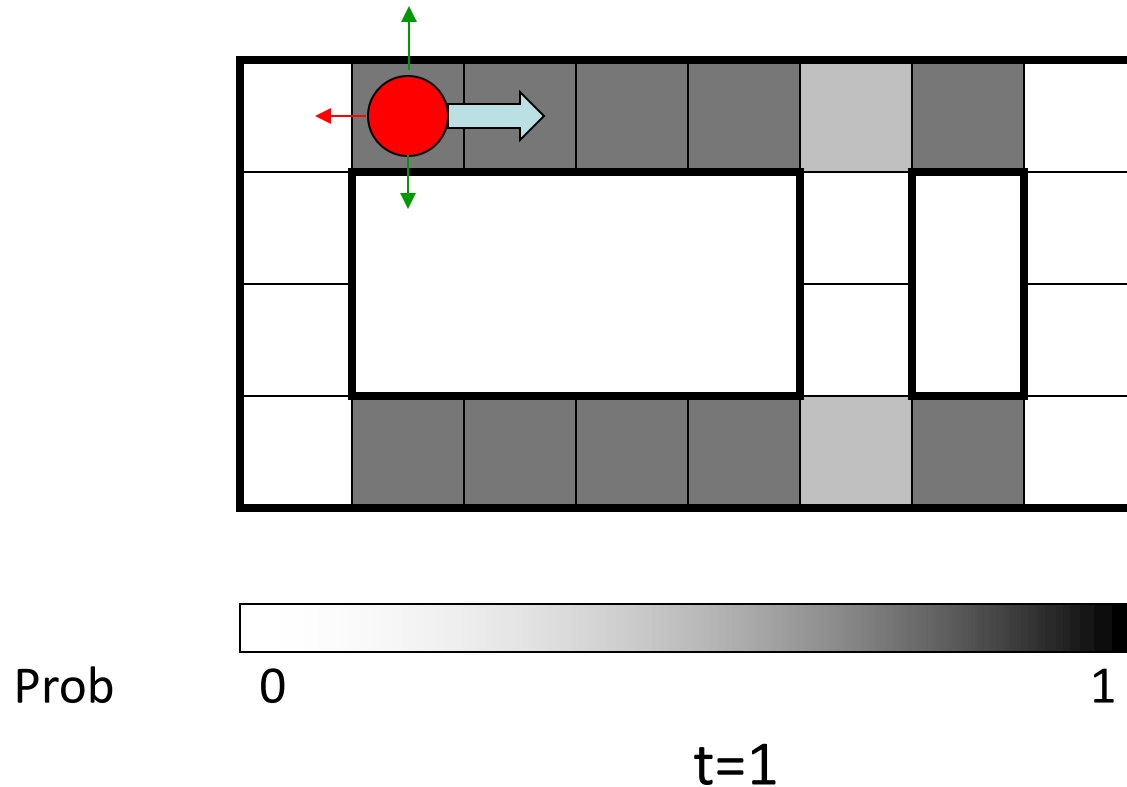
Example: Robot Localization

Prob   0   1

t=0
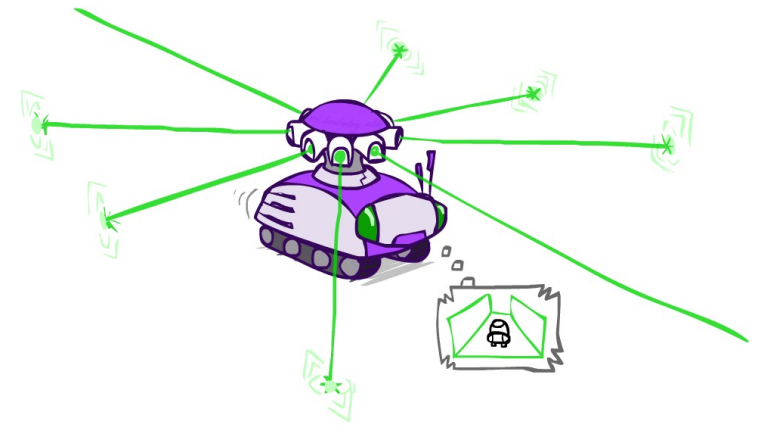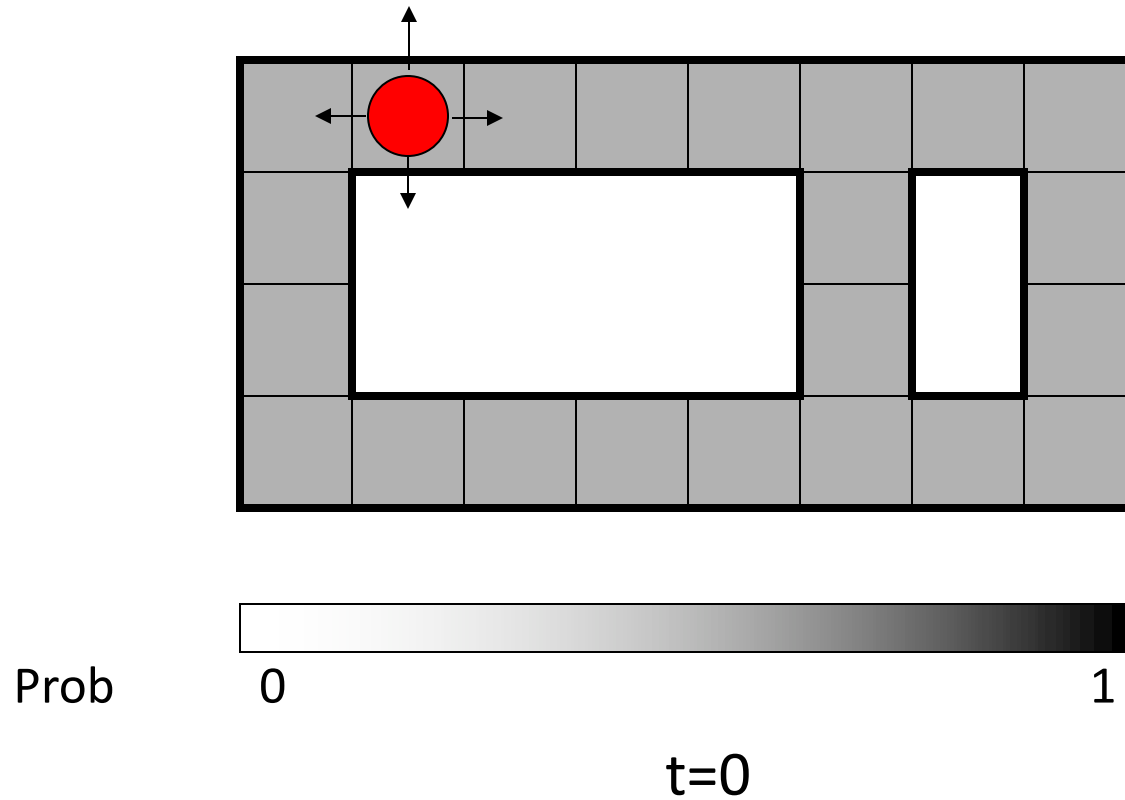
Sensor model: can read in which directions there is a wall,
never more than 1 mistake

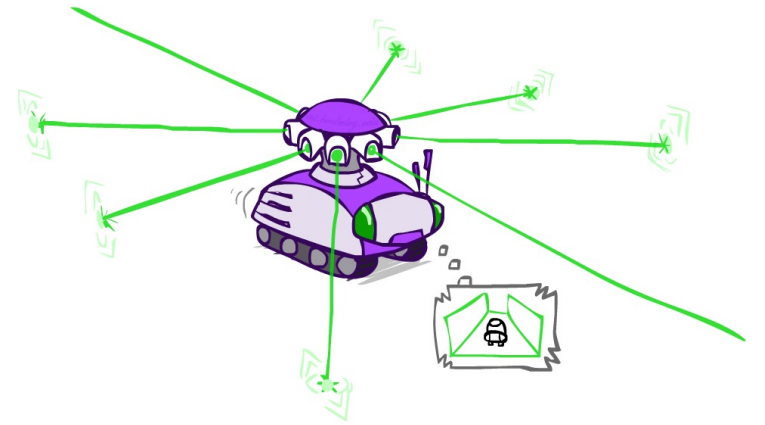Motion model: may not execute action with small prob.

Example: Robot Localization

Prob   0   1

t=1

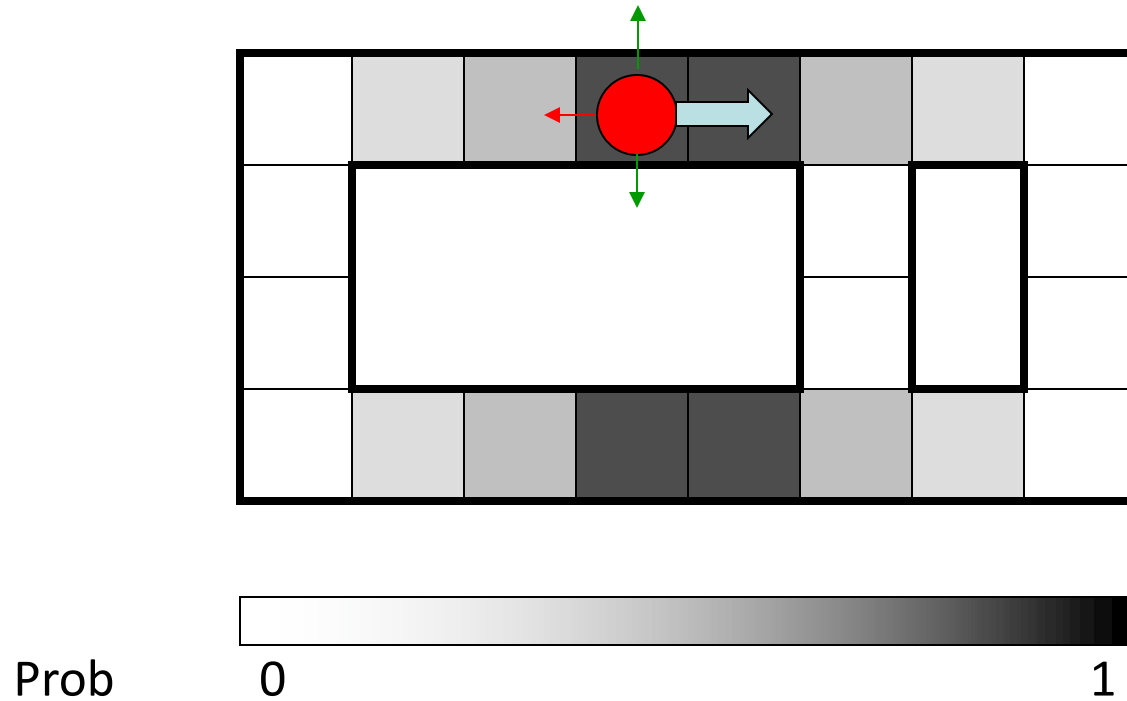Lighter grey: was possible to get the reading, but less likely b/c required 1 mistake
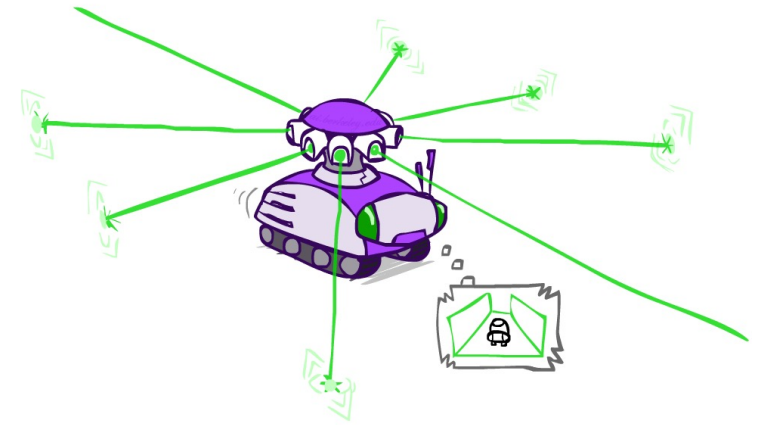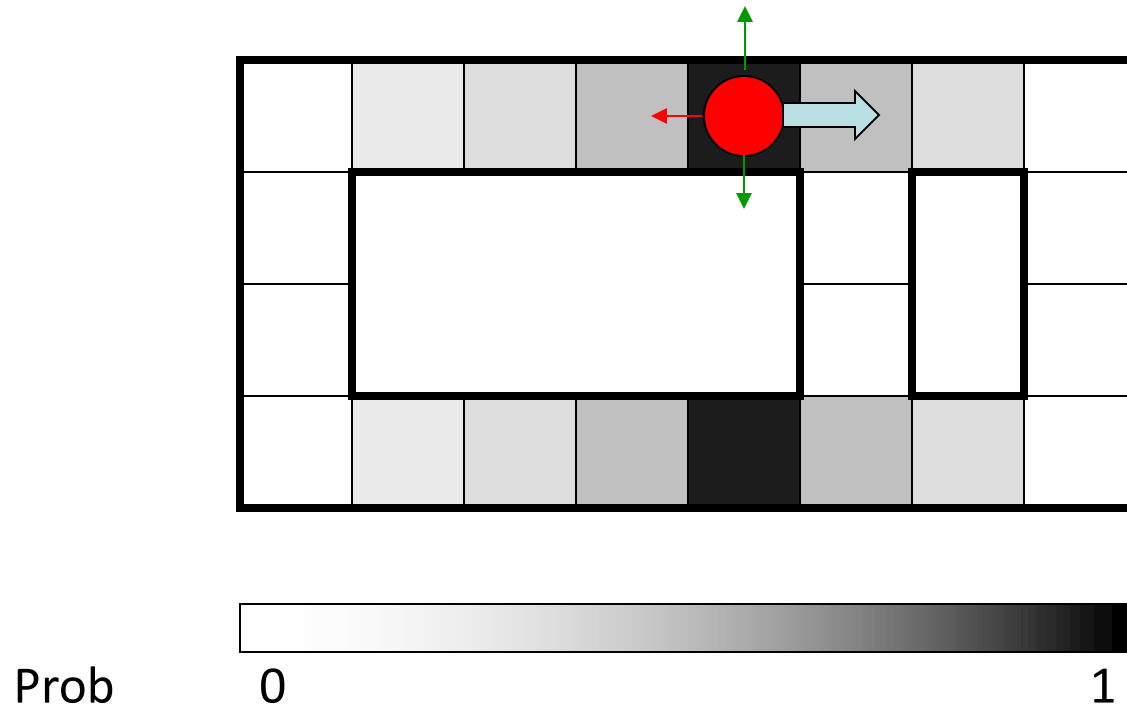
Example: Robot Localization

Prob   0   1

t=2

Prob   0   1
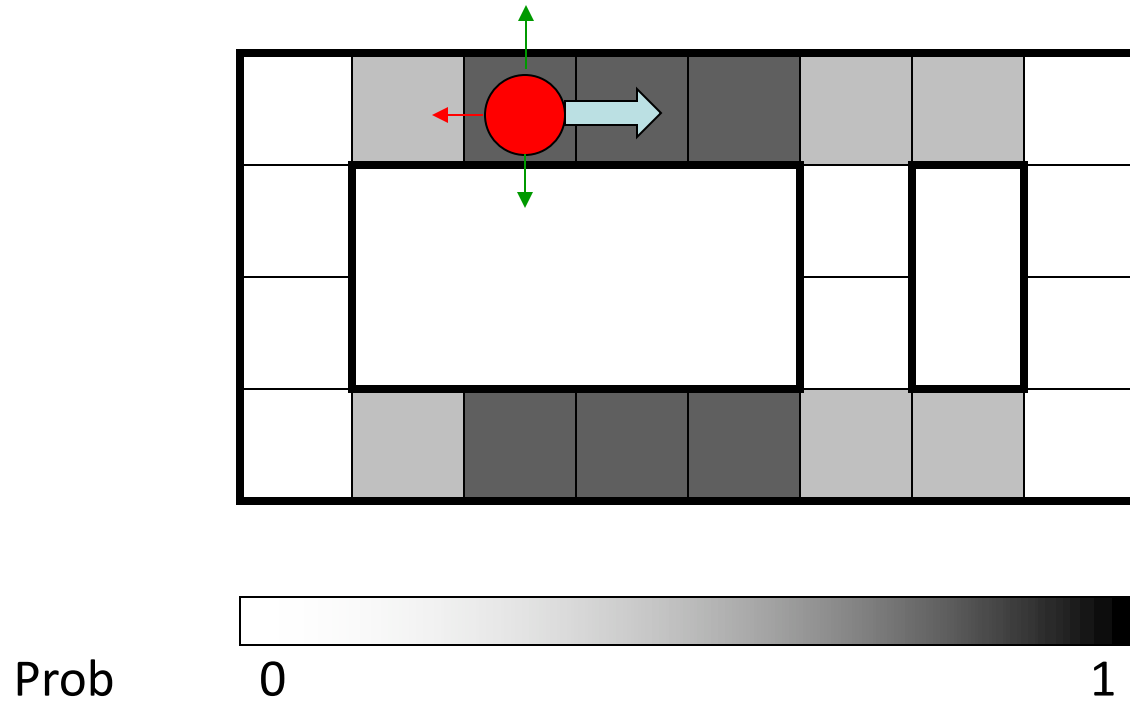
t=3

Example: Robot Localization

Prob   0   1

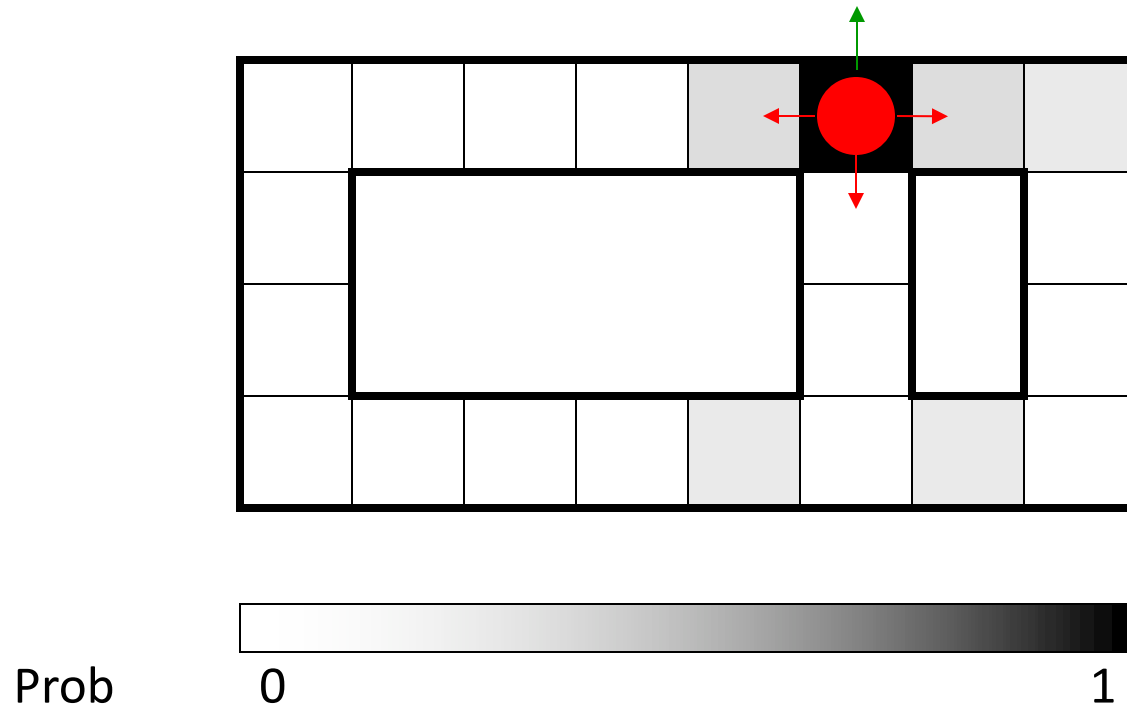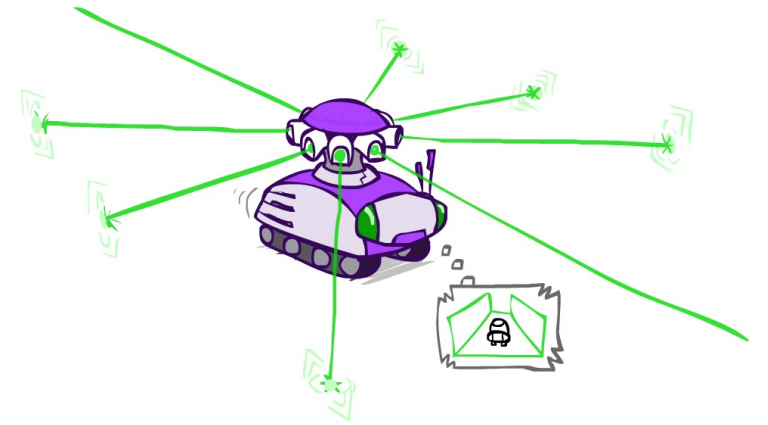t=4

Example: Robot Localization

Prob   0   1
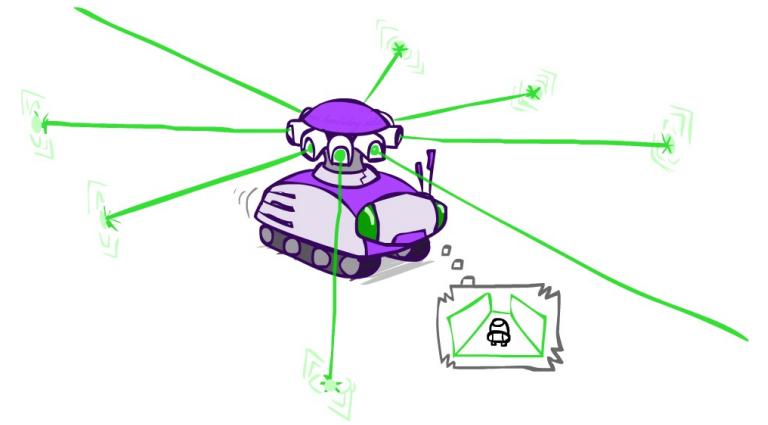
t=5

Inference: Find State Given Evidence

o We are given evidence at each time and want to know

$$B_t(X) = P(X_t | e_{1:t})$$

o Idea: start with $P(X_1)$ and derive B_t in terms of B_{t-1}

  o equivalently, derive B_{t+1} in terms of B_t