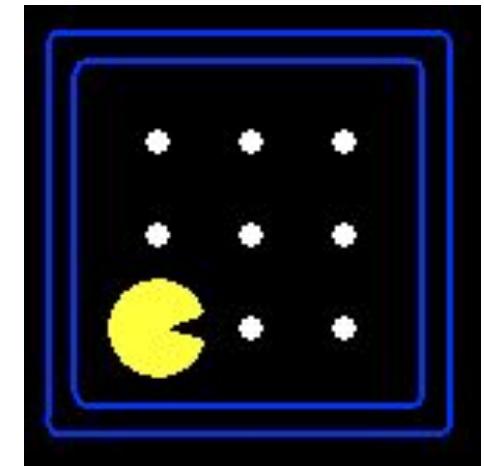


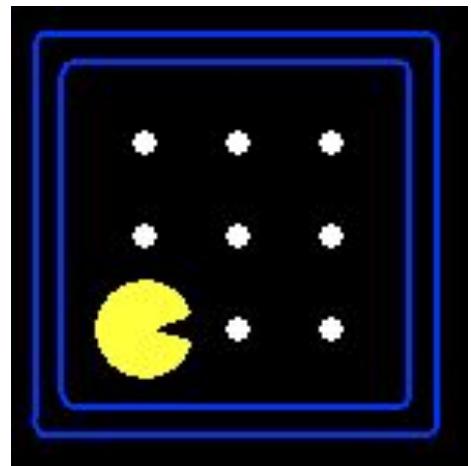
Reminder: Partially observable Pacman

- Pacman knows the map but perceives just wall/gap to NSEW
- Formulation: ***what variables do we need?***
 - Wall locations
 - Wall_0,0 there is a wall at [0,0]
 - Wall_0,1 there is a wall at [0,1], etc. (N symbols for N locations)
 - Percepts
 - ~~Blocked_W (blocked by wall to my West) etc.~~
 - Blocked_W_0 (blocked by wall to my West at time 0) etc. ($4T$ symbols for T time steps)
 - Actions
 - W_0 (Pacman moves West at time 0), E_0 etc. ($4T$ symbols)
 - Pacman's location
 - At_0,0_0 (Pacman is at [0,0] at time 0), At_0,1_0 etc. (NT symbols)



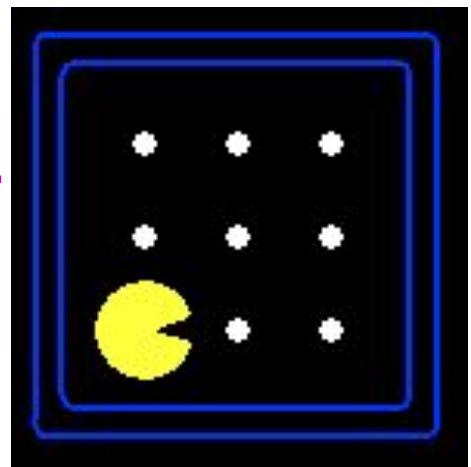
Pacman's knowledge base: Map

- Pacman knows where the walls are:
 - $\text{Wall}_{0,0} \wedge \text{Wall}_{0,1} \wedge \text{Wall}_{0,2} \wedge \text{Wall}_{0,3} \wedge \text{Wall}_{0,4} \wedge \text{Wall}_{1,4}$
 $\wedge \dots$
- Pacman knows where the walls aren't!
 - $\neg\text{Wall}_{1,1} \wedge \neg\text{Wall}_{1,2} \wedge \neg\text{Wall}_{1,3} \wedge \neg\text{Wall}_{2,1} \wedge \neg\text{Wall}_{2,2} \wedge \dots$



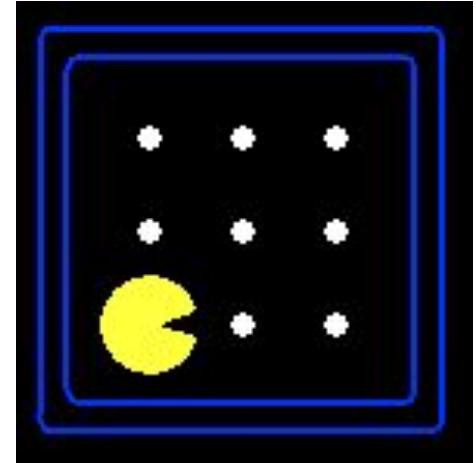
Pacman's knowledge base: Initial state

- Pacman doesn't know where he is!
- But he knows he's somewhere!
 - $\text{At_1,1_0} \vee \text{At_1,2_0} \vee \text{At_1,3_0} \vee \text{At_2,1_0} \vee \dots$
- And he knows he's not where the walls are!
 - $\neg\text{At_0,0_0} \wedge \neg\text{At_0,1_0} \wedge \neg\text{At_0,2_0} \wedge \dots$
- And he knows he's not in two places at once!
 - $\neg(\text{At_1,1_0} \wedge \text{At_1,2_0}) \wedge \neg(\text{At_1,1_0} \wedge \text{At_1,3_0}) \wedge \dots$



Pacman's knowledge base: Sensor model

- State facts about how Pacman's percepts arise...
 - $\langle \text{Percept variable at } t \rangle \Leftrightarrow \langle \text{some condition on world at } t \rangle$
- Pacman perceives a wall to the West at time t
if and only if he is in x,y and there is a wall at $x-1,y$
 - $\text{Blocked_W_0} \Leftrightarrow ((\text{At}_{-1,1_0} \wedge \text{Wall}_{0,1}) \vee (\text{At}_{-1,2_0} \wedge \text{Wall}_{0,2}) \vee (\text{At}_{-1,3_0} \wedge \text{Wall}_{0,3}) \vee \dots)$
- $4T$ sentences, each of size $O(N)$
- Note: these are valid for any map



Pacman's knowledge base: Transition model

- How does each ***state variable*** at each time gets its value?
 - Here we care about location variables, e.g., At_3,3_17
- A state variable X gets its value according to a ***successor-state axiom***
 - $X_t \Leftrightarrow [X_{t-1} \wedge \neg(\text{some action } t-1 \text{ made it false})] \vee [\neg X_{t-1} \wedge (\text{some action } t-1 \text{ made it true})]$
- For Pacman location:
 - $\text{At_3,3_17} \Leftrightarrow [\text{At_3,3_16} \wedge \neg((\neg \text{Wall_3,4} \wedge \text{N_16}) \vee (\neg \text{Wall_4,3} \wedge \text{E_16}) \vee \dots)]$
 $\vee [\neg \text{At_3,3_16} \wedge ((\text{At_3,2_16} \wedge \neg \text{Wall_3,3} \wedge \text{N_16}) \vee (\text{At_2,3_16} \wedge \neg \text{Wall_3,3} \wedge \text{N_16}) \vee \dots)]$

How many sentences?

- Vast majority of KB occupied by $O(NT)$ transition model sentences
 - Each about 10 lines of text
 - $N=200, T=400 \Rightarrow \sim 800,000$ lines of text, or 20,000 pages
- This is because propositional logic has limited expressive power
- Are we really going to write 20,000 pages of logic sentences???
- No, but your code will generate all those sentences!
- (In first-order logic, we need $O(1)$ transition model sentences)

A knowledge-based agent

```
function KB-AGENT(percept) returns an action
  persistent: KB, a knowledge base
    t, an integer, initially 0
    TELL(KB, MAKE-PERCEPT-SENTENCE(percept, t))
    action ← ASK(KB, MAKE-ACTION-QUERY(t))
    TELL(KB, MAKE-ACTION-SENTENCE(action, t))
    t←t+1
  return action
```

Some reasoning tasks

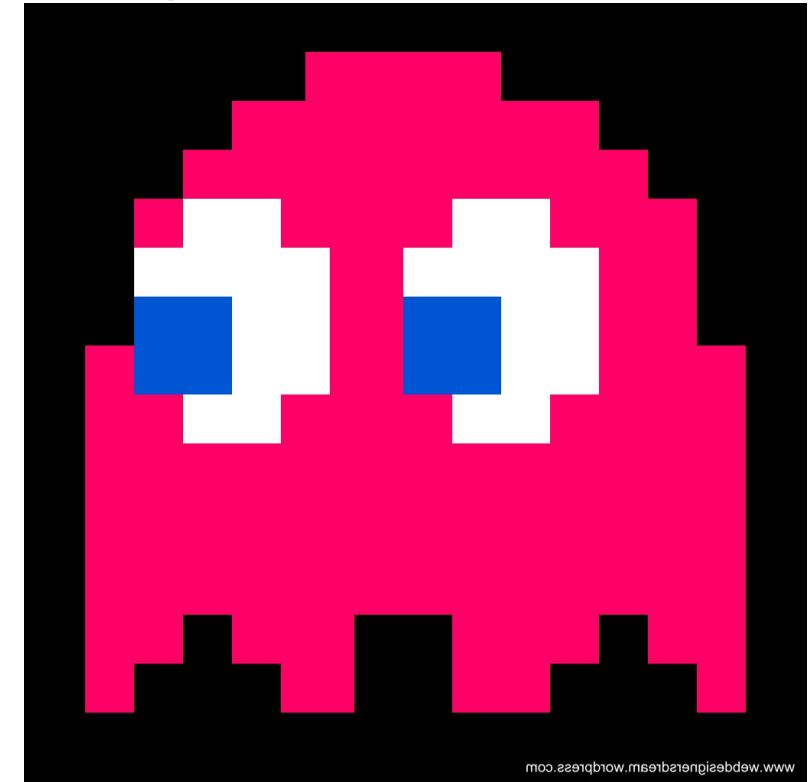
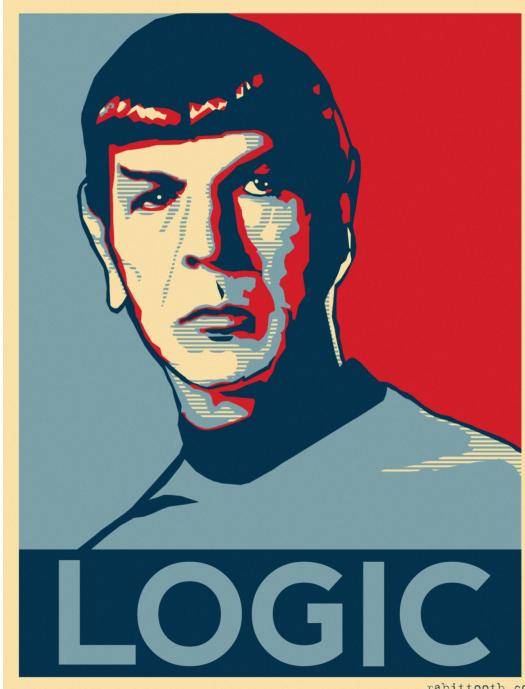
- Localization with a map and local sensing:
 - Given an initial KB, plus a sequence of percepts and actions, where am I?
- Mapping with a location sensor:
 - Given an initial KB, plus a sequence of percepts and actions, what is the map?
- Simultaneous localization and mapping:
 - Given ..., where am I and what is the map?
- Planning:
 - Given ..., what action sequence is guaranteed to reach the goal?
- **ALL OF THESE USE THE SAME KB AND THE SAME ALGORITHM!!**

Summary

- One possible agent architecture: knowledge + inference
- Logics provide a formal way to encode knowledge
 - A logic is defined by: syntax, set of possible worlds, truth condition
- A simple KB for Pacman covers the initial state, sensor model, and transition model
- Logical inference computes entailment relations among sentences, enabling a wide range of tasks to be solved

CS 188: Artificial Intelligence

Inference in Propositional Logic



Instructors: Stuart Russell and Dawn Song

University of California, Berkeley

Inference (reminder)

- Method 1: ***model-checking***
 - For every possible world, if α is true make sure that is β true too
- Method 2: ***theorem-proving***
 - Search for a sequence of proof steps (applications of ***inference rules***) leading from α to β
- ***Sound*** algorithm: everything it claims to prove is in fact entailed
- ***Complete*** algorithm: every that is entailed can be proved

Simple theorem proving: Forward chaining

- Forward chaining applies Modus Ponens to generate new facts:
 - **Given** $X_1 \wedge X_2 \wedge \dots \wedge X_n \Rightarrow Y$ and X_1, X_2, \dots, X_n , **infer** Y
- Forward chaining keeps applying this rule, adding new facts, until nothing more can be added
- Requires KB to contain only **definite clauses**:
 - (Conjunction of symbols) \Rightarrow symbol; or
 - A single symbol (note that X is equivalent to $\text{True} \Rightarrow X$)
- Runs in **linear** time using two simple tricks:
 - Each symbol X_i knows which rules it appears in
 - Each rule keeps count of how many of its premises are not yet satisfied

Forward chaining algorithm: Details

function PL-FC-ENTAILS?(KB, q) **returns** true or false

count \leftarrow a table, where count[c] is the number of symbols in c's premise

inferred \leftarrow a table, where inferred[s] is initially false for all s

agenda \leftarrow a queue of symbols, initially symbols known to be true in KB

while agenda is not empty **do**

 p \leftarrow Pop(agenda)

if p = q **then return** true

if inferred[p] = false **then**

 inferred[p] \leftarrow true

for each clause c in KB where p is in c.premise **do**

 decrement count[c]

if count[c] = 0 **then** add c.conclusion to agenda

return false

Properties of forward chaining

- Theorem: FC is sound and complete for definite-clause KBs
- Soundness: follows from soundness of Modus Ponens (easy to check)
- Completeness proof:
 1. FC reaches a fixed point where no new atomic sentences are derived
 2. Consider the final set of known-to-be-true symbols as a model m (other ones false)
 3. Every clause in the original KB is true in m

Proof: Suppose a clause $a_1 \wedge \dots \wedge a_k \Rightarrow b$ is false in m
Then $a_1 \wedge \dots \wedge a_k$ is true in m and b is false in m
Therefore the algorithm has not reached a fixed point!
 4. Hence m is a model of KB
 5. If $\text{KB} \models q$, q is true in every model of KB, including m

Resolution (briefly)

- The resolution inference rule takes two implication sentences (of a particular form) and infers a new implication sentence:
- Example: $A \wedge B \wedge C \Rightarrow U \vee v$

$$\frac{D \wedge E \wedge U \Rightarrow x \vee y}{A \wedge B \wedge C \wedge D \wedge E \Rightarrow v \vee x \vee y}$$

- Resolution is complete for propositional logic
- Exponential time in the worst case

Satisfiability and entailment

- A sentence is **satisfiable** if it is true in at least one world
- Suppose we have a hyper-efficient SAT solver (WARNING: NP-COMPLETE 😈 😈 😈); how can we use it to test entailment?
 - $\alpha \models \beta$
 - iff $\alpha \Rightarrow \beta$ is true in all worlds
 - iff $\neg(\alpha \Rightarrow \beta)$ is false in all worlds
 - iff $\alpha \wedge \neg\beta$ is false in all worlds, i.e., unsatisfiable
- So, add the **negated** conclusion to what you know, test for (un)satisfiability; also known as **reductio ad absurdum**
- Efficient SAT solvers operate on **conjunctive normal form**

Conjunctive normal form (CNF)

- Every sentence can be expressed as a **conjunction** of **clauses**.
Replace biconditional by two implications **clauses**
- Each clause is a **disjunction** of **literals**.
Replace $\alpha \Rightarrow \beta$ by $\neg\alpha \vee \beta$
- Each literal is a symbol or a negated symbol.
Distribute \vee over \wedge
- Conversion to CNF by a sequence of standard transformations:
 - $\text{At_1,1_0} \Rightarrow (\text{Wall_0,1} \Leftrightarrow \text{Blocked_W_0})$
 - $\text{At_1,1_0} \Rightarrow ((\text{Wall_0,1} \Rightarrow \text{Blocked_W_0}) \wedge (\text{Blocked_W_0} \Rightarrow \text{Wall_0,1}))$
 - $\neg\text{At_1,1_0} \vee ((\neg\text{Wall_0,1} \vee \text{Blocked_W_0}) \wedge (\neg\text{Blocked_W_0} \vee \text{Wall_0,1}))$
 - $(\neg\text{At_1,1_0} \vee \neg\text{Wall_0,1} \vee \text{Blocked_W_0}) \wedge$
 - $(\neg\text{At_1,1_0} \vee \neg\text{Blocked_W_0} \vee \text{Wall_0,1})$

Efficient SAT solvers

- DPLL (Davis-Putnam-Logemann-Loveland) is the core of modern solvers
- Recursive depth-first search over models with some extras:
 - ***Early termination***: stop if
 - all clauses are satisfied; e.g., $(A \vee B) \wedge (A \vee \neg C)$ is satisfied by $\{A=\text{true}\}$
 - any clause is falsified; e.g., $(A \vee B) \wedge (A \vee \neg C)$ is satisfied by $\{A=\text{false}, B=\text{false}\}$
 - ***Pure literals***: if all occurrences of a symbol in as-yet-unsatisfied clauses have the same sign, then give the symbol that value
 - E.g., A is pure and positive in $(A \vee B) \wedge (A \vee \neg C) \wedge (C \vee \neg B)$ so set it to **true**
 - ***Unit clauses***: if a clause is left with a single literal, set symbol to satisfy clause
 - E.g., if $A=\text{false}$, $(A \vee B) \wedge (A \vee \neg C)$ becomes $(\text{false} \vee B) \wedge (\text{false} \vee \neg C)$, i.e. $(B) \wedge (\neg C)$
 - Satisfying the unit clauses often leads to further propagation, new unit clauses, etc.

DPLL algorithm

function DPLL(*clauses,symbols,model*) **returns** true or false

if every clause in *clauses* is true in *model* **then return** true

if some clause in *clauses* is false in *model* **then return** false

P,value \leftarrow FIND-PURE-SYMBOL(*symbols,clauses,model*)

if *P* is non-null **then return** DPLL(*clauses, symbols–P, model U {P=value}*)

P,value \leftarrow FIND-UNIT-CLAUSE(*clauses,model*)

if *P* is non-null **then return** DPLL(*clauses, symbols–P, model U {P=value}*)

P \leftarrow First(*symbols*); *rest* \leftarrow Rest(*symbols*)

return or(DPLL(*clauses,rest,model U {P=true}*),
DPLL(*clauses,rest,model U {P=false}*))

Efficiency

- Naïve implementation of DPLL: solve ~100 variables
- Extras:
 - Smart variable and value ordering
 - Divide and conquer
 - Caching unsolvable subcases as extra clauses to avoid redoing them
 - Cool indexing and incremental recomputation tricks so that every step of the DPLL algorithm is efficient (typically O(1))
 - Index of clauses in which each variable appears +ve/-ve
 - Keep track number of satisfied clauses, update when variables assigned
 - Keep track of number of remaining literals in each clause
- Real implementation of DPLL: solve ~100000000 variables

SAT solvers in practice

- Circuit verification: does this VLSI circuit compute the right answer?
- Software verification: does this program compute the right answer?
- Software synthesis: what program computes the right answer?
- Protocol verification: can this security protocol be broken?
- Protocol synthesis: what protocol is secure for this task?
- Lots of combinatorial problems: what is the solution?
- Planning: ***how can I eat all the dots???***

Summary

- Inference in propositional logic:
 - Inference algorithms determine whether $\alpha \models \beta$
 - Theorem provers apply inference rules to construct proofs
 - Model checkers enumerate models to establish entailment directly
 - Forward chaining is sound, complete, and linear-time for definite clauses
 - DPLL enumerates possible models via recursive depth-first search
 - Even though propositional logic KBs are often very large, modern solvers (usually based on DPLL) are usually very efficient in practice