



# Variable elimination: The basic ideas

- Move summations inwards as far as possible
  - $P(B \mid j, m) = \alpha \sum_{e,a} P(B) P(e) P(a \mid B, e) P(j \mid a) P(m \mid a)$
  - $= \alpha P(B) \sum_e P(e) \sum_a P(a \mid B, e) P(j \mid a) P(m \mid a)$
- Do the calculation from the inside out
  - I.e., sum over  $a$  first, then sum over  $e$
  - Problem:  $P(a \mid B, e)$  isn't a single number, it's a bunch of different numbers depending on the values of  $B$  and  $e$
  - Solution: use arrays of numbers (of various dimensions) with appropriate operations on them; these are called **factors**



# Operation 1: Pointwise product

- First basic operation: **pointwise product** of factors (similar to a **database join**, **not** matrix multiply!)
  - New factor has **union** of variables of the two original factors
  - Each entry is the product of the corresponding entries from the original factors
- Example:  $P(A,J) \times P(A,M) = P(A,J,M)$



$P(A,J)$

A \ J	true	false
true	0.09	0.01
false	0.045	0.855

$\times$

$P(A,M)$

A \ M	true	false
true	0.07	0.03
false	0.009	0.891

$=$

$P(A,J,M)$

J \ M	true	false	
J \ M	true	false	
true			18 A=false
false		.0003	A=true

# Operation 2: Summing out a variable

- Second basic operation: **summing out** (or eliminating) a variable from a factor
  - Shrinks a factor to a smaller one
- Example:  $\sum_j P(A, J) = P(A, j) + P(A, \neg j) = P(A)$

$P(A, J)$


A \ J	true	false
true	0.09	0.01
false	0.045	0.855

Sum out  $J$




$P(A)$

	0.1
true	0.1
false	0.9




# Summing out from a product of factors

- Project the factors each way first, then sum the products
- Example:  $\sum_a P(a|B,e) \times P(j|a) \times P(m|a)$
- $= P(a|B,e) \times P(j|a) \times P(m|a) +$
- $P(\neg a|B,e) \times P(j|\neg a) \times P(m|\neg a)$

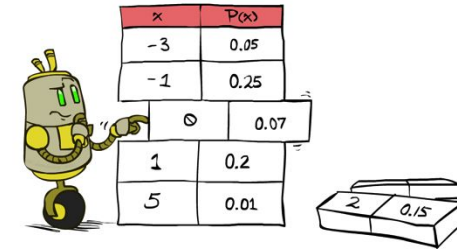


# Variable Elimination




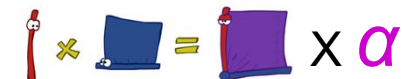
# Variable Elimination

- Query:  $P(Q|e)$
- Start with initial factors:
  - Local CPTs (but instantiated by evidence)
- For each hidden variable  $H_j$ 
  - Sum out  $H_j$  from the product of all factors mentioning  $H_j$
- Join all remaining factors and normalize



x	P(x)
-3	0.05
-1	0.25
0	0.07
1	0.2
5	0.01





$$\text{stick figure} \times \text{blue square} = \text{purple square} \times \alpha$$

# Example

Query  $P(B \mid j, m)$

$P(B)$	$P(E)$	$P(A \mid B, E)$	$P(j \mid A)$	$P(m \mid A)$
--------	--------	------------------	---------------	---------------



Choose  $A$

$P(A \mid B, E)$	$\Rightarrow$	$\times$	$\Rightarrow$	$\Sigma$	$\Rightarrow$	$P(j, m \mid B, E)$
$P(j \mid A)$						
$P(m \mid A)$						

$P(B)$	$P(E)$	$P(j, m \mid B, E)$
--------	--------	---------------------

# Example

$P(B)$	$P(E)$	$P(j,m B,E)$
--------	--------	--------------


Choose  $E$

$$\begin{array}{c} P(E) \\ P(j,m|B,E) \end{array} \xrightarrow{\times} \xrightarrow{\Sigma} P(j,m|B)$$

$P(B)$	$P(j,m B)$
--------	------------

Finish with  $B$

$$\begin{array}{c} P(B) \\ P(j,m|B) \end{array} \xrightarrow{\times} P(j,m,B) \xrightarrow{\text{Normalize}} P(B | j,m)$$






# Order matters

- Order the terms  $Z, A, B, C, D$

- $P(D) = \alpha \sum_{z,a,b,c} P(z) P(a|z) P(b|z) P(c|z) P(D|z)$
- $= \alpha \sum_z P(z) \sum_a P(a|z) \sum_b P(b|z) \sum_c P(c|z) P(D|z)$
- Largest factor has 2 variables ( $D, Z$ )

- Order the terms  $A, B, C, D, Z$

- $P(D) = \alpha \sum_{a,b,c,z} P(a|z) P(b|z) P(c|z) P(D|z) P(z)$
- $= \alpha \sum_a \sum_b \sum_c \sum_z P(a|z) P(b|z) P(c|z) P(D|z) P(z)$
- Largest factor has 4 variables ( $A, B, C, D$ )
- In general, with  $n$  leaves, factor of size  $2^n$




# VE: Computational and Space Complexity

---

- The computational and space complexity of variable elimination is determined by the largest factor (and it's space that kills you)
- The elimination ordering can greatly affect the size of the largest factor.
  - E.g., previous slide's example  $2^n$  vs. 2
- Does there always exist an ordering that only results in small factors?
  - No!





# Worst Case Complexity? Reduction from SAT



- Variables:  $W, X, Y, Z$
- CNF clauses:
  - $C_1 = W \vee X \vee Y$
  - $C_2 = Y \vee Z \vee \neg W$
  - $C_3 = X \vee Y \vee \neg Z$
- Sentence  $S = C_1 \wedge C_2 \wedge C_3$
- $P(S) > 0$  iff  $S$  is satisfiable
  - $\Rightarrow$  **NP-hard**
- $P(S) = K \times 0.5^n$  where  $K$  is the number of satisfying assignments for clauses
  - $\Rightarrow$  **#P-hard**


# Polytrees

- A polytree is a directed graph with no undirected cycles
- For poly-trees the complexity of variable elimination is **linear in the network size** if you eliminate from the leaves towards the roots



# CS 188: Artificial Intelligence

## Bayes Nets: Approximate Inference



Instructors: Stuart Russell and Dawn Song

University of California, Berkeley


# Sampling

- Basic idea

- Draw  $N$  samples from a *sampling distribution*  $S$
- Compute an approximate posterior probability
- Show this converges to the true probability  $P$

- Why sample?

- Often very fast to get a decent approximate answer
- The algorithms are very simple and general (easy to apply to fancy models)
- They require very little memory ( $O(n)$ )
- They can be applied to large models, whereas exact algorithms blow up




# Example

- Suppose you have two agent programs **A** and **B** for Monopoly
- What is the probability that **A** wins?
  - Method 1:
    - Let  $s$  be a sequence of dice rolls and Chance and Community Chest cards
    - Given  $s$ , the outcome  $V(s)$  is determined (1 for a win, 0 for a loss)
    - Probability that **A** wins is  $\sum_s P(s) V(s)$
    - Problem: infinitely many sequences  $s$  !
  - Method 2:
    - Sample  $N$  sequences from  $P(s)$  , play  $N$  games (maybe 100)
    - Probability that **A** wins is roughly  $1/N \sum_i V(s_i)$  i.e., fraction of wins in the sample

# Sampling basics: discrete (*categorical*) distribution

- To simulate a biased d-sided coin:

- Step 1: Get sample  $u$  from uniform distribution over  $[0, 1)$ 
  - E.g. `random()` in python
- Step 2: Convert this sample  $u$  into an outcome for the given distribution by associating each outcome  $x$  with a  $P(x)$ -sized sub-interval of  $[0,1)$



- Example


$C$	$P(C)$
red	0.6
green	0.1
blue	0.3

$0.0 \leq u < 0.6, \rightarrow C=\text{red}$

$0.6 \leq u < 0.7, \rightarrow C=\text{green}$

$0.7 \leq u < 1.0, \rightarrow C=\text{blue}$

- If `random()` returns  $u = 0.83$ , then the sample is  $C = \text{blue}$
- E.g, after sampling 8 times:






# Sampling in Bayes Nets


---

- Prior Sampling
- Rejection Sampling
- Likelihood Weighting
- Gibbs Sampling

# Prior Sampling




# Prior Sampling



# Prior Sampling

- For  $i=1, 2, \dots, n$  (in topological order)
  - Sample  $X_i$  from  $P(X_i \mid \text{parents}(X_i))$
- Return  $(x_1, x_2, \dots, x_n)$



# Prior Sampling

- This process generates samples with probability:

$$S_{PS}(x_1, \dots, x_n) = \prod_i P(x_i \mid \text{parents}(X_i)) = P(x_1, \dots, x_n)$$

...i.e. the BN's joint probability

- Let the number of samples of an event be  $N_{PS}(x_1, \dots, x_n)$
- Estimate from  $N$  samples is  $Q_N(x_1, \dots, x_n) = N_{PS}(x_1, \dots, x_n)/N$
- Then  $\lim_{N \rightarrow \infty} Q_N(x_1, \dots, x_n) = \lim_{N \rightarrow \infty} N_{PS}(x_1, \dots, x_n)/N$   
     $= S_{PS}(x_1, \dots, x_n)$   
     $= P(x_1, \dots, x_n)$
- I.e., the sampling procedure is **consistent**

# Example

- We'll get a bunch of samples from the BN:

$C, \neg S, r, w$

$C, S, r, w$


$\neg C, S, r, \neg w$

$C, \neg S, r, w$


$\neg C, \neg S, \neg r, w$

- If we want to know  $P(W)$

- We have counts  $\langle w:4, \neg w:1 \rangle$
- Normalize to get  $P(W) = \langle w:0.8, \neg w:0.2 \rangle$
- This will get closer to the true distribution with more samples
- Can estimate anything else, too
  - E.g., for query  $P(C \mid r, w)$  use  $P(C \mid r, w) = \alpha P(C, r, w)$




# Rejection Sampling



# Rejection Sampling

- A simple modification of prior sampling for conditional probabilities
- Let's say we want  $P(C \mid r, w)$
- Count the  $C$  outcomes, but ignore (reject) samples that don't have  $R=\text{true}$ ,  $W=\text{true}$ 
  - This is called **rejection sampling**
  - It is also consistent for conditional probabilities (i.e., correct in the limit)




$c, \neg s, r, w$   
 ~~$c, s, \neg r$~~   
 ~~$\neg c, s, r, \neg w$~~   
 ~~$c, \neg s, \neg r$~~   
 $\neg c, \neg s, r, w$




# Rejection Sampling

- Input: evidence  $e_1, \dots, e_k$
- For  $i=1, 2, \dots, n$ 
  - Sample  $x_i$  from  $P(x_i \mid \text{parents}(x_i))$
  - If  $x_i$  not consistent with evidence
    - Reject: Return, and no sample is generated in this cycle
- Return  $(x_1, x_2, \dots, x_n)$




# Likelihood Weighting




# Likelihood Weighting

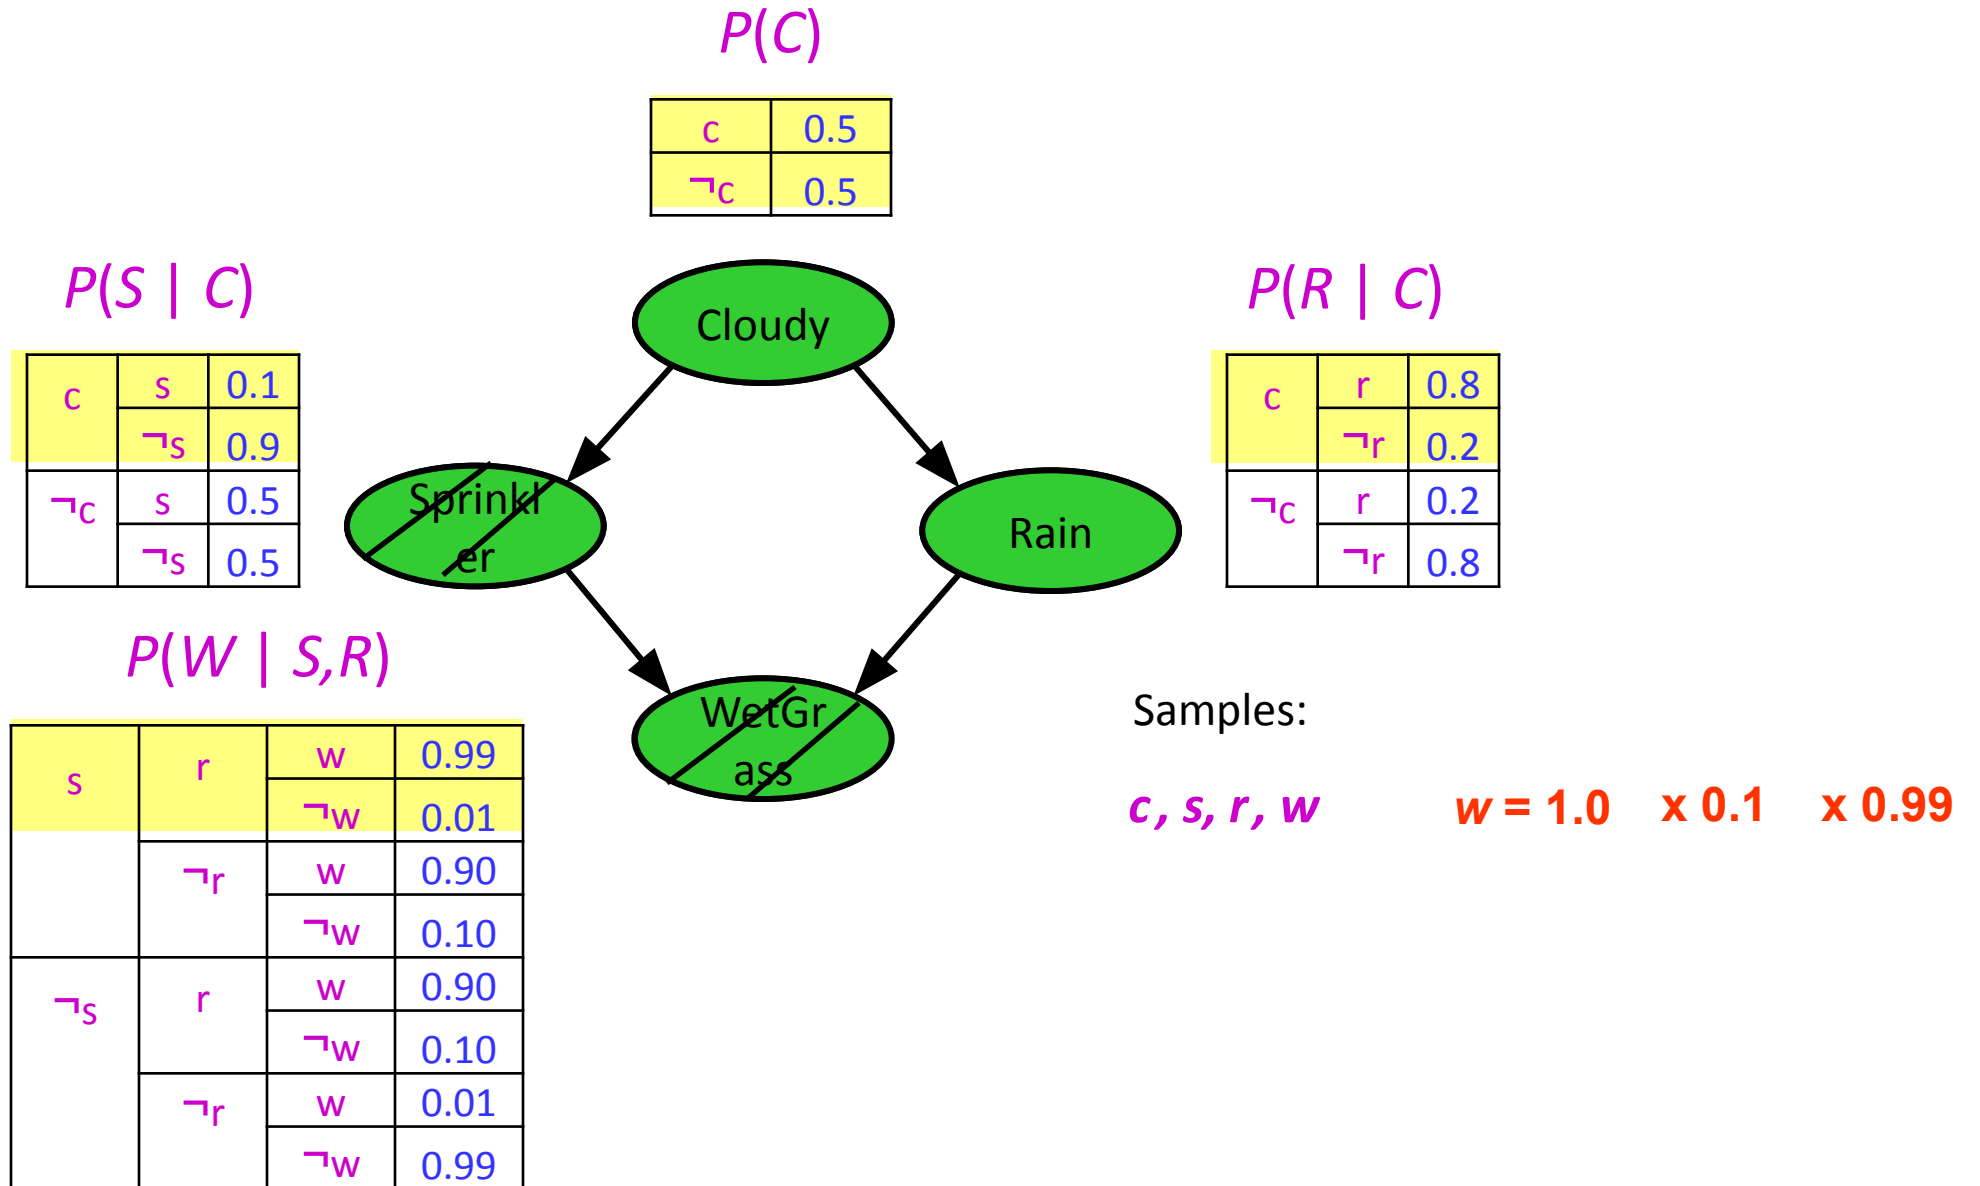
- Problem with rejection sampling:
  - If evidence is unlikely, rejects lots of samples
  - Evidence not exploited as you sample
  - Consider  $P(\text{Shape} | \text{Color}=\text{blue})$



- Idea: fix evidence variables, sample the rest
  - Problem: sample distribution not consistent!
  - Solution: **weight** each sample by probability of evidence variables given parents




# Likelihood Weighting



# Likelihood Weighting

- Input: evidence  $e_1, \dots, e_k$
- $w = 1.0$
- for  $i=1, 2, \dots, n$ 
  - if  $X_i$  is an evidence variable
    - $x_i = \text{observed value}_i \text{ for } X_i$
    - Set  $w = w * P(x_i \mid \text{parents}(X_i))$
  - else
    - Sample  $x_i$  from  $P(X_i \mid \text{parents}(X_i))$
- return  $(x_1, x_2, \dots, x_n), w$



# Likelihood Weighting

- Sampling distribution if  $\mathbf{z}$  sampled and  $\mathbf{e}$  fixed evidence

$$S_{WS}(\mathbf{z}, \mathbf{e}) = \prod_j P(z_j \mid \text{parents}(Z_j))$$


- Now, samples have weights

$$w(\mathbf{z}, \mathbf{e}) = \prod_k P(e_k \mid \text{parents}(E_k))$$

- Together, weighted sampling distribution is consistent


$$\begin{aligned} S_{WS}(\mathbf{z}, \mathbf{e}) \cdot w(\mathbf{z}, \mathbf{e}) &= \prod_j P(z_j \mid \text{parents}(Z_j)) \prod_k P(e_k \mid \text{parents}(E_k)) \\ &= P(\mathbf{z}, \mathbf{e}) \end{aligned}$$

- Likelihood weighting is an example of **importance sampling**
  - Would like to estimate some quantity based on samples from  $P$
  - $P$  is hard to sample from, so use  $Q$  instead
  - Weight each sample  $x$  by  $P(x)/Q(x)$



# Likelihood Weighting

- Likelihood weighting is good
  - All samples are used
  - The values of **downstream** variables are influenced by **upstream** evidence



- Likelihood weighting still has weaknesses
  - The values of **upstream** variables are unaffected by **downstream** evidence
    - E.g., suppose evidence is a video of a traffic accident
  - With evidence in  $k$  leaf nodes, weights will be  $O(2^{-k})$
  - With high probability, one lucky sample will have much larger weight than the others, dominating the result
- We would like each variable to “see” **all** the evidence!



# Quiz

- Suppose I perform a random walk on a graph, following the arcs out of a node *uniformly at random*. In the infinite limit, what fraction of time do I spend at each node?
  - Consider these two examples:

