



CS 188: Artificial Intelligence

Dynamic Bayes Nets and Particle Filters



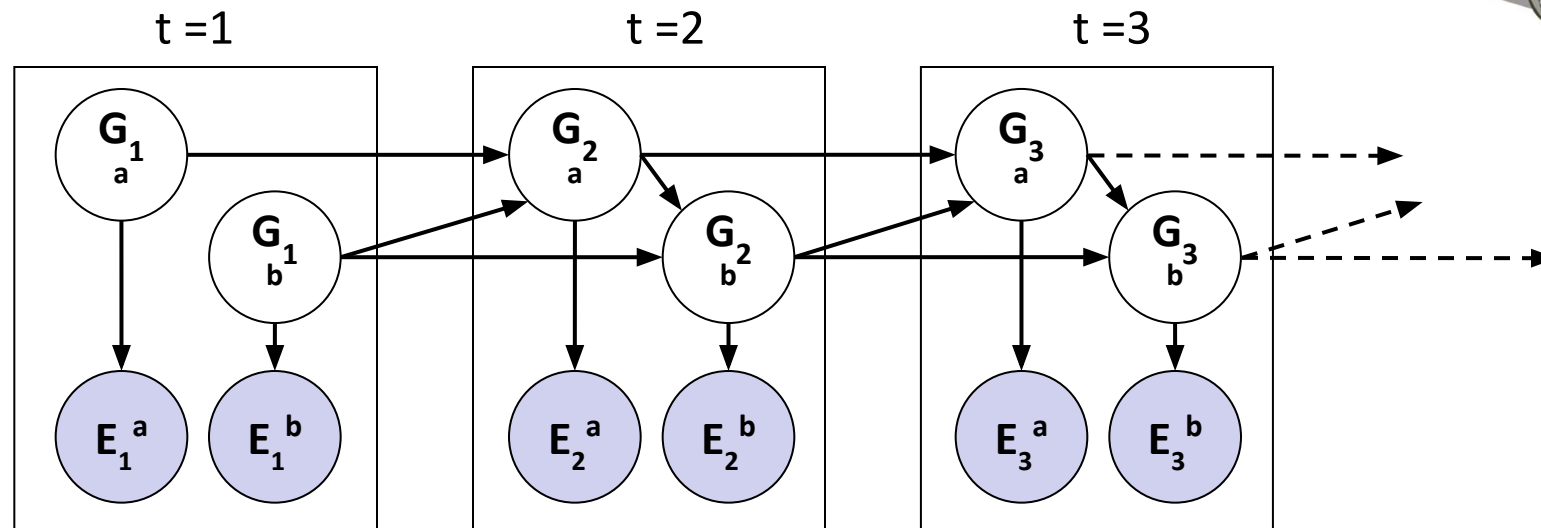

Instructor: Stuart Russell and Dawn Song
University of California, Berkeley

Dynamic Bayes Nets




Dynamic Bayes Nets (DBNs)

- We want to track multiple variables over time, using multiple sources of evidence
- Idea: Repeat a fixed Bayes net structure at each time
- Variables at time t can have parents at time $t-1$



DBNs and HMMs


- Every HMM is a single-variable DBN
- Every discrete DBN is an HMM
 - HMM state is Cartesian product of DBN state variables




- Sparse dependencies \Rightarrow exponentially fewer parameters in DBN
 - E.g., 20 state variables, 3 parents each;
DBN has $20 \times 2^3 = 160$ parameters, HMM has $2^{20} \times 2^{20} \approx 10^{12}$ parameters

Exact Inference in DBNs

- Variable elimination applies to dynamic Bayes nets
- Offline: “unroll” the network for T time steps, then eliminate variables to find $P(X_T | e_{1:T})$



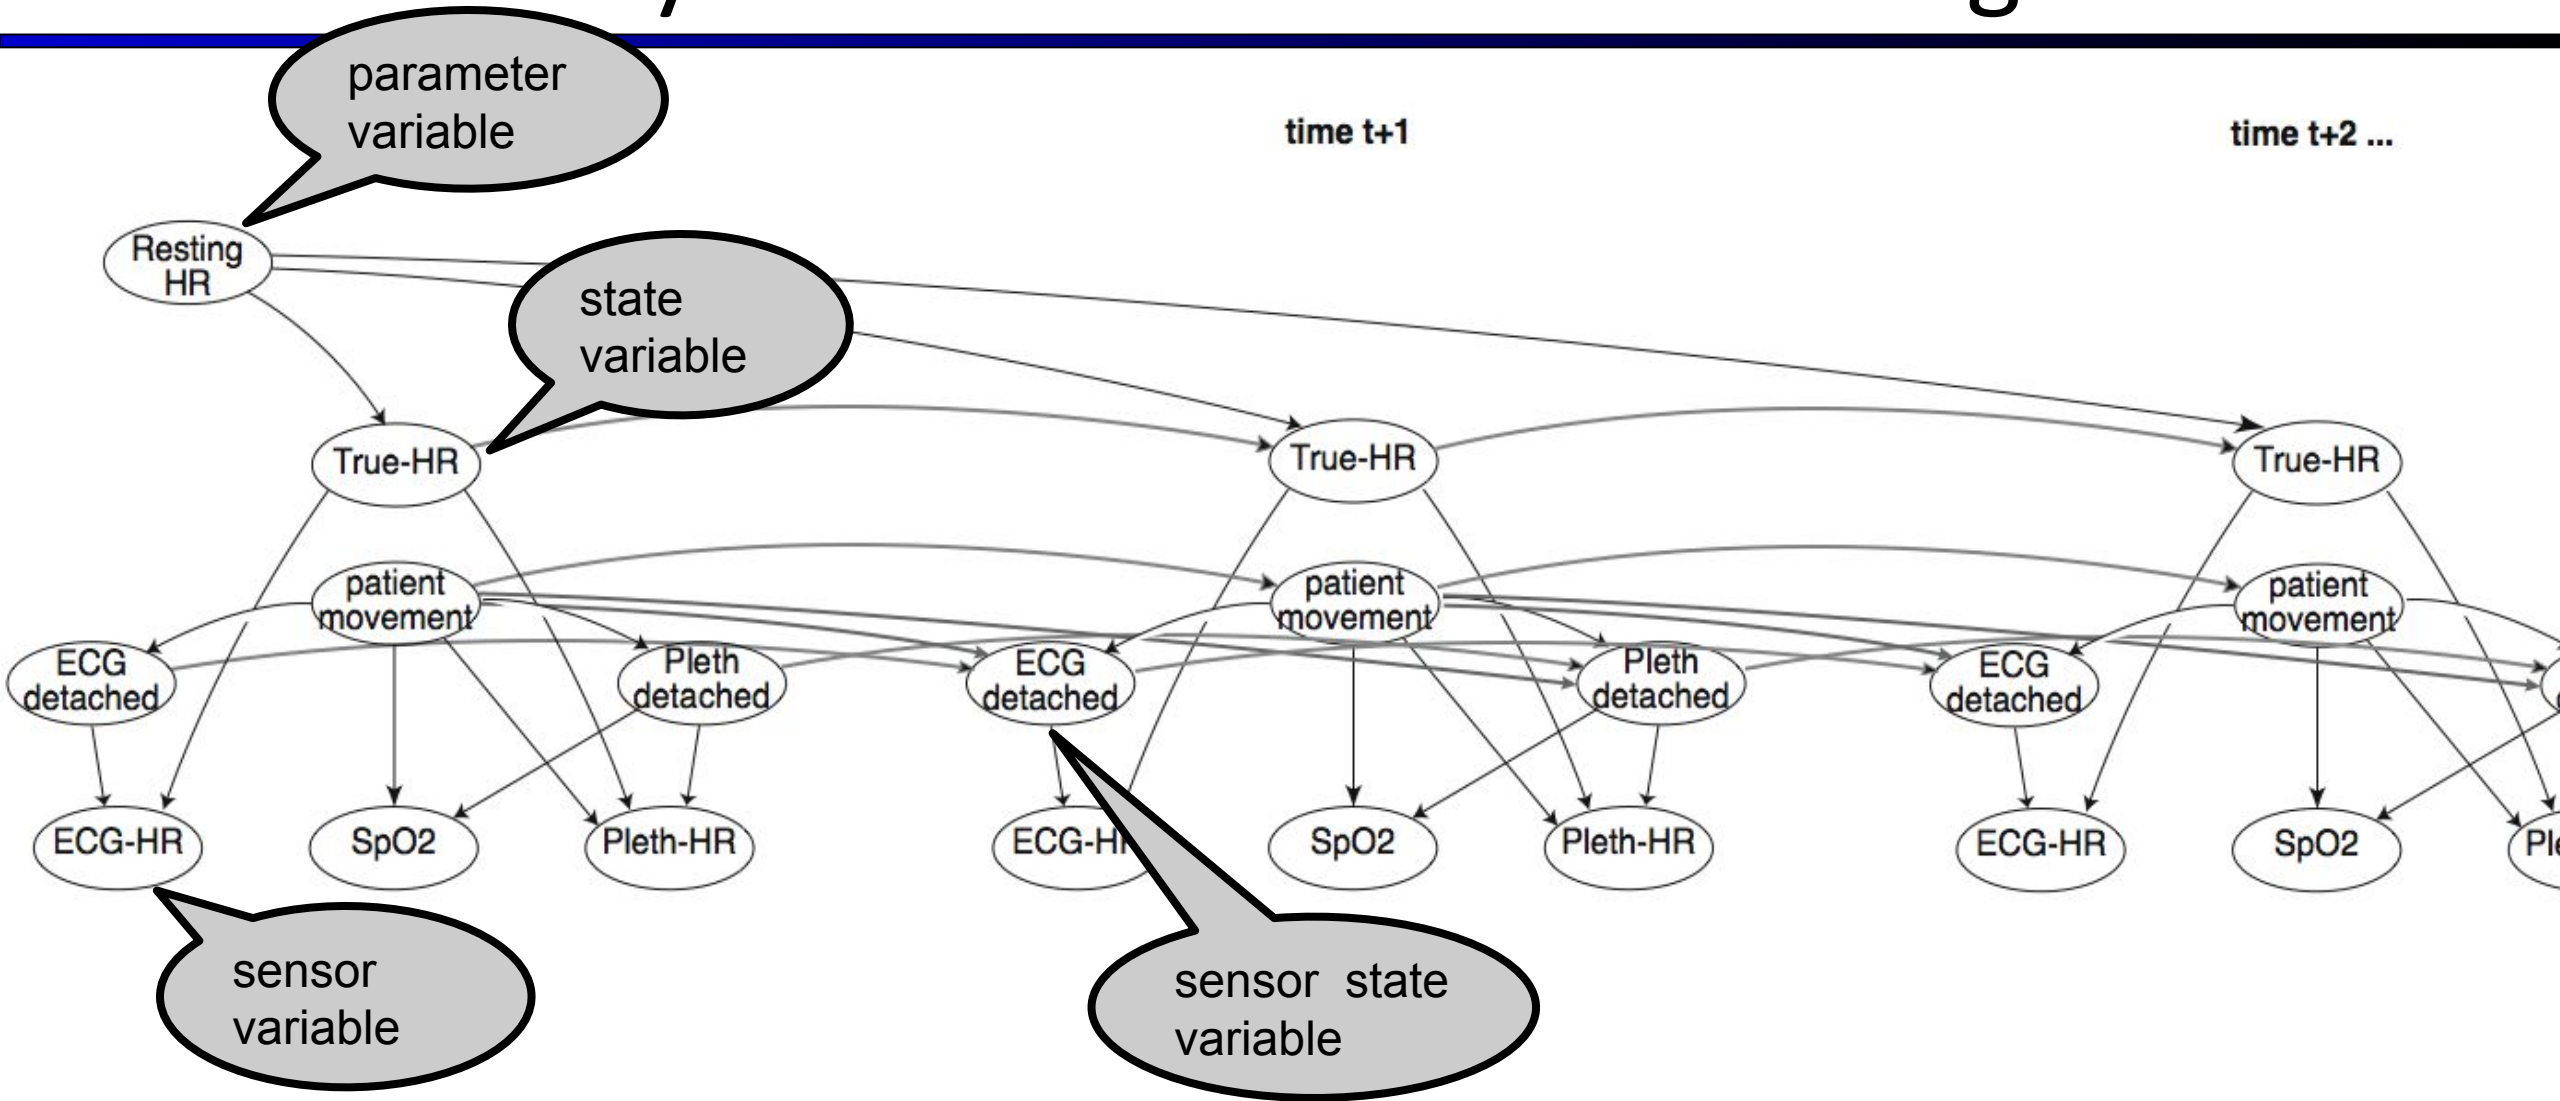
- Online: eliminate all variables from the previous time step; store factors for current time only
- Problem: largest factor contains all variables for current time (plus a few more)




Application: ICU monitoring

- ***State***: variables describing physiological state of patient
- ***Evidence***: values obtained from monitoring devices
- ***Transition model***: physiological dynamics, sensor dynamics
- ***Query variables***: pathophysiological conditions (a.k.a. bad things)


Toy DBN: heart rate monitoring




The enhanced heart-rate DBN's inferences on data from a healthy 40-year-old




ICU data: 22 variables, 1min ave







Particle Filtering




We need a new algorithm!

- When $|X|$ is more than 10^6 or so (e.g., 3 ghosts in a 10x20 world), exact inference becomes infeasible
- Likelihood weighting fails completely – number of samples needed grows *exponentially* with T



We need a new idea!





- The problem: sample state trajectories go off into low-probability regions, ignoring the evidence; too few “reasonable” samples
- Solution: kill the bad ones, make more of the good ones
- This way the population of samples stays in the high-probability region
- This is called *resampling* or survival of the fittest

Particle Filtering


- Represent belief state by a set of samples
 - Samples are called *particles*
 - Time per step is linear in the number of samples
 - But: number needed may be large
- This is how robot localization works in practice

0	0.1	0
0	0	0.2
0	0.2	0.5



Representation: Particles

- Our representation of $P(X)$ is now a list of $N \ll |X|$ particles
- $P(x)$ approximated by number of particles with value x
 - So, many x may have $P(x) = 0$!
 - More particles => more accuracy (cf. frequency histograms)
 - Usually we want a **low-dimensional** marginal
 - E.g., “Where is ghost 1?” rather than “Are ghosts 1,2,3 in [2,6], [5,6], and [8,11]?”



Particles:


(3,3)
(2,3)
(3,3)
(3,2)
(3,3)
(3,2)
(1,2)
(3,3)
(3,3)
(2,3)

Particle Filtering: Prediction step

- Particle j in state $x_t^{(j)}$ samples a new state directly from the transition model:
 - $x_{t+1}^{(j)} \sim P(X_{t+1} | x_t^{(j)})$
 - Here, most samples move clockwise, but some move in another direction or stay in place


Particles:

(3,3)
(2,3)
(3,3)
(3,2)
(3,3)
(3,2)
(1,2)
(3,3)
(3,3)
(2,3)



Particles:

(3,2)
(2,3)
(3,2)
(3,1)
(3,3)
(3,2)
(1,3)
(2,3)
(3,2)
(2,2)




Particle Filtering: Update step

- After observing e_{t+1} :
 - As in likelihood weighting, weight each sample based on the evidence
 - $w^{(j)} = P(e_{t+1} | x_{t+1}^{(j)})$
 - Normalize the weights: particles that fit the data better get higher weights, others get lower weights


Particles:

(3,2)
(2,3)
(3,2)
(3,1)
(3,3)
(3,2)
(1,3)
(2,3)
(3,2)
(2,2)




Particles:

(3,2) $w=.9$
(2,3) $w=.2$
(3,2) $w=.9$
(3,1) $w=.4$
(3,3) $w=.4$
(3,2) $w=.9$
(1,3) $w=.1$
(2,3) $w=.2$
(3,2) $w=.9$
(2,2) $w=.4$



Particle Filtering: Resample

- Rather than tracking weighted samples, we *resample*
- N times, we choose from our weighted sample distribution (i.e., draw with replacement)
- Now the update is complete for this time step, continue with the next one (with weights reset to 1)




Particles:

(3,2) $w=.9$
(2,3) $w=.2$
(3,2) $w=.9$
(3,1) $w=.4$
(3,3) $w=.4$
(3,2) $w=.9$
(1,3) $w=.1$
(2,3) $w=.2$
(3,2) $w=.9$
(2,2) $w=.4$


(New) Particles:

(3,2)
(2,2)
(3,2)
(2,3)
(3,3)
(3,2)
(1,3)
(2,3)
(3,2)
(3,2)




Summary: Particle Filtering

- Particles: track samples of states rather than an explicit distribution





Particle filtering on umbrella model



Robot Localization

- In robot localization:
 - We know the map, but not the robot's position
 - Observations may be vectors of range finder readings
 - State space and readings are typically continuous so we cannot usually represent or compute an exact posterior
 - Particle filtering is a main technique




Particle Filter Localization (Sonar)




Robot Mapping


- SLAM: Simultaneous Localization And Mapping
 - Robot does not know map or location
 - State $x_t^{(j)}$ consists of position+orientation, map!
 - (Each map usually inferred exactly given sampled position+orientation sequence)



DP-SLAM, Ron Parr




Particle Filter SLAM – Video 1



[Demo: PARTICLES-SLAM-mapping1-new.avi]

Particle Filter SLAM – Video 2



[Demo: PARTICLES-SLAM-fastslam.avi]