

CS 188: Artificial Intelligence

Markov Decision Processes



Instructor: Stuart Russell and Dawn Song

University of California, Berkeley

Recap: Decision Networks

- What is a decision network?

Recap: Decision Networks

- Decision network = Bayes net + Actions + Utilities




- Action nodes** (rectangles, cannot have parents, will have value fixed by algorithm)



- Utility nodes** (diamond, depends on action and chance nodes)

- Decision network represents a decision problem, containing all the information needed to for the agent to decide

What types of decisions?



Recap: Decision Networks

- Decision network = Bayes net + Actions + Utilities




- Action nodes** (rectangles, cannot have parents, will have value fixed by algorithm)



- Utility nodes** (diamond, depends on action and chance nodes)

- Decision network represents a decision problem, containing all the information needed to for the agent to decide

- What action to take given evidence e
 - Decision algorithm
 - Fix evidence e
 - For each possible action a
 - Fix action node to a
 - Compute posterior $P(W|e,a)$ for parents W of U
 - Compute expected utility $\sum_w P(w|e,a) U(a,w)$
 - Return action with highest expected utility



Recap: Decision Networks

- Decision network = Bayes net + Actions + Utilities




- Action nodes** (rectangles, cannot have parents, will have value fixed by algorithm)



- Utility nodes** (diamond, depends on action and chance nodes)

- Decision network represents a decision problem, containing all the information needed to for the agent to decide
 - What action to take given evidence e
 - Decision algorithm
 - Value of information
 - $VPI(E_i | e) = [\sum_{e_i} P(e_i | e) \max_a EU(a | e_i, e)] - \max_a EU(a | e)$



Decisions with unknown preferences

- In reality the assumption that we can write down our exact preferences for the machine to optimize is false
- A machine optimizing the wrong preferences causes problems




I'm sorry, Dave, I'm afraid I can't do that



Decisions with unknown preferences

- In reality the assumption that we can write down our exact preferences for the machine to optimize is false
- A machine optimizing the wrong preferences causes problems
- A machine that is explicitly uncertain about the human's preferences will defer to the human (e.g., allow itself to be switched off)

Off-switch problem (example)



$$EU(\text{act}) = +10$$

$$EU(\text{wait}) = (0.4 * 0) + (0.6 * 30) = +18$$


Off-switch problem (general proof)

- $EU(act) = \int_{-\infty}^{+\infty} P(u) \cdot u \, du = \int_{-\infty}^0 P(u) \cdot u \, du + \int_0^{+\infty} P(u) \cdot u \, du$
- $EU(wait) = \int_{-\infty}^0 P(u) \cdot 0 \, du + \int_0^{+\infty} P(u) \cdot u \, du$
- Obviously $\int_{-\infty}^0 P(u) \cdot u \, du \leq \int_{-\infty}^0 P(u) \cdot 0 \, du$
- Hence $EU(act) \leq EU(wait)$

Sequential decisions under uncertainty

So far, decision problem is one-shot --- concerning only one action

Sequential decision problem: agent's utility depends on a sequence of actions



Markov Decision Process (MDP)

- Environment history: $[s_0, a_0, s_1, a_1, \dots, s_t]$
- “Markov” generally means that given the present state, the future and the past are independent
- For Markov decision processes, “Markov” means action outcomes depend only on the current state

$$P(S_{t+1} = s' | S_t = s_t, A_t = a_t, S_{t-1} = s_{t-1}, A_{t-1}, \dots, S_0 = s_0)$$

=

$$P(S_{t+1} = s' | S_t = s_t, A_t = a_t)$$


- This is just like search, where the successor function could only depend on the current state (not the history)



Andrey Markov
(1856-1922)

Markov Decision Process (MDP)


- An MDP is defined by:
 - A set of states $s \in S$
 - A set of actions $a \in A$
 - A transition model $T(s, a, s')$
 - Probability that a from s leads to s' , i.e., $P(s' | s, a)$
 - A reward function $R(s, a, s')$ for each transition
 - A start state
 - Possibly a terminal state (or absorbing state)
 - Utility function which is additive (discounted) rewards



- MDPs are fully observable but probabilistic search problems


Example: Grid World

- A maze-like problem
 - The agent lives in a grid
 - Walls block the agent's path
- Noisy movement: actions do not always go as planned
 - 80% of the time, the action North takes the agent North (if there is no wall there)
 - 10% of the time, North takes the agent West; 10% East
 - If there is a wall in the direction the agent would have been taken, the agent stays put
- The agent receives rewards each time step
 - Small “living” reward r each step (can be negative)
 - Big rewards come at the end (good or bad)
- Goal: maximize sum of rewards




Policies

- A policy π gives an action for each state, $\pi: S \rightarrow A$
- In deterministic single-agent search problems, we wanted an optimal **plan**, or sequence of actions, from start to a goal
- For MDPs, we want an optimal **policy** $\pi^*: S \rightarrow A$
 - An optimal policy maximizes expected utility
 - An explicit policy defines a reflex agent




Optimal policy for $r > 0$



$r > 0$

Optimal policy for $r > 0$



$r > 0$

Sample Optimal Policies

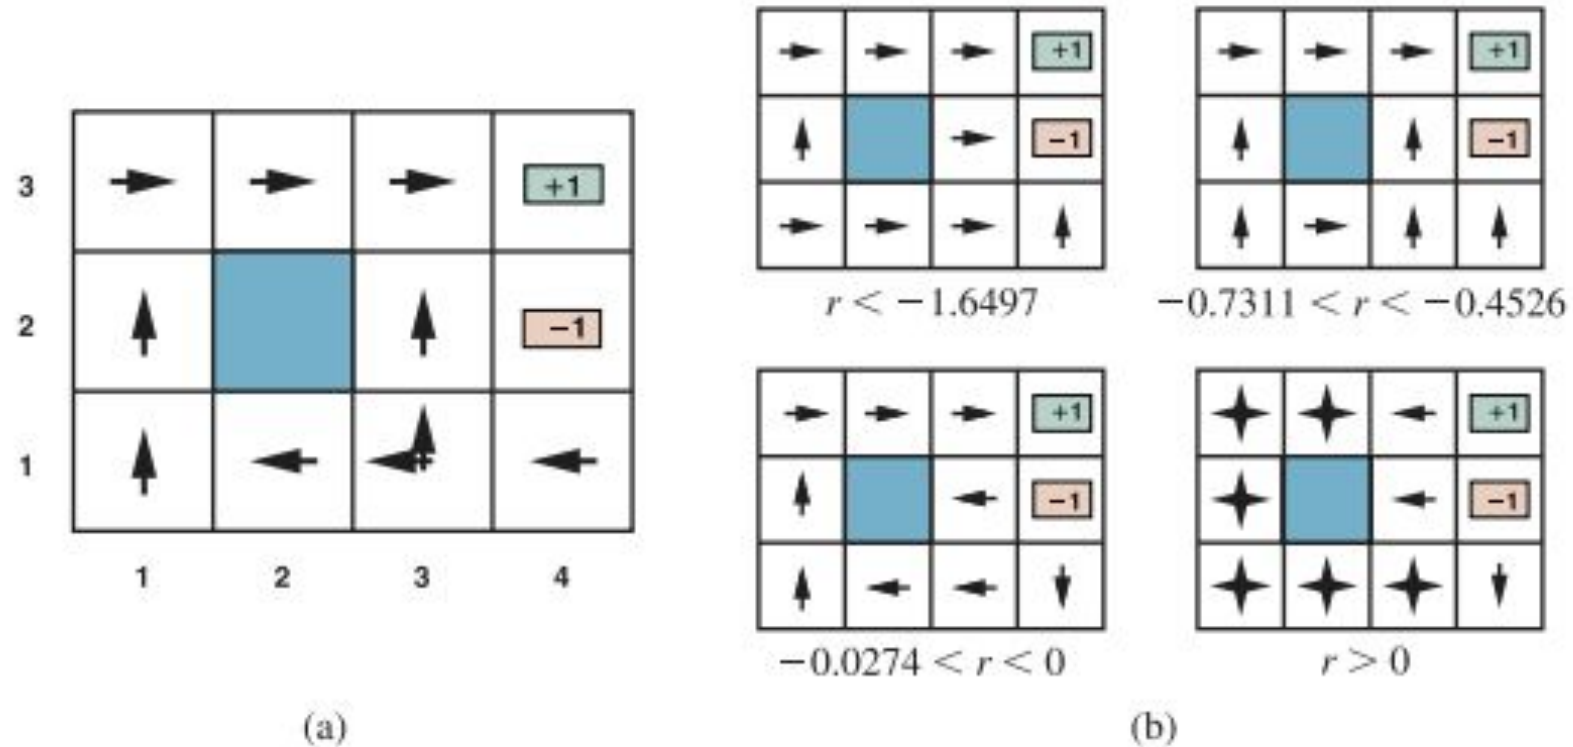




Figure 17.2 (a) The optimal policies for the stochastic environment with $r = -0.04$ for transitions between nonterminal states. There are two policies because in state (3,1) both *Left* and *Up* are optimal. (b) Optimal policies for four different ranges of r .

Utilities of Sequences



Utilities of Sequences

- What preferences should an agent have over reward sequences?
- More or less? $[1, 2, 2]$ or $[2, 3, 4]$
- Now or later? $[0, 0, 1]$ or $[1, 0, 0]$



Stationary Preferences

- Theorem: if we assume **stationary preferences**:


$$[s_0, a_0, s_1, a_1, s_2, \dots] > [s'_0, a'_0, s'_1, a'_1, s'_2, \dots], \quad s_0 = s'_0, a_0 = a'_0, \text{ and } s_1 = s'_1$$
$$\Leftrightarrow [s_1, a_1, s_2, \dots] > [s'_1, a'_1, s'_2, \dots]$$

then there is only one way to define utilities:

- **Additive discounted utility:**

$$U_h([s_0, a_0, s_1, a_1, s_2, \dots]) = R(s_0, a_0, s_1) + \gamma R(s_1, a_1, s_2) + \gamma^2 R(s_2, a_2, s_3) + \dots$$

where $\gamma \in [0, 1]$ is the **discount factor**



Discounting



Worth r now



Worth γr next step




Worth $\gamma^2 r$ in two steps

- Discounting conveniently solves the problem of infinite reward streams!
 - Geometric series: $1 + \gamma + \gamma^2 + \dots = 1/(1 - \gamma)$
 - Assume rewards bounded by $\pm R_{\max}$
 - Then $r_0 + \gamma r_1 + \gamma^2 r_2 + \dots$ is bounded by $\pm R_{\max}/(1 - \gamma)$
- (Another solution: environment contains a **terminal state**; **and** agent reaches it with probability 1)

Quiz: Discounting

- Given:



- Actions: East, West, and Exit (only available in exit states a, e)
- Transitions: deterministic

- Quiz 1: For $\gamma = 1$, what is the optimal policy?



- Quiz 2: For $\gamma = 0.1$, what is the optimal policy?



- Quiz 3: For which γ are West and East equally good when in state d?

The utility of a policy

- Executing a policy π from any state s_0 generates a sequence

$s_0, \pi(s_0), s_1, \pi(s_1), s_2, \dots$

- This corresponds to a sequence of rewards

$R(s_0, \pi(s_0), s_1), R(s_1, \pi(s_1), s_2), \dots$


- This reward sequence happens with probability

$P(s_1 | s_0, \pi(s_0)) \times P(s_2 | s_1, \pi(s_1)) \times \dots$

- The value (expected utility) of π in s_0 is written $U^\pi(s_0)$

- It's the sum over all possible state sequences of
(discounted sum of rewards) \times (probability of state sequence)

$$U^\pi(s_0) = E \left[\sum_{t=0}^{\infty} \gamma^t R(s_t, \pi(s_t), s_{t+1}) \right]$$



Optimal Quantities

- The optimal policy:

$\pi^*(s)$ = optimal action from state s

Gives highest $U^\pi(s)$ for any π


- The value (utility) of a state s :

$U^*(s) = U^{\pi^*}(s)$ = expected utility starting in s and acting optimally

- The value (utility) of a q-state (s,a) :

$Q^*(s,a)$ = expected utility of taking action a in state s and (thereafter) acting optimally

$U^*(s) = \max_a Q^*(s,a)$



Bellman equations (Shapley, 1953)

- The value/utility of a state is
 - The expected reward for the next transition plus the discounted value/utility of the next state, assuming the agent chooses the optimal action
- Hence we have a recursive definition of value (Bellman equation):

$$U(s) = \max_{a \in A(s)} \sum_{s'} P(s' | s, a) [R(s, a, s') + \gamma U(s')]$$

- Similarly, Bellman equation for Q-functions

$$\begin{aligned} Q(s, a) &= \sum_{s'} P(s' | s, a) [R(s, a, s') + \gamma U(s')] \\ &= \sum_{s'} P(s' | s, a) [R(s, a, s') + \gamma \max_{a'} Q(s', a')] \end{aligned}$$