
EECS 16A Designing Information Devices and Systems I

Spring 2021 Homework 3

This homework is due February 12, 2021, at 23:59.

Self-grades are due February 16, 2021, at 23:59.

Submission Format

Your homework submission should consist of **one** file.

- `hw3.pdf`: A single PDF file that contains all of your answers (any handwritten answers should be scanned) as well as your IPython notebook saved as a PDF.
If you do not attach a PDF “printout” of your IPython notebook, you will not receive credit for problems that involve coding. Make sure that your results and your plots are visible. Assign the IPython printout to the correct problem(s) on Gradescope.

Submit each file to its respective assignment on Gradescope.

The first question on this homework gives you a chance to practice the zoom proctoring setup for exams. You may want to start with this question

1. Exam Policy and Practice

Please read through the entirety of the [EECS 16A exam proctoring policies](#) carefully before proceeding. This question is designed to familiarize you with how the exam will be run and help you setup and practice.

- (a) After reading through the proctoring policies carefully, please answer the following questions.
 - i. If you experience no disruptions during the exam, how many total minutes do you have for scanning and submission? What if there is a disruption?
 - ii. Are you required to record locally during the exam? How much space should you have available on your computer for a local recording?
 - iii. How should you contact the course staff in case of an emergency situation during the exam?
- (b) Please configure your Zoom link.
 - i. Fill out the following [Google form](#) to submit the Zoom link you will be using. You must use this Zoom link for this assignment as well as for the exams. If you wish, you may use your Personal Meeting Room link and set your Personal Meeting ID as your default on all devices (desktop + laptop + phone) you will be using for the exams.
 - ii. Ensure that anyone can join your Zoom link and that there is no waiting room for your Zoom meeting. Try to do this by entering the meeting on one device that is logged in to your official Berkeley Zoom account and then entering the meeting on another device that is logged into some other Zoom account. If you are able to join automatically, then your Zoom link is *joinable*. If you are not put into a waiting room, then your Zoom meeting will not have a waiting room. (You might want to have a member of your study group try this out with you if you don't already have two Zoom accounts.)
- (c) You will now practice a Zoom recording. You should use this recording to work through a homework problem or other study material to simulate the actual circumstances of the final exam.

- i. Start the Zoom call for the link you provided above. Turn on your microphone and recording device (webcam, phone camera). You may turn off your speaker. Please share your entire desktop (not just a particular window). Your video should be visible on the desktop and at maximum size. Please refer to the image below.
- ii. Start recording via Zoom. You may record locally or on the cloud (see the exam policy document for details).
- iii. Hold your Cal ID or any photo ID (as detailed in the exam policy document) next to your face and record yourself saying your name and SID number into the webcam. Both your face and your entire Cal ID should be visible in the video. We should be able to read your name and SID. This step should take **at least 3 seconds**. See Fig. 1.

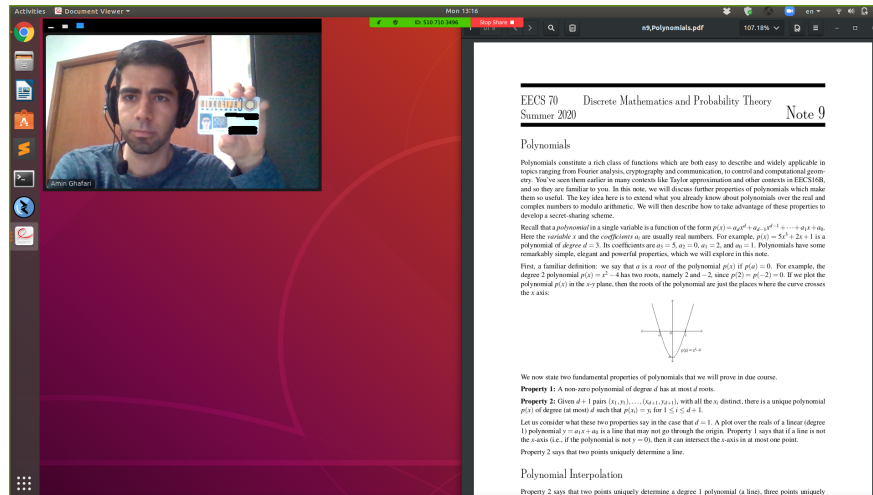


Figure 1: ID card demonstration. Do not black out your SID and name.

- iv. Turn your recording device (webcam, phone) around 360° **slowly** so that we can see your entire room clearly. There should be no uncovered screens anywhere in the room during your exam. Only select members of course staff will be able to see your videos (both for this assignment and for the actual exams).
- v. Position your recording device in such a way that we can see your workspace and your hands. It is perfectly fine if your face is not visible at this point. If you are not using your phone to record your workspace, then it should be visible in the recording, face down. See Figure 2. **Think about how you want to set this up during this test run. On the actual exam, you will want to use the computer to see the exam itself, so make sure this works for you. If you are using a laptop's built-in webcam to record and also see the exam, make sure you have a setup that works for you. Think about how you might position the laptop. You may also consider using your phone or an external webcam on a stand to record your workspace. We want you to iron out these details ahead of the actual exam so that the exam itself has no additional stress due to proctoring logistics. Please contact STEP if you would like to request an external webcam.**

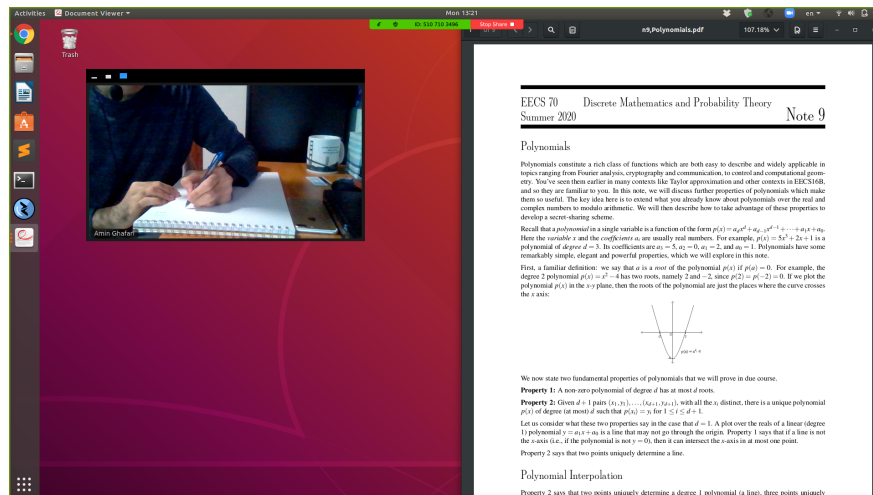


Figure 2: Demonstration of taking your exam. Your setup should look like this while you are taking the exam.

- vi. Your microphone should be on at all times. The recording should also include the time on your desktop at all times.
 - vii. Record for a full two and a half hours. We recommend you use this time to work through a homework problem or other study material for the course to simulate an exam-like environment. (We also want to make sure that your computer can handle the video encoding task if you are doing a local recording.)
 - viii. After two and a half hours, make sure your face is back in your recording and then stop the recording. In the actual exam, prior to ending the recording you will scan and submit your exam to Gradescope.
 - ix. After stopping the recording, check your recording to confirm that it contains your video as well as your desktop throughout its duration. Upload your video to Google drive (if saved locally) and submit a link to the video using this [Google form](#). Please name your file using the studentname-exam naming format described in the proctoring instructions. **Please make sure that the link sharing option is set to "Anybody at UC Berkeley with the link can view."** Again, we will not be sharing these links except with a small subset of course staff.
- (d) A Midterm Google document should be shared with you with instructions for the midterm a few days before the midterm. More details will be made available closer to the exam date.

Link for policy:

<https://docs.google.com/document/d/10pnWwxyZ40nlpbCM4aOYTxxOjc36sIQaMx9m8zyaR8w/edit?usp=sharing>

Form to submit Zoom link:

<https://forms.gle/P53yLv0xnW8j672W9>

Form to submit video link:

<https://forms.gle/21MaHrjyTEVMBenQ8>

- 2. Reading Assignment** For this homework, please read Notes 3-5. Notes 3 and 4 will give you an overview of linear independence, span and an introduction to thinking about and writing proofs. Note 5 will give you an overview of pump systems and transition matrices. You are always welcome and encouraged to read beyond this as well. Please write a few sentences about :

- (a) how you can use the strategies in the notes to tackle proof questions
- (b) how we can interpret the rows and columns of a state transition matrix

3. Multiply the Matrices

Learning Objective: Practice evaluating matrix-matrix multiplication.

- (a) We have two matrices **A** and **B**, where **A** is a 3×2 matrix and **B** is a 2×4 matrix. Would the multiplication **AB** be a valid operation? If yes, what do you expect the dimensions of **AB** to be?
- (b) Compute **AB** by hand, where **A** and **B** are given by

$$\mathbf{A} = \begin{bmatrix} 1 & 0 \\ 2 & 1 \\ 0 & 1 \end{bmatrix}, \text{ and } \mathbf{B} = \begin{bmatrix} 1 & 2 & -1 & 0 \\ -3 & 0 & 2 & -1 \end{bmatrix}$$

Compute **BA** too if the operation is valid. If it is invalid, explain why. Make sure you show the work for your calculations.

- (c) Now let us assume $\mathbf{A} \in \mathbb{R}^{2 \times n}$ is a **new matrix with 2 rows**, which are given by the **transposes** of column vectors \vec{r}_1, \vec{r}_2 i.e.

$$\mathbf{A} = \begin{bmatrix} - & \vec{r}_1^T & - \\ - & \vec{r}_2^T & - \end{bmatrix} \quad \text{where,} \quad \vec{r}_1 = \begin{bmatrix} r_{11} \\ r_{12} \\ \vdots \\ r_{1n} \end{bmatrix}, \text{ and } \vec{r}_2 = \begin{bmatrix} r_{21} \\ r_{22} \\ \vdots \\ r_{2n} \end{bmatrix}.$$

$\mathbf{B} \in \mathbb{R}^{n \times 3}$ is a **new matrix with 3 columns**, which are called \vec{c}_1, \vec{c}_2 , and \vec{c}_3 , i.e.

$$\mathbf{B} = \begin{bmatrix} | & | & | \\ \vec{c}_1 & \vec{c}_2 & \vec{c}_3 \\ | & | & | \end{bmatrix} \quad \text{where,} \quad \vec{c}_1 = \begin{bmatrix} c_{11} \\ c_{12} \\ \vdots \\ c_{1n} \end{bmatrix}, \quad \vec{c}_2 = \begin{bmatrix} c_{21} \\ c_{22} \\ \vdots \\ c_{2n} \end{bmatrix}, \text{ and } \vec{c}_3 = \begin{bmatrix} c_{31} \\ c_{32} \\ \vdots \\ c_{3n} \end{bmatrix}$$

Now show that:

$$\mathbf{AB} = \begin{bmatrix} \vec{r}_1^T \vec{c}_1 & \vec{r}_1^T \vec{c}_2 & \vec{r}_1^T \vec{c}_3 \\ \vec{r}_2^T \vec{c}_1 & \vec{r}_2^T \vec{c}_2 & \vec{r}_2^T \vec{c}_3 \end{bmatrix}$$

if **AB** is a valid operation.

4. Linear Dependence in a Square Matrix

Learning Objective: This is an opportunity to practice applying proof techniques. This question is specifically focused on linear dependence of rows and columns in a square matrix.

Let A be a square $n \times n$ matrix, (i.e. both the columns and rows are vectors in \mathbb{R}^n). Suppose we are told that the columns of A are linearly dependent. Prove, then, that the rows of A must also be linearly dependent. You can use the following conclusion in your proof:

If Gaussian elimination is applied to a matrix A , and the resulting matrix (in reduced row echelon form) has at least one row of all zeros, this means that the rows of A are linearly dependent.

(Hint: Can you use the linear dependence of the columns to say something about the number of solutions to $A\vec{x} = \vec{0}$? How does the number of solutions relate to the result of Gaussian elimination?)

5. Kinematic Model for a Simple Car

Learning Objective: Many real world systems are not actually linear and have more complex behaviors. However, linear models can approximate non-linear systems under certain conditions.

Building a self-driving car first requires understanding the basic motions of a car. In this problem, we will explore how to model the motion of a car.

There are several models that we can use to model the motion of a car. Assume we use a kinematic model, described in the following four equations and Figure 3.

$$x[k+1] = x[k] + v[k] \cos(\theta[k]) \Delta t \quad (1)$$

$$y[k+1] = y[k] + v[k] \sin(\theta[k]) \Delta t \quad (2)$$

$$\theta[k+1] = \theta[k] + \frac{v[k]}{L} \tan(\phi[k]) \Delta t \quad (3)$$

$$v[k+1] = v[k] + a[k] \Delta t \quad (4)$$




Figure 3: Vehicle Kinematic Model

where

- k , a nonnegative integer, indicates the time step at which we measure each variable (e.g. $v[k]$ is the speed at time step k and $v[k+1]$ is the speed at the following time step)
- $x[k]$ and $y[k]$ denote the coordinates of the vehicle (meters)
- $\theta[k]$ denotes the heading of the vehicle, or the angle with respect to the x-axis (radians)
- $v[k]$ is the speed of the car (meters per second)
- $a[k]$ is the acceleration of the car (meters per second squared)
- $\phi[k]$ is the steering angle input we command (radians)
- Δt is a constant measuring the time difference (in seconds) between time steps $k+1$ and k
- L is a constant and is the length of the car (in meters)

For this problem, let L be 1.0 meter and Δt be 0.1 seconds.

The variables $x[k], y[k], \theta[k], v[k]$ describe the **state** of the car at time step k . The state captures all the information needed to fully determine the current position, speed, and heading of the car. The **inputs** at

time step k are $a[k]$ and $\phi[k]$. These are provided by the driver. The current value of these inputs, along with the current state of the vehicle, will determine the state of the vehicle at the next time step.

We note that the problem is nonlinear, due to the sine, cosine and tangent functions, as well as terms including the product of states and inputs.

The purpose of this problem is to show that we can approximate a nonlinear model with a simple linear model and do reasonably well. This is why, despite many systems being nonlinear, linear algebra tools are widely used in practice.

For Parts (b) - (d), fill out the corresponding sections in prob3.ipynb.

- (a) We assume that the car has a small heading, θ , which is a **very small but nonzero** value, and that the steering angle ϕ is also **very small but nonzero**. In this case, we could use the following approximations:

$$\sin(\alpha) \approx 0,$$

$$\cos(\alpha) \approx 1,$$

$$\tan(\alpha) \approx 0.$$

where α is the small angle of interest, and \approx means “approximately equal to”. Here, we use a very simple approximation for small angles; in later classes, you may learn better approximations.

Draw, by hand, graphs of $\sin(\alpha)$ and $\cos(\alpha)$, for α ranging from $-\pi$ to π . Using these graphs can you justify the approximation we are making for small values of α ?

- (b) Applying the approximation described in the previous part, write down a system of linear equations that approximates the nonlinear vehicle model given above in Equations (1) to (4). In particular, find the 4×4 matrix \mathbf{A} and 4×2 matrix \mathbf{B} that satisfy the equation given below.

$$\begin{bmatrix} x[k+1] \\ y[k+1] \\ \theta[k+1] \\ v[k+1] \end{bmatrix} = \mathbf{A} \begin{bmatrix} x[k] \\ y[k] \\ \theta[k] \\ v[k] \end{bmatrix} + \mathbf{B} \begin{bmatrix} a[k] \\ \phi[k] \end{bmatrix}$$

Hint: Start with simplifying Equations (1) to (4), using the $\sin(\alpha)$, $\cos(\alpha)$, and $\tan(\alpha)$ approximations from part (a) only to approximate nonlinear terms. Do NOT replace ϕ , θ with 0 unless it is for the $\sin(\alpha)$, $\cos(\alpha)$, and $\tan(\alpha)$ approximations.

- (c) Suppose we drive the car so that the direction of travel is aligned with the x-axis, and we are driving nearly straight, i.e. the steering angle is $\phi[k] = 0.0001$ radians. (Driving exactly straight would have the steering angle $\phi[k] = 0$ radians.) The initial state and input are:

$$\begin{bmatrix} x[0] \\ y[0] \\ \theta[0] \\ v[0] \end{bmatrix} = \begin{bmatrix} 5.0 \\ 10.0 \\ 0.0 \\ 2.0 \end{bmatrix}$$

$$\begin{bmatrix} a[k] \\ \phi[k] \end{bmatrix} = \begin{bmatrix} 1.0 \\ 0.0001 \end{bmatrix}$$

You can use these values in the IPython notebook to compare how the nonlinear system evolves in comparison to the linear approximation that you made. The IPython notebook simulates the car for ten time steps. Are the trajectories similar or very different? Why?

- (d) Now suppose we drive the vehicle from the same starting state, but we turn left instead of going straight, i.e. the steering angle is $\phi[k] = 0.5$ radians. The initial state and input are:

$$\begin{bmatrix} x[0] \\ y[0] \\ \theta[0] \\ v[0] \end{bmatrix} = \begin{bmatrix} 5.0 \\ 10.0 \\ 0.0 \\ 2.0 \end{bmatrix}$$

$$\begin{bmatrix} a[k] \\ \phi[k] \end{bmatrix} = \begin{bmatrix} 1.0 \\ 0.5 \end{bmatrix}$$

You can use these values in the IPython notebook to compare how the nonlinear system evolves in comparison to the linear approximation that you made. The IPython notebook simulates the car for ten time steps. Are the trajectories similar or very different? Why?

6. Image Stitching

Learning Objective: This problem is similar to one that students might experience in an upper division computer vision course. Our goal is to give students a flavor of the power of tools from fundamental linear algebra and their wide range of applications.

Often, when people take pictures of a large object, they are constrained by the field of vision of the camera. This means that they have two options to capture the entire object:

- Stand as far away as they need to include the entire object in the camera's field of view (clearly, we do not want to do this as it reduces the amount of detail in the image)
- (This is more exciting) Take several pictures of different parts of the object and stitch them together like a jigsaw puzzle.

We are going to explore the second option in this problem. Daniel, who is a professional photographer, wants to construct an image by using "image stitching". Unfortunately, Daniel took some of the pictures from different angles as well as from different positions and distances from the object. While processing these pictures, Daniel lost information about the positions and orientations from which the pictures were taken. Luckily, you and your friend Marcela, with your wealth of newly acquired knowledge about vectors and matrices, can help him!

You and Marcela are designing an iPhone app that stitches photographs together into one larger image. Marcela has already written an algorithm that finds common points in overlapping images. **It's your job to figure out how to stitch the images together using Marcela's common points to reconstruct the larger image.**

We will use vectors to represent the common points which are related by a linear transformation. Your idea is to find this linear transformation. For this you will use a single matrix, \mathbf{R} , and a vector, \vec{t} , that transforms every common point in one image to their corresponding point in the other image. Once you find \mathbf{R} and \vec{t} you will be able to transform one image so that it lines up with the other image.

Suppose $\vec{p} = \begin{bmatrix} p_x \\ p_y \end{bmatrix}$ is a point in one image, which is transformed to $\vec{q} = \begin{bmatrix} q_x \\ q_y \end{bmatrix}$ is the corresponding point in the other image (i.e., they represent the same object in the scene). For example, Fig. 4 shows how the points $\vec{p}_1, \vec{p}_2 \dots$ in the right image are transformed to points $\vec{q}_1, \vec{q}_2 \dots$ on the left image. You write down the following relationship between \vec{p} and \vec{q} .




Figure 4: Two images to be stitched together with pairs of matching points labeled.

$$\begin{bmatrix} q_x \\ q_y \end{bmatrix} = \underbrace{\begin{bmatrix} r_{xx} & r_{xy} \\ r_{yx} & r_{yy} \end{bmatrix}}_{\mathbf{R}} \begin{bmatrix} p_x \\ p_y \end{bmatrix} + \underbrace{\begin{bmatrix} t_x \\ t_y \end{bmatrix}}_{\vec{t}} \quad (5)$$

This problem focuses on finding the unknowns (i.e. the components of \mathbf{R} and \vec{t}), so that you will be able to stitch the image together.

- (a) To understand how the matrix \mathbf{R} and vector \vec{t} transforms any vector representing a point on a image, Consider this equation similar to Equation (5),

$$\vec{v} = \begin{bmatrix} 2 & 2 \\ -2 & 2 \end{bmatrix} \vec{u} + \vec{w} = \vec{v}_1 + \vec{w}. \quad (6)$$

Use $\vec{w} = \begin{bmatrix} 0 \\ 1 \end{bmatrix}$, $\vec{u} = \begin{bmatrix} 1 \\ 1 \end{bmatrix}$ for this part.

We want to find out what geometric transformation(s) can be applied on \vec{u} to give \vec{v} .

Step 1: Find out how $\begin{bmatrix} 2 & 2 \\ -2 & 2 \end{bmatrix}$ is transforming \vec{u} . Evaluate $\vec{v}_1 = \begin{bmatrix} 2 & 2 \\ -2 & 2 \end{bmatrix} \vec{u}$.

What **geometric transformation(s)** might be applied to \vec{u} to get \vec{v}_1 ? Choose the option that answers the question and explain your choice.

- (i) Rotation
- (ii) Scaling
- (iii) Shifting/Translation
- (iv) Rotation and Scaling

Drawing the vectors \vec{u} , and \vec{v}_1 in two dimensions on a single plot might help you to visualize the transformations. You can also look into the corresponding demo in the IPython notebook prob3.ipynb.

Step 2: Find out $\vec{v} = \vec{v}_1 + \vec{w}$. Find out how **addition of \vec{w} is geometrically transforming \vec{v}_1** . Choose the option that answers the question and explain your choice.

- (i) Rotation
- (ii) Scaling

(iii) Shifting/ Translation.

Drawing the vectors \vec{v} , \vec{w} , and \vec{v}_1 in two dimensions on a single plot might help you to visualize the transformations. You can also look into the corresponding demo in the IPython notebook `prob3.ipynb`.

(b) Multiply Equation (5) out into **two scalar linear equations**.

- (i) What are the known values and what are the unknowns in each equation?
- (ii) How many unknowns are there?
- (iii) How many independent equations do you need to solve for all the unknowns?
- (iv) How many pairs of common points \vec{p} and \vec{q} will you need in order to write down a system of equations that you can use to solve for the unknowns? **Hint:** Remember that each pair of \vec{p} and \vec{q} will give you **two** different linear equations.

- (c) Write out a system of linear equations that you can use to solve for $\vec{\alpha} = \begin{bmatrix} r_{xx} \\ r_{xy} \\ r_{yx} \\ r_{yy} \\ t_x \\ t_y \end{bmatrix}$. Assume that all four

pairs of points from Fig. 4 are labeled as:

$$\vec{q}_1 = \begin{bmatrix} q_{1x} \\ q_{1y} \end{bmatrix}, \vec{p}_1 = \begin{bmatrix} p_{1x} \\ p_{1y} \end{bmatrix} \quad \vec{q}_2 = \begin{bmatrix} q_{2x} \\ q_{2y} \end{bmatrix}, \vec{p}_2 = \begin{bmatrix} p_{2x} \\ p_{2y} \end{bmatrix} \quad \vec{q}_3 = \begin{bmatrix} q_{3x} \\ q_{3y} \end{bmatrix}, \vec{p}_3 = \begin{bmatrix} p_{3x} \\ p_{3y} \end{bmatrix} \quad \vec{q}_4 = \begin{bmatrix} q_{4x} \\ q_{4y} \end{bmatrix}, \vec{p}_4 = \begin{bmatrix} p_{4x} \\ p_{4y} \end{bmatrix}.$$

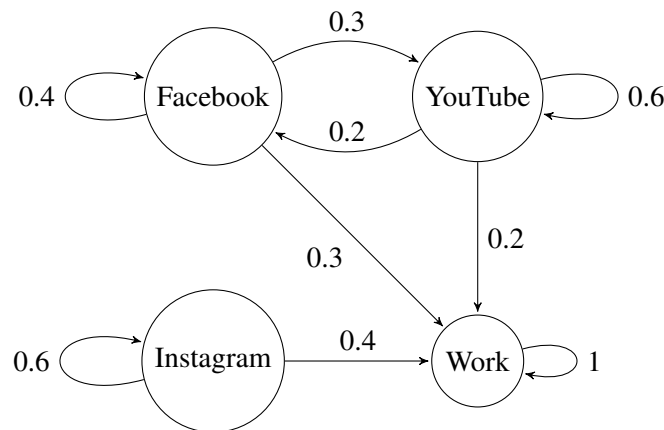
Now think of your answer to Part b(iv). How many pairs of these points would you need to solve for $\vec{\alpha}$. Choose **just enough** equations required to solve for $\vec{\alpha}$ and express these linear equations using **matrix-vector form**.

- (d) In the IPython notebook `prob3.ipynb`, you will have a chance to test out your solution. Plug in the values that you are given for p_x , p_y , q_x , and q_y for each pair of points into your system of equations to solve for the matrix, \mathbf{R} , and vector, \vec{t} . The notebook will solve the system of equations, apply your transformation to the second image, and show you if your stitching algorithm works. **You are NOT responsible for understanding the image stitching code or Marcela's algorithm.**

7. Social Media

Learning Objective: Practice setting up transition matrices from a diagram and understand how to compute subsequent states of the system.

As a tech-savvy Berkeley student, the distractions of social media are always calling you away from productive stuff like homework for your classes. You're curious—are you the only one who spends hours switching between Facebook or YouTube? How do other students manage to get stuff done and balance pursuing Insta-fame? You conduct an experiment, collect some data, and notice Berkeley students tend to follow a pattern of behavior similar to the figure below. So, for example, if 100 students are on Facebook, in the next timestep, 30 of them will click on a link and move to YouTube.



- (a) Let us define $x_F[n]$ as the number of students on Facebook at timestep n , $x_Y[n]$ as the number of students on YouTube at timestep n , $x_I[n]$ as the number of students on Instagram at timestep n , and $x_W[n]$ as the number of students working at timestep n . Let the state vector be: $\vec{x}[n] = \begin{bmatrix} x_F[n] \\ x_Y[n] \\ x_I[n] \\ x_W[n] \end{bmatrix}$. Derive the corresponding transition matrix.
- (b) There are 1500 of you in the class. Suppose on a given Friday evening (the day when HW is due), there are 700 EECS16A students on Facebook, 450 on YouTube, 200 on Instagram, and 150 actually doing work. In the next timestep, how many people will be doing each activity? In other words, after you apply the matrix once to reach the next timestep, what is the state vector?
- (c) Compute the sum of each column in the state transition matrix. What is the interpretation of this?


8. [PRACTICE/OPTIONAL] Splitting The Tips

Learning Objective: This problem showcases how you can understand general systems of equations by looking at simpler examples. In particular, see if you can generalize your intuition from the case of 5 and 6 guests to a general number of guests.


A number of guests gather around a table for a dinner. Between every adjacent pair of guests, there is a jar for tips. When everyone has finished eating, each person places half their tip in the jar to their left and half in the jar to their right. Suppose you can only see the amount of tips in each jar after everyone has left. Can you deduce the amount that each individual tipped?

Note: For this question, if we assume that tips are positive, then we need to introduce additional constraint that would make the system of equations no longer linear. We are going to ignore this constraint and assume that negative tips are acceptable.

- (a) Suppose six guests (represented by green circles) sit around a hexagonal table and there are six jars of tips (represented by black dots). If we know the amount of tip in each jar, J_1 to J_6 , can we determine each individual's tip amount, G_1 to G_6 ? If yes, explain why by examining the relationship between the jar values, J_1 to J_6 , and guest tips, G_1 to G_6 . If not, give two different assignments of G_1 to G_6 that will result in the same J_1 to J_6 .



- (b) Now let's consider five guests around a pentagonal table, G_1 to G_5 , and we can see the amount of tips in the five jars, J_1 to J_5 . In this new setting can you figure out each guest's tip values, G_1 to G_5 ?



- (c) *This part will challenge you to further reason about and generalize the results you obtained in parts a and b.*

If n is the total number of guests sitting around a table, for which values of n can you figure out everyone's tip? You do not have to rigorously prove your answer. (**Hint:** consider what is different between parts a and b.)

9. Homework Process and Study Group

Who did you work with on this homework? List names and student ID's. (In case you met people at homework party or in office hours, you can also just describe the group.) How did you work on this homework? If you worked in your study group, explain what role each student played for the meetings this week.