

EECS 16A Designing Information Devices and Systems I Homework 13

Spring 2021

This homework is due April 30, 2021, at 23:59.

Self-grades are due May 3, 2021, at 23:59.

Submission Format

Your homework submission should consist of **one** file.

- `hw13.pdf`: A single PDF file that contains all of your answers (any handwritten answers should be scanned) as well as your IPython notebook saved as a PDF.
If you do not attach a PDF “printout” of your IPython notebook, you will not receive credit for problems that involve coding. Make sure that your results and your plots are visible. Assign the IPython printout to the correct problem(s) on Gradescope.

Submit the file to the appropriate assignment on Gradescope.


1. Reading Assignment

Review Note 23 (Least Squares).

2. Noise Canceling Headphones

Learning Goal: This problem will employ least squares method to minimize the effect of additive noise through headphones.

In this problem, we will explore a common design for noise cancellation using noise-canceling headphones as an example application. We will work with the model shown in the figure below.



A music signal is generated at a speaker and transmitted to the listener's ear. If there is noise in the environment (*e.g.* other people's voices, a train going by), this noise signal will be superimposed (*i.e.* added) on the music signal and the listener will hear both. In order to cancel the noise, we will try **to record the noise and subtract it directly from the transmitted signal** with the hope that we can achieve perfect cancellation of everything but the music. Since our system is imperfect, we'll have to solve a **least squares problem**.

The gain blocks marked by γ (**Greek "gamma"**) represent **scalar multiplication**, and we will assume that they can take on any real number, positive or negative.

Assume we have an **additive noise signal** noted by \vec{n} , which we want to cancel. This represents the extra noise coming from the train.

$$\vec{n} = \begin{bmatrix} n_1 \\ n_2 \\ n_3 \\ n_4 \\ n_5 \end{bmatrix}.$$

Let us assume the music signal is represented by

$$\vec{m} = \begin{bmatrix} m_1 \\ m_2 \\ m_3 \\ m_4 \\ m_5 \end{bmatrix}.$$

In absence of any noise cancellation, the user will hear:

$$\vec{s} = \vec{m} + \vec{n}.$$

However, we want to cancel this noise. For this we use three microphones to record this noise, Mic A, Mic B, and Mic C. However, they cannot perfectly record the noise \vec{n} and have erroneous recordings. Let \vec{r}_A , \vec{r}_B , and \vec{r}_C represent the **noise that each microphone picks up**:

$$\vec{r}_A = \begin{bmatrix} a_1 \\ a_2 \\ a_3 \\ a_4 \\ a_5 \end{bmatrix}, \vec{r}_B = \begin{bmatrix} b_1 \\ b_2 \\ b_3 \\ b_4 \\ b_5 \end{bmatrix}, \vec{r}_C = \begin{bmatrix} c_1 \\ c_2 \\ c_3 \\ c_4 \\ c_5 \end{bmatrix}.$$

We can arrange the recordings into a matrix \mathbf{R} and the microphone gains, γ , into a vector $\vec{\gamma}$ to get:

$$\mathbf{R} = [\vec{r}_A \quad \vec{r}_B \quad \vec{r}_C] = \begin{bmatrix} a_1 & b_1 & c_1 \\ a_2 & b_2 & c_2 \\ a_3 & b_3 & c_3 \\ a_4 & b_4 & c_4 \\ a_5 & b_5 & c_5 \end{bmatrix}, \vec{\gamma} = \begin{bmatrix} \gamma_A \\ \gamma_B \\ \gamma_C \end{bmatrix}.$$

We want to choose a $\vec{\gamma}$ such that we minimize the impact of the noise.

The listener's ear will hear a total signal of:

$$\vec{s} = \vec{m} + \mathbf{R}\vec{\gamma} + \vec{n},$$

$$\begin{bmatrix} s_1 \\ s_2 \\ s_3 \\ s_4 \\ s_5 \end{bmatrix} = \begin{bmatrix} m_1 \\ m_2 \\ m_3 \\ m_4 \\ m_5 \end{bmatrix} + \begin{bmatrix} a_1 & b_1 & c_1 \\ a_2 & b_2 & c_2 \\ a_3 & b_3 & c_3 \\ a_4 & b_4 & c_4 \\ a_5 & b_5 & c_5 \end{bmatrix} \begin{bmatrix} \gamma_A \\ \gamma_B \\ \gamma_C \end{bmatrix} + \begin{bmatrix} n_1 \\ n_2 \\ n_3 \\ n_4 \\ n_5 \end{bmatrix}.$$

This problem focuses on how to choose $\vec{\gamma}$ so as to have the signal \vec{s} be as close to \vec{m} as possible.

- (a) Ideally, we would want to have a signal at the ear that matches the original music signal \vec{m} perfectly. In reality, this is not possible, so we will aim to minimize the effect of the noise. What quantity would we need to minimize to make sure this happens? Write your answer in terms of the matrix \mathbf{R} , the vector of mic gains $\vec{\gamma}$, and the noise vector \vec{n} . *Hint: Your answer should be the norm of something.*
- (b) We can solve the noise minimization problem by the least squares method. In effect, if we have a problem, $\min_{\vec{x}} \|\mathbf{A}\vec{x} - \vec{b}\|$, then the \vec{x} that solves this problem is,

$$\vec{x} = (\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T \vec{b}. \quad (1)$$

Implement this least squares solution in the IPython Notebook helper function `doLeastSquares`.

Remember that $\min_{\vec{x}} \|\mathbf{A}\vec{x} + \vec{b}\| = \min_{\vec{x}} \|\mathbf{A}\vec{x} - (-\vec{b})\|$.

- (c) In order to verify the least squares function, we want to compute an example by hand. For this part only, we will be using a system with 2 microphones instead of 3 and a signal of length 3. For the given \vec{n} and the recordings, \vec{r}_A , \vec{r}_B , below, find the γ 's that minimize the effect of noise. You must show your work but you may check your answer with the helper function.

$$\vec{n} = \begin{bmatrix} 0.20 \\ 0.10 \\ 0.50 \end{bmatrix}, \vec{r}_A = \begin{bmatrix} 0.50 \\ 0 \\ 0.50 \end{bmatrix}, \vec{r}_B = \begin{bmatrix} 0 \\ 0.10 \\ 0 \end{bmatrix}.$$

- (d) For the given \vec{n} and the recordings, \vec{r}_A , \vec{r}_B , \vec{r}_C , below, find the γ 's that minimize the effect of noise. Use the helper function `doLeastSquares` you made to solve this.

$$\vec{n} = \begin{bmatrix} 0.10 \\ 0.37 \\ -0.45 \\ 0.068 \\ 0.036 \end{bmatrix}, \vec{r}_A = \begin{bmatrix} 0 \\ 0.11 \\ -0.31 \\ -0.012 \\ -0.018 \end{bmatrix}, \vec{r}_B = \begin{bmatrix} 0 \\ 0.22 \\ -0.20 \\ 0.080 \\ 0.056 \end{bmatrix}, \vec{r}_C = \begin{bmatrix} 0 \\ 0.37 \\ -0.44 \\ 0.065 \\ 0.038 \end{bmatrix}.$$

The next few questions can be answered in the IPython notebook by running the associated cells.

- (e) Follow the instructions in the IPython notebook to load a **music signal** and some **noise signals**. Listen to the music signal `music_y` and the two noise signals `noise1_y` and `noise2_y`. Which ones are full of static and which ones are not?
Also listen to the **noisy music signal** `noisyMusic`, generated by adding the music signal `music_y` and the first noise signal `noise1_y`. What do you hear?
- (f) For this part we assume the noise signal \vec{n} is given by `noise1_y` only. \mathbf{R} , i.e. the matrix of recorded noise from 3 microphones, is simulated with the `recordAmbientNoise` function. Calculate a vector $\vec{\gamma}$ that minimizes the effect of noise using `doLeastSquares`.

- (g) Use your answers from the previous part to find the **noise-cancellation signal** $\mathbf{R}\tilde{\gamma}$. Add this noise-cancellation signal ($\mathbf{R}\tilde{\gamma}$), the music signal \tilde{m} (`music_y`), and the noise signal \tilde{n} (`noise1_y`), in order to generate a **noise-cancelled signal**. Play the noisy signal from part (e) and the noise-cancelled signal. Can you hear a difference?
- (h) **[Optional]** Try adding the other noise signal `noise2_y` to the music signal to generate a new noisy signal. Don't re-calculate new values for $\tilde{\gamma}$ i.e. don't solve the least squares problem again. Recalculate \mathbf{R} and find the new noise-cancellation signal $\mathbf{R}\tilde{\gamma}$ using $\tilde{\gamma}$ from part (g). Add this noise-cancellation signal to your noisy signal. Comment on the quality of the resulting noise-cancelled signal. Is it perfect or are there artifacts?

Note: We want to find out how the $\tilde{\gamma}$ we calculated work for any type of noise.

3. Image Analysis

Learning Goal: This problem introduces a method of fitting a non-linear model through a set of measured data points using the least squares method.

Applications in medical imaging often require an analysis of images based on the image's pixels. For instance, we might want to count the number of cells in a given biological sample. One way to do this is to take a picture of the cells and use the pixels to determine their locations and how many there are. Automatic detection of shape is useful in image classification as well (e.g. consider a robot trying to find out autonomously where a mug is in its field of vision).

Let us focus back on the medical imaging scenario. You are interested in finding the exact position and shape of a cell in an image. You will do this by finding the **equation of the circle or ellipse** that bounds the cell relative to a given coordinate system in the image. Your collaborator uses edge detection techniques to find a **bunch of points that are approximately along the edge of the cell**. We assume that the origin of the coordinate system is in the center of the image with standard axes (x, y) and your collaborator gives you the following points that approximately bound the cell:

$(0.3, -0.69), (0.5, 0.87), (0.9, -0.86), (1, 0.88), (1.2, -0.82), (1.5, 0.64), (1.8, 0)$.

Recall that an equation of the form

$$a_1x^2 + b_1xy + c_1y^2 + d_1x + e_1y = 1$$

can be used to represent an ellipse (if $b_1^2 - 4a_1c_1 < 0$), and an equation of the form

$$a_1(x^2 + y^2) + d_1x + e_1y = 1$$

is a circle if $d_1^2 + e_1^2 + 4a_1 > 0$. Notice that the circle has fewer parameters.

You don't need to consider these constraints in your least squares setup, but you are encouraged to check whether your least squares solutions satisfy these constraints.

- (a) How can you find the equation of a *circle* that surrounds the cell by fitting the data points? First, provide a setup and formulate a minimization problem to do this, i.e. a least squares problem minimizing the squared error $\|\mathbf{A}\vec{v} - \vec{b}\|^2$, where you attempt to find the **unknown coefficients** a_1 , d_1 , and e_1 from your data points. Here your unknown vector $\vec{v} = \begin{bmatrix} a_1 \\ d_1 \\ e_1 \end{bmatrix}$. *Hint: The quantities $(x^2 + y^2)$, x , and y can be thought of as known values calculated from your data points.*

You do not need to simplify the numerical values for the matrix elements; just writing out the matrix with numerical expressions will suffice.

- (b) How can you find the equation of an ellipse (instead of a circle) that surrounds the cell? Provide a setup and formulate a minimization problem similar to that in part (a). Now the unknown vector \vec{v} will be different from the earlier parts.

You do not need to simplify the numerical values for the matrix elements; just writing out the numerical expressions will suffice.

- (c) In the IPython notebook, write a short program that uses least-squares to fit a circle to the given points. A helper function `plot_circle` is provided. What is $\frac{\|\vec{e}\|}{N}$, where $\vec{e} = \mathbf{A}\vec{v} - \vec{b}$ and N is the number of data points? Plot your points and the best fit circle in IPython.
- (d) In the IPython notebook, write a short program that uses least-squares to fit an ellipse to the given points. A helper function `plot_ellipse` is provided. What is $\frac{\|\vec{e}\|}{N}$, where $\vec{e} = \mathbf{A}\vec{v} - \vec{b}$ and N is the number of data points? Now the unknown vector \vec{v} will be different from the earlier parts. Plot your points and the best fit ellipse in IPython. How does this error compare to the one in the previous subpart? Which technique is better, and why?

4. Trilateration With Noise!

Learning Goal: This problem will help to understand how noise affects the accuracy of trilateration and consistency of the corresponding system of equations.

In this question, we will explore how various types of noise affect the quality of triangulating a point on the 2D plane to see when trilateration works well and when it does not.

First, we will remind ourselves of the fundamental equations underlying trilateration.

- (a) There are four beacons at the known coordinates $(x_1, y_1), (x_2, y_2), (x_3, y_3), (x_4, y_4)$. You are located at some unknown coordinate (x, y) that you want to determine. The distance between your location and each of the four beacons are d_1 through d_4 , respectively. Write down one equation for each beacon that relates the coordinates to the distances using the Pythagorean Theorem.
- (b) Unfortunately, the above system of equations is nonlinear, so we can't use least squares or Gaussian Elimination to solve it. We will use the technique discussed in lecture to obtain a system of linear equations. In particular, we can subtract the first of the above equations (involving x_1, y_1 and d_1) from the other three to obtain three linear equations (cancel out the nonlinear terms). Write down these three linear equations.
- (c) Combine the three equations in the above system into a single matrix equation of the form

$$\mathbf{A} \begin{bmatrix} x \\ y \end{bmatrix} = \vec{b}.$$

- (d) Now, go to the IPython notebook. In the notebook we are given three possible sets of measurements for the distances of each beacon from the receiver:
- `ideal_distances`: the ideal set of measurements, the true distances of our receiver to the beacons. $d_1 = d_2 = d_3 = d_4 = 5$.
 - `imperfect_distances`: imperfect measurements. $d_1 = 5.5, d_2 = 4.5, d_3 = 5, d_4 = 5$.
 - `one_bad_distances`: mostly perfect measurements, but d_1 is a very bad measurement. $d_1 = 6.5$ and $d_2 = d_3 = d_4 = 5$.

Plot the graph illustrating the case when the receiver has received `ideal_distances` and visually solve for the position of the observer (x, y) . What is the coordinate?

- (e) You will now set up the above linear system using IPython. Fill in each element of the matrix \mathbf{A} that you found in part (c).
- (f) Similarly, fill in the entries of \vec{b} from part (c) in the `make_b` function.
- (g) Now, you should be able to plot the estimated position of (x,y) using the supplied code for the `ideal_distances` observations. Modify the code to estimate (x,y) for `imperfect_distances` and `one_bad_distances`, and comment on the results.

In particular, for `one_bad_distances` would you intuitively have chosen the same point that our trilateration solution did knowing that only one measurement was bad?

5. Constrained Least Squares Optimization

In this problem, we will guide you through solving the following optimization problem:

Consider a matrix $\mathbf{A} \in \mathbb{R}^{M \times N}$ where $M > N$ and all N columns are linearly independent. Determine a unit vector \vec{x} that minimizes $\|\mathbf{A}\vec{x}\|$, where $\|\cdot\|$ denotes the norm—that is,

$$\|\mathbf{A}\vec{x}\|^2 \triangleq \langle \mathbf{A}\vec{x}, \mathbf{A}\vec{x} \rangle = (\mathbf{A}\vec{x})^T \mathbf{A}\vec{x} = \vec{x}^T \mathbf{A}^T \mathbf{A} \vec{x}.$$

This is equivalent to solving the following optimization problem:

$$\min_{\vec{x}} \|\mathbf{A}\vec{x}\|^2 \quad \text{subject to the constraint} \quad \|\vec{x}\|^2 = 1.$$

This task may *seem* like solving a standard least squares problem $\mathbf{A}\vec{x} = \vec{b}$ where $\vec{b} = \vec{0}$, but it is different. As an example, notice $\vec{x} = \vec{0}$ is *not* a valid solution to our problem because the norm of the zero vector does not equal one. Our optimization problem is a least squares problem with a constraint—hence the term *constrained least squares optimization*. The constraint can be visualized as limiting the vector \vec{x} to lie on a unit circle (radius of the circle is one) if $N = 2$ and on a unit sphere if $N = 3$.

Let $(\lambda_1, \vec{v}_1), \dots, (\lambda_N, \vec{v}_N)$ denote the eigenpairs (i.e., eigenvalue/eigenvector pairs) of $\mathbf{A}^T \mathbf{A}$. Assume that the eigenvalues are all real, distinct and indexed in an descending fashion—that is,

$$\lambda_1 > \dots > \lambda_N.$$

Assume, too, that each eigenvector has been normalized to have unit length—that is, $\|\vec{v}_k\| = 1$ for all $k \in \{1, \dots, N\}$.

- (a) Show that $\lambda_N > 0$, i.e. all the eigenvalues are strictly positive.
Hint: Consider $\|\mathbf{A}\vec{v}_N\|^2$. Write $\|\mathbf{A}\vec{v}_N\|^2$ in the matrix multiplication form, and use the eigenpair equation $\mathbf{A}^T \mathbf{A} \vec{v}_N = \lambda_N \vec{v}_N$.
- (b) Consider two eigenpairs (λ_k, \vec{v}_k) and $(\lambda_\ell, \vec{v}_\ell)$ corresponding to distinct eigenvalues of $\mathbf{A}^T \mathbf{A}$ —that is, $\lambda_k \neq \lambda_\ell$. Prove that the corresponding eigenvectors \vec{v}_k and \vec{v}_ℓ are orthogonal: $\vec{v}_k \perp \vec{v}_\ell$.

To help you get started, consider the two equations

$$\mathbf{A}^T \mathbf{A} \vec{v}_k = \lambda_k \vec{v}_k \tag{2}$$

and

$$\vec{v}_\ell^T \mathbf{A}^T \mathbf{A} = \lambda_\ell \vec{v}_\ell^T. \tag{3}$$

The second equation can be derived by taking the transpose of both sides in the eigenvalue equation $\mathbf{A}^T \mathbf{A} \vec{v}_\ell = \lambda_\ell \vec{v}_\ell$. Premultiply Equation 2 with \vec{v}_ℓ^T , postmultiply Equation 3 with \vec{v}_k , compare the two, and explain how one may then infer that \vec{v}_k and \vec{v}_ℓ are orthogonal, i.e. $\langle \vec{v}_k, \vec{v}_\ell \rangle = 0$.

Premultiplication by a vector \vec{x} means multiplying an expression by \vec{x} on the left. For example, pre-multiplying the matrix \mathbf{A} by \vec{x} gives $\vec{x}\mathbf{A}$. Postmultiplication means multiplying on the right. Remember that in general matrix-vector multiplication is not commutative, so those operations are not identical.

- (c) The results of part (b) imply that the N eigenvectors of $\mathbf{A}^T\mathbf{A}$ are mutually orthogonal. A basis formed by vectors that are both (1) mutually orthogonal and (2) have unit length is called an orthonormal basis. Since the eigenvalues of $\mathbf{A}^T\mathbf{A}$ are distinct and have a norm of one, the eigenvectors form an orthonormal basis in \mathbb{R}^N . This means that we can express an arbitrary vector $\vec{x} \in \mathbb{R}^N$ as a linear combination of the eigenvectors $\vec{v}_1, \dots, \vec{v}_N$, as follows:

$$\vec{x} = \sum_{n=1}^N \alpha_n \vec{v}_n.$$

- i. Determine the n^{th} coefficient α_n in terms of \vec{x} and one or more of the eigenvectors $\vec{v}_1, \dots, \vec{v}_N$.
- ii. Suppose \vec{x} is a unit-length vector (i.e., a unit vector) in \mathbb{R}^N . Show that

$$\sum_{n=1}^N \alpha_n^2 = 1$$

where the α_n 's are the coefficients of \vec{x} in the basis defined earlier.

- (d) Now express $\|\mathbf{A}\vec{x}\|^2$ in terms of $\{\alpha_1, \alpha_2 \dots \alpha_N\}$, $\{\lambda_1, \lambda_2 \dots \lambda_N\}$, and $\{\vec{v}_1, \vec{v}_2 \dots \vec{v}_N\}$, and find an expression for \vec{x} such that $\|\mathbf{A}\vec{x}\|^2$ is minimized. Do *not* use any tool from calculus to solve this problem, so avoid differentiation.

Hint: After expressing $\|\mathbf{A}\vec{x}\|^2$ in terms of $\{\alpha_1, \alpha_2 \dots \alpha_N\}$, $\{\lambda_1, \lambda_2 \dots \lambda_N\}$, and $\{\vec{v}_1, \vec{v}_2 \dots \vec{v}_N\}$, which variable are you minimizing over?

6. (OPTIONAL, PRACTICE) Labeling Patients Using Gene Expression Data

Learning Goal: This problem aims to design a predictive model using the least squares method on a set of training data and test the efficacy of the model using a set of test data.

Least squares techniques are useful for many different kinds of prediction problems. Numerous researchers have extensively further developed the core ideas that we have learned in class. These ideas are commonly used in machine learning for finance, healthcare, advertising, image processing, and many other fields. Here, we'll explore how least squares can be used for classification of data in a medical context.

Gene expression data of patients, along with other factors such as height, weight, age, and family history, are often used to predict the likelihood that a patient might develop a certain disease. This data can be combined into a vector that describes each patient. This vector is often referred to as a feature vector.

Many scientific studies examine mice to understand how gene expression relates to diabetes in humans. Studies have shown that the expression of the *tomosinB* and *tsA* genes are correlated to the onset of diabetes in mice. How can we predict whether or not a mouse will develop diabetes based on data about this expression as well as other factors of the mouse? We will use some (fake) data to explore this.

We are given feature vectors for each mouse as:

$$\begin{bmatrix} \text{age} \\ \text{weight} \\ \text{tomosinB} \\ \text{tsA} \\ \text{chnA} \end{bmatrix}$$

Age and weight in the vector above are represented by real numbers, while the presence or absence of the expression of the genes *tomosinB*, *tsA*, and *chnA* is captured by $+1$ and -1 respectively. For example, the vector $[2 \ 20 \ 1 \ -1 \ -1]^T$ means a 2 month old mouse, that weighs 20 grams, expresses the genes *tomosinB*, but not *tsA* or *chnA*.

We would like the following expression to be **positive if the mouse has diabetes and negative if the mouse does not have diabetes**:

$$f(\text{age, weight, tomosinB, tsA, chnA}) = \alpha_1(\text{age}) + \alpha_2(\text{weight}) + \alpha_3(\text{tomosinB}) + \alpha_4(\text{tsA}) + \alpha_5(\text{chnA}). \quad (4)$$

- (a) We wish to set up a linear model for the problem in the format $\mathbf{A}\vec{x} = \vec{b}$. Here, \vec{b} will be a vector with $+1, -1$ entries where a 1 represents that the mouse is diabetic and -1 represents that the mouse is not diabetic. The feature vectors of each mouse will be included in each row of the matrix \mathbf{A} . For example, if the first row of \mathbf{A} contains the data of mouse #1 and the first entry of \vec{b} is $+1$, that means the mouse #1 has diabetes.

Set up the problem by writing \mathbf{A} , \vec{x} , and \vec{b} in terms of the variables in the feature vectors, α_i , and any other variables you define. What are your unknowns?

- (b) Training data is data that is used to develop your model. The matrix \mathbf{A} and vector \vec{b} , provided in `gene_data_train.npy` and `diabetes_train.npy` respectively, represent the (fake) *training* data. Use the data to find the optimal model parameters $\alpha_1, \dots, \alpha_5$ for the given data set. Find the optimal parameter values using least squares method and the provided IPython notebook.

- (c) Now it is time to use the model you have developed to make some predictions! It is interesting to note here that we are not looking for a real number to model whether each mouse has diabetes or not; we are looking for a **binary label**. Therefore, we will use the **sign of the expression from Equation (4)** to assign a ± 1 value to each mouse. Each mouse characteristics are represented by each row of `gene_data_test.npy`.

Predict whether each mouse with the characteristics in the *test* data set `gene_data_test.npy` will get diabetes. There are four mice/ rows in the test data set. Calculate the ± 1 prediction vector \vec{b}_{calc} that shows the prediction whether a mouse will be diabetic or not. Observe the ± 1 vector \vec{b}_{test} from `diabetes_test.npy` that indicates whether or not the mice *actually* have diabetes.

What is the prediction accuracy (number of correct predictions divided by total number of predictions) of your model?

7. (OPTIONAL, PRACTICE) How Much Is Too Much?

When discussing circuits in this course, we only talked about resistor I - V curves. There are many other devices that can be found in nature that do not have linear I - V relations. Instead, I is some general function of V , that is $I = f(V)$. Often times, the function describing the I - V relationship is not known beforehand. The function f is assumed to be a polynomial, and the parameters of f (the coefficients for every power of V) are computed using least squares.

Throughout this problem, we are provided with \vec{x} , a set of voltage measurements, and \vec{y} , a set of current measurements.

- (a) Let's first try to model a resistor I - V curve. Run the code in the attached IPython notebook. What is the degree of the polynomial that fits an ideal resistor I - V curve? Play around the with degree in the IPython notebook and observe the best fit polynomial's shape. Is the noise influencing the higher degree polynomials?

- (b) The attached IPython notebook provides functions `data_matrix`, `leastSquares` that allow you to fit polynomials of different degrees to the data provided. We also provide a function `cost` that computes the squared error of the fit. In the attached IPython notebook, plot the cost of various degree polynomials fitting to the measured I - V data points for a resistor using the given `cost` function. The `cost` function returns $\|\vec{y} - \mathbf{A}\vec{f}\|^2$, i.e. the squared magnitude of the error vector.

$$\mathbf{A} = \begin{bmatrix} 1 & x_1 & x_1^2 & \dots \\ 1 & x_2 & x_2^2 & \dots \\ 1 & x_3 & x_3^2 & \dots \\ \vdots & & & \ddots \end{bmatrix}$$

As seen above \mathbf{A} is the appropriately sized matrix containing powers of the elements of \vec{x} and the vector \vec{f} contains entries f_n that are the coefficients for the n th power of the elements of \vec{x} . Comment on the shape of the “Cost vs. Degree” graph. Do we want to choose a best fit polynomial of degree greater than one if the cost is lower than the polynomial of degree one? Should we choose the degree of the polynomial based on this graph? This question is meant to make you think, do not worry too much about getting a precise right answer here.

8. Homework Process and Study Group

Who did you work with on this homework? List names and student ID’s. (In case you met people at homework party or in office hours, you can also just describe the group.) How did you work on this homework? If you worked in your study group, explain what role each student played for the meetings this week.